PRODUCT REQUIREMENTS DOCUMENT

prd-forge

A domain expert, or a product manager who speaks for one, with no product management training, reaches an approved PRD in which every section is graded with no blocking defect and every buildability gap is closed or owned before anyone estimates.

VERSION	0.13.0
STATUS	draft
OWNER	Richard Hightower
DOMAIN	product management tooling
CREATED	2026-10-23
UPDATED	2026-11-01
GOLDEN THREAD	J-001
STAKEHOLDERS	Product managers who are, or speak for, the domain expert; Engineering leads who build from the resulting PRD; Sponsors who approve the PRD for a release. For this PRD the approver is Richard Hightower, the owner of the place where it is kept, named under FR-021.; Delivery leads who run story-mapping workshops

Executive summary

PROBLEM	A domain expert knows exactly what their product must do, but what they can write down is a feature list with no named users, no pass-or-fail criteria and the technology already chosen. The team cannot build from it, and the expert is asked the same questions again by each engineer who tries.
USERS	Product managers who are, or speak for, the domain expert, the engineering leads who build from what they write, the sponsors who approve it, and the delivery leads who derive a backlog from it.
WHY NOW	In October 2025, 98% of 117 product managers surveyed said they use AI at work, and only 39% had received job-specific AI training (General Assembly, 6 Nov 2025). A PRD that reads well but cannot be built from costs more than no PRD at all, because it is believed.
BET	A domain expert, or a product manager who speaks for one, with no product management training, reaches an approved PRD in which every section is graded with no blocking defect and every buildability gap is closed or owned before anyone estimates.
OBJECTIVE	O-001
MEASURE	Share of the 5 pilot PRDs that pass the buildability review on first submission, where a pass is the verdict `buildable`, or `buildable with gaps` with every gap closed or owned. The pilot PRDs are the first 5 written with release 1 by people other than the product's author, between REL-001, on 2027-03-31, and 2027-09-30. Target: 80%, at least 4 of 5, by 2027-09-30, provisional until Q-001 sets the baseline, by 2027-02-28.

Contents

§	SECTION	STATUS	GRADE	BLOCKING
1	Problem and opportunity	draft	83/100	—
2	Target customer and market	draft	75/100	—
3	Business objectives and success metrics	draft	88/100	—
4	Users and personas	draft	82/100	—
5	User journeys	draft	94/100	—
5	Scenarios	draft	89/100	—
6	Scope	draft	88/100	—
7	Functional requirements	draft	88/100	—
7	Wireframes	draft	91/100	—
8	Non-functional requirements	draft	89/100	—

§	SECTION	STATUS	GRADE	BLOCKING
9	Constraints	draft	95/100	—
10	Decision rights and trust rules	draft	84/100	—
11	Assumptions, risks and dependencies	draft	96/100	—
12	Open questions	draft	88/100	—
13	Release plan	draft	83/100	—
13	Roadmap	draft	92/100	—
14	Glossary	draft	85/100	—
15	Change log and approvals	draft	75/100	—

Approvers

- Richard Hightower, Sponsor

In scope

- IN-001: State the idea and the domain's vocabulary once, and see every section of the set and its state
- IN-002: Be taught each section before its interview, be asked one or two questions at a time with follow-ups, end a section early with its gaps recorded, and stop being asked once the section would meet the gate
- IN-003: Record no invented domain fact, and convert a stated technology into the need beneath it, kept for the engineering lead
- IN-004: Read the draft as prose with its grade and every blocking defect, refute a finding, move to any section, refresh to see a change, and be told when the set states one item two ways
- IN-005: Answer one question about the weakest part, see the section re-graded against the same rubric, and resume after a break without repeating an answer
- IN-006: Get a buildability review of the whole set, read as an engineer would, and see which release-line stages no requirement covers and which requirements cannot be tested
- IN-007: Let the expert name each approver, give the sponsor one document with the summary, the scope, the grades, and the open questions, bring the sponsor's comments on an exported copy back into the set with a response to each, and let the expert record the sponsor's decision, given outside the product, with any conditions, against a named version

Out of scope

- OUT-001: Deriving features, epics and stories from the approved PRD - journey J-002
- OUT-002: More than one expert working on the same set at the same time
- OUT-004: An engineering lead reading the set and recording their own confirmation that it can be built from
- OUT-003: Any view of the document set other than reading the documents themselves and stepping through them beside the conversation
- OUT-005: A reviewer commenting on a section inside the set, not on an exported copy

Undecided

- Q-001: What share of specifications reaching a development team today pass a buildability review on first submission? Decider: Richard Hightower. By: 2027-02-28.
- Q-002: How many working days does it take today from the first question put to the domain expert about a product to an approved specification? Decider: Richard Hightower. By: 2027-02-28.
- Q-003: What share of `must` functional requirements in today's specifications can an engineer who was not in the interview write a test from? Decider: Richard Hightower. By: 2027-02-28.
- Q-004: How many organizations are in each segment in section 2? Decider: Richard Hightower. By: 2027-06-30.
- Q-005: Given what three sponsors say about a set shown with and without grades, does FR-013 stay `must`, with each section's grade shown to the sponsor? Decider: Richard Hightower. By: 2027-01-31.
- Q-006: How many sections does an expert complete in one sitting before stopping? Decider: Richard Hightower. By: 2027-02-28.
- Q-007: Does writing the requirements send an expert back to the journeys? Decider: Richard Hightower. By: 2027-02-28.
- Q-008: Does an engineering lead read a PRD that names no technology as disciplined or as incomplete? Decider: Richard Hightower. By: 2027-02-28.
- Q-009: Which version of the rubric and the eleven derivability checks does release 1 grade against? Decider: Richard Hightower. By: 2027-01-31.
- Q-010: Must an answer the expert gave survive the end of the session before its section is drafted, given A-005, that domain content stays in the expert's own session and in the documents? Decider: Richard Hightower. By: 2027-01-31.
- Q-012: Which segment in section 2 converts first? Decider: Richard Hightower. By: 2027-06-30.
- Q-013: What day rate does a buyer pay to contract a product manager for the writing of one set? Decider: Richard Hightower. By: 2027-06-30.
- Q-014: Given how experts describe a loop-back in D-002's sittings, does ST-004 keep its step that sends the expert back to an earlier section? Decider: Richard Hightower. By: 2027-02-28.
- Q-015: Who owns the rubric and the eleven derivability checks? Decider: Richard Hightower. By: 2027-01-15.
- Q-016: Must a grade show which version of the rubric it was made against? Decider: Richard Hightower. By: 2027-01-31.
- Q-017: Whom does the decider record as owner of NFR-006's accessibility, given that the product's output is read in places the non-goals in section 6 say we do not own? Decider: Richard Hightower. By: 2027-02-28.
- Q-018: What does each problem in section 1 cost, as one counted case per problem? Decider: Richard Hightower. By: 2027-01-31.
- Q-020: Which section do experts stop at most often? Decider: Richard Hightower. By: 2027-02-28.
- Q-021: Whom does the decider record as owner of NFR-012's access control, given that a set is stored where the organization already stores documents? Decider: Richard Hightower. By: 2027-02-28.
- Q-022: If D-002 or D-003 reports after 2027-01-31, does REL-001 move or ship as an experiment? Decider: Richard Hightower. By: 2027-01-31.
- Q-023: Do delivery leads run a story-mapping workshop about once per release and want a backbone agreed before it, as PER-004 assumes? Decider: Richard Hightower. By: 2027-05-31.

Blocking defects

None.

About this document

A domain expert knows exactly what their product must do, but what they can write down is a feature list with no named users, no pass-or-fail criteria and the technology already chosen. The team cannot build from it, and the expert is asked the same questions again by each engineer who tries. The users are product managers who are, or speak for, the domain expert, the engineering leads who build from what they write, the sponsors who approve it, and the delivery leads who derive a backlog from it. In October 2025, 98% of 117 product managers surveyed said they use AI at work, and only 39% had received job-specific AI training (General Assembly, 6 Nov 2025). A PRD that reads well but cannot be built from costs more than no PRD at all, because it is believed.

The bet is one outcome: the expert, with no product management training, reaches an approved PRD in which every section is graded with no blocking defect and every buildability gap is closed or owned before anyone estimates. We will know it worked when at least 4 of the 5 pilot PRDs, 80%, pass the buildability review on first submission by 2027-09-30. The pilot PRDs are the first 5 written with release 1, between REL-001 on 2027-03-31 and 2027-09-30, by people other than the product's author. The target is provisional until Q-001 sets the baseline, by 2027-02-28.

Contents

SECTION	DOCUMENT
1	Problem and opportunity
2	Target customer and market
3	Business objectives and success metrics
4	Users and personas
5	User journeys
5	Scenarios
6	Scope
7	Functional requirements
7	Wireframes
8	Non-functional requirements
9	Constraints
10	Decision rights and trust rules
11	Assumptions, risks and dependencies
12	Open questions
13	Release plan
13	Roadmap

SECTION	DOCUMENT
14	Glossary
15	Change log and approvals

How this document set is numbered

A PRD has **fifteen numbered sections** and this set has **eighteen section documents**, because section 5 holds two of them, journeys and scenarios, section 7 holds two, functional requirements and wireframes, and section 13 holds two, the release plan and the roadmap. The two section 7 files share the prefix `08`. The roadmap is `18-roadmap.md`, but it is read after the release plan, in section 13. From section 5 onward, the file prefix and the section number differ: `07-scope.md` is section 6, and `16-change-log.md` is section 15. **Every cross-reference in this set is by section number, never by file number.** Three independent reviewers misread this, so it is stated here rather than left to be inferred.

A note on this PRD

This is prd-forge's PRD, written with prd-forge's own conventions. The requirements in section 7 describe only what a person does and sees, and section 15 records what each version changed.

SECTION 1

Problem and opportunity

DRAFT 83/100 NO BLOCKING DEFECTS

ID	PROBLEM	WHO
P-001	What the expert writes is a feature list with no user and no testable outcome	Product manager or domain expert
P-002	The steps between requirements are never written, so paths have gaps	Engineering lead
P-003	Buildability is invisible until the team starts building	Sponsor
P-004	The expert does not know what a section is for, so they fill it badly	Product manager or domain expert

The four compound. P-004 means the expert does not know what good looks like; P-001 is what they produce as a result; P-002 is the specific gap a build trips over; and P-003 is why none of it is caught before it is expensive. P-002 is the engineering lead's problem, and the delivery lead, PER-004, meets it too, at the story-mapping workshop.

Evidence note. Each problem now has one counted case, with its source and date. The counts for P-001 to P-003 come from one engagement, this product's own PRD, written and graded with prd-forge. The count for P-004 comes from two public templates. One engagement shows that each problem occurs. It does not show how often. The other evidence items describe a pattern and are not measured. Each why-now claim cites the General Assembly survey of 117 product managers, run in October 2025. No cost has a number yet, and Q-018 asks for one per problem. All three objectives in section 3 still have no baseline. See Q-001, Q-002, and Q-003.

P-001

STATEMENT
A domain expert can describe what their product must do out loud, but what they write down is a feature list: no named user, no outcome anyone could check, and a technology already chosen for reasons they cannot defend.

WHO
Product manager who is, or speaks for, the domain expert

PERSONA
PER-001

EVIDENCE
- Claim: In this PRD at version 0.1.0, the grading pass found 2 of 14 functional requirements, FR-003 and FR-004, with no outcome a person could see. Both were fixed at version 0.2.0; Source: prd-forge repository, commits 5ba59cb (version 0.1.0) and 5d4d936 (version 0.2.0); As of: 2026-10-26
- A team that receives a feature list re-derives the user and the acceptance criteria itself before it can estimate.
- Experts asked to write requirements reach for the words of the software they have used, because that is the only requirements language they have been exposed to.

WHY NOW
In October 2025, 98% of 117 product managers surveyed said they use AI at work (General Assembly, "AI is everywhere in product management, but skills gaps and 'shadow AI' are holding teams back", 6 Nov 2025), which means drafts written with AI help arrive fluent. Fluent and unbuildable is a worse failure than rough and unbuildable, because nobody queries it.

COST OF INACTION
The gap is closed by an engineer guessing, during build, at what the expert meant, and the guess is discovered at review or in production. Q-018 asks for one counted cost.

STATUS QUO
The expert writes a feature list from a template, and an engineer re-derives the user and the criteria before estimating.

DESIRED OUTCOME
Every requirement names a user and an outcome that someone else can check.

P-002

STATEMENT	Requirements are written one at a time, so the steps between them are never written at all, and nobody notices until the team tries to assemble a working path through the product.
WHO	Engineering lead building from the PRD
PERSONA	PER-002
EVIDENCE	• Claim: At version 0.1.0, the buildability review of this PRD failed 4 of 11 derivability checks and named 4 gates an engineer would have to invent; Source: prd-forge repository, the commit messages of 5ba59cb (version 0.1.0) and 5d4d936 (version 0.2.0); As of: 2026-10-26 • A specification with no user journey gives a story-mapping workshop no backbone, so features and epics cannot be derived from it at all. • Gaps surface as mid-iteration questions about what happens between two requirements that both looked complete.
WHY NOW	In the same survey, 44% of product managers use AI for backlog grooming and QA support (General Assembly, "AI is everywhere in product management, but skills gaps and 'shadow AI' are holding teams back", 6 Nov 2025), which means more backlog items are written with AI help, and each one needs the path through the product that it serves.
COST OF INACTION	The release line is drawn across a feature list rather than a journey, and a release ships features that no user can complete anything with. Q-018 asks for one counted cost.
STATUS QUO	Gaps between requirements surface as questions in the middle of an iteration.
DESIRED OUTCOME	Every step a user must take to reach the goal is written down, and every gap is closed or owned before build.

P-003

STATEMENT	Nobody can tell a buildable specification from a readable one until the team has already started building from it.
WHO	Sponsor approving the PRD
PERSONA	PER-003
EVIDENCE	• Claim: At version 0.8.0, this PRD passed every schema and cross-reference check with 0 errors and 0 warnings. The grading pass found 23 blocking defects in 8 of 18 documents; Source: prd-forge repository, commit 6f77bc0 (version 0.8.0); As of: 2026-10-27 • Approval is given on readability and completeness of headings, because those are the only properties an approver can assess unaided. • The first real assessment of buildability happens when an engineer estimates, which is after approval and after the plan was made.
WHY NOW	In the same survey, only 39% of product managers had received comprehensive, job-specific AI training (General Assembly, "AI is everywhere in product management, but skills gaps and 'shadow AI' are holding teams back", 6 Nov 2025), which means 61% of product managers judge AI drafts without that training, and a fluent draft is no evidence that it can be built.
COST OF INACTION	Plans and commitments are made on an approved document whose defects are found by the people asked to deliver against them. Q-018 asks for one counted cost.
STATUS QUO	The sponsor approves on readability, and buildability is first assessed when an engineer estimates.
DESIRED OUTCOME	The sponsor knows, before deciding, which parts of the PRD a team could not build from and what is still undecided.

P-004

STATEMENT	An expert filling in a section of a PRD does not know what the section is for, so they fill it with what they have seen elsewhere: adjectives, feature names, and headings copied from a template.
WHO	Product manager who is, or speaks for, the domain expert
PERSONA	PER-001
EVIDENCE	• Claim: Of 20 sections in two public PRD templates, all 20 say what to write, and 0 say what makes an entry good or how a reader checks it. The Aha! Agile PRD template has 6 sections, and the Product School PRD template has 14; Source: Aha! Agile PRD template, https://www.aha.io/roadmapping/guide/requirements-management/what-is-a-good-product-requirements-document-template, and Product School PRD template, https://productschool.com/blog/product-strategy/product-template-requirements-document-prd; As of: 2026-10-26 • The craft exists in books and in experienced product managers, and neither is at hand at the moment the expert is typing.
WHY NOW	In the same survey, 47% of product managers want to prototype and validate products without engineering, and 38% can (General Assembly, "AI is everywhere in product management, but skills gaps and 'shadow AI' are holding teams back", 6 Nov 2025), which means more of the specifying falls on the person with the domain knowledge, with no specialist between them and the page.
COST OF INACTION	Each expert reinvents the form badly, and the organization's PRDs cannot be compared to one another or improved as a body of work. Q-018 asks for one counted cost.
STATUS QUO	The expert fills a template whose headings say what to write but not what makes an entry good.
DESIRED OUTCOME	The expert knows what each section is for when they write it.

SECTION 2

Target customer and market

DRAFT 75/100 NO BLOCKING DEFECTS

ID	SEGMENT	BUYER	USER
S-001	Consultancies and advisory firms turning client domain knowledge into PRDs	Partner or director of the consultancy, the budget holder, as the sponsor	The consultant with the domain knowledge, as the product manager
S-002	Regulated-industry organizations where practitioners specify their own internal software	Director of product or transformation, the budget holder, as the sponsor	The practitioner, as the product manager or domain expert

Neither segment is sized yet. Both count organizations, and Q-004 asks for the count and its source.

COMPETITOR	KIND
ChatPRD	Direct
A blank PRD template	Indirect
A general-purpose assistant conversation	Indirect
A product manager contracted without the domain	Indirect
A product manager who already holds the domain	Indirect
The status quo	Status quo

Product teams entering a domain they do not know are a non-goal in section 6. They already employ the product skill and want the domain, which this product does not supply.

Why we win

The expert confirmed each sponsor's words on 2026-10-27.

- **S-001, the consultancy partner:** "Before we hand the PRD to the client's engineers, I can see which parts they could not build from, and my senior people bill fewer hours fixing it later." The partner weighs it against the consultant's billable hours.
- **S-002, the director:** "Before I approve the budget, I can see which parts of the PRD my team could not build from." The director weighs it against contracting a product manager without the domain, whose day rate is Q-013.

In S-002 the organization works under regulation, so its PRDs may cite regulations. The set states no regulation the expert did not supply (FR-006, TR-005), which is what lets such a PRD be circulated.

Honest weakness. Against a product manager who already holds the domain, we have no quality edge, only availability. Any cost edge is unproven until Q-013. A skeptical buyer is right to test that.

Segments

S-001 · Consultancies and advisory firms turning client domain knowledge into PRDs

BUYER	Partner or director of the consultancy, the budget holder, as the sponsor
BUYER PERSONA	PER-003
USER	The consultant with the domain knowledge, as the product manager
USER PERSONAS	• PER-001
SIZE	Not yet measured; see Q-004.
SIZE UNIT	organizations
WILLINGNESS TO PAY	Already pays for the consultant's time; the comparison is billable hours saved, not a software budget.
ADOPTION BARRIERS	• We may trust a machine's grade less than our own partner review (Q-005). • We must approve client knowledge going into the assistant session our consultants already use before anyone works that way (A-004).

S-002 · Regulated-industry organizations where practitioners specify their own internal software

BUYER	Director of product or transformation, the budget holder, as the sponsor
BUYER PERSONA	PER-003
USER	The practitioner, as the product manager or domain expert
USER PERSONAS	• PER-001
SIZE	Not yet measured; see Q-004.
SIZE UNIT	organizations
WILLINGNESS TO PAY	Compared against contracting a product manager for the project. The day rate is Q-013.
ADOPTION BARRIERS	• We may trust a machine's grade less than our own review (Q-005). • Our engineering leads may send back a PRD that names no technology (A-003). • We cannot circulate a PRD that cites a regulation our own expert did not supply (FR-006).

Competitors

ChatPRD, an AI tool that drafts PRDs from a chat

KIND	direct
OFFER	A structured PRD drafted fast from little input, from a chat. Source: ChatPRD docs, https://www.chatprd.ai/docs/writing-your-first-prd, accessed 2026-10-27.
OUR EDGE	Before the budget is committed, the sponsor sees in one document each section's grade and every blocking defect, with the section it is in (FR-013). The sponsor can open any set written with this product and look for the same list in a ChatPRD draft.

A blank PRD template in the team's document tool

KIND	indirect
OFFER	The section headings, immediately, free, in a tool everyone already has.
OUR EDGE	When the PRD is done, each fact the expert could not supply is listed with a named person and a date to settle it, and the PRD states no figure or citation that nobody can stand behind.

A general-purpose assistant conversation with no structure

KIND	indirect
OFFER	Fluent prose on any section, instantly, with no setup at all.
OUR EDGE	Before buying, the buyer can open any set written with this product, such as this one, and see which sections are not ready, with the quoted reason for each. A general-purpose conversation shows no such list. Writing that reads well is the failure P-003 names.

Contracting a product manager without the domain for the project

KIND	indirect
OFFER	Genuine product judgement, and someone accountable for the outcome.
OUR EDGE	The buyer gets a PRD whose domain facts come from the person who holds them, rather than facts filled in by someone without the domain. Any cost advantage is unproven until Q-013 gives the day rate. The buyer can check the PRD for each fact the expert could not supply: each one has a named decider and a date.

Contracting a product manager who already holds the domain

KIND	indirect
OFFER	Product judgement and the domain in one person, accountable for the outcome.
OUR EDGE	No quality edge. The buyer gets availability only: the work starts when their own domain expert can start, with no search for a person who holds both. Any cost advantage is unproven until Q-013 gives the day rate.

The status quo: a smart practitioner, a document, and a willing engineer

KIND	status-quo
OFFER	No tool, no cost, no adoption, and it does sometimes work.
OUR EDGE	At handover, an engineer who was not in the room gets each fact the PRD does not yet know, with a named person and a date to settle it, so they can estimate without first interviewing the author.

SECTION 3

Business objectives and success metrics

DRAFT 88/100 NO BLOCKING DEFECTS

ID	OBJECTIVE	METRIC	BASELINE	TARGET
O-001	Buildable without re-derivation	Share of the 5 pilot PRDs that pass the buildability review first time	unknown (Q-001)	80%, at least 4 of 5
O-002	Days to an approved PRD	Median working days, first interview question to approval, for pilot PRDs with one journey	unknown (Q-002)	under 5
O-003	Requirements testable as written	Share of `must` requirements an outside engineer can test	unknown (Q-003)	95%

Why there are no baselines

All three baselines are unknown, and each has a matching open question with a decider and a date. This is deliberate and it is the honest state: this product has not run against a real engagement yet. Each target is provisional until its question is answered, by 2027-02-28, and section 15 records the moment each baseline arrives and each target is either confirmed or changed.

The pilot PRDs are the first 5 PRDs written with release 1 by people other than the product's author, between REL-001 on 2027-03-31 and 2027-09-30. A PRD passes the buildability review when the verdict is `buildable`, or `buildable with gaps` with every gap closed or owned. For O-003, a `must` requirement counts as testable when at least two of three engineering leads, scoring independently, say a test can be written from it.

A reader should push on this. An objective with an invented baseline is worse than one with an admitted gap, because the invented one gets planned against.

What would move these without this product

Worth asking, and recorded here because the answer is not nothing: a single experienced product manager embedded with the expert would move all three. The bet is on availability and cost, not on beating that person. See section 2.

Objectives

O-001

OBJECTIVE	A PRD handed to a development team can be built from without the team re-deriving it first.
PROBLEMS	• P-001 • P-002 • P-003
METRIC	Share of the pilot PRDs that pass the buildability review, the eleven derivability checks, on first submission. The pilot PRDs are the first 5 PRDs written with release 1 by people other than the product's author, between REL-001, on 2027-03-31, and 2027-09-30. A PRD passes when the verdict is `buildable`, or `buildable with gaps` with every gap closed or owned.
MEASUREMENT	The verdict of the buildability review on each pilot PRD at its first submission.
INDICATOR	lagging
BASELINE	unknown - not measured; see Q-001.
BASELINE QUESTION	Q-001
TARGET	80% of the pilot PRDs, at least 4 of 5, pass on first submission. Provisional until Q-001 is answered, by 2027-02-28.
BY	2027-09-30
OWNER	Richard Hightower

O-002 · A domain expert, or a product manager who speaks for one, reaches an approved PRD in days.

PROBLEMS	• P-004
METRIC	Median elapsed working days, across the pilot PRDs that have one journey, from the first interview question to an approval recorded in section 15.
MEASUREMENT	The date of the first interview question and the date of the approval recorded in section 15, for each pilot PRD that has one journey.
INDICATOR	lagging
BASELINE	unknown - not measured; see Q-002.
BASELINE QUESTION	Q-002
TARGET	Under 5 working days for a single-journey product. Provisional until Q-002 is answered, by 2027-02-28.
BY	2027-09-30
OWNER	Richard Hightower

O-003

OBJECTIVE	Requirements are testable as written, so a team can close them without asking the author what they meant.
PROBLEMS	• P-001 • P-004
METRIC	Share of `must` functional requirements whose acceptance criteria an engineer who was not in the interview judges sufficient to write a test from.
MEASUREMENT	Three engineering leads who were not in the interview score each `must` requirement of each pilot PRD independently, as Q-003 does. A requirement counts as testable when at least two of the three say a test can be written from its acceptance criteria.
INDICATOR	leading
BASELINE	unknown - not measured; see Q-003.
BASELINE QUESTION	Q-003
TARGET	95% of `must` requirements. Provisional until Q-003 is answered, by 2027-02-28.
BY	2027-09-30
OWNER	Richard Hightower

SECTION 4

Users and personas

DRAFT 82/100 NO BLOCKING DEFECTS

ID	ROLE	PER PRD	FREQUENCY	DECIDES
PER-001	Senior product manager in a consultancy or an in-house product team: the domain expert, or a proxy who speaks for one	One; 2 PRDs at once and 6 a year per organization	3 working sittings a week of 1 to 2 hours, for 1 to 3 weeks per PRD	What the product must do, scope, the domain's terms, when an interview ends, drafts, section states, whether to accept or refute a grade, domain facts, whether the domain check may research outside the session, and when the set is offered; not the PRD itself where a sponsor exists
PER-002	Engineering lead (staff or principal engineer) on the development team	One or two	At estimation and at each release planning	How it is built, and whether a set can be built from, with its gaps (TR-006)
PER-003	Sponsor: the budget holder (director or partner), who is also the approver	One	Once per version, about 6 a year	Approves the PRD for a release, with or without conditions, or rejects it (TR-001)
PER-004	Delivery lead (delivery manager or product owner) on the development team	One	Once per release, about 6 a year	The backlog's shape and order, and whether a backlog can be mapped at all

No persona is external. Under a client contract, in segment S-001, the sponsor, the engineering lead, and the delivery lead all work for the consultancy. The client's staff only supply domain knowledge to the expert, and they do not use the product.

PER-001 is the primary persona

PER-001 is the product manager who writes the PRD: either the domain expert, or a product manager who speaks for one. The set calls this person the expert, because they hold the domain knowledge in the interview. Everything is designed around someone with deep domain knowledge, no product training, and a burst of intense use followed by months away. Two consequences run through the rest of this PRD: nothing may assume prior knowledge of product vocabulary, and the work must survive being put down for a week and picked back up. See FR-015 and NFR-004.

Who is deliberately not a persona

No mechanism is a persona. Whatever delivers the teaching, the interview, and the grading is how the product reaches PER-001, and the set keeps it out of the requirements: section 6 makes a separate application a non-goal, and section 11 records the session as dependency D-001. What matters to this PRD is that PER-001 sees a grade they did not produce themselves.

> **PER-001**
>
> NAME — Product manager (senior: manager to principal), in a consultancy or an in-house product team: the domain expert, or a proxy who speaks for one
>
> SHORT NAME — DomainExpert
>
> PRIMARY — yes

POPULATION
One per PRD. Per organization, 2 PRDs run at once and 6 start in a year.

EXTERNAL
no

PROBLEMS
- P-001
- P-004

BASIS
Assumed, not researched. Section 1's evidence note says the evidence comes from one engagement. The ten observed sittings in D-002 test it.

CONTEXT
Works inside the session they already use, as section 11 records (D-001), in bursts weeks apart (NFR-004).

GOALS
- Get what is in my head into a form the development team can build from
- Stop being asked the same question by four different engineers
- Not be the reason the project is late

PAINS
- I know what it must do, and I cannot write it the way they need it
- I am told my requirements are not testable and I do not know what that means
- I reach for the names of software I have used, because that is the only requirements language I have

FREQUENCY
Three working sittings a week of 1 to 2 hours each, for one to three weeks per PRD, about 6 to 9 sittings per PRD, then rarely until the next one.

AUTHORITY
Decides what the product must do and what is out of scope; decides the domain's own terms; decides when a section's interview is finished and when a section may be left at draft; accepts or corrects a drafted section; sets a section to final; accepts or refutes a grade; decides whether the set may be offered for approval; decides whether the domain check may research the set's domain outside the session; and is the only one who may state a domain fact as true. May not approve the PRD for a release on their own where a sponsor exists.

PROFICIENCY
Either deep in the domain with no product management training, or a product manager who speaks for such an expert. The product assumes the first: no familiarity with MoSCoW, Given/When/Then, or the word "persona".

PER-002

NAME
Engineering lead (staff or principal engineer) on the development team, who must build from the PRD

SHORT NAME
EngineerLead

POPULATION
One or two per PRD, so 2 to 4 at once in an organization running 2 PRDs at once.

EXTERNAL
no

PROBLEMS
- P-001
- P-002

BASIS
Assumed, not researched, as section 1's evidence note says. The three engineering leads in D-003 test it.

CONTEXT
Reads the set cold at estimation and at each release planning, without having been in the interview.

GOALS
- Estimate without first interviewing the author
- Know which decisions are mine and which have been made for me
- Find out now, not in iteration three, what the document does not say

PAINS
- Specifications that name a technology and so pre-empt my decision badly
- Requirements I cannot write a test for
- Gaps between two requirements that both looked complete

FREQUENCY
Reads each PRD at estimation and at each release planning, so at least 12 reads a year in an organization that starts 6 PRDs a year.

AUTHORITY	Decides how it is built. Judges whether a set can be built from and reports the gaps, under TR-006, and records a verdict that a set is not buildable when it is not.
PROFICIENCY	Expert in delivery, usually knows nothing about the expert's domain.

PER-003 · Sponsor: the budget holder (director or partner), who is also the approver

SHORT NAME	Sponsor
POPULATION	One per PRD, and one organization starts 6 PRDs a year.
EXTERNAL	no
PROBLEMS	• P-003
BASIS	Assumed, not researched, as section 1's evidence note says. Q-005 tests it with three sponsors.
CONTEXT	Reads the root document or an exported copy outside the product, once per version (ST-006).
GOALS	• Be confident enough to commit budget and a date • Know what is still undecided, and who is deciding it
PAINS	• I can only judge whether it reads well • I find out it was not buildable after I have committed to the plan
FREQUENCY	Once per version offered, about 6 versions a year per organization, plus any conditions.
AUTHORITY	Approves the PRD for a release; may approve with conditions; may reject.
PROFICIENCY	Senior in the business, not in product or engineering.

PER-004

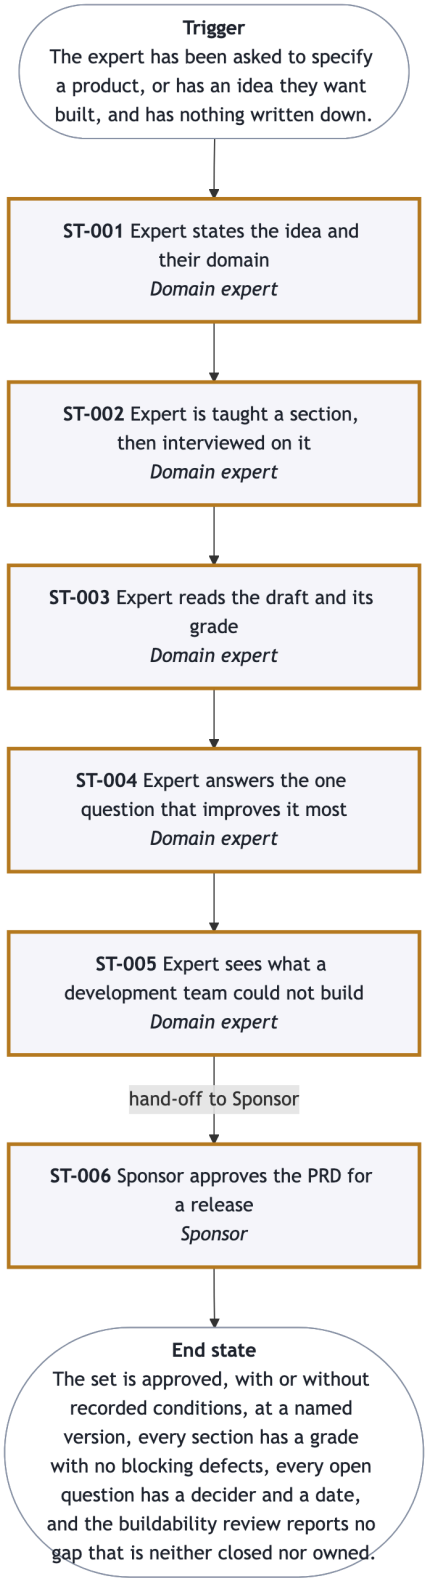
NAME	Delivery lead (delivery manager or product owner) on the development team, who runs the story-mapping workshop
SHORT NAME	DeliveryLead
POPULATION	One per PRD, and one organization starts 6 PRDs a year.
EXTERNAL	no
PROBLEMS	• P-002
BASIS	Assumed, not researched, as section 1's evidence note says. Q-023 tests it with three delivery leads.
CONTEXT	Runs the story-mapping workshop from the approved set (ST-007, ST-008).
GOALS	• Walk into the workshop with a backbone already agreed • Trace every story back to something the expert actually said
PAINS	• No journey means no backbone, so the workshop invents one in the room • Stories that trace to nothing, so nobody can say why they exist
FREQUENCY	Once per release, about 6 workshops a year per organization, plus refinement.
AUTHORITY	Decides the backlog's shape and order, and decides whether a backlog can be mapped from a set at all; does not decide scope.
PROFICIENCY	Expert in agile delivery; depends entirely on the PRD for the domain.

SECTION 5

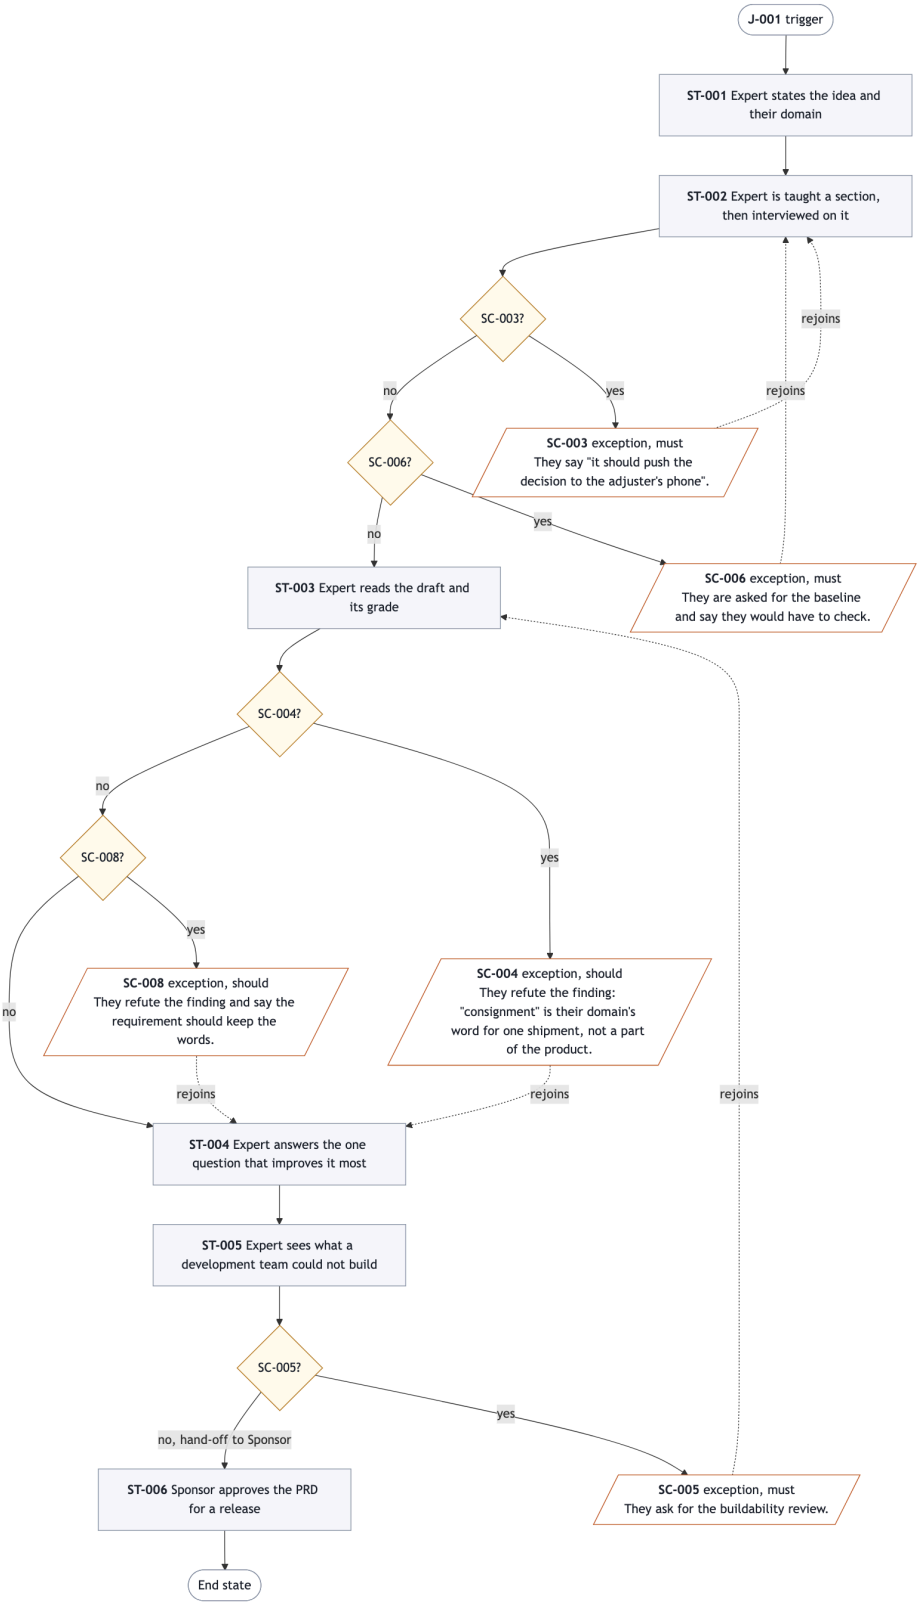
User journeys

DRAFT 94/100 NO BLOCKING DEFECTS

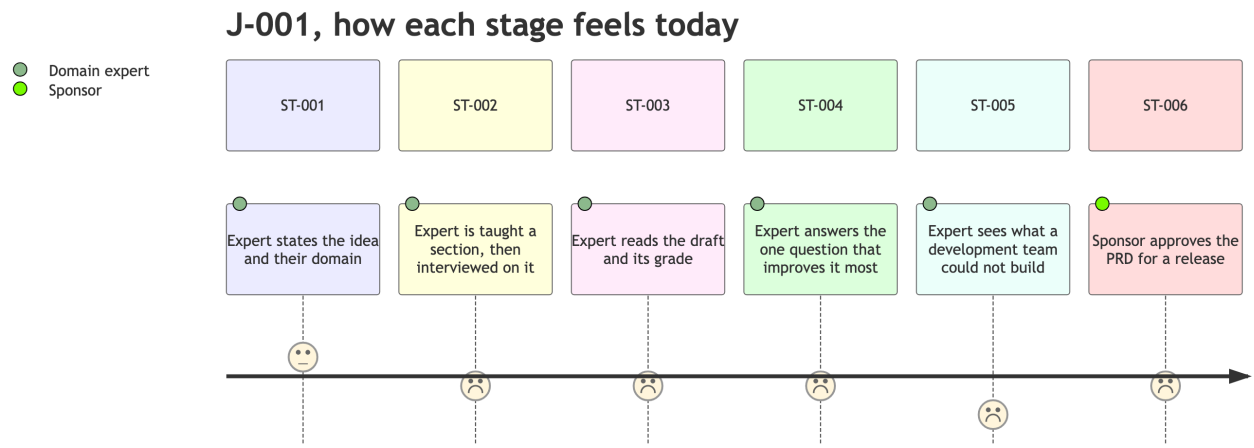
J-001 - A domain expert reaches an approved PRD whose gaps are known (golden thread)



J-001 journey diagram



J-001 overview with its main alternate and exception paths



J-001 satisfaction

Persona PER-001. Trigger: they have been asked to specify a product and have nothing written. Stages ST-001 to ST-006.

STAGE	ACTOR	DECISION
ST-001 States the idea and the domain	Domain expert	What this is, and the domain's own words
ST-002 Taught a section, then interviewed	Domain expert	Have the questions reached the end of what they know
ST-003 Reads the draft and its grade	Domain expert	Accept, correct, or refute the grade
ST-004 Answers the one question that matters most	Domain expert	Is this section done enough to leave
ST-005 Sees what a team could not build	Domain expert	May the set be offered to someone who will commit
ST-006 Sponsor approves for a release, outside the product	Sponsor	Approve, approve with conditions, or reject

ST-002 to ST-004 are a loop, run once per section. The loop is the whole product: the expert is taught, asked, shown a draft, shown a grade, and asked one more thing.

End state: approved, with or without recorded conditions, at a named version, every section graded with nothing blocking, every open question with a decider and a date, and no buildability gap left that is neither closed nor owned.

In release 1, ST-006 happens outside the product. The sponsor reads the root document or an exported copy and gives a decision. The expert records it in section 15 under FR-014, with the sponsor's name, role, and date (section 6).

How each stage feels today

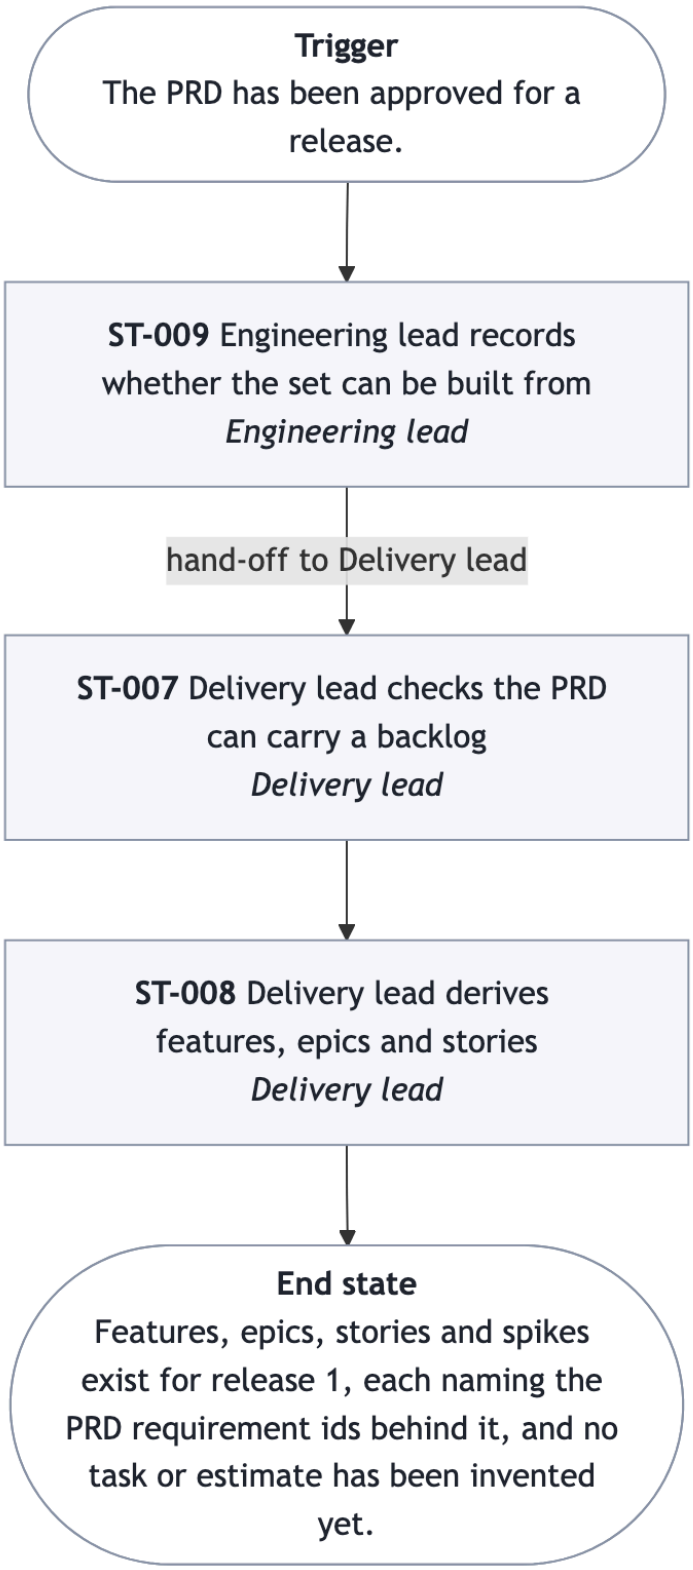
Each stage is scored from 1 (bad) to 5 (good) for how it feels today, without this product. The scores are the expert's.

STAGE	SCORE	PAIN POINT BEHIND IT
ST-001	3	None. Stating the idea comes naturally.
ST-002	2	Headings filled with adjectives, and software names in place of needs
ST-003	2	Cannot tell whether what was written is good
ST-004	2	Put down for a week and cannot remember where they were

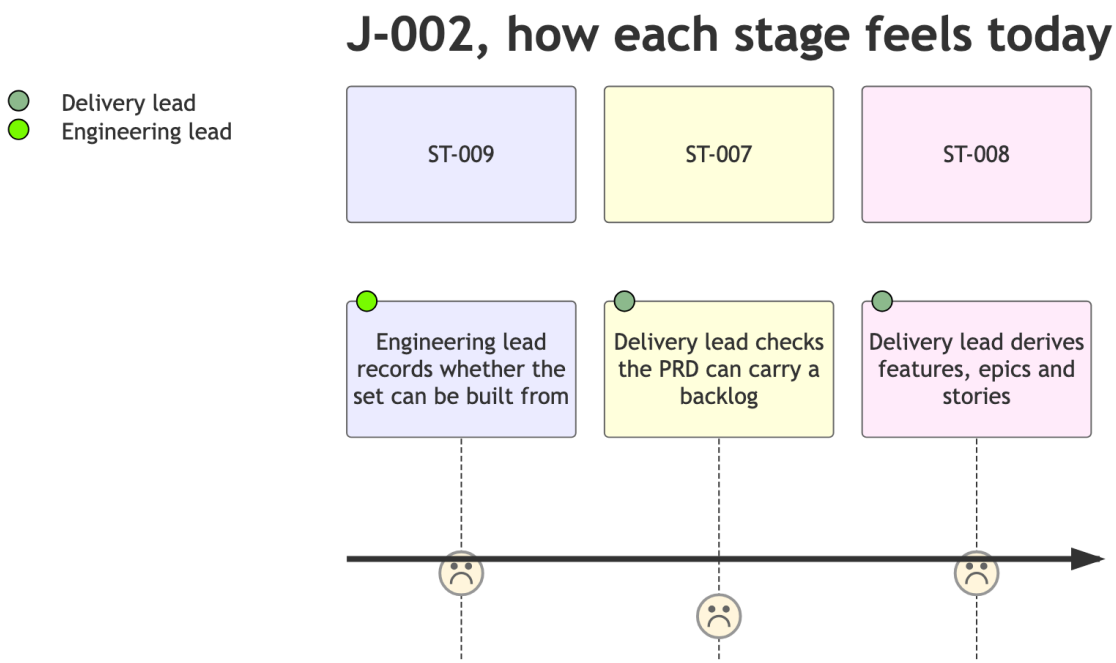
STAGE	SCORE	PAIN POINT BEHIND IT
ST-005	1	Nobody says the requirements are untestable until a team tries to build them
ST-006	2	The sponsor can only judge whether it reads well
ST-009	2	Gaps between two requirements that both looked complete
ST-007	1	The workshop invents a backbone in the room
ST-008	2	Stories trace to nothing, and constraints arrive late

The goal deliberately does not claim an engineering lead's confirmation. No stage in ST-001 to ST-006 puts one in the actor seat, so claiming it would promise something the release line cannot deliver. A lead's own recorded confirmation is OUT-004, deferred to release 2 with J-002, where it is stage ST-009.

J-002 - A delivery lead derives a backlog (not the golden thread)



J-002 journey diagram



J-002 satisfaction

Persona PER-004, stages ST-009, ST-007, and ST-008. In ST-009 the engineering lead, PER-002, reads the report of the buildability review and records their own verdict under TR-006, before the delivery lead takes over. ST-009 was added at 0.8.1, after ST-008, and ids are never renumbered. Release 1 does not have to complete this journey (see section 6), but the PRD must carry what it needs, which is what ST-007 checks.

What these journeys deliberately do not say

No stage says where the teaching text comes from, how a grade is computed, or what reviews the set in ST-005. Those are mechanisms. Every stage above names a person, what they do and see, and what they decide.

J-001

GOAL	An approved PRD whose buildability gaps have been found and either closed or given a decider and a date, before anyone estimates it.
PERSONA	PER-001
TRIGGER	The expert has been asked to specify a product, or has an idea they want built, and has nothing written down.
GOLDEN THREAD	yes
OBJECTIVES	• O-001 • O-002 • O-003

STAGES

Expert states the idea and their domain

ID	ST-001
ACTOR	Domain expert
SATISFACTION	3
ACCEPTANCE	• Given an expert with an idea and nothing written down, when they state the idea, the domain, and three words an outsider would misread, then they see a document set with every section listed and marked as not yet started.
STEPS	• Says what they want built, in their own words, however roughly • Is asked what the business domain is and what the unit of work is called, such as a load, a claim, or a case • Names three words their team uses that an outsider would misread • Sees that a set of documents now exists, every section listed and marked as not yet started
INPUTS	• The idea, the domain, and the vocabulary
OUTPUTS	• A document set with every section visible and honestly marked empty
DECISION	What this product is, and what the domain's own words are

Expert is taught a section, then interviewed on it

ID	ST-002
ACTOR	Domain expert
SATISFACTION	2
ACCEPTANCE	• Given a section not yet interviewed, when the expert reads what it is for and answers its questions one or two at a time, then the interview ends when the question list is finished, when they are told the section already meets the gate, or when they say they have nothing more to add, and they see the section drafted.
STEPS	• Reads what the section is, why it matters, and what belongs in it, in under ten lines, with an example from their own domain • Is asked one or two questions at a time, in their own vocabulary • Is asked follow-ups where an answer was thin • Says when they have nothing more to add
INPUTS	• Their knowledge of the domain, in answer to each question
OUTPUTS	• The answers the section is drafted from
DECISION	Whether the questions have reached the end of what they know

Expert reads the draft and its grade

ID	ST-003
ACTOR	Domain expert
SATISFACTION	2
ACCEPTANCE	• Given a drafted and graded section, when the expert opens it, then they read it as prose and see its score, each blocking defect, and its weakest part.

STEPS	• Reads the drafted section as prose, not as a form • Sees a score, how many blocking defects there are, and what each one is • Sees which part of the section is weakest, and why • Decides whether the grade is fair
INPUTS	• The drafted section and its grade
OUTPUTS	• A judgement on whether the draft says what they meant
DECISION	Accept the draft, correct it, or refute the grade

Expert answers the one question that improves it most

ID	ST-004
ACTOR	Domain expert
SATISFACTION	2
ACCEPTANCE	• Given one question about the weakest part of a section, when the expert answers it, then they see the section re-graded against the same rubric. • Given one question whose premise the expert says is wrong, when they say why, then the finding behind the question is checked against the document as a refutation under FR-016, and they see whether it was dropped. • Given an expert returning after a break, when they say they want to carry on, then they resume at the first section not yet at review, and no question they already answered is put to them again.
STEPS	• Reads a single question, not a list of twelve • Answers it, or says why the premise is wrong • Sees the section re-graded against the same rubric • Moves to the next section, or goes back to an earlier one when this section showed it was incomplete • Stops when they need to, and later resumes at the first section not yet at review
INPUTS	• One answer
OUTPUTS	• A section at a grade the expert accepts, and a set that is one section further on
DECISION	Whether this section is done enough to leave
HANDOFF	Back into ST-002 for the next section, until every section has been through this loop.

Expert sees what a development team could not build

ID	ST-005
ACTOR	Domain expert
SATISFACTION	1

ACCEPTANCE	• Given a set with every section written at least once, when the expert asks for a buildability review, then they see what an engineer could not build and which release-line stages have no requirement behind them. • Given a gap the buildability review reports, when the expert does not close it, then they see it recorded as an open question with a decider and a date.
STEPS	• Asks for a buildability review of the whole set, read as an engineer would • Sees what an engineer could not build, and what they would have to guess • Sees which release-line stages have no requirement behind them • Closes each gap, or records it as an open question with a decider and a date • Offers the set to the sponsor for approval
INPUTS	• Every section written at least once, so there is something to review
OUTPUTS	• A list of what is not buildable, and a set with each gap closed or owned
DECISION	May the set be offered to someone who will commit to it
HANDOFF	Sponsor, who receives the set offered for approval as its root document or an exported copy.

Sponsor approves the PRD for a release

ID	ST-006
ACTOR	Sponsor
SATISFACTION	2
ACCEPTANCE	• Given a set offered for approval, when the sponsor approves, approves with conditions, or rejects it, and the expert records the decision, then the sponsor sees the decision, their name, their role, and the date recorded against the set's version.
STEPS	• Reads the root document or an exported copy, outside the product • Sees the summary, the scope, what is still open, each section's grade, and whether anything is still blocking • Approves, approves with conditions naming what the condition is, or rejects, and tells the expert • Sees the expert record the decision in section 15 against this version
INPUTS	• The set offered for approval, as its root document or an exported copy
OUTPUTS	• An approved PRD at a named version, with the approval recorded
DECISION	Approve, approve with conditions, or reject
HANDOFF	Engineering lead, to record a verdict on the approved set in ST-009

PAIN POINTS

PAIN	I do not know what the section is for, so I fill the heading with adjectives
STAGE	ST-002
NEED	Teach me the section before you ask me about it, in my own domain
PAIN	Nobody tells me my requirements are untestable until a team tries to build them
STAGE	ST-005
NEED	Show me, before anyone builds, what an engineer cannot build
PAIN	I reach for the names of software I have used
STAGE	ST-002
NEED	Convert my technology talk into the need underneath it, without arguing
PAIN	I put this down for a week and cannot remember where I was
STAGE	ST-004
NEED	Let me stop and resume without starting over
PAIN	I cannot tell whether what I have written is good
STAGE	ST-003
NEED	Tell me where it is weak, against a stated standard, and ask me one thing
PAIN	As the sponsor, I can only judge whether it reads well
STAGE	ST-006
NEED	Show the sponsor each grade and every blocking defect before they decide

NEEDS

- Teach me the section before you ask me about it, in my own domain
- Tell me where it is weak, against a stated standard, and ask me one thing
- Convert my technology talk into the need underneath it, without arguing
- Let me stop and resume without starting over
- Show me, before anyone builds, what an engineer cannot build
- Show the sponsor each grade and every blocking defect before they decide

END STATE

The set is approved, with or without recorded conditions, at a named version, every section has a grade with no blocking defects, every open question has a decider and a date, and the buildability review reports no gap that is neither closed nor owned.

J-002 · A backlog whose every story traces to a requirement the expert actually stated.

PERSONA

PER-004

TRIGGER

The PRD has been approved for a release.

OBJECTIVES

- O-001

STAGES

Engineering lead records whether the set can be built from

ID	ST-009
ACTOR	Engineering lead
SATISFACTION	2
ACCEPTANCE	• Given an approved set and the report of its buildability review, when the engineering lead records a verdict that differs from the report's, then they see their verdict recorded as the verdict, with the report kept alongside it. • Given an approved set and the report of its buildability review, when the engineering lead records a verdict that agrees with the report's, then they see their verdict recorded with the report. • Given a gap in the set, when the engineering lead records it, then they see the document quoted for it.
STEPS	• Reads the report of the buildability review, its verdict, and the gaps it quotes • Checks each reported gap against the document it quotes • Records their own verdict on whether the set can be built from
INPUTS	• An approved set and the report of its buildability review
OUTPUTS	• The engineering lead's own recorded verdict, with the report kept alongside it
DECISION	Can a team build from this set
HANDOFF	The delivery lead, to check the PRD can carry a backlog

Delivery lead checks the PRD can carry a backlog

ID	ST-007
ACTOR	Delivery lead
SATISFACTION	1
ACCEPTANCE	• Given an approved set and the release it is mapped for, when the delivery lead checks it, then they see whether a golden-thread journey exists and what would block deriving a backlog. • Given an approved set with no journey, when the delivery lead checks it, then they see a no-go, with the reason that a backlog cannot be derived, and the workshop waits.
STEPS	• Sees whether a journey exists and which one is the golden thread • Sees which release-line stages have requirements behind them • Sees what would block deriving a backlog, before the workshop
INPUTS	• An approved set, and the release it is being mapped for
OUTPUTS	• A go or no-go on holding the story-mapping workshop
DECISION	Is there a backbone to map against

Delivery lead derives features, epics and stories

ID	ST-008
ACTOR	Delivery lead
SATISFACTION	2

ACCEPTANCE	• Given an approved set, when the delivery lead derives a backlog from it, then they see one feature per journey stage, epics scoped by the release line, and a story per requirement, each naming the id it came from.
STEPS	• Sees one feature per journey stage, with the requirement ids it carries • Sees epics scoped by the release line • Sees a story per requirement, in the persona's own words, with the requirement's criteria as its acceptance criteria • Sees spikes for the open questions that block the release line • Sees each non-functional requirement and constraint that applies to a story as a definition-of-done item on it
INPUTS	• The approved set
OUTPUTS	• A backlog whose every level names what it came from
DECISION	Which epics come first, and why
HANDOFF	The development team, for iteration planning

PAIN POINTS

PAIN	The workshop invents a backbone in the room because the PRD has none
STAGE	ST-007
NEED	A journey to map against, agreed before the workshop

PAIN	Stories trace to nothing, so nobody can say why they exist
STAGE	ST-008
NEED	Every story carrying the requirement id it came from

PAIN	Non-functional requirements and constraints are forgotten until late
STAGE	ST-008
NEED	Constraints arriving as definition-of-done items, not as surprises

PAIN	As the engineering lead, I find gaps between two requirements that both looked complete
STAGE	ST-009
NEED	Every gap quoted from the document, before the team estimates

NEEDS

- A journey to map against, agreed before the workshop
- Every story carrying the requirement id it came from
- Constraints arriving as definition-of-done items, not as surprises
- Every gap quoted from the document, before the team estimates

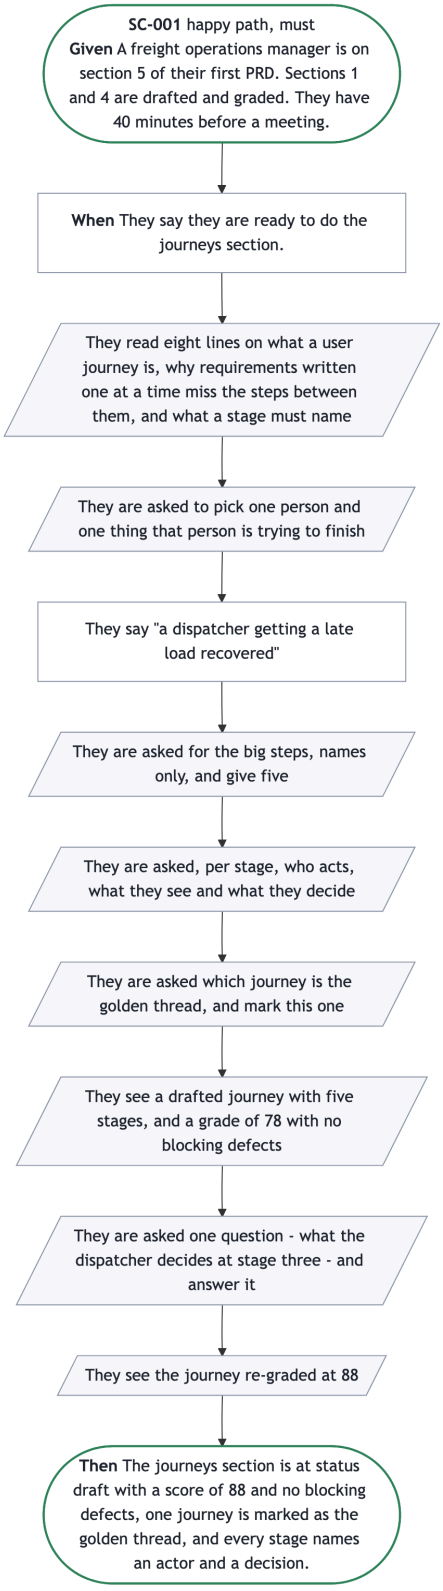
END STATE

Features, epics, stories and spikes exist for release 1, each naming the PRD requirement ids behind it, and no task or estimate has been invented yet.

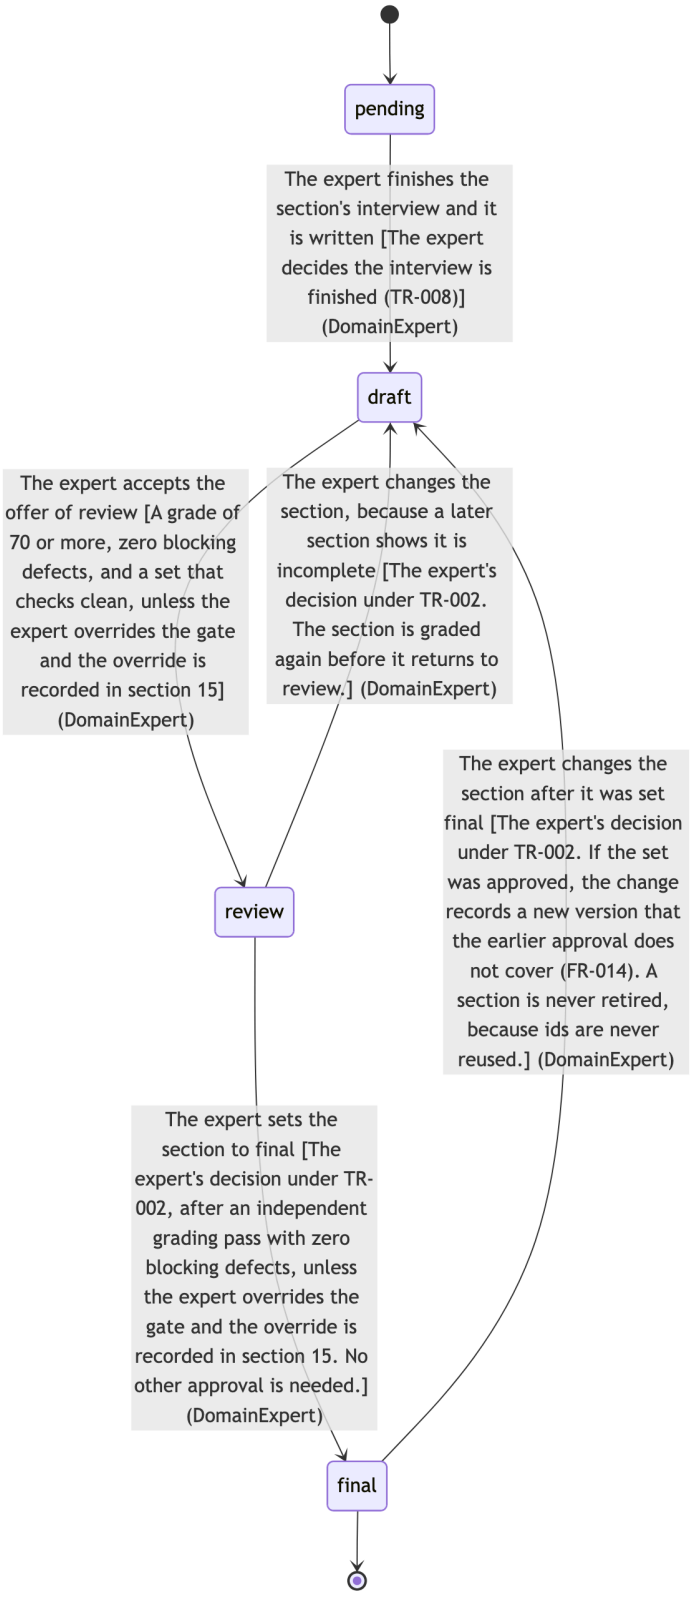
SECTION 5

Scenarios

DRAFT 89/100 NO BLOCKING DEFECTS

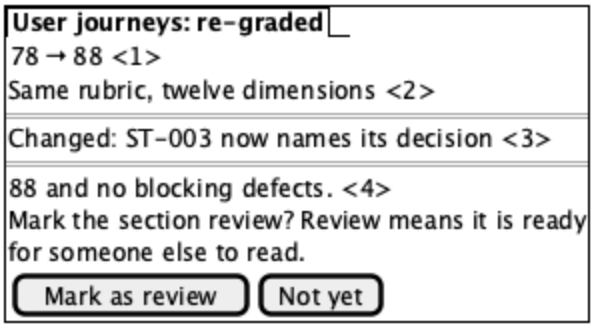


SC-001 scenario

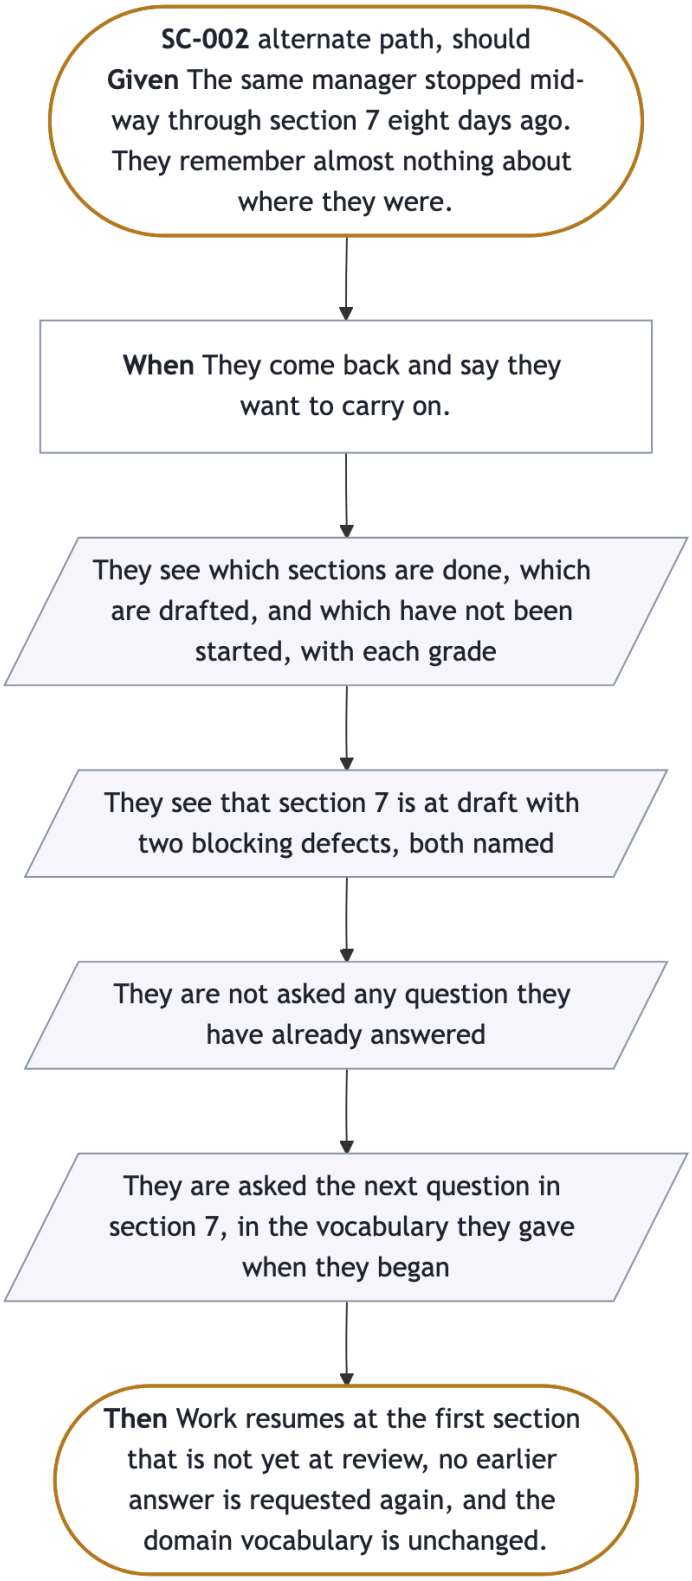


LC-001 lifecycle of a section of the PRD

WF-008 A section re-graded after one answer

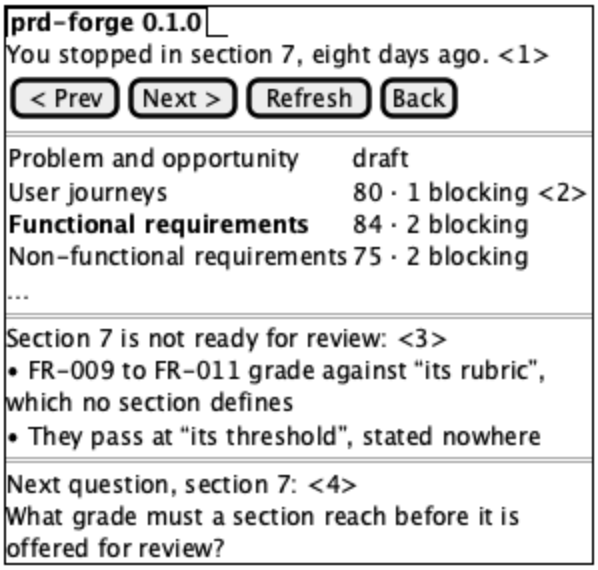


WF-008 drawing

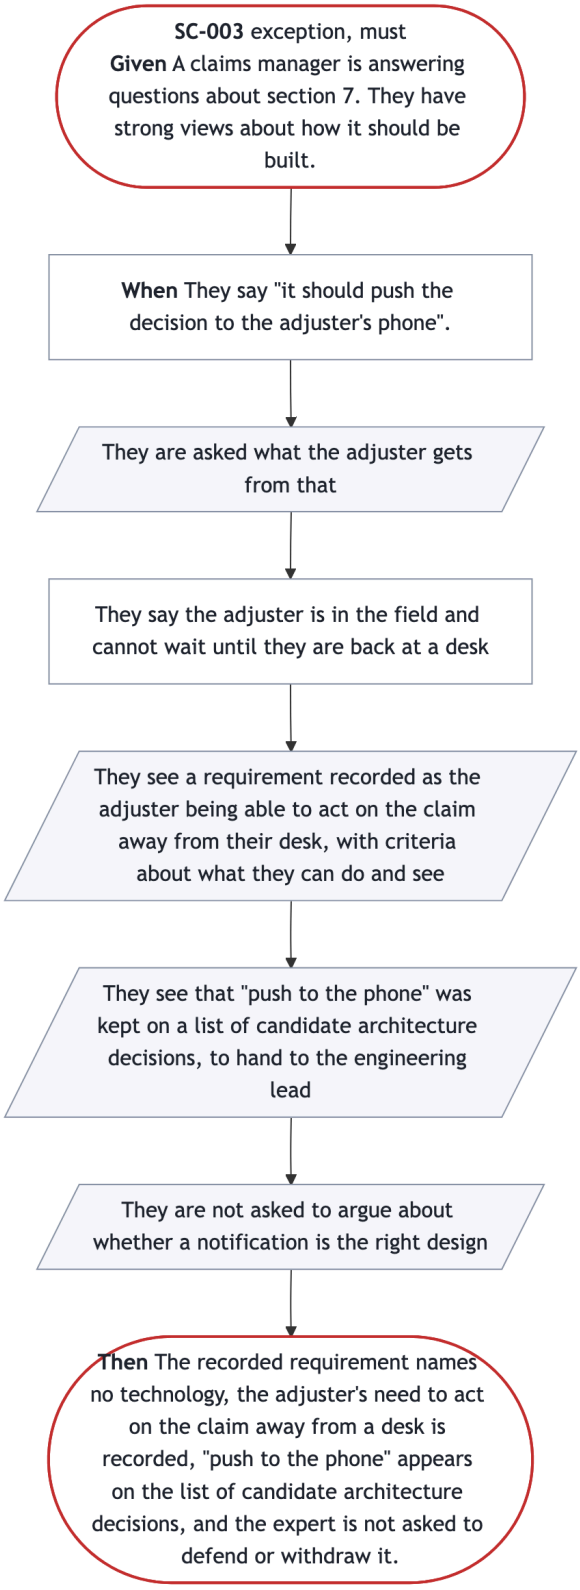


SC-002 scenario

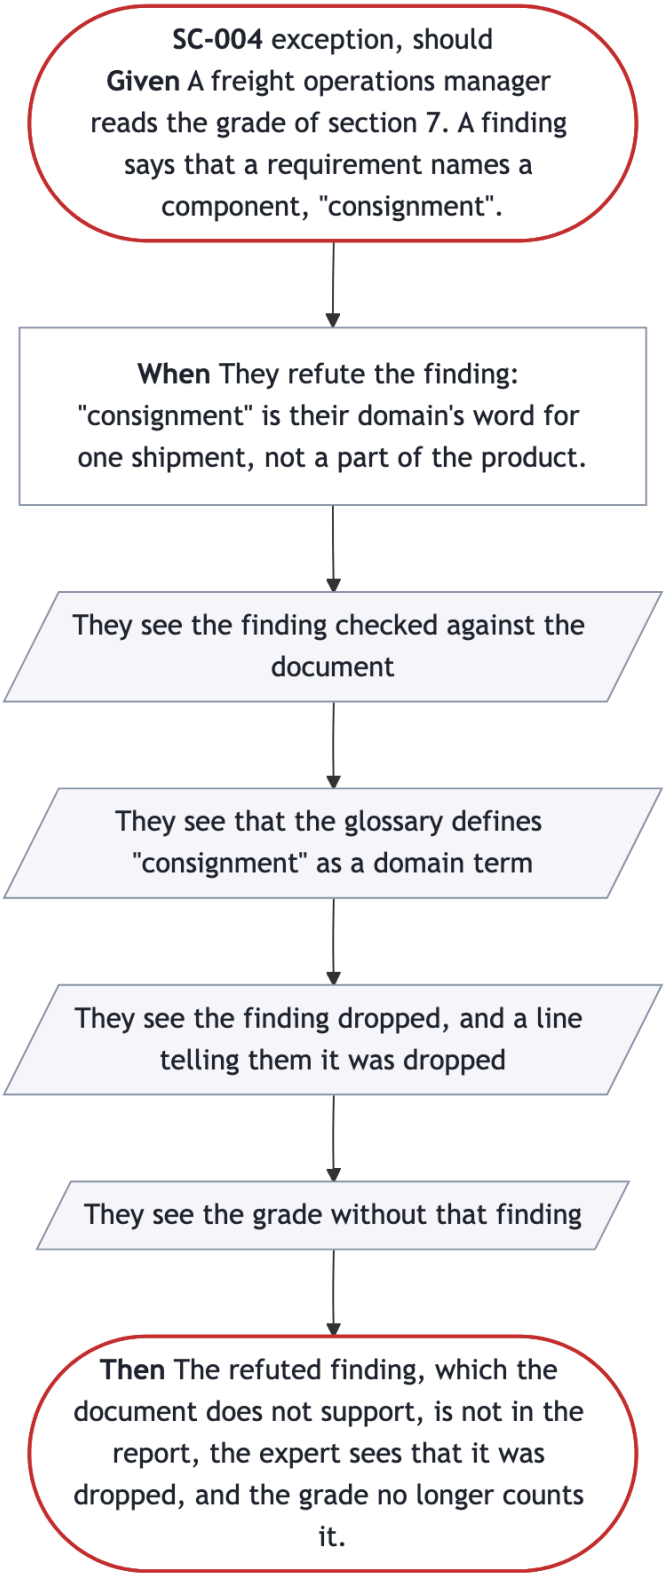
WF-010 Resuming a set after a week away



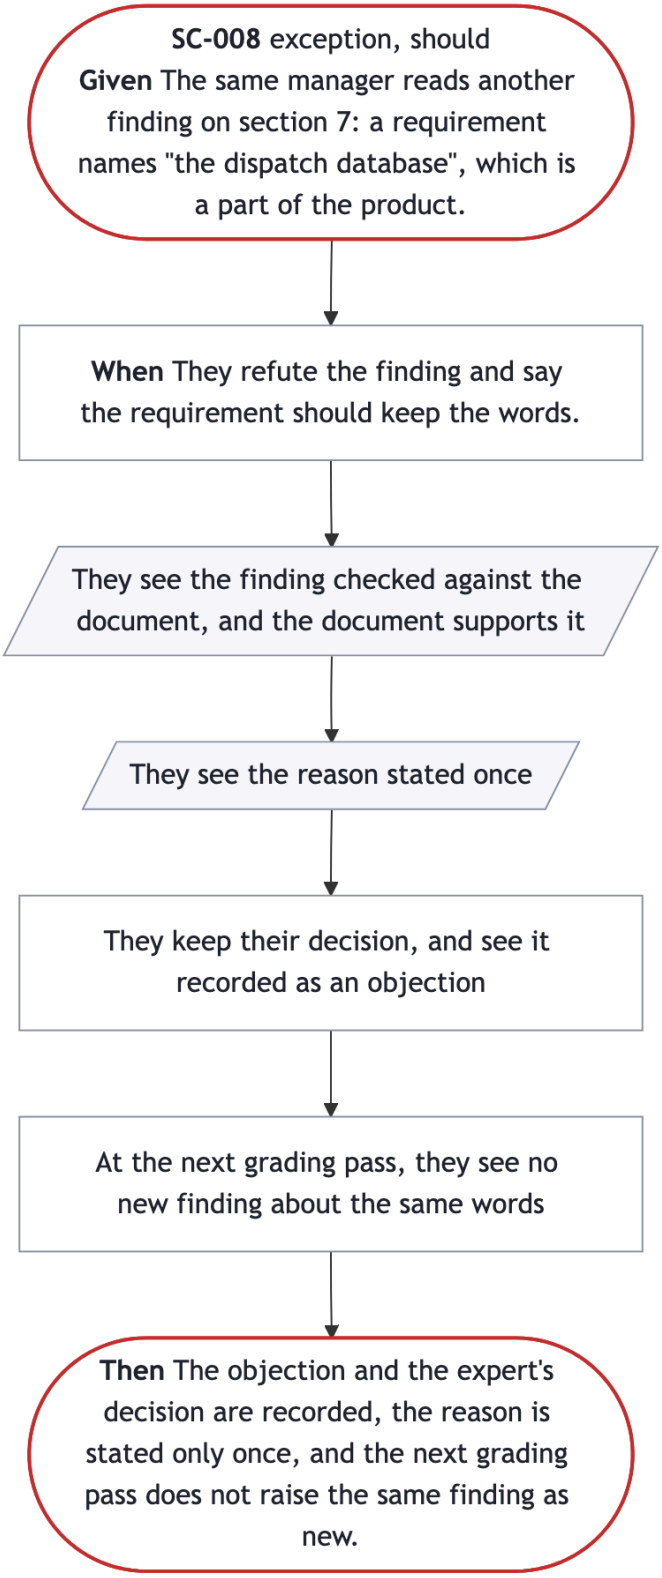
WF-010 drawing



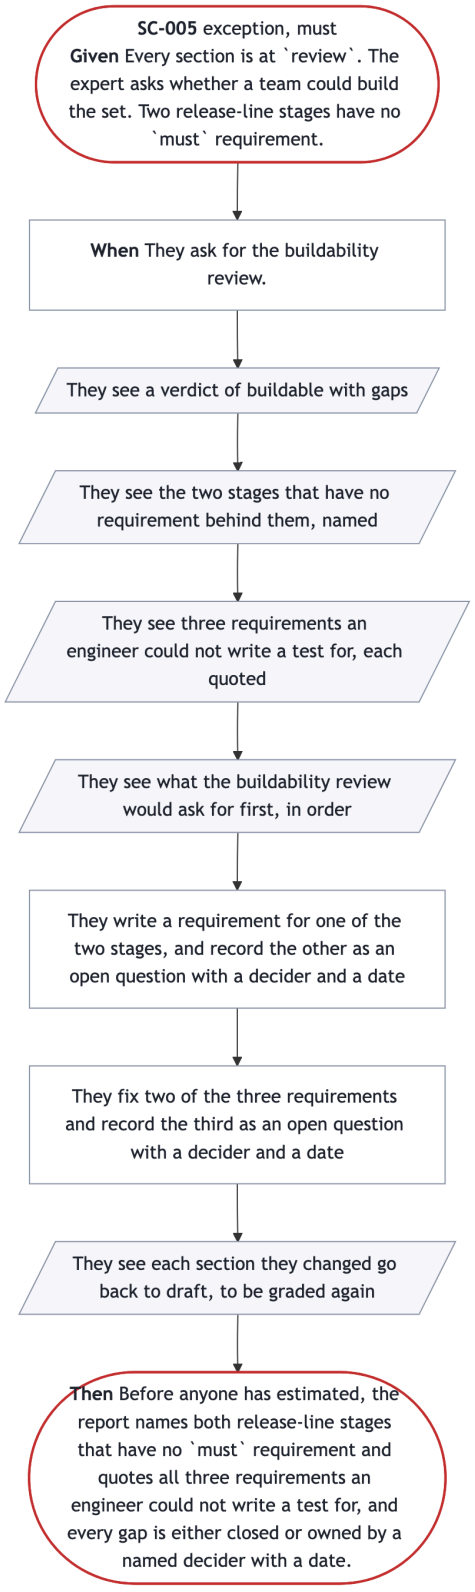
SC-003 scenario



SC-004 scenario



SC-008 scenario



SC-005 scenario

WF-009 A buildability review that finds gaps

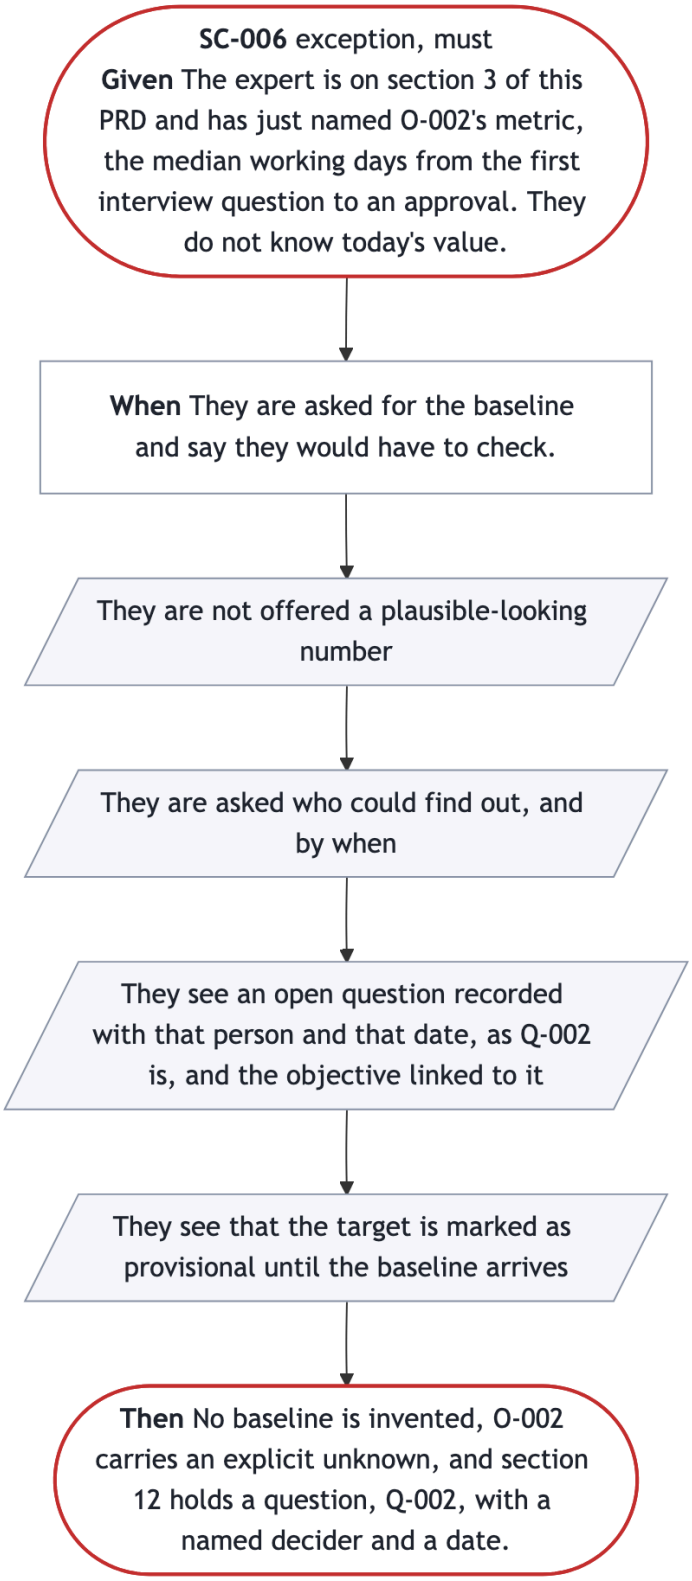
What a team could not build from this set
Buildable with gaps. <1>
✗ ST-004 no must requirement <2> ✗ ST-005 no must requirement ✓ ST-001, ST-002, ST-003, ST-006
Each gap becomes a requirement, or an open question with a decider and a date. <3>

WF-009 drawing

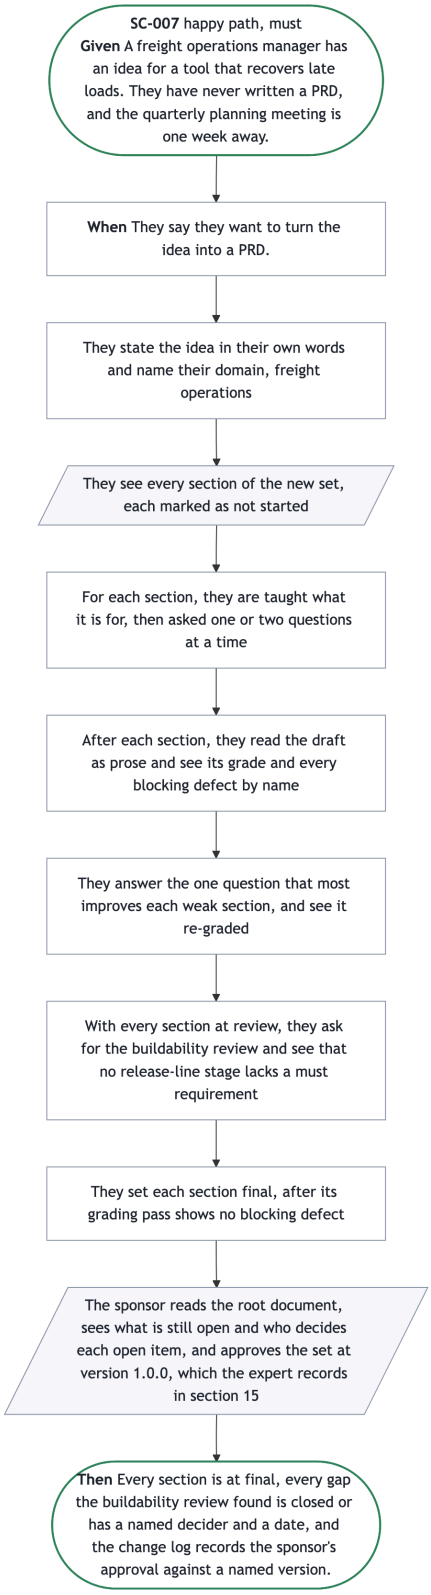
WF-005 What a team could not build from this set

What a team could not build from this set
prd-forge PRD 0.4.0. Buildable with gaps. <1>
Release line <2> ✓ ST-001 FR-001, FR-002, FR-015 ✓ ST-003 FR-008, FR-009, FR-016, FR-017, FR-018 ✓ ST-006 FR-013, FR-014 …
Cannot write a test from <3> 1. FR-004: “two ways differing in meaning", no judge 2. FR-014: no test for a rejection
Blocked by an open question <4> Q-009 blocks FR-012: which rubric version applies. Richard Hightower, by 2027-01-31
Ask for first <5> 1. Who judges FR-004's two meanings

WF-005 drawing

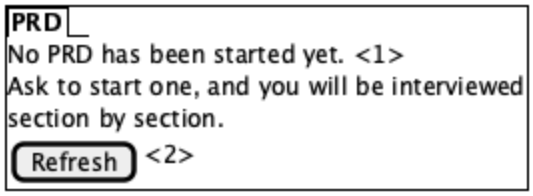


SC-006 scenario



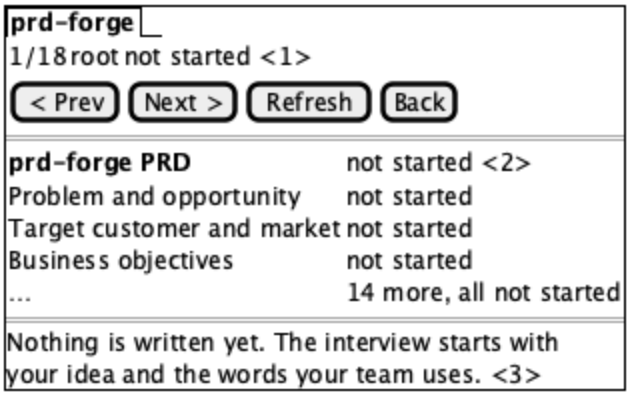
SC-007 scenario

WF-007 Before any set exists



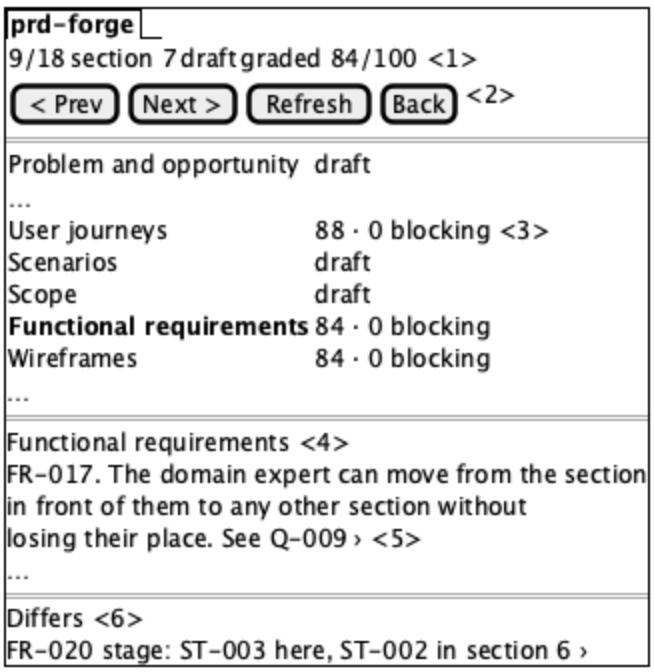
WF-007 drawing

WF-002 A new set, before the first interview



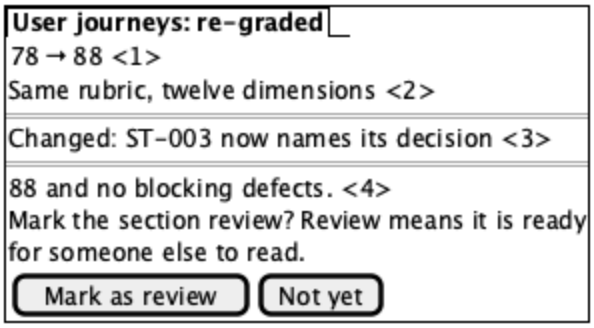
WF-002 drawing

WF-001 The set with one section open



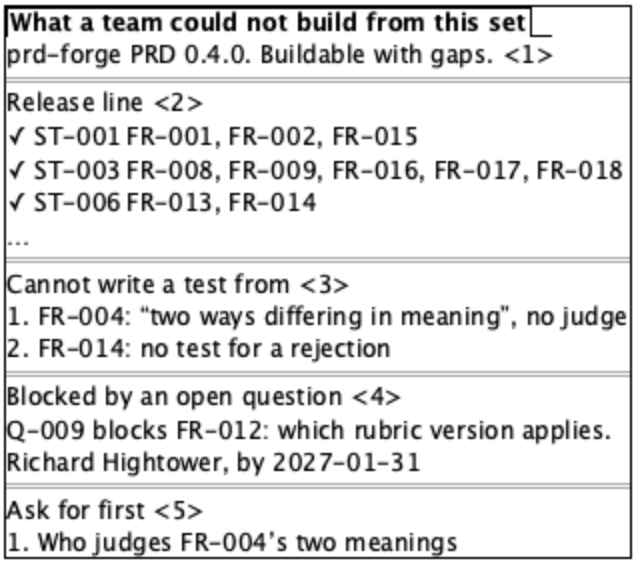
WF-001 drawing

WF-008 A section re-graded after one answer

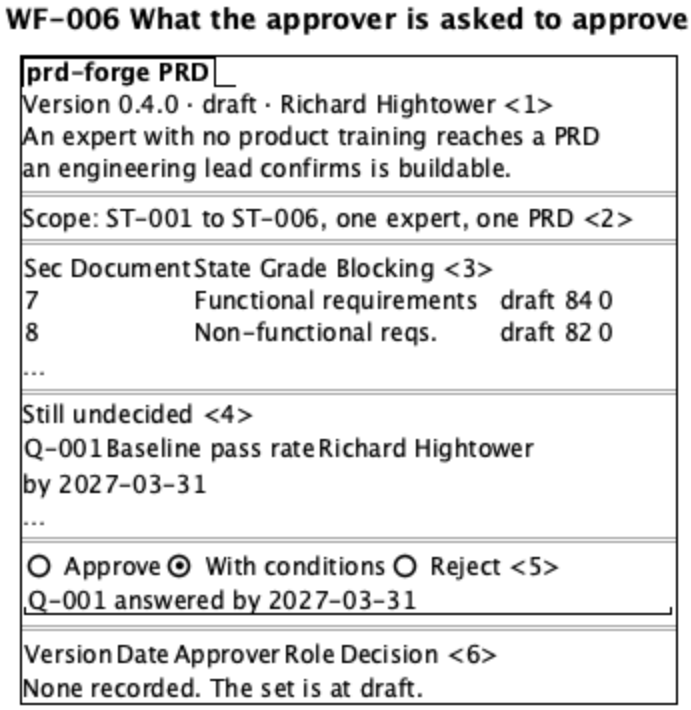


WF-008 drawing

WF-005 What a team could not build from this set



WF-005 drawing



WF-006 drawing

ID	STAGE	VARIANT	PRIORITY	WHAT IT COVERS
SC-001	ST-002	happy	must	A section taught, interviewed, drafted, graded, improved
SC-002	ST-004	alternate	should	Resuming after eight days away
SC-003	ST-002	exception	must	The expert states a technology
SC-004	ST-003	exception	should	The expert refutes a finding that the document does not support
SC-005	ST-005	exception	must	The buildability review finds gaps
SC-006	ST-002	exception	must	The expert has no baseline for a metric
SC-007	ST-001 to ST-006	happy	must	One idea, from the first sentence to an approved set
SC-008	ST-003	exception	should	The expert refutes a finding that the document supports

SC-007 is the end-to-end acceptance test for release 1: it follows the golden thread through every release-line stage. SC-001 is the same path inside one section. SC-005 is the expert's own run of the buildability review, which is why its persona is PER-001 and not the engineering lead: no engineering lead is in release 1 at all, per OUT-004.

The priorities follow the requirements. SC-002, SC-004, and SC-008 are `should`, because FR-015 and FR-016, which they exercise, are `should`.

SC-001 and SC-007 are first use, and SC-002 is a return after time away. No scenario has two experts on one set at once, which OUT-002 defers to a later release. A sponsor reading the set while the expert changes it is covered by FR-014: the change records a new version that the earlier approval does not cover.

The five exceptions are the ones that decide whether this product is trustworthy. SC-003 and SC-006 are where a PRD quietly acquires a technology choice or an invented fact. SC-004 and SC-008 are where the grading either earns the

expert's trust or loses it: a grading pass that cannot be refuted is one nobody believes. SC-005 is the moment the product is supposed to pay for itself.

Scenarios

SC-001

CONTEXT	A freight operations manager is on section 5 of their first PRD. Sections 1 and 4 are drafted and graded. They have 40 minutes before a meeting.
JOURNEY	J-001
STAGE	ST-002
STAGES	• ST-002 • ST-003 • ST-004
PERSONA	PER-001
VARIANT	happy
PRIORITY	must
SITUATION	first-use
SHOWN IN	screen
LIFECYCLES	• LC-001
WIREFRAMES	• WF-008
TRIGGER	They say they are ready to do the journeys section.
COURSE	• They read eight lines on what a user journey is, why requirements written one at a time miss the steps between them, and what a stage must name • They are asked to pick one person and one thing that person is trying to finish • They say "a dispatcher getting a late load recovered" • They are asked for the big steps, names only, and give five • They are asked, per stage, who acts, what they see and what they decide • They are asked which journey is the golden thread, and mark this one • They see a drafted journey with five stages, and a grade of 78 with no blocking defects • They are asked one question - what the dispatcher decides at stage three - and answer it • They see the journey re-graded at 88
EXPECTED OUTCOME	The journeys section is at status draft with a score of 88 and no blocking defects, one journey is marked as the golden thread, and every stage names an actor and a decision.

SC-002

CONTEXT
The same manager stopped mid-way through section 7 eight days ago. They remember almost nothing about where they were.

BRANCHES FROM

SCENARIO	SC-007
STEP	5

JOURNEY J-001

STAGE ST-004

REJOINS ST-004

PERSONA PER-001

VARIANT alternate

PRIORITY should

SITUATION return

SHOWN IN screen

LIFECYCLES
- LC-001

WIREFRAMES
- WF-010

TRIGGER
They come back and say they want to carry on.

COURSE
- They see which sections are done, which are drafted, and which have not been started, with each grade
- They see that section 7 is at draft with two blocking defects, both named
- They are not asked any question they have already answered
- They are asked the next question in section 7, in the vocabulary they gave when they began

EXPECTED OUTCOME
Work resumes at the first section that is not yet at review, no earlier answer is requested again, and the domain vocabulary is unchanged.

SC-003

CONTEXT
A claims manager is answering questions about section 7. They have strong views about how it should be built.

BRANCHES FROM

SCENARIO	SC-007
STEP	3

JOURNEY J-001

STAGE ST-002

REJOINS ST-002

PERSONA PER-001

VARIANT exception

PRIORITY must

SHOWN IN conversation

TRIGGER
They say "it should push the decision to the adjuster's phone".

COURSE

- They are asked what the adjuster gets from that
- They say the adjuster is in the field and cannot wait until they are back at a desk
- They see a requirement recorded as the adjuster being able to act on the claim away from their desk, with criteria about what they can do and see
- They see that "push to the phone" was kept on a list of candidate architecture decisions, to hand to the engineering lead
- They are not asked to argue about whether a notification is the right design

EXPECTED OUTCOME

The recorded requirement names no technology, the adjuster's need to act on the claim away from a desk is recorded, "push to the phone" appears on the list of candidate architecture decisions, and the expert is not asked to defend or withdraw it.

SC-004

CONTEXT

A freight operations manager reads the grade of section 7. A finding says that a requirement names a component, "consignment".

BRANCHES FROM

SCENARIO	SC-007
STEP	4

JOURNEY

J-001

STAGE

ST-003

REJOINS

ST-004

PERSONA

PER-001

VARIANT

exception

PRIORITY

should

SHOWN IN

conversation

TRIGGER

They refute the finding: "consignment" is their domain's word for one shipment, not a part of the product.

COURSE

- They see the finding checked against the document
- They see that the glossary defines "consignment" as a domain term
- They see the finding dropped, and a line telling them it was dropped
- They see the grade without that finding

EXPECTED OUTCOME

The refuted finding, which the document does not support, is not in the report, the expert sees that it was dropped, and the grade no longer counts it.

SC-008

CONTEXT
The same manager reads another finding on section 7: a requirement names "the dispatch database", which is a part of the product.

BRANCHES FROM

SCENARIO	SC-007
STEP	4

JOURNEY J-001

STAGE ST-003

REJOINS ST-004

PERSONA PER-001

VARIANT exception

PRIORITY should

SHOWN IN conversation

TRIGGER
They refute the finding and say the requirement should keep the words.

COURSE

- They see the finding checked against the document, and the document supports it
- They see the reason stated once
- They keep their decision, and see it recorded as an objection
- At the next grading pass, they see no new finding about the same words

EXPECTED OUTCOME
The objection and the expert's decision are recorded, the reason is stated only once, and the next grading pass does not raise the same finding as new.

SC-005

CONTEXT
Every section is at `review`. The expert asks whether a team could build the set. Two release-line stages have no `must` requirement.

BRANCHES FROM

SCENARIO	SC-007
STEP	6

JOURNEY J-001

STAGE ST-005

REJOINS ST-003

PERSONA PER-001

VARIANT exception

PRIORITY must

SHOWN IN report

WIREFRAMES

- WF-009
- WF-005

TRIGGER
They ask for the buildability review.

COURSE

- They see a verdict of buildable with gaps
- They see the two stages that have no requirement behind them, named
- They see three requirements an engineer could not write a test for, each quoted
- They see what the buildability review would ask for first, in order
- They write a requirement for one of the two stages, and record the other as an open question with a decider and a date
- They fix two of the three requirements and record the third as an open question with a decider and a date
- They see each section they changed go back to draft, to be graded again

EXPECTED OUTCOME

Before anyone has estimated, the report names both release-line stages that have no `must` requirement and quotes all three requirements an engineer could not write a test for, and every gap is either closed or owned by a named decider with a date.

SC-006

CONTEXT

The expert is on section 3 of this PRD and has just named O-002's metric, the median working days from the first interview question to an approval. They do not know today's value.

BRANCHES FROM

SCENARIO	SC-007
STEP	3

JOURNEY

J-001

STAGE

ST-002

REJOINS

ST-002

PERSONA

PER-001

VARIANT

exception

PRIORITY

must

SHOWN IN

conversation

TRIGGER

They are asked for the baseline and say they would have to check.

COURSE

- They are not offered a plausible-looking number
- They are asked who could find out, and by when
- They see an open question recorded with that person and that date, as Q-002 is, and the objective linked to it
- They see that the target is marked as provisional until the baseline arrives

EXPECTED OUTCOME

No baseline is invented, O-002 carries an explicit unknown, and section 12 holds a question, Q-002, with a named decider and a date.

SC-007

CONTEXT

A freight operations manager has an idea for a tool that recovers late loads. They have never written a PRD, and the quarterly planning meeting is one week away.

JOURNEY

J-001

STAGE

ST-001

STAGES	• ST-001 • ST-002 • ST-003 • ST-004 • ST-005 • ST-006
PERSONA	PER-001
VARIANT	happy
PRIORITY	must
SITUATION	first-use
SHOWN IN	screen
LIFECYCLES	• LC-001
WIREFRAMES	• WF-007 • WF-002 • WF-001 • WF-008 • WF-005 • WF-006
TRIGGER	They say they want to turn the idea into a PRD.
COURSE	• They state the idea in their own words and name their domain, freight operations • They see every section of the new set, each marked as not started • For each section, they are taught what it is for, then asked one or two questions at a time • After each section, they read the draft as prose and see its grade and every blocking defect by name • They answer the one question that most improves each weak section, and see it re-graded • With every section at review, they ask for the buildability review and see that no release-line stage lacks a must requirement • They set each section final, after its grading pass shows no blocking defect • The sponsor reads the root document, sees what is still open and who decides each open item, and approves the set at version 1.0.0, which the expert records in section 15
EXPECTED OUTCOME	Every section is at final, every gap the buildability review found is closed or has a named decider and a date, and the change log records the sponsor's approval against a named version.

Lifecycles

LC-001 · a section of the PRD

STATES

pending

MEANING — Listed in the set and not yet interviewed.

draft

MEANING — Interviewed and written, and graded or waiting for a grade.

review

MEANING — Ready for someone other than the expert to read.

final

MEANING — Set final by the expert under TR-002.

INITIAL — pending

FINAL

- final

TRANSITIONS

FROM	pending
TO	draft
WHEN	The expert finishes the section's interview and it is written
RULE	The expert decides the interview is finished (TR-008)
BY	PER-001

FROM	draft
TO	review
WHEN	The expert accepts the offer of review
RULE	A grade of 70 or more, zero blocking defects, and a set that checks clean, unless the expert overrides the gate and the override is recorded in section 15
BY	PER-001

FROM	review
TO	final
WHEN	The expert sets the section to final
RULE	The expert's decision under TR-002, after an independent grading pass with zero blocking defects, unless the expert overrides the gate and the override is recorded in section 15. No other approval is needed.
BY	PER-001

FROM	review
TO	draft
WHEN	The expert changes the section, because a later section shows it is incomplete
RULE	The expert's decision under TR-002. The section is graded again before it returns to review.
BY	PER-001

FROM	final
TO	draft
WHEN	The expert changes the section after it was set final
RULE	The expert's decision under TR-002. If the set was approved, the change records a new version that the earlier approval does not cover (FR-014). A section is never retired, because ids are never reused.
BY	PER-001

SECTION 6

Scope

DRAFT 88/100 NO BLOCKING DEFECTS

Release line

ST-001 through ST-006 for one expert, one PRD, one golden-thread journey. The expert reaches J-001's end state: the set approved at a named version, every section graded with no blocking defects, every open question with a decider and a date, and no buildability gap that is neither closed nor owned.

ST-006 happens outside the product in REL-001. The sponsor reads the root document or an exported copy and gives a decision. Comments the sponsor leaves in an exported copy come back under FR-023. The expert records the decision in section 15 under FR-014, with the sponsor's name, role, and date.

ID	IN SCOPE	STAGE	REQUIREMENTS
IN-001	State the idea and vocabulary once, and see every section and its state	ST-001	FR-001, FR-002
IN-002	Be taught each section, be asked one or two questions at a time, end early with gaps recorded, and stop once it would meet the gate	ST-002	FR-003, FR-004, FR-005, FR-019
IN-003	Record no invented fact, and convert a stated technology into the need beneath it	ST-002	FR-006, FR-007
IN-004	Read the draft with its grade, refute a finding, move to any section, refresh, and see one item stated two ways	ST-003	FR-008, FR-009, FR-016, FR-017, FR-018, FR-020
IN-005	Answer one question, see the section re-graded, and resume without repeating an answer	ST-004	FR-010, FR-011, FR-015
IN-006	See which release-line stages no requirement covers and which requirements cannot be tested	ST-005	FR-012
IN-007	Name each approver, give the sponsor one document, bring back the sponsor's comments on it, and let the expert record their decision against a named version	ST-006	FR-013, FR-014, FR-021, FR-022, FR-023

Out of scope for REL-001

ID	OUT OF SCOPE	REVISIT
OUT-001	Deriving the backlog (J-002)	REL-002; the PRD already carries what it needs
OUT-002	Two experts on one set at once	When three real engagements have each reached an approved set that passed the buildability review on first submission
OUT-003	Any view beyond reading the documents and stepping through them beside the conversation	When two of the three sponsors in Q-005's spike say the root document is not enough
OUT-004	An engineering lead's own recorded confirmation	REL-002, with J-002

ID	OUT OF SCOPE	REVISIT
OUT-005	A reviewer's comments inside the set, not on an exported copy	REL-003, when OUT-002's condition holds; until then, comments come back from an exported copy under FR-022 and FR-023

Non-goals

NON-GOAL	WHY NOT
Supplying domain knowledge	The expert has it; the product supplies product management
Deciding scope, priority, or whether to build	The expert decides; the product shows the consequences
A backlog's tasks, subtasks, estimates, or sizing	They come from engineering design, not from a PRD
Replacing the document store, issue tracker, or approval system, or being the system of record	The organization already stores and approves documents, and the PRD stays readable with no part of this product running (NFR-004)
A separate application to open, install, or log into	The expert works in short bursts, and a second place to open would split the work (D-001)
Certifying a PRD	A grade is evidence for the expert's judgement, never a warrant

The first is the one that defines the product: **we do not supply domain knowledge.** Someone without the domain is not served by this and should not buy it. The last is the one that protects its users: **a grade is evidence, not a warrant.** A high score on a PRD for the wrong product is still the wrong product.

Two non-goals were constraints C-002 and C-001 in section 9 until 0.8.1. Product teams new to a domain, once segment S-003 in section 2, fall under the first non-goal: the product never supplies the domain.

In scope

IN-001

SUMMARY: State the idea and the domain's vocabulary once, and see every section of the set and its state

JOURNEY STAGE: ST-001

REQUIREMENTS:
- FR-001
- FR-002

IN-002

SUMMARY: Be taught each section before its interview, be asked one or two questions at a time with follow-ups, end a section early with its gaps recorded, and stop being asked once the section would meet the gate

JOURNEY STAGE: ST-002

REQUIREMENTS:
- FR-003
- FR-004
- FR-005
- FR-019

IN-003

SUMMARY Record no invented domain fact, and convert a stated technology into the need beneath it, kept for the engineering lead

JOURNEY STAGE ST-002

REQUIREMENTS

- FR-006
- FR-007

IN-004

SUMMARY Read the draft as prose with its grade and every blocking defect, refute a finding, move to any section, refresh to see a change, and be told when the set states one item two ways

JOURNEY STAGE ST-003

REQUIREMENTS

- FR-008
- FR-009
- FR-016
- FR-017
- FR-018
- FR-020

IN-005

SUMMARY Answer one question about the weakest part, see the section re-graded against the same rubric, and resume after a break without repeating an answer

JOURNEY STAGE ST-004

REQUIREMENTS

- FR-010
- FR-011
- FR-015

IN-006

SUMMARY Get a buildability review of the whole set, read as an engineer would, and see which release-line stages no requirement covers and which requirements cannot be tested

JOURNEY STAGE ST-005

REQUIREMENTS

- FR-012

IN-007

SUMMARY Let the expert name each approver, give the sponsor one document with the summary, the scope, the grades, and the open questions, bring the sponsor's comments on an exported copy back into the set with a response to each, and let the expert record the sponsor's decision, given outside the product, with any conditions, against a named version

JOURNEY STAGE ST-006

REQUIREMENTS

- FR-013
- FR-014
- FR-021
- FR-022
- FR-023

Out of scope

OUT-001 · Deriving features, epics and stories from the approved PRD - journey J-002

REVISIT — REL-002. The PRD must already carry what J-002 needs, which FR-012 checks, and nothing is lost by deferring the derivation itself.

OUT-002 · More than one expert working on the same set at the same time

REVISIT — Once the single-expert path is proven: three real engagements have each reached an approved set that passed the buildability review on first submission.

OUT-004

SUMMARY — An engineering lead reading the set and recording their own confirmation that it can be built from

REVISIT — REL-002, with J-002. REL-001 gives the expert the buildability review (FR-012), and the gaps are found before anyone estimates. A named lead's own confirmation is a second person's time, which REL-001 does not ask for.

OUT-003

SUMMARY — Any view of the document set other than reading the documents themselves and stepping through them beside the conversation

REVISIT — When at least two of the three sponsors in Q-005's spike say the root document is not enough to decide. Until then the sponsor reads one document, the root document, which FR-013 says carries everything they need.

OUT-005 · A reviewer commenting on a section inside the set, not on an exported copy

REVISIT — REL-003, when OUT-002's condition holds: three real engagements have each reached an approved set that passed the buildability review on first submission. Until then, the sponsor comments on an exported copy, and the comments come back under FR-022 and FR-023.

SECTION 7

Functional requirements

DRAFT 88/100 NO BLOCKING DEFECTS

ID	USER	CAPABILITY	PRIORITY	STAGE
FR-001	PER-001	State the idea; establish the vocabulary once	must	ST-001
FR-002	PER-001	See the whole set's state at a glance	should	ST-001
FR-003	PER-001	Be taught the section before being asked	must	ST-002
FR-004	PER-001	Be asked one or two questions at a time	must	ST-002
FR-005	PER-001	End a section and see what was left behind	should	ST-002
FR-006	PER-001	Unknown facts become open questions, never inventions	must	ST-002
FR-007	PER-001	Technology talk converted; preference preserved	must	ST-002
FR-008	PER-001	Read the draft as prose	should	ST-003
FR-009	PER-001	See score, blocking defects, weakest part	must	ST-003
FR-010	PER-001	Be asked one question, not twelve	must	ST-004
FR-011	PER-001	See the re-grade against the same standard	should	ST-004
FR-012	PER-001	See what a team could not build from it	must	ST-005
FR-013	PER-003	See what is approved and what is open, on one page	must	ST-006
FR-014	PER-003	Approve, with conditions, against a version	must	ST-006
FR-015	PER-001	Stop and resume without being asked anything twice	should	ST-004
FR-016	PER-001	Refute a finding and have it withdrawn	should	ST-003
FR-017	PER-001	Move to any section without losing their place	should	ST-003
FR-018	PER-001	See the set as it now stands	should	ST-003
FR-019	PER-001	Stop being asked once the section would meet the gate	should	ST-002
FR-020	PER-001	Be told when the set states one item two ways	should	ST-003
FR-021	PER-001	Name each approver before offering the set	must	ST-006
FR-022	PER-003	Comment on an exported copy, and see each comment's response	should	ST-006
FR-023	PER-001	Bring an exported copy's review comments into the set, and respond to each	should	ST-006

Every release-line stage has at least one `must` requirement. Eleven of the 23 are `must`: without any one of them, REL-001 does not do what it promises. The twelve `should` requirements make the work faster or easier to follow, and the

expert can still reach an approved set without them. The expert set this split at version 0.8.0, and set FR-021 to `must` at version 0.9.0.

Five of the `must` requirements wait on an open question in section 12. FR-004, FR-009, FR-010, and FR-012 wait on Q-009, which version of the rubric and the derivability checks release 1 grades against. The same four also wait on Q-015, who owns the rubric and the checks, because Q-009 cannot fix a version until Q-015 names its owner. FR-010 also waits on Q-007 and Q-014, whether writing the requirements sends the expert back to the journeys and how that feels, because both block ST-004. FR-013 waits on Q-005, whether a sponsor accepts a PRD that carries a grade. FR-014 no longer waits: Q-011 is answered, and in release 1 the expert records the sponsor's decision under FR-014. Three `should` requirements also wait: FR-011 and FR-019 on Q-009, and FR-015 on Q-010.

FR-015 and FR-016 were split out of FR-002 and FR-009 at version 0.2.0, because each has a failure mode of its own: a build could deliver perfect set-state visibility and still re-ask three answered questions, and could report a grade perfectly and still defend a finding the document does not support. Ids are never renumbered, so they sit at the end rather than in stage order.

FR-019 and FR-020 were split out of FR-004 and FR-008 at version 0.9.0, for the same reason. The early stop and the check across the set each have a failure mode of their own, and each needs its own priority. The expert set both to `should`. FR-019 also waits on Q-009, because the early stop applies the rubric.

FR-021 was added at version 0.9.0 on the expert's decision. A set records who will approve it before it is offered, so an approval names a person, not only the sponsor's role in section 10.

FR-022 was added at version 0.10.0 on the expert's decision, as a `should` in REL-001. A sponsor reads an exported copy outside the product at ST-006 and comments in it. FR-023 brings those comments back, and under FR-022 the next copy carries each response and what it changed. WF-011 shows that copy. A comment still without a response when the sponsor approves does not block the approval: the expert records it as a condition under FR-014, as the expert decided.

FR-023 was split out of FR-022 at version 0.11.0, as FR-015 and FR-016 were. FR-022 is what the sponsor sees, and FR-023 is what the expert does to bring the comments in and answer them. Each has a failure mode of its own: a build could import every comment and never show the sponsor a response, or show responses to comments it lost. The approval rule moved into FR-014, which records every condition, and the sponsor sets each condition's date.

A judgement call worth challenging

FR-002, FR-008 and FR-013 refer to "the set", "the section", "a table row" and "the root document". A reader may read those as naming an output format and therefore failing test 5. The judgement taken here is that a written PRD is the product's subject, not its delivery technology, so these words name the artifact rather than the mechanism. What is deliberately absent everywhere: the assistant the product runs inside, how grading is performed, who or what reviews the set, and what the documents are stored as. Those are in section 6 as non-goals, in section 11 as dependencies, or nowhere.

FR-001, FR-009, FR-013 and FR-014 each read as more than one capability, and a grading pass has called them out for it. They are left whole deliberately: each is one act with several parts, and a build delivering half of one would be visibly unfinished rather than arguably done. FR-015 and FR-016 were split precisely because they were not like that, each had a failure mode the other half could hide.

FR-001 · State the idea in their own words and establish the domain's vocabulary once

USER	PER-001
OUTCOME	Every question they are asked afterwards uses their unit of work and their terms, so they are never translating.
ACCEPTANCE	• Given an expert who says the unit of work is a "load", when any later question refers to that unit, then it says "load" and not "item", "record" or "entity". • Given three words the expert says an outsider would misread, when the first section is drafted, then each of those words appears in the glossary with the expert's own meaning. • Given an expert who has not stated a domain, when they ask to begin a section, then they are asked for the domain and the unit of work first.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-001
OBJECTIVE	O-003

FR-002 · See every section of the set and what state each one is in, at any point

USER	PER-001
OUTCOME	They can see at a glance what is done, what is weak and what has not been started, without opening any section.
ACCEPTANCE	• Given a set where four section documents are drafted and thirteen are pending, when the expert asks where things stand, then they see all seventeen section documents with each one's state, and the pending ones are shown as not started rather than omitted. • Given a section that has been graded, when the expert sees the set's state, then that section shows its score and how many blocking defects it has. • Given a set where no section has been started, when the expert asks where things stand, then every section is listed and shown as not started.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-001
OBJECTIVE	O-002

FR-003

CAPABILITY	Be taught what a section is, why it matters and what belongs in it, before being asked anything about it
USER	PER-001
OUTCOME	Before the first question of a section reaches them, what the section is for and what belongs in it is in front of them, in their own domain's terms.
ACCEPTANCE	• Given a section the expert has not done before, when they begin it, then they are first told what it is, why it matters and what belongs in it, in no more than ten lines. • Given a stated domain, when a section is taught, then the example used is in that domain and not in another. • Given an expert who says they already know the section, when they ask to skip the teaching, then the questions begin and the teaching is not forced.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-003

FR-004 · Be asked one or two questions at a time, never a questionnaire

USER	PER-001
OUTCOME	No message puts more than two questions to them, so each answer they give addresses one thing.
ACCEPTANCE	• Given any section's interview, when the expert is asked for input, then at most two questions are put to them at once. • Given an answer that leaves something the section needs still unanswered, or that could be written into the section in two ways differing in meaning, when it is received, then a follow-up is asked about that answer before the next topic is opened.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-003

FR-005 · Say that they have nothing more to add, and have that end the section's interview

USER	PER-001
OUTCOME	They are never trapped in an interview, and what they chose to leave incomplete is recorded rather than hidden.
ACCEPTANCE	• Given an expert who says they are done with a section, when the section does not yet meet the gate, then it is written anyway at draft, and what is missing is recorded as an open question or named plainly to them. • Given an expert who ends a section early, when they are told what was left behind, then the statement names the specific gap, not just a score.
PRIORITY	should
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-002

FR-006

CAPABILITY	Have a fact they do not know recorded as an open question with a decider and a date, rather than filled in for them
USER	PER-001
OUTCOME	Nothing in their PRD is a number or a citation that nobody can stand behind.
ACCEPTANCE	• Given an expert who says they would have to check a figure, when the section is drafted, then no figure is written in its place and an open question records who could find out and by when. • Given an objective whose baseline is unknown, when the section is drafted, then the baseline says so explicitly and an open question exists for it. • Given a regulation or standard the expert has not named, when a constraint is drafted, then no regulation number appears that the expert did not supply.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-001

FR-007

CAPABILITY	State a technology they want and have the need beneath it recorded instead, with their preference preserved for the engineering lead
USER	PER-001
OUTCOME	Their requirement names no technology, their reasoning is not lost, and they do not feel overruled.
ACCEPTANCE	• Given an expert who names a technology, when they are asked what the user gets from it and answer, then the recorded requirement states that outcome and names no technology, component, vendor or screen. • Given a technology the expert stated, when the set is handed over, then that preference appears on a list of candidate architecture decisions for the engineering lead. • Given an expert who insists the technology belongs in the requirement, when they have been told once why it does not, then their decision is recorded and is not argued again at the next grading pass.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-001

FR-008 · Read the drafted section as prose rather than as a filled-in form

USER	PER-001
OUTCOME	They can judge from the prose alone whether it says what they meant.
ACCEPTANCE	• Given a drafted section, when the expert reads it, then every item is shown in a sentence or in a table row under named column headings, and no item is shown only as a bare list of names and values.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-003

FR-009 · See a score, every blocking defect by name, and which part of the section is weakest

USER	PER-001
OUTCOME	They know how good the section is against a stated standard, and what specifically is wrong, rather than being told it looks good.
ACCEPTANCE	• Given a graded section, when the expert reads the grade, then they see a score, a count of blocking defects, and each blocking defect stated with where it is. • Given a section scoring 70 or more with one blocking defect, when the grade is reported, then the section is not offered as ready, and the defect is named. • Given a graded section, when the expert reads the grade, then the weakest part is named and why it is weak is stated in no more than three lines. • Given a grade that has just been reported, when the expert reads it, then they are asked to accept it or refute it, under TR-014.
PRIORITY	must
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-001

FR-010

CAPABILITY	Be asked the single question that would most improve the section, not a list of everything wrong with it
USER	PER-001
OUTCOME	They can act immediately, and the section improves one answer at a time instead of stalling under a list.
ACCEPTANCE	• Given a graded section scoring under 70, or carrying any blocking defect, when the expert is asked to improve it, then exactly one question is put to them. • Given that one question, when it is asked, then it addresses the weakest part that was just named, and not a different part.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-004
OBJECTIVE	O-003

FR-011 · See the section re-graded against the same standard after answering

USER	PER-001
OUTCOME	They can see whether their answer actually improved the section, and by how much.
ACCEPTANCE	• Given an answered question, when the section is re-graded, then it is scored against the same rubric as before and both scores are visible. • Given a re-grade, when the score has risen, then the document has changed; a score never rises without a change to the document. • Given a section scoring 70 or more with no blocking defects, in a set that checks clean, when it is re-graded, then the expert is offered the state `review`, and is told that `review` means the section is ready for someone else to read.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-004
OBJECTIVE	O-003

FR-012

CAPABILITY	See what a development team would not be able to build from the set, before anyone has estimated it
USER	PER-001
OUTCOME	They learn where their own PRD fails while it is still cheap to fix, rather than at iteration three.
ACCEPTANCE	• Given a set in which every section has been written at least once, when the expert asks for a buildability review, then they see a verdict and, for each release-line stage, whether a requirement covers it or it is a gap. • Given a requirement whose acceptance criteria an engineer who was not in the interview judges insufficient to write a test from, when the buildability review reports it, then the requirement is named and the document is quoted. • Given an open question that blocks a `must` requirement, when the buildability review runs, then that requirement's wait on the question is reported. • Given a buildability review finding, when it is reported to the expert, then it has been checked against the document and anything not supported by the document has been dropped.
PRIORITY	must
SHOWN IN	report
JOURNEY STAGE	ST-005
OBJECTIVE	O-001

FR-013

CAPABILITY	See, in the set's own root document, what is being approved, what is still undecided and who owns each undecided thing
USER	PER-003
OUTCOME	They can commit budget and a date knowing what they are committing to and what is still open.
ACCEPTANCE	• Given a set offered for approval, when the approver opens its root document, then that one document carries the summary, the scope, every section's state and grade, and every open question with its decider and date, without the approver opening any other document. • Given a set with any blocking defect anywhere, when the approver opens it, then they see that, with the sections it is in.
PRIORITY	must
SHOWN IN	report
JOURNEY STAGE	ST-006
OBJECTIVE	O-001

FR-014

CAPABILITY	Approve, approve with conditions, or reject, and have that recorded against a named version
USER	PER-003
OUTCOME	What was approved, when, by whom and on what conditions is answerable later without asking anyone.

ACCEPTANCE

- Given an approver who approves with conditions, when they state the conditions, then each condition names the open question, requirement, or review comment it depends on and a date.
- Given a recorded approval, when the set is read later, then the version approved, the approver, their role, the date and the decision are all present.
- Given an approved set that is then changed, when the change is written, then a new version is recorded and the earlier approval is not represented as covering it.
- Given an approver who rejects the set, when the expert records it, then the rejection, its reason, the approver, and the date are recorded against the version, and the set reads as not approved.
- Given a sponsor who approves while a review comment has no response, when the expert records the approval, then each such comment is recorded as a condition on it, with the comment's id and the date the sponsor sets for it, and only the sponsor clears it.

PRIORITY: must

SHOWN IN: report

JOURNEY STAGE: ST-006

OBJECTIVE: O-001

FR-015 · Stop in the middle of a set and resume days later without being asked anything twice

USER: PER-001

OUTCOME: They pick up at the first unfinished section, and nothing they have already told the product is asked for again.

ACCEPTANCE

- Given an expert returning after eight days, when they say they want to carry on, then work resumes at the first section not yet at `review`.
- Given an expert who has already stated their domain and vocabulary, when they resume, then they are not asked for either again and the vocabulary in use is unchanged.
- Given an expert who answered a question in a section before stopping, when they resume that section, then that question is not put to them again.

PRIORITY: should

SHOWN IN: conversation

JOURNEY STAGE: ST-004

OBJECTIVE: O-002

FR-016 · Refute a grading finding by pointing at the document, and have it withdrawn

USER: PER-001

OUTCOME: They are told the finding was dropped, and it does not reappear at the next grading pass.

ACCEPTANCE

- Given a finding the expert refutes, when it is checked against the document and the document does not support it, then it is dropped and the expert is told it was dropped.
- Given a finding the document does support, when the expert refutes it, then the reason is stated once and their decision is recorded.
- Given a decision the expert reaffirmed after an objection, when the section is graded again, then the objection is not raised as a new finding.

PRIORITY: should

SHOWN IN: conversation

JOURNEY STAGE: ST-003

OBJECTIVE: O-001

FR-017

CAPABILITY	Move from the section in front of them to any other section of the set without losing their place
USER	PER-001
OUTCOME	They reach the section they want in one step, and can always return to the view they were on before.
ACCEPTANCE	• Given a section in front of them, when they ask for the next or the previous section, then the section after or before it in reading order is shown, and at the first or the last section the view does not change. • Given a section that cites an id another section defines, such as Q-001, when they follow that id, then the section that defines it is shown and the id is named above it. • Given they have moved from a section to three other things in the set, one after another, when they ask to go back three times, then they see each earlier thing again, the most recent first, and end at the section they started from.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-002

FR-018 · See the set as it now stands after it was changed outside the view in front of them

USER	PER-001
OUTCOME	After they ask to refresh, what they read is the set as it now stands.
ACCEPTANCE	• Given a section changed after they opened the set, when they ask to refresh, then within 10 seconds that section shows its new text and its new state. • Given a section that is graded again while the set is open, when they ask to refresh, then within 10 seconds the new grade is shown. • Given no set existed when they looked, when a set is written and they ask to refresh, then every section of it is listed.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-002

FR-019 · Stop being asked once the section would already meet the gate

USER	PER-001
OUTCOME	The interview ends when the section has what it needs, so they are not asked questions it no longer needs.
ACCEPTANCE	• Given an expert whose answers would already score the section 70 or more with no blocking defects, when further questions remain unasked on the question list, then the interview stops rather than continuing to the end of it. • Given an expert whose answers would score the section under 70, or with a blocking defect, when questions remain on the question list, then the next question is put to them. • Given an answer that FR-004 says needs a follow-up, when the answers would already score the section 70 or more with no blocking defects, then the follow-up is asked first, and the interview stops only after it is answered.
PRIORITY	should
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-003

FR-020 · Be told when the set states one item two different ways

USER	PER-001
OUTCOME	They see every id, priority or link that two places in the set state differently, before anyone builds from either.
ACCEPTANCE	• Given an id, priority or link that two places in the set state differently, when the section is read, then the expert is told of the difference rather than it being silently resolved. • Given a set in which every id, priority and link is stated the same way wherever it appears, when the section is read, then no difference is reported.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-003

FR-021 · Name each approver of a set, by name and role, before offering it for approval

USER	PER-001
OUTCOME	The set records who will approve it, and nobody else's decision counts as an approval.
ACCEPTANCE	• Given a set the expert is about to offer for approval, when they name an approver, then the set records that approver's name and role, and the root document shows them. • Given a set with no approver named, when the expert asks to offer it for approval, then it is not offered, and they are told an approver must be named first. • Given a decision from someone the set does not name as an approver, when the expert records it, then it is not recorded as an approval.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-006
OBJECTIVE	O-002

FR-022

CAPABILITY	Comment on an exported copy of the set without using the product, and see each comment's response in the next copy
USER	PER-003
OUTCOME	A sponsor who never opens the product sees every comment they left in the next copy: with its status, accepted, declined, or answered, its response, and what changed, or marked open while it has no response.
ACCEPTANCE	• Given a sponsor who comments on an item in an exported copy, when they read the next exported copy, then they see the comment against that item with their name, the date, and the text they marked. • Given a comment the expert responded to, when the sponsor reads the next exported copy, then they see the comment, its status, the response, and the ids it changed. • Given a comment with no response yet, when the sponsor reads the next exported copy, then they see the comment marked open.
PRIORITY	should
SHOWN IN	report
JOURNEY STAGE	ST-006
OBJECTIVE	O-002

FR-023 · Bring the review comments in an exported copy into the set, and record a response to each

USER	PER-001
OUTCOME	Every review comment left in an exported copy is in the set once, against the item it marks, and none is lost in an inbox.
ACCEPTANCE	• Given an exported copy with review comments, when the expert brings them into the set, then each one is recorded against the item it marks, with the name of the person who left it, the date, and the text they marked. • Given a copy whose comments were already brought in, when the expert brings it in again, then no comment is recorded twice. • Given a comment that marks no text, when the expert brings the copy's comments into the set, then it is recorded against the section it sits in. • Given a comment in the set, when the expert records a response, then the comment records its status, accepted, declined, or answered, the response, and the ids that changed.
PRIORITY	should
SHOWN IN	conversation
JOURNEY STAGE	ST-006
OBJECTIVE	O-002

SECTION 7

Wireframes

DRAFT 91/100 NO BLOCKING DEFECTS

Every requirement in section 7 whose outcome a person meets on a screen or in a report is drawn here as a text sketch. Each sketch shows what the person sees and what they can do there, with the set's own content. None of them decides how anything looks or how it is built.

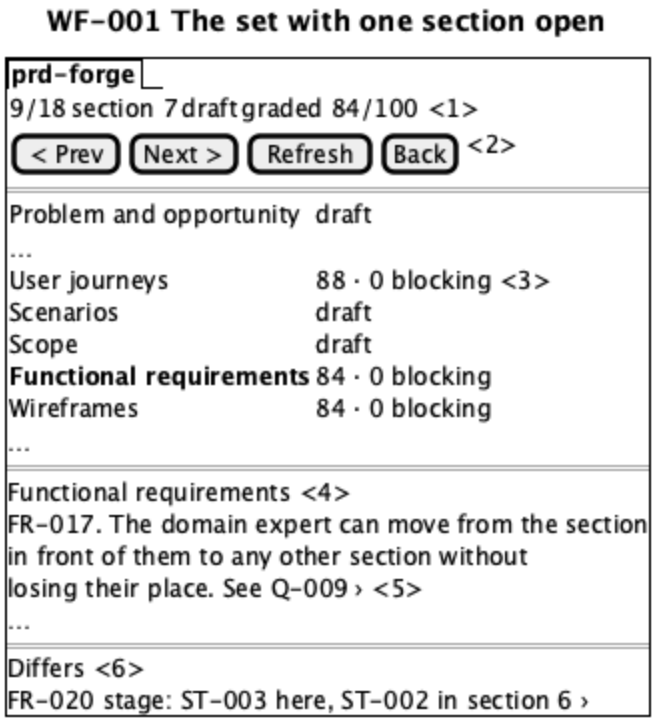
Five sketches follow the domain expert through the golden thread: a new set with nothing started, the set with a section open, why a section got its grade, the rubric behind it, and a re-grade. Two reports decide whether the set is buildable: the buildability review and the approval pack. A third, the review comments with their responses, is what the sponsor reads in the next exported copy. Three smaller sketches draw the other states: no set at all, which FR-018's third criterion names; the review that finds gaps, which SC-005 names; and the return after a week away, which SC-002 names.

The numbered marks in each sketch are explained, in order, by its `shows` list. Each sketch is also drawn for the printed PRD.

WF-001 The set with one section open

Illustration, not a design. Illustrates FR-002, FR-008, FR-017, FR-018, and FR-020. PER-001 sees it at ST-003.

The expert reads a drafted section with the whole set beside it. Every graded section shows its score and its blocking count, so what is weak is visible without opening anything. When two places in the set state one id, priority, or link differently, mark 6 lists each difference with both values. The difference drawn is an example, not a fact of this set.

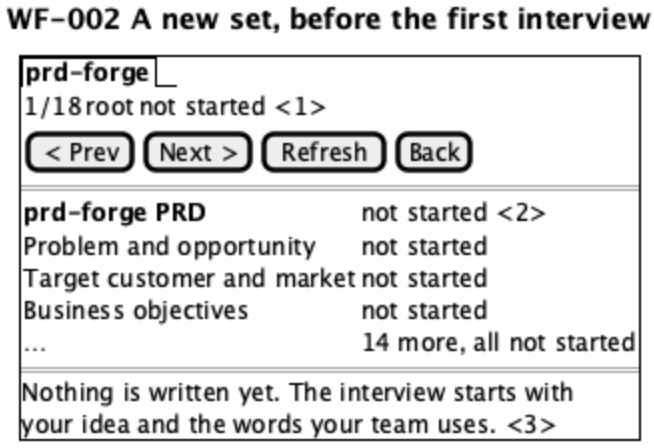


WF-001 drawing

WF-002 A new set, before the first interview

Illustration, not a design. Illustrates FR-002, FR-017, and FR-018. PER-001 sees it at ST-001.

At ST-001 the set exists and nothing in it is written. Every section is listed and marked not started, so the expert sees the whole shape of the work before the first question.

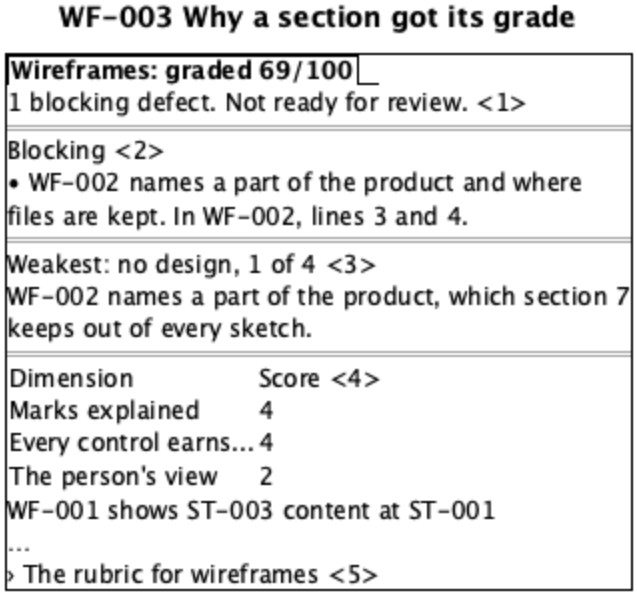


WF-002 drawing

WF-003 Why a section got its grade

Illustration, not a design. Illustrates FR-009 and FR-017. PER-001 sees it at ST-003.

A score alone does not tell the expert what to fix. This view names each blocking defect and where it is, and names the weakest part, quoting the section.



WF-003 drawing

WF-004 The standard a section is graded against

Illustration, not a design. Illustrates FR-009 and FR-017. PER-001 sees it at ST-003.

A grade against a stated standard is only fair if the expert can read the standard. This is the same rubric the graders used.

WF-004 The standard a section is graded against

Rubric for functional requirements Each dimension scores 0, 2 or 4. <1>			
Dimension	0	2	4 <2>
Atomicity	Several in one	Mostly single capability	One
Coverage	Stages with none	Most covered	Every stage has a must
...			
Blocking defects <3> • A requirement that names a technology ...			

WF-004 drawing

WF-005 What a team could not build from this set

Illustration, not a design. Illustrates FR-012. PER-001 sees it at ST-005.

The values are prd-forge's own at version 0.4.0, a snapshot. The expert reads this before anyone has estimated anything, while a fix is still a sentence. Every finding quotes the set, so the expert can check it rather than trust it.

WF-005 What a team could not build from this set

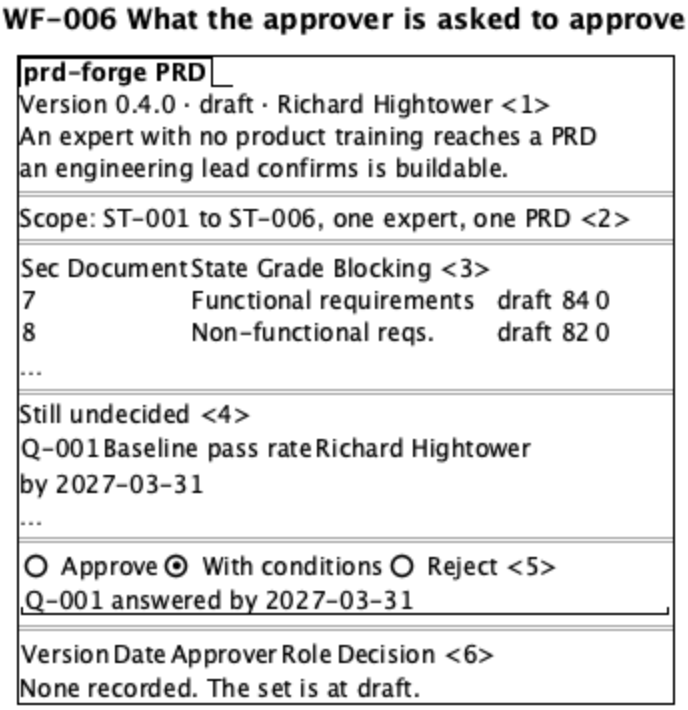
What a team could not build from this set prd-forge PRD 0.4.0. Buildable with gaps. <1>
Release line <2> ✓ ST-001 FR-001, FR-002, FR-015 ✓ ST-003 FR-008, FR-009, FR-016, FR-017, FR-018 ✓ ST-006 FR-013, FR-014 ...
Cannot write a test from <3> 1. FR-004: "two ways differing in meaning", no judge 2. FR-014: no test for a rejection
Blocked by an open question <4> Q-009 blocks FR-012: which rubric version applies. Richard Hightower, by 2027-01-31
Ask for first <5> 1. Who judges FR-004's two meanings

WF-005 drawing

WF-006 What the approver is asked to approve

Illustration, not a design. Illustrates FR-013 and FR-014. PER-003 sees it at ST-006.

The values are prd-forge's own at version 0.4.0, a snapshot. The approver reads one document, not eighteen. It says what they are committing to, what is still open and who owns it, and where their decision is recorded.

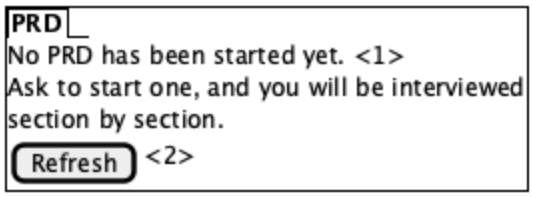


WF-006 drawing

WF-007 Before any set exists

Illustration, not a design. Illustrates FR-018. PER-001 sees it at ST-001.

Before the first interview there is no set to list. The view says so plainly, and a refresh lists every section once the set is written.

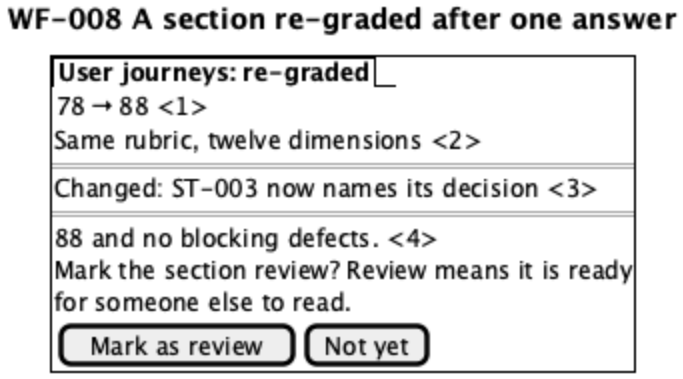


WF-007 drawing

WF-008 A section re-graded after one answer

Illustration, not a design. Illustrates FR-011. PER-001 sees it at ST-004.

After one answer the section is graded again against the same rubric. Both scores are visible, and at 70 or more with no blocking defects the expert is offered review. The values are SC-001's.

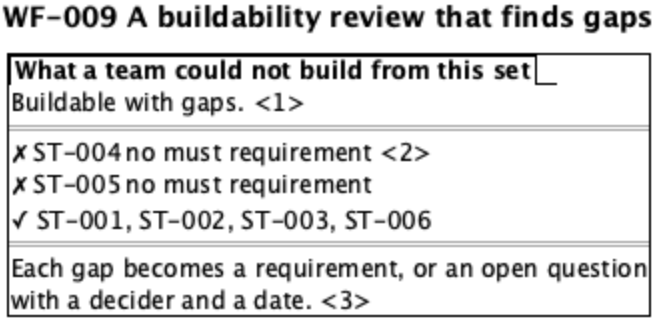


WF-008 drawing

WF-009 A buildability review that finds gaps

Illustration, not a design. Illustrates FR-012. PER-001 sees it at ST-005.

SC-005 is the buildability review that finds work missing: two release-line stages have no must requirement behind them. The gaps are named first, before anything else in the buildability review.

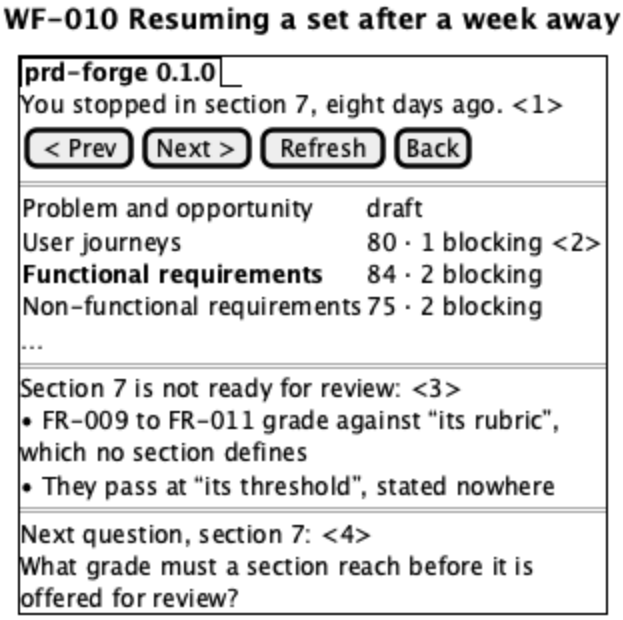


WF-009 drawing

WF-010 Resuming a set after a week away

Illustration, not a design. Illustrates FR-002, FR-009, FR-015, FR-017, and FR-018. PER-001 sees it at ST-004.

SC-002 is the expert who returns after eight days and remembers nothing. The values are prd-forge's own at version 0.1.0, when section 7 scored 84 with two blocking defects: a grade of 70 or more that is still not offered as ready.

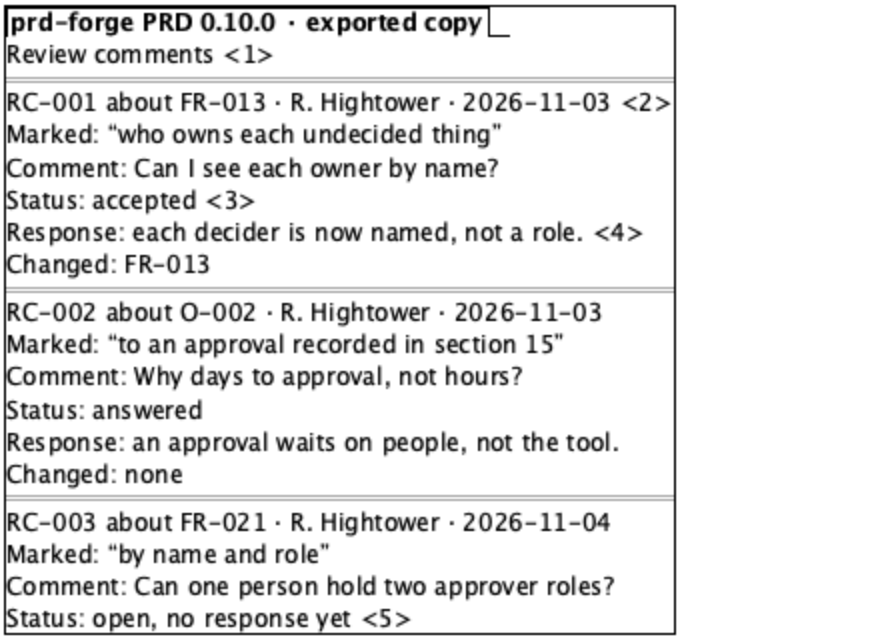


WF-010 drawing

WF-011 Review comments with their responses in the next exported copy

Illustration, not a design. Illustrates FR-022. PER-003 sees it at ST-006.

The sponsor commented on an earlier copy, and this is the copy they read next. Each marked text is the item's own wording. The comments, responses, and dates are invented, to show the form.



WF-011 drawing

WF-001 · The set with one section open

TYPE screen

PERSONA PER-001

JOURNEY STAGE ST-003

ILLUSTRATES

- FR-002
- FR-008
- FR-017
- FR-018
- FR-020

SHOWS

- Where they are in the set, the section's state, and its grade
- Next, previous, refresh, and back to the view before
- Each section's grade and blocking count; opening it shows why
- The section itself, as prose
- An ID that another section defines opens that section
- Each place where the set states one id, priority, or link two ways, with both values, and each place can be followed

WF-002 · A new set, before the first interview

TYPE screen

PERSONA PER-001

JOURNEY STAGE ST-001

ILLUSTRATES

- FR-002
- FR-017
- FR-018

SHOWS

- Where they are, and that nothing is started yet
- Every section listed, each marked not started
- What happens first

WF-003 · Why a section got its grade

TYPE screen

PERSONA PER-001

JOURNEY STAGE ST-003

ILLUSTRATES

- FR-009
- FR-017

SHOWS

- The score, and that the section is not ready
- Each blocking defect, and where it is
- The weakest part, and why, in three lines or fewer
- Each dimension's score, with the evidence for it
- The standard the grade was given against

WF-004 · The standard a section is graded against

TYPE	screen
PERSONA	PER-001
JOURNEY STAGE	ST-003
ILLUSTRATES	• FR-009 • FR-017
SHOWS	• How a dimension becomes a grade • What 0, 2 and 4 look like for each dimension • What stops a section at any grade

WF-005 · What a team could not build from this set

TYPE	report
PERSONA	PER-001
JOURNEY STAGE	ST-005
ILLUSTRATES	• FR-012
SHOWS	• Which version of the set the buildability review read, and the verdict • Each release-line stage, covered or a gap • Each requirement an engineer could not test, quoted • An open question that blocks a must, its decider and date • What the buildability review would ask for first

WF-006 · What the approver is asked to approve

TYPE	report
PERSONA	PER-003
JOURNEY STAGE	ST-006
ILLUSTRATES	• FR-013 • FR-014
SHOWS	• The version, its state, its owner, and what it is for • What release 1 covers • Every section with its state and grade, and what blocks • What is still undecided, who decides it, and by when • The decision, with any condition tied to a question • Who decided what, in which role, on which version

WF-007 · Before any set exists

TYPE	screen
PERSONA	PER-001
JOURNEY STAGE	ST-001
STATE OF	WF-001
ILLUSTRATES	• FR-018
SHOWS	• That nothing has been started • Look again once a set exists

WF-008 · A section re-graded after one answer

TYPE	screen
PERSONA	PER-001
JOURNEY STAGE	ST-004
ILLUSTRATES	• FR-011
SHOWS	• The new score beside the old one • The same standard both times • What changed in the section • The offer of review, and what review means

WF-009 · A buildability review that finds gaps

TYPE	report
PERSONA	PER-001
JOURNEY STAGE	ST-005
STATE OF	WF-005
ILLUSTRATES	• FR-012
SHOWS	• The verdict • The release-line stages with no requirement, named • That each gap must become a requirement or a question

WF-010 · Resuming a set after a week away

TYPE	screen
PERSONA	PER-001
JOURNEY STAGE	ST-004
STATE OF	WF-001
ILLUSTRATES	• FR-002 • FR-009 • FR-015 • FR-017 • FR-018
SHOWS	• Where they stopped • Every section's grade and blocking count • Each blocking defect in the section, named • The next question, not one already answered

WF-011 · Review comments with their responses in the next exported copy

TYPE	report
PERSONA	PER-003
JOURNEY STAGE	ST-006
ILLUSTRATES	• FR-022
SHOWS	• The comments, together in the copy • Each comment, with what it is about, who left it, when, and the text they marked • Its status, accepted, declined, or answered • The response, and what changed because of it • A comment with no response yet, marked open

SECTION 8

Non-functional requirements

DRAFT 89/100 NO BLOCKING DEFECTS

ID	CATEGORY	TARGET
NFR-001	performance	Next question under 20s at p95
NFR-002	performance	Section grade under 90s at p95
NFR-003	scale	120 requirements, 10 journeys x 10 stages, 200 terms
NFR-004	durability	100% of a set readable with no part of the product running
NFR-005	usability	90% of sections completed without asking what a term means
NFR-006	accessibility	Zero WCAG 2.2 AA violations on the golden-thread path
NFR-007	observability	100% of grades reconstructable after the session
NFR-008	privacy	Zero destinations for the set's text or the expert's answers; domain research sends only standard and term names, and can be turned off
NFR-009	supportability	100% of older sets readable and gradable within a major product release
NFR-010	retention	100% retrievable from the organisation's own store, no expiry imposed by us
NFR-011	availability	Zero answers lost per interruption; resume within 1 minute
NFR-012	security	Zero read paths the organisation's own access rules do not govern

Where these came from

NFR-001 and NFR-002 come from PER-001's frequency: a burst of intense use. A 20-second wait is tolerable once; across 60 questions it is 20 minutes of dead time and the expert stops.

NFR-004 is the one that matters most and is cheapest to get wrong. PER-001 works in bursts weeks apart. If the set is only readable through this product, then a version change, a tool change or an expired session destroys weeks of an expert's thinking. The target is deliberately absolute.

NFR-005, NFR-006, NFR-007, NFR-008, and NFR-012 need people rather than instruments to measure. NFR-005 is the one most likely to be skipped, and it is the direct test of whether P-004 was solved.

NFR-006 covers the approver too, not only the expert: PER-003 reads this once and must not need a mouse to do it.

There is no concurrency target, and that is a decision rather than an omission. Under D-001 the product reaches one expert inside a session they are already in, so the session, not the product, carries simultaneous users; NFR-003 therefore reads scale as the size of one set, which is the dimension we do control. If D-001 or the separate-application non-goal in section 6 ever changes, concurrency comes back and this paragraph goes.

NFR-006 and NFR-012 set targets on things the non-goals in section 6 say we do not own: what the session renders, and where the customer keeps its documents. They are kept as `must` because a target nobody owns is how accessibility and access control get lost. Who owns them is Q-017 for NFR-006 and Q-021 for NFR-012.

NFR-010, NFR-011 and NFR-012 were added at version 0.2.0. Retention, availability and security had simply not been reached, and C-003, now assumption A-005, stated a retention rule that no measurable requirement stood behind. NFR-010 is the one that answers A-005's second half: the product retains nothing of its own, so the only retention question is how long the organization keeps the documents, and the answer must be that we impose no expiry on them.

NFR-011 waits on Q-010, whether an answer must survive the end of the session before its section is drafted, and NFR-007 waits on Q-016, whether a grade must show the rubric version it was made against.

Requirements

NFR-001

STATEMENT	An expert answering a question does not lose their train of thought waiting for the next one.
CATEGORY	performance
METRIC	95th and 99th percentile elapsed time from the expert submitting an answer to the next question or the drafted section appearing.
TARGET	under 20 seconds at p95, and under 60 seconds at p99
MEASUREMENT	At least 300 timed observations per release, taken across five scripted runs of all seventeen section documents on a set already holding 60 requirements, from the expert's own side of the conversation rather than from inside the product.
PRIORITY	must
APPLIES TO	• FR-004 • FR-010

NFR-002 · A grading pass does not make the expert wait so long that they stop asking for one.

CATEGORY	performance
METRIC	95th percentile elapsed time from the moment a section's draft is written to its grade being reported to the expert, measured from the expert's own side of the conversation, per grading pass on the section just drafted.
TARGET	under 90 seconds at p95 for a single section
MEASUREMENT	At least 200 grading passes per release, spread across all seventeen section documents and across sets of 20, 60 and 120 requirements.
PRIORITY	must
APPLIES TO	• FR-009 • FR-011

NFR-003 · A set stays workable at the size real products reach.

CATEGORY	scale
METRIC	The size of one set at which every requirement in section 7 still meets its acceptance criteria and every target in this section still holds.
TARGET	120 functional requirements, 10 journeys of 10 stages, and 200 glossary terms
MEASUREMENT	A generated set at those sizes, run through every function in section 7, per release.
PRIORITY	must

NFR-004 · Work survives the expert stopping, and survives the tool they were using going away.

CATEGORY	durability
METRIC	Share of a set's content recoverable after the authoring session ends, with no part of the product running.
TARGET	100%, readable without this product
MEASUREMENT	End a session mid-section, then read the whole set with nothing but a text editor, once per release.
PRIORITY	must
APPLIES TO	• all

NFR-005

STATEMENT	An expert with no product management training can complete a section without asking anyone what a word means.
CATEGORY	usability
METRIC	Share of sections a first-time expert completes without asking for the meaning of a product management term.
TARGET	90% of sections, measured over 10 first-time experts
MEASUREMENT	Observed sittings with 10 first-time experts before release, questions to the observer counted and categorised.
PRIORITY	must

NFR-006

STATEMENT	Everything the expert and the approver read is usable with a screen reader and without a mouse.
CATEGORY	accessibility
METRIC	Conformance to WCAG 2.2 AA of everything the product puts in front of the expert or the approver along the golden-thread journey.
TARGET	zero level-A or level-AA violations on the golden-thread path
MEASUREMENT	One manual screen-reader pass of the golden-thread journey per release, plus an automated check in the release pipeline of everything that journey shows.
PRIORITY	must

NFR-007 · An expert can tell why a section received the score it did, after the fact.

CATEGORY	observability
METRIC	Share of grades for which the score per rubric dimension and the quotation behind each can be recovered after the session.
TARGET	100% for any section at review or better
MEASUREMENT	Re-read a graded set a week later and reconstruct every score, once per release.
PRIORITY	should
APPLIES TO	• FR-009

NFR-008

STATEMENT	What the expert says about their domain stays within whatever boundary their organization already works in; this product adds no new place for it to go. The one outbound request is the domain check's research, on by default, which carries only the names of standards, terms of art, and the domain. An expert can turn it off for one set under TR-015, and an organization's own policy can forbid it for every set.
CATEGORY	privacy
METRIC	Number of destinations beyond the boundary A-005 names that the set's text or the expert's answers reach. An outbound query counts only when it holds more than the names of standards, terms of art, and the domain.
TARGET	zero
MEASUREMENT	Review of every outbound destination per release, recorded against A-005, and a review of the queries the domain check sent while grading the reference set, per release.
PRIORITY	must

NFR-009

STATEMENT	A set written by an earlier release of the product is still readable and still gradable by a later release.
CATEGORY	supportability
METRIC	Share of sets written by product release N that product release N+1 reads and grades without the expert rewriting anything.
TARGET	100% within a major product release
MEASUREMENT	Keep one set per product release and run the current product release against all of them, per product release.
PRIORITY	should

NFR-010

STATEMENT	A set remains readable and retrievable for as long as the organization keeps the PRDs it builds from, without this product being involved.
CATEGORY	retention
METRIC	Share of a set retrievable and readable from the organization's own document store, with no part of this product present.
TARGET	100%, for as long as the organization retains the set; no expiry imposed by us
MEASUREMENT	Retrieve a set written one release ago from the document store of a test organization, set up for the measurement because no customer organization exists at release, and read it with a text editor, once per release.
PRIORITY	must
APPLIES TO	• all

NFR-011

STATEMENT	An interview in progress survives the product being unavailable; no session loses an answer the expert has already given.
CATEGORY	availability
METRIC	Answers lost per interruption, and time to resume after one. An interruption is the product becoming unavailable in the middle of a section. The resume clock starts when the product is available again and the expert asks to carry on.
TARGET	Zero answers lost, on every interruption; resume within 1 minute at p95 and within 3 minutes at p99
MEASUREMENT	Interrupt a session mid-section 100 times per release, spread across all seventeen section documents, and count answers lost and seconds to resume.
PRIORITY	must
APPLIES TO	• FR-015

NFR-012

STATEMENT	Nobody but the people the organization has already given access to its documents can read a set.
CATEGORY	security
METRIC	Number of ways to read a set that the organization's own document access rules do not govern.
TARGET	zero
MEASUREMENT	Enumerate every read path to a set per release, in a test organization set up for the measurement because no customer organization exists at release, and check each against that organization's access rules. Recorded against the system-of-record non-goal in section 6 and A-005.
PRIORITY	must

SECTION 9

Constraints

DRAFT 95/100 NO BLOCKING DEFECTS

This section holds no constraints. No rule imposed from outside the team binds prd-forge itself. Each rule that this section once held was a decision of our own, a belief about customers, or the product's chosen response to a risk. The team could change each one by agreeing among themselves, so each one now lives where that kind of item belongs. The ids C-001 to C-004 are not reused.

An expert who uses prd-forge can still be bound by a professional code, a contract, or a firm policy. Those constraints belong in that expert's own PRD, not in this one.

Where the rules live now

At version 0.8.1, C-001 to C-003 left this section.

- C-001, no separate application for the expert to open, install, or log into, is a non-goal in section 6. The session it needs is dependency D-001 in section 11.
- C-002, the output is plain documents and the product is not the system of record, is a non-goal in section 6.
- C-003, domain content stays in the expert's session and in the documents, is assumption A-005 in section 11.

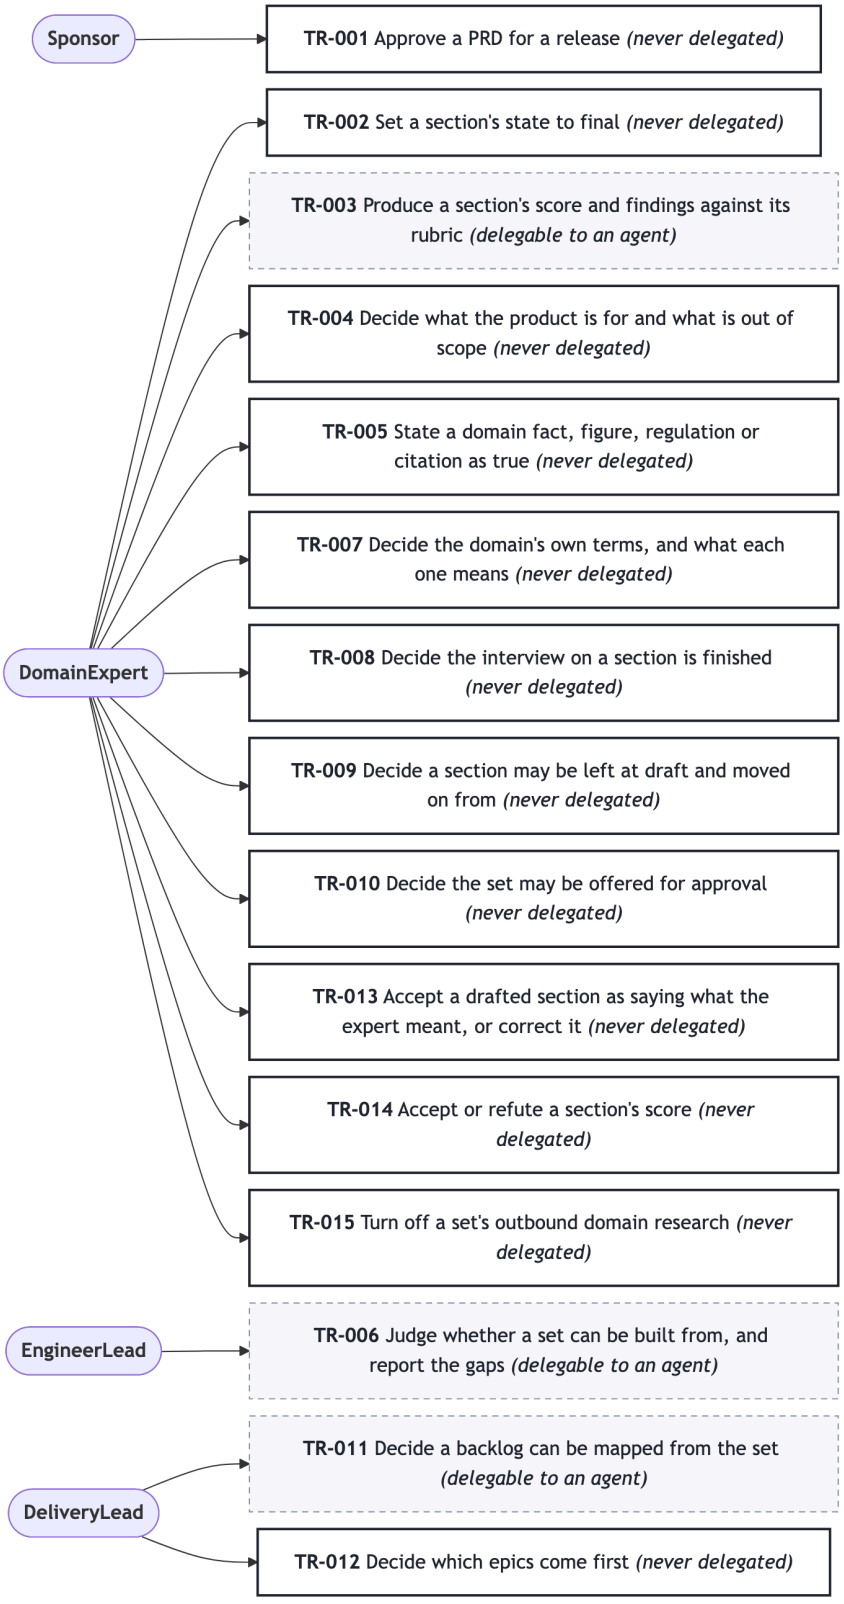
At version 0.9.0, C-004 left this section.

- C-004, no regulation, standard, clause, or statistic is stated as fact unless the expert supplied it or it was checked against its source, is trust rule TR-005 in section 10 and requirement FR-006 in section 7. TR-005 makes the expert the only one who can state such a fact as true, and it is never delegated. FR-006 records a fact the expert does not know as an open question with a decider and a date.

SECTION 10

Decision rights and trust rules

DRAFT 84/100 NO BLOCKING DEFECTS



Who decides what

ID	DECISION	DECIDER	DELEGABLE TO AN AGENT	NEVER DELEGATED
TR-001	Approve a PRD for a release	PER-003	no	yes
TR-002	Set a section to final	PER-001	no	yes
TR-003	Produce a section's score and findings	PER-001	**yes**	no
TR-004	Decide scope	PER-001	no	yes
TR-005	State a domain fact as true	PER-001	no	yes
TR-006	Judge buildability, and report the gaps	PER-002	**yes**	no
TR-007	Decide the domain's own terms	PER-001	no	yes
TR-008	Decide a section's interview is finished	PER-001	no	yes
TR-009	Decide a section may be left at draft	PER-001	no	yes
TR-010	Decide the set may be offered for approval	PER-001	no	yes
TR-011	Decide a backlog can be mapped	PER-004	**yes**	no
TR-012	Decide which epics come first	PER-004	no	yes
TR-013	Accept or correct a drafted section	PER-001	no	yes
TR-014	Accept or refute a section's score	PER-001	no	yes
TR-015	Turn off a set's outbound domain research	PER-001	no	yes

Every decision named in a stage of section 5 appears here. TR-007 to TR-013 were added at version 0.2.0 to close that: five release-line stage decisions and both of J-002's had no rule, which meant gates a development team would have had to invent. TR-013 covers the one in the middle of the loop, the expert accepting that a draft says what they meant, which is distinct from scoring it (TR-003), from accepting or refuting the score (TR-014), from leaving it at draft (TR-009) and from setting it final (TR-002).

TR-015 was added at version 0.12.0. It settles a conflict between the domain check and NFR-008: the domain check researches the domain by default, because a citation it cannot confirm is worth less to the expert. The expert can turn it off for one set, and an organization's own policy can forbid it for every set. NFR-008 and A-005 now state that default.

TR-001's decider is the sponsor's role, PER-003. FR-021 names the person: before a set is offered for approval, the expert names each approver by name and role, and only a named approver's decision counts as an approval.

The shape of the trust boundary

Only three of fifteen are delegable, and all three are **assessments**, not decisions: TR-003 produces a document's score, which the expert then accepts or refutes under TR-014, TR-006 judges whether it can be built from, and TR-011 judges whether a backlog can be mapped from it. All three produce evidence for a person. None changes a document, and none can set a state.

Everything that commits anyone to anything, such as approving, finalising, choosing scope, or asserting a fact, is never delegated. TR-005 is the strictest: a fact nobody can stand behind is not written at all, by anyone.

The delegable assessments all carry the same two obligations: **quote the document for every claim**, and **drop any finding the document does not support**. An assessment that cannot be refuted by pointing at the text is not evidence, it is an assertion, and the expert will correctly stop believing it. See R-002 in section 11.

TR-001 · Approve a PRD for a release

STAGES	• ST-006
APPROVER	PER-003
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The summary, the scope, every section's state and grade, every blocking defect, and every open question with its decider and date.
ESCALATION	No substitute approver, and approval waits. Where no sponsor exists at all, such as a solo expert or an internal tool with no PER-003, the expert approves their own set under this rule, shown the same evidence, and the change log records that no separate approver existed. A set is never treated as approved because its score is high. A rejection goes back to the expert with its reason, and the set reads as not approved until a later version is approved.
ACCEPTANCE	• Given a set with any blocking defect, when approval is sought, then the approver is shown the defects and the set is not presented as fit to approve. • Given a set with no recorded approval, when anyone reads its state, then it reads as not approved, whatever its scores are. • Given an engagement with no sponsor, when the expert approves their own set, then the approval records that no separate approver existed. • Given an approver who rejects the set, when the decision is recorded, then the rejection carries its reason, and the set reads as not approved. • Given an approver who approves with conditions, when the approval is recorded, then each condition names the open question, requirement, or review comment it depends on and a date, and the set reads as approved-with-conditions until every condition is cleared by that same approver.

TR-002 · Set a section's state to final

STAGES	• ST-004
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The section's grade, its blocking defect count, and the result of validating it.
ESCALATION	The expert may override the gate; the override and the score at the time are recorded in the change log. Nobody may set a section final on the expert's behalf while they are away: it waits.
ACCEPTANCE	• Given a section with one or more blocking defects, when final is requested, then it is not set without an explicit override from the expert. • Given an override, when it is applied, then the change log records that it was an override and the score at the time.

TR-003 · Produce a section's score and findings against its rubric

STAGES	• ST-003
APPROVER	PER-001
DELEGABLE TO AN AGENT	yes
NEVER DELEGATED	no
AGENT ROLE	proposes
EVIDENCE REQUIRED	The rubric's dimensions and a quotation from the document for every score below full marks.
ESCALATION	The score runs by itself after each draft, and nobody has to start it. The score does not count until the expert accepts it under TR-014, and a finding the document does not support is dropped (FR-016).
ACCEPTANCE	• Given a section that has just been drafted, when the draft is written, then its score and findings are produced without anyone asking for them. • Given a score below full marks on a dimension, when it is reported, then a quotation from the document is given for it.

TR-004 · Decide what the product is for and what is out of scope

STAGES	• ST-001
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	What the consequence of the choice is, including any release-line stage it would leave uncovered.
ESCALATION	None, and no substitute. The expert decides; while they are away, scope does not change. Where the choice weakens the PRD the objection is stated once and the decision is recorded with the dissent noted.
ACCEPTANCE	• Given an expert decision that leaves a release-line stage uncovered, when it is recorded, then the consequence is stated to the expert once and the decision stands. • Given a decision the expert reaffirmed after an objection, when later grading passes run, then the objection is not raised again as a new finding.

TR-005 · State a domain fact, figure, regulation or citation as true

APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The expert's own statement of it, or a source that has been checked.
ESCALATION	Where neither exists the fact is not written at all; it becomes an open question with a decider and a date. Nobody may supply the fact on the expert's behalf while they are away.
ACCEPTANCE	• Given a figure or citation the expert did not supply and that has not been verified, when a section is drafted, then it does not appear in the document. • Given a citation that was checked and could not be confirmed, when it is reported, then it is marked unverified rather than stated as fact.

TR-006 · Judge whether a set can be built from, and report the gaps

STAGES	• ST-009
APPROVER	PER-002
DELEGABLE TO AN AGENT	yes
NEVER DELEGATED	no
AGENT ROLE	proposes
EVIDENCE REQUIRED	The release line, the requirements covering each stage, and a quotation for every gap claimed.
ESCALATION	The engineering lead's own judgement overrides any reported verdict; the report is evidence for them, not a substitute for them. In release 1 no engineering lead is in the loop, so the report goes to the expert, who decides what to do with it under TR-010.
ACCEPTANCE	• Given a reported verdict, when the engineering lead disagrees, then their judgement is recorded as the verdict and the report is kept alongside it. • Given a reported gap, when it is stated, then the document is quoted for it.

TR-007 · Decide the domain's own terms, and what each one means

STAGES	• ST-001
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	Where the term is already used in the set, and any standard-setter's definition that differs from the expert's.
ESCALATION	None, and no substitute. Terms are the expert's; while they are away the vocabulary does not change.
ACCEPTANCE	• Given a term the expert has defined, when any later question or document uses that idea, then it uses the expert's word and not a synonym. • Given a standard-setter's definition that differs, when the term is recorded, then both are stated and the expert says which the set uses.

TR-008 · Decide the interview on a section is finished

STAGES	• ST-002
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	What the section would score as it stands, and what is still missing from it.
ESCALATION	None. The expert ends the interview; nobody ends it for them, and no question list runs past the point they call it done.
ACCEPTANCE	• Given an expert who says they have nothing more to add, when they say it, then no further question is put to them in that section. • Given an interview ended early, when the section is written, then what is missing is named to the expert, not hidden behind a score.

TR-009 · Decide a section may be left at draft and moved on from

STAGES	• ST-004
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The section's score, its blocking defects, and which later sections depend on what is missing.
ESCALATION	None, and no substitute. While the expert is away the section stays where it is.
ACCEPTANCE	• Given a section below 70 or carrying a blocking defect, when the expert chooses to move on, then it is written at `draft` and what was left behind is recorded as an open question or named plainly to them.

TR-010 · Decide the set may be offered for approval

STAGES	• ST-005
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	Every section's state and grade, every blocking defect, and the buildability review's gaps with, for each, whether it is closed or owned.
ESCALATION	None, and no substitute. The expert offers the set; nobody offers it for them.
ACCEPTANCE	• Given a buildability gap that is neither closed nor given a decider and a date, when the expert offers the set for approval, then that gap is stated to them first. • Given a set offered for approval, when the approver opens it, then what the expert was shown under this rule is what the approver sees.

TR-011 · Decide a backlog can be mapped from the set

STAGES	• ST-007
APPROVER	PER-004
DELEGABLE TO AN AGENT	yes
NEVER DELEGATED	no
AGENT ROLE	proposes
EVIDENCE REQUIRED	Whether a golden thread exists, and which release-line stages have requirements behind them.
ESCALATION	The delivery lead's own judgement stands over any reported verdict, and an override is recorded with its reason. Where they are unavailable the workshop waits rather than proceeding without a backbone.
ACCEPTANCE	• Given a set with no journey, when the check runs, then it reports that a backlog cannot be derived, and names that as the reason. • Given a reported verdict the delivery lead disagrees with, when they record their own, then their verdict stands and the override and its reason are recorded beside the report.

TR-012 · Decide which epics come first

STAGES	• ST-008
APPROVER	PER-004
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The release line, the dependencies in section 11, and the objective each epic moves.
ESCALATION	The delivery lead may re-order within the release line at any time; order never changes scope, which stays with TR-004. While they are away the order stands as last recorded, and nobody re-orders on their behalf.
ACCEPTANCE	• Given an ordering, when it is recorded, then each decision names whether value, dependency or the release line drove it.

TR-013 · Accept a drafted section as saying what the expert meant, or correct it

STAGES	• ST-003
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The drafted section in readable form, and its grade, so the expert judges the words and not the score.
ESCALATION	None, and no substitute. An unaccepted draft stays at `draft`; nobody accepts a section's wording on the expert's behalf while they are away.
ACCEPTANCE	• Given a drafted section the expert has not accepted, when the set's state is read, then that section is not shown as ready for anyone else to read. • Given an expert who corrects a drafted section, when the correction is applied, then it is their words that are kept, and the section is graded again before it advances.

TR-014 · Accept or refute a section's score

STAGES	• ST-003
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The score, each finding with the quotation behind it, and the section's weakest part. The product asks as soon as the score is produced, so the expert's only action is to accept or to refute.
ESCALATION	None, and no substitute. While the expert is away the score waits, unaccepted, and the section does not advance. Nobody accepts a score on the expert's behalf.
ACCEPTANCE	• Given a score that has just been produced, when the expert sees it, then they are asked whether they accept it or refute it, and need take no other action. • Given an expert who accepts a score, when they accept it, then the acceptance is recorded and the section may advance on that score. • Given an expert who refutes a finding, when they give their reason, then the refutation and the reason are recorded, and the finding is checked against the document under FR-016.

TR-015 · Turn off a set's outbound domain research

STAGES	• ST-003
APPROVER	PER-001
DELEGABLE TO AN AGENT	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	What the domain check sends out by default: the names of standards, terms of art, and the domain, and never the set's text or the expert's answers.
ESCALATION	While the expert has not decided, research stays on, carrying only the names of standards, terms of art, and the domain. Nobody turns it off or on for one set on the expert's behalf. An organization's own policy that forbids outbound research is a rule from outside the product, not a decision in it: the domain check obeys that rule for every set, and no decision of the expert overrides it.
ACCEPTANCE	• Given a set whose expert has turned off its outbound domain research, when a section is graded, then the domain check makes no outbound request, and marks each citation it cannot confirm from the set as unverified. • Given a set whose expert has made no choice, when a section is graded, then the domain check may research the domain's standards and terms, and no query it sends holds the set's text or the expert's answers. • Given an organization whose policy forbids outbound research, when any set is graded, then the domain check makes no outbound request, whatever the expert has decided for the set.

SECTION 11

Assumptions, risks and dependencies

DRAFT 96/100 NO BLOCKING DEFECTS

Assumptions

ID	ASSUMPTION	IF FALSE
A-001	An expert will answer 50-100 questions across a few sittings	The method fails; sections 5-7 change completely
A-002	An expert will act on a grade they did not produce	The grading is decoration
A-003	An engineering lead accepts a PRD naming no technology	Test 5 becomes a liability
A-004	Organizations permit drafting inside an assistant session	Section 6's separate-application non-goal and A-005 have to change, and FR-002, FR-015, and FR-018 with them
A-005	Organizations accept that domain content stays in the session and the documents	NFR-008, NFR-010, and NFR-012 have to change

A-001 and A-002 are the two the product stands on, and neither has been tested. Both are validated by the same ten observed sittings (D-002), which makes D-002 the single most load-bearing dependency in this PRD.

A-005 was constraint C-003 in section 9 until 0.8.1. It moved here because it is a belief about what customers accept, which a customer can falsify, and not a rule imposed from outside.

Risks

ID	RISK	L	I	RESPONSE	MITIGATED BY	STATUS	OWNER
R-001	High score becomes false assurance	high	high	reduce	2027-01-31	open	Richard Hightower
R-002	Unsupported findings destroy trust in all findings	high	high	reduce	2027-01-31	open	Richard Hightower
R-003	A beautifully graded PRD for the wrong product	medium	high	reduce	none	open	Richard Hightower
R-004	Interview too long, and sets abandoned	medium	medium	reduce	none	open	Richard Hightower
R-005	A version change strands an expert's set	low	high	avoid	none	open	Richard Hightower

R-001 and R-002 are high and high, and both mitigations must be in place by 2027-01-31, two months before REL-001.

R-003 is not closed. The mitigations listed reduce it and do not solve it. Nothing in this product checks whether the product being specified should exist. This section states it rather than hide it behind a mitigation that sounds adequate.

Dependencies

ID	NEED	FROM	BY	STATUS	OWNER
D-001	An assistant session able to carry the interview	The assistant platform the customer already uses	2027-01-31	delivered	Richard Hightower
D-002	10 observed first-time experts	Early customers and the professional network	2027-01-31	committed	Richard Hightower
D-003	3 engineering leads to read a set cold	Early customers	2027-01-31	committed	Richard Hightower

D-001 is delivered: this PRD was written and graded with prd-forge inside an assistant session, which carried the interview and read and wrote the documents.

D-002 and D-003 are both committed and both due by 2027-01-31, two months before REL-001's target in section 13. On 2026-11-02, the parties behind both committed in writing to report by that date, which answers Q-019. Richard Hightower, the owner of both, holds the written commitments. The set does not name the parties. The questions they answer are due by 2027-02-28, so REL-001 can use the results. Q-022 decides what happens to REL-001 if either party reports late.

Assumptions

A-001

ASSUMPTION	A domain expert will answer fifty to a hundred questions about their own field across a few sittings, if each one is clearly worth answering.
IF FALSE	The whole method fails. The product would have to produce a usable partial PRD from perhaps fifteen answers, which changes sections 5, 6 and 7 completely and makes the grading thresholds meaningless.
VALIDATED BY	Observed sittings with 10 first-time experts, counting where they stop, the same sittings NFR-005 uses (D-002).
OWNER	Richard Hightower
CATEGORY	usability
EVIDENCE	none
VALIDATE BY	2027-02-28

A-002 · An expert will act on a grade they did not produce, rather than dismissing it.

IF FALSE	The grading is decoration, and O-001 to O-003 and FR-009 to FR-011 measure something nobody uses. The product reduces to a structured interview.
VALIDATED BY	In the same observed sittings, whether the expert answers the one question or moves on regardless (D-002).
OWNER	Richard Hightower
CATEGORY	value
EVIDENCE	some
VALIDATE BY	2027-02-28

A-003

ASSUMPTION	An engineering lead will accept a PRD that names no technology, rather than reading the absence as incompleteness.
IF FALSE	Test 5 becomes a liability rather than a discipline, and the candidate architecture decisions list in FR-007 has to become a first-class output rather than a hand-over note.
VALIDATED BY	Three engineering leads reading a completed set cold and saying what they would ask for first (D-003).
OWNER	Richard Hightower
CATEGORY	usability
EVIDENCE	none
VALIDATE BY	2027-02-28

A-004 · Organizations will let a PRD for their product be drafted inside an assistant session.

IF FALSE	The separate-application non-goal in section 6 and A-005 are unacceptable as written and the delivery model has to change, which invalidates FR-002, FR-015, and FR-018, because each assumes the expert's own session.
VALIDATED BY	The information security review of the first three client contracts.
OWNER	Richard Hightower
CATEGORY	viability
EVIDENCE	none

A-005

ASSUMPTION	Organizations accept that the domain content an expert enters stays in the expert's own session and in the documents. This product sends it nowhere else and keeps it nowhere else. They also accept the domain check's research, on by default, which sends only the names of standards, terms of art, and the domain, or they turn it off under TR-015.
IF FALSE	An organization whose information policy needs more than this cannot adopt the product as written, and NFR-008, NFR-010, and NFR-012, which are measured against this boundary, have to change.
VALIDATED BY	The information security review of the first three client contracts, the same review that tests A-004.
OWNER	Richard Hightower
CATEGORY	viability
EVIDENCE	none
VALIDATE BY	2027-03-15

Risks

R-001

RISK	The expert accepts a weak PRD because the score was high. The grade becomes false assurance, which is worse than no grade.
LIKELIHOOD	high
IMPACT	high
MITIGATION	Blocking defects are counted and reported separately from the score and no score clears them (FR-009); the non-goal in section 6 says explicitly that a grade is evidence, not a warrant; FR-012 puts a buildability verdict in front of the expert before approval.
OWNER	Richard Hightower
TRIGGER	An expert approves a set with a high score while its buildability review still reports a gap that is neither closed nor owned.
RESPONSE	reduce
MITIGATION BY	2027-01-31
STATUS	open

R-002

RISK	The grading reports findings that are not in the document. The expert catches one, stops trusting all of them, and the grading becomes noise they skip.
LIKELIHOOD	high
IMPACT	high
MITIGATION	Every finding must quote the document, and anything unsupported is dropped before the expert sees it (FR-016, TR-003, TR-006); a refuted finding the document does not support is acknowledged as dropped rather than defended (SC-004).
OWNER	Richard Hightower
TRIGGER	An expert refutes a finding and the document does not support it.
RESPONSE	reduce
MITIGATION BY	2027-01-31
STATUS	open

R-003

RISK	A fluent, well-structured, well-graded PRD is produced for the wrong product, and the structure makes the error harder to see rather than easier.
LIKELIHOOD	medium
IMPACT	high
MITIGATION	Section 3 asks what would move the objectives without this product at all; section 1 requires evidence per problem and names what is unmeasured. Neither is sufficient, and this risk is not closed. A reader should say so.
OWNER	Richard Hightower
TRIGGER	A sponsor or an engineering lead says that an approved set specifies the wrong product.
RESPONSE	reduce
STATUS	open

R-004

RISK	The interview is too long to finish, so sets are abandoned half-written and the product is remembered as the thing that asked a hundred questions.
LIKELIHOOD	medium
IMPACT	medium
MITIGATION	FR-005 lets the expert end any section and records what was left; FR-015 makes resumption cheap; the working order in the method, which Q-007 tests, puts the sections that unlock the others first, and an abandoned set is still useful.
OWNER	Richard Hightower
TRIGGER	An expert stops a set with sections still pending and does not resume it (Q-006 and Q-020 count this).
RESPONSE	reduce
STATUS	open

R-005

RISK	A set written by one version cannot be read or graded by a later one, and an expert loses weeks of thinking to an upgrade.
LIKELIHOOD	low
IMPACT	high
MITIGATION	NFR-004 makes every set readable with no part of the product running; NFR-009 requires older sets to stay gradable within a major version.
OWNER	Richard Hightower
TRIGGER	A set written by one version fails to read or grade in a later version with the same major number.
RESPONSE	avoid
STATUS	open

Dependencies

D-001

NEED	An assistant session the expert is already working in, able to carry a long interview and read and write the documents.
FROM	The assistant platform the customer already uses
BY	2027-01-31
STATUS	delivered
OWNER	Richard Hightower
BLOCKS	• ST-001 • ST-002

D-002 · Ten first-time domain experts willing to be observed completing a PRD.

FROM	Early customers and the professional network
BY	2027-01-31
STATUS	committed
OWNER	Richard Hightower
BLOCKS	• A-001 • A-002 • NFR-005 • REL-001

D-003

NEED	Three engineering leads willing to read a completed set cold and report what they would ask for first.
FROM	Early customers
BY	2027-01-31
STATUS	committed
OWNER	Richard Hightower
BLOCKS	• A-003 • O-001 • REL-001

SECTION 12

Open questions

DRAFT 88/100 NO BLOCKING DEFECTS

ID	QUESTION	DECIDER	BY	BLOCKS
Q-001	Today's first-pass buildability rate	R. Hightower	2027-02-28	O-001
Q-002	Today's days from the first question to an approval	R. Hightower	2027-02-28	O-002
Q-003	Today's share of testable `must` requirements	R. Hightower	2027-02-28	O-003
Q-004	How many organizations are in each segment	R. Hightower	2027-06-30	S-001, S-002
Q-005	Does FR-013 stay `must`, given what sponsors say about grades	R. Hightower	2027-01-31	FR-013, TR-001, OUT-003, PER-003
Q-006	Sections completed per sitting	R. Hightower	2027-02-28	R-004
Q-007	Does writing the requirements send an expert back to the journeys	R. Hightower	2027-02-28	ST-004, FR-010, R-004
Q-008	Disciplined or incomplete, to an engineering lead	R. Hightower	2027-02-28	A-003
Q-009	Which version of the rubric and the checks release 1 grades against	R. Hightower	2027-01-31	FR-004, FR-009, FR-010, FR-011, FR-012, FR-019
Q-010	Must an answer survive the session before its section is drafted	R. Hightower	2027-01-31	FR-015, NFR-011
Q-012	Which segment converts first	R. Hightower	2027-06-30	S-001, S-002
Q-013	The day rate a buyer pays for a contracted product manager	R. Hightower	2027-06-30	S-002
Q-014	Does ST-004 keep its loop-back, given what experts say	R. Hightower	2027-02-28	ST-004, FR-010
Q-015	Who owns the rubric and the checks	R. Hightower	2027-01-15	FR-004, FR-009 to FR-012, FR-019
Q-016	Must a grade show the rubric version it was made against	R. Hightower	2027-01-31	NFR-007
Q-017	Who owns NFR-006's accessibility	R. Hightower	2027-02-28	NFR-006

ID	QUESTION	DECIDER	BY	BLOCKS
Q-018	One counted cost per problem in section 1	R. Hightower	2027-01-31	P-001 to P-004
Q-020	Which section experts stop at most	R. Hightower	2027-02-28	R-004
Q-021	Who owns NFR-012's access control	R. Hightower	2027-02-28	NFR-012
Q-022	Does REL-001 move or ship as an experiment if D-002 or D-003 is late	R. Hightower	2027-01-31	REL-001
Q-023	Does the delivery lead persona hold	R. Hightower	2027-05-31	PER-004

The ones that matter most

Q-001 to Q-003 are the baselines for all three objectives. They are due by 2027-02-28, a month before REL-001, so its exit criteria are fixed before it ships. Until they are answered, section 3's targets are provisional. Nothing should be committed to a plan on the strength of them.

Q-015, Q-009, and Q-010 block the most. Q-015, who owns the rubric and the checks, is due first, by 2027-01-15, because Q-009 cannot fix a version without an owner. It blocks the same six requirements as Q-009. Q-009 blocks four `must` requirements, FR-004, FR-009, FR-010, and FR-012, and two `should` requirements, FR-011 and FR-019. Those six requirements sit on four release-line stages, ST-002 to ST-005, which can be scaffolded and not finished until Q-009 is answered. Q-010 blocks FR-015, a `should`, and NFR-011. It is the one where guessing breaks assumption A-005 rather than costing rework.

Q-005 is the one that could invalidate the product's central idea. If showing a sponsor a grade reduces their confidence rather than raising it, then FR-013 is wrong and the grading becomes an author's private tool rather than part of the approval. It takes two days, and it should be asked early.

Q-022 is the schedule. D-002 and D-003 report by 2027-01-31, and the questions they answer are due by 2027-02-28, a month before REL-001. Q-019 is answered: on 2026-11-02 both parties committed in writing, so REL-001 keeps 2027-03-31. Q-022 is now the only open question on REL-001's date. It decides what happens if either dependency reports late.

Q-008 tests the fifth of the five tests in the glossary: a requirement names no technology. If engineering leads read the absence of technology as incompleteness, the discipline this whole product enforces is a liability, and A-003 fails.

Candidate architecture decisions, not questions

Three things the set needs are design choices, not product questions. They go to the engineering lead as candidate architecture decisions, as FR-007 hands over a technology the expert names:

- Where the rubric and the eleven derivability checks live. Q-009 decides only which version release 1 grades against.
- How a grade records the rubric version it was made against. Q-016 decides only whether it must.
- Whether an interview in progress is a document of the set, and where an answer sits before its section is drafted. Q-010 decides only whether the answer must survive the end of the session.

Honest note

Every decider here is the same person. For a one-person product that is accurate. For anything larger it is a defect, and the first thing to fix when there is a second person.

Q-001

QUESTION	What share of specifications reaching a development team today pass a buildability review on first submission?
DECIDER	Richard Hightower
BY	2027-02-28
BLOCKS	• O-001
SPIKE	Run the eleven derivability checks against ten specifications collected from real client projects done without this product, and count. Time-boxed to 3 days. It depends on Q-009, because the version of the checks must be fixed first.

Q-002

QUESTION	How many working days does it take today from the first question put to the domain expert about a product to an approved specification?
DECIDER	Richard Hightower
BY	2027-02-28
BLOCKS	• O-002
SPIKE	Reconstruct the timeline from the first question put to the expert to the approval, for five recent projects done without this product. Time-boxed to 2 days.

Q-003

QUESTION	What share of `must` functional requirements in today's specifications can an engineer who was not in the interview write a test from?
DECIDER	Richard Hightower
BY	2027-02-28
BLOCKS	• O-003
SPIKE	Three engineering leads score the requirements of the same ten specifications from Q-001, independently. Time-boxed to 3 days.

Q-004 · How many organizations are in each segment in section 2?

DECIDER	Richard Hightower
BY	2027-06-30
BLOCKS	• S-001 • S-002
SPIKE	Name the register or source per segment, and count. Time-boxed to 3 days.

Q-005

QUESTION — Given what three sponsors say about a set shown with and without grades, does FR-013 stay `must`, with each section's grade shown to the sponsor?

DECIDER — Richard Hightower

CONSULTED — the first three sponsors

BY — 2027-01-31

BLOCKS
- FR-013
- TR-001
- OUT-003
- PER-003

SPIKE — Show three sponsors the same set with and without grades and ask which they would approve. FR-013 stays `must` if at least 2 of the 3 approve the graded set at least as readily as the ungraded one. Otherwise FR-013 becomes `should`. Time-boxed to 2 days.

Q-006 · How many sections does an expert complete in one sitting before stopping?

DECIDER — Richard Hightower

BY — 2027-02-28

BLOCKS
- R-004

SPIKE — Count in D-002's ten observed sittings, held by 2027-01-31. Time-boxed to 1 day after the sittings.

Q-007 · Does writing the requirements send an expert back to the journeys?

DECIDER — Richard Hightower

BY — 2027-02-28

BLOCKS
- ST-004
- FR-010
- R-004

SPIKE — Count loop-backs in D-002's ten observed sittings, held by 2027-01-31. The answer is yes if at least 3 of the 10 sittings loop back. Time-boxed to 1 day after the sittings.

Q-008

QUESTION — Does an engineering lead read a PRD that names no technology as disciplined or as incomplete?

DECIDER — Richard Hightower

CONSULTED — the three engineering leads from D-003

BY — 2027-02-28

BLOCKS
- A-003

SPIKE — Ask each lead, after reading a set cold, whether anything was missing and whether they wanted a technology named. A-003 holds if no more than 1 of the 3 leads says a named technology was missing. Time-boxed to 1 day.

Q-009

QUESTION Which version of the rubric and the eleven derivability checks does release 1 grade against?

DECIDER Richard Hightower

BY 2027-01-31

BLOCKS

- FR-004
- FR-009
- FR-010
- FR-011
- FR-012
- FR-019

SPIKE List the rubric dimensions of each section and the eleven checks as they stand, and have the decider approve that list as release 1's version. Time-boxed to 1 day. Until it is answered, ST-002 to ST-005 can be scaffolded and not finished.

Q-010

QUESTION Must an answer the expert gave survive the end of the session before its section is drafted, given A-005, that domain content stays in the expert's own session and in the documents?

DECIDER Richard Hightower

BY 2027-01-31

BLOCKS

- FR-015
- NFR-011

SPIKE List what an expert must find when they come back after a break (FR-015), and check each item against A-005's boundary. The answer is yes if any item the expert must find is an answer not yet drafted. Time-boxed to 2 days. Guessing this one breaks assumption A-005, which customers rely on, rather than merely costing rework.

Q-011

QUESTION How does an approver record a decision, given that D-001 puts the product inside the expert's session and OUT-003 excludes any view of our own?

DECIDER Richard Hightower

BY 2027-02-28

BLOCKS

- FR-014

ANSWER Outside the product, in release 1. The sponsor reads the root document or an exported copy and gives a decision. The expert records it in section 15 under FR-014, with the sponsor's name, role, and date. Section 6 records ST-006 as outside the product.

ANSWERED 2026-10-27

Q-012 · Which segment in section 2 converts first?

DECIDER Richard Hightower

BY 2027-06-30

BLOCKS

- S-001
- S-002

SPIKE For the first ten buyers spoken to in each segment, count how many buy by 2027-06-30, and compare the two counts. Time-boxed to 1 day at that date.

Q-013 · What day rate does a buyer pay to contract a product manager for the writing of one set?

DECIDER Richard Hightower

BY 2027-06-30

BLOCKS
- S-002

SPIKE Ask the buyers in the first ten conversations what they paid the last time. Time-boxed to 1 day.

Q-014

QUESTION Given how experts describe a loop-back in D-002's sittings, does ST-004 keep its step that sends the expert back to an earlier section?

DECIDER Richard Hightower

BY 2027-02-28

BLOCKS
- ST-004
- FR-010

SPIKE Ask each expert after a loop-back, in D-002's ten observed sittings, held by 2027-01-31. ST-004 keeps the step if at least half of the experts who loop back say it helped. Time-boxed to 1 day after the sittings.

Q-015 · Who owns the rubric and the eleven derivability checks?

DECIDER Richard Hightower

BY 2027-01-15

BLOCKS
- FR-004
- FR-009
- FR-010
- FR-011
- FR-012
- FR-019

SPIKE Name the person accountable for the content of the rubric and the checks, and confirm with that person that they accept it, before Q-009 fixes a version. Time-boxed to 1 day.

Q-016 · Must a grade show which version of the rubric it was made against?

DECIDER Richard Hightower

BY 2027-01-31

BLOCKS
- NFR-007

SPIKE Regrade one section of this set against an earlier rubric, and ask three readers whether they need to tell the two grades apart. The answer is yes if at least 2 of the 3 say they need to. Time-boxed to 1 day.

Q-017

QUESTION Whom does the decider record as owner of NFR-006's accessibility, given that the product's output is read in places the non-goals in section 6 say we do not own?

DECIDER Richard Hightower

BY 2027-02-28

BLOCKS
- NFR-006

SPIKE List each place the product's output is read and who controls it, and get that party's written confirmation that they own it. The decider then records the owner. Time-boxed to 2 days.

Q-018 · What does each problem in section 1 cost, as one counted case per problem?

DECIDER Richard Hightower

BY 2027-01-31

BLOCKS
- P-001
- P-002
- P-003
- P-004

SPIKE For one recent project done without this product, count the engineer days spent re-deriving the specification, the gaps found during build, and the plans changed after approval. Time-boxed to 2 days.

Q-019

QUESTION Will both parties behind D-002 and D-003 commit in writing to report by 2027-01-31, so that REL-001 keeps 2027-03-31?

DECIDER Richard Hightower

BY 2027-01-31

BLOCKS
- REL-001

SPIKE Ask each party behind D-002 and D-003 for a written date. The answer is yes if both written dates are 2027-01-31 or earlier, and no otherwise, in which case Q-022 decides. Time-boxed to 1 day.

ANSWER Yes. The parties behind D-002 and D-003 both committed in writing to report by 2027-01-31, so REL-001 keeps 2027-03-31. Q-022 still decides what happens if either reports late.

ANSWERED 2026-11-02

Q-020 · Which section do experts stop at most often?

DECIDER Richard Hightower

BY 2027-02-28

BLOCKS
- R-004

SPIKE Note the section at each stop in D-002's ten observed sittings, held by 2027-01-31. Time-boxed to 1 day after the sittings.

Q-021

QUESTION Whom does the decider record as owner of NFR-012's access control, given that a set is stored where the organization already stores documents?

DECIDER Richard Hightower

BY 2027-02-28

BLOCKS
- NFR-012

SPIKE List every path by which someone can read a set and who grants access to it, and get that party's written confirmation that they own it. The decider then records the owner. Time-boxed to 2 days.

Q-022 · If D-002 or D-003 reports after 2027-01-31, does REL-001 move or ship as an experiment?

DECIDER Richard Hightower

BY 2027-01-31

BLOCKS

- REL-001

SPIKE For each outcome, list which of REL-001's exit criteria can still be judged with the late results, and choose the outcome that leaves more of them judgeable. Time-boxed to 1 day.

Q-023

QUESTION Do delivery leads run a story-mapping workshop about once per release and want a backbone agreed before it, as PER-004 assumes?

DECIDER Richard Hightower

BY 2027-05-31

BLOCKS

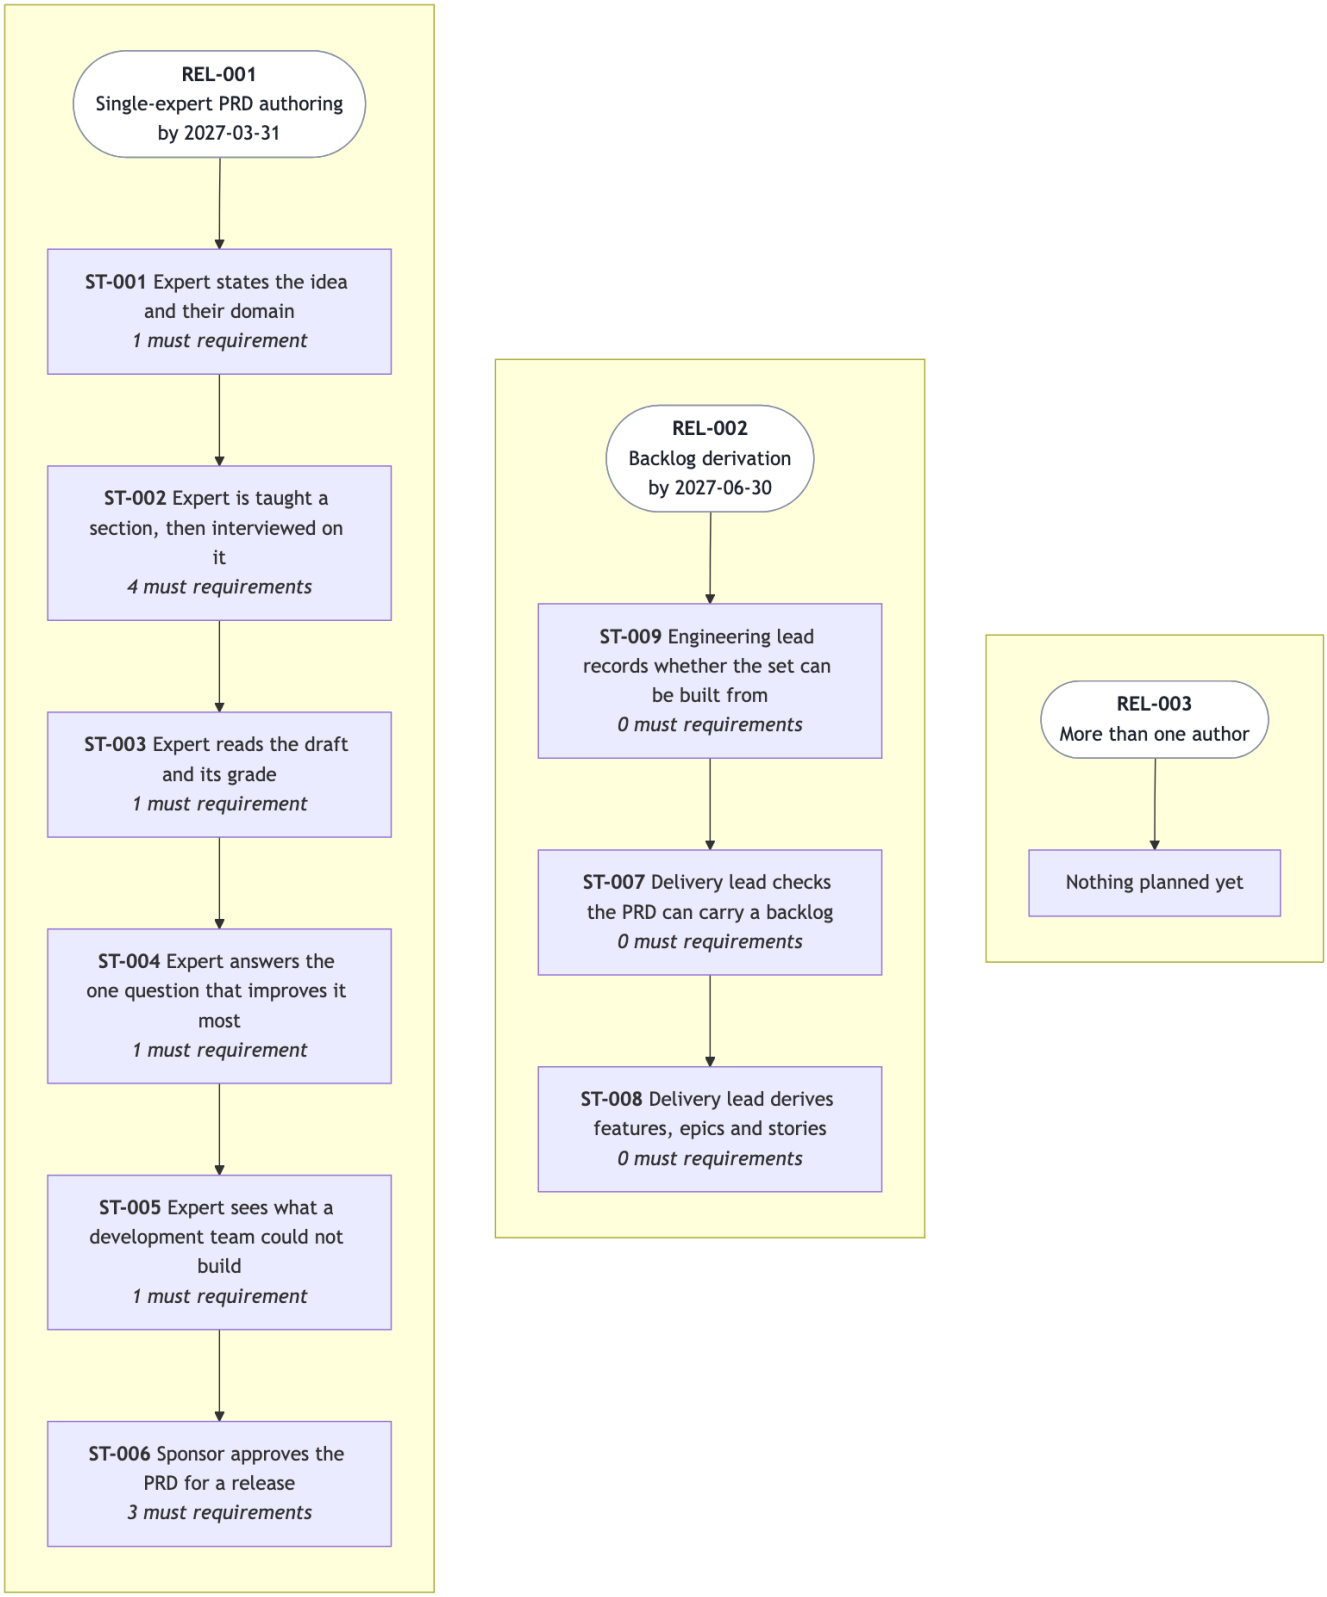
- PER-004

SPIKE Ask three delivery leads how often they run the workshop and what they need agreed before it. PER-004 holds if at least 2 of the 3 run the workshop about once per release and want a backbone agreed before it. Time-boxed to 2 days.

SECTION 13

Release plan

DRAFT 83/100 NO BLOCKING DEFECTS



The release line

ID	RELEASE	STAGES	TARGET
REL-001	Single-expert PRD authoring	ST-001..ST-006	2027-03-31
REL-002	Backlog derivation	ST-009, ST-007, ST-008	2027-06-30
REL-003	More than one author	none	none

Why this order

REL-001 is exactly the release line in section 6. It is the smallest release in which anyone gets anything: an expert reaching an approved PRD with its buildability review.

REL-002 is deferred rather than cut because the PRD REL-001 produces already has to carry everything a backlog needs, and FR-012 checks precisely that. Deferring the derivation costs nothing; deferring what it needs would have cost a rewrite. REL-002 does not wait for REL-001's exit thresholds, which the pilots can take until 2027-09-30 to show. It starts when the first pilot PRD written with REL-001 is approved, because the derivation needs one real approved PRD as its input. Richard Hightower decided this on 2026-11-02.

REL-002 links to O-001 and REL-003 to O-002, but neither release can move those numbers as section 3 measures them. O-001 is judged before a backlog exists, and O-002 counts single-journey pilot PRDs written with REL-001. The exit criteria of REL-002 and REL-003 carry their own targets for now. Each becomes an objective in section 3 when its release has requirements that can reference it.

REL-003 waits for evidence that one author's path works at all. It starts when OUT-002's condition holds: three real engagements have each reached an approved set that passed the buildability review on first submission. Building for two authors before one is proven would be building on A-001 and A-002 before either is tested. A sponsor can already comment on an exported copy in REL-001, under FR-022. REL-003 adds comments inside the set, by a reviewer who does not hold the pen, which OUT-005 defers.

How you know each release worked

REL-001 exits when the pilot PRDs meet the three objective targets in section 3: at least 4 of the 5 pilot PRDs, 80%, pass the buildability review on first submission, the median reaches approval in under 5 working days, and 95% of `must` requirements are testable as written. REL-002 exits when every derived story names its requirement, a delivery lead runs a workshop from the backlog without re-deriving it, the delivery lead accepts the derived backlog within 2 working days for at least 2 of the first 3 approved pilot PRDs, and an engineering lead records a verdict on each of those 3. REL-003 exits when the median of the first 3 sets written by two experts reaches approval within 10 working days.

Both target dates are indicative, not committed. D-002 and D-003 have not reported yet, and Q-022 is still open. The schedule problem below says why.

The schedule problem, stated plainly

REL-001 targets 2027-03-31. D-002 and D-003 in section 11, the ten observed experts and the three engineering leads, are due by 2027-01-31. Both are *committed*: on 2026-11-02 their parties committed in writing, which answers Q-019. The questions they answer are due by 2027-02-28, a month before REL-001, so REL-001 can use what they find. If either reports late, REL-001 moves or ships knowingly as an experiment. Q-022 decides which, and Richard Hightower decides it by 2027-01-31.

REL-002 carries no requirements yet. Section 7 holds none for ST-007, ST-008, or ST-009, and they are written when REL-002 is planned.

Not in this section, deliberately

No tasks, no estimates, no sizing, no team assignments, and no technology milestone. Those come from engineering design and iteration planning.

REL-001 · Single-expert PRD authoring

OUTCOMES

- An expert with no product training, a domain expert or a product manager who speaks for one, can take one product idea to an approved PRD. Only the sponsor's reading, comments, and decision at ST-006 happen outside the product.
- The expert can see, before anyone estimates, which parts of their own PRD a development team could not build from.
- A sponsor can see what they are approving, what is still open, and who decides each open question, in one document.

JOURNEY STAGES

- ST-001
- ST-002
- ST-003
- ST-004
- ST-005
- ST-006

REQUIREMENTS

- FR-001
- FR-002
- FR-003
- FR-004
- FR-005
- FR-006
- FR-007
- FR-008
- FR-009
- FR-010
- FR-011
- FR-012
- FR-013
- FR-014
- FR-015
- FR-016
- FR-017
- FR-018
- FR-019
- FR-020
- FR-021
- FR-022
- FR-023

OBJECTIVES

- O-001
- O-002
- O-003

DEPENDS ON

- D-002
- D-003
- Q-022
- R-001
- R-002

EXIT CRITERIA

- At least 4 of the 5 pilot PRDs, 80%, pass the buildability review on first submission, where a pass is the verdict `buildable`, or `buildable with gaps` with every gap closed or owned (O-001). The pilot PRDs are as O-001 defines them: the first 5 PRDs written with release 1 by people other than the product's author, between 2027-03-31 and 2027-09-30.
- The median pilot PRD reaches a recorded approval in under 5 working days from the first interview question (O-002).
- An engineer who was not in the interview judges 95% of the pilot PRDs' `must` requirements testable as written (O-003).

TARGET DATE

2027-03-31

DATE STATUS

indicative

REL-002 · Backlog derivation

OUTCOMES

- A delivery lead can walk into a story-mapping workshop with features and a release line already derived from the approved PRD.
- A delivery lead sees every story name the requirement id it came from, so nobody has to ask why it exists.
- An engineering lead can record their own verdict on whether an approved set can be built from, with the report of its buildability review kept beside it.

JOURNEY STAGES

- ST-009
- ST-007
- ST-008

OBJECTIVES

- O-001

DEPENDS ON

- Q-023

EXIT CRITERIA

- Every story derived from a pilot PRD names the requirement id it came from.
- A delivery lead runs a story-mapping workshop from a derived backlog without re-deriving any feature from the PRD.
- For at least 2 of the first 3 approved pilot PRDs, the delivery lead accepts the derived backlog within 2 working days of the approval.
- An engineering lead records a verdict on each of the first 3 approved pilot PRDs, with its buildability review kept beside it.

TARGET DATE

2027-06-30

DATE STATUS

indicative

REL-003 · More than one author

OUTCOMES

- Two experts can work on different sections of the same set without overwriting each other.
- A reviewer, such as the engineering lead (PER-002), can comment on a section inside the set, not on an exported copy, without being able to change it.

OBJECTIVES

- O-002

EXIT CRITERIA

- The median of the first 3 sets written by two experts reaches a recorded approval within 10 working days of the first interview question.

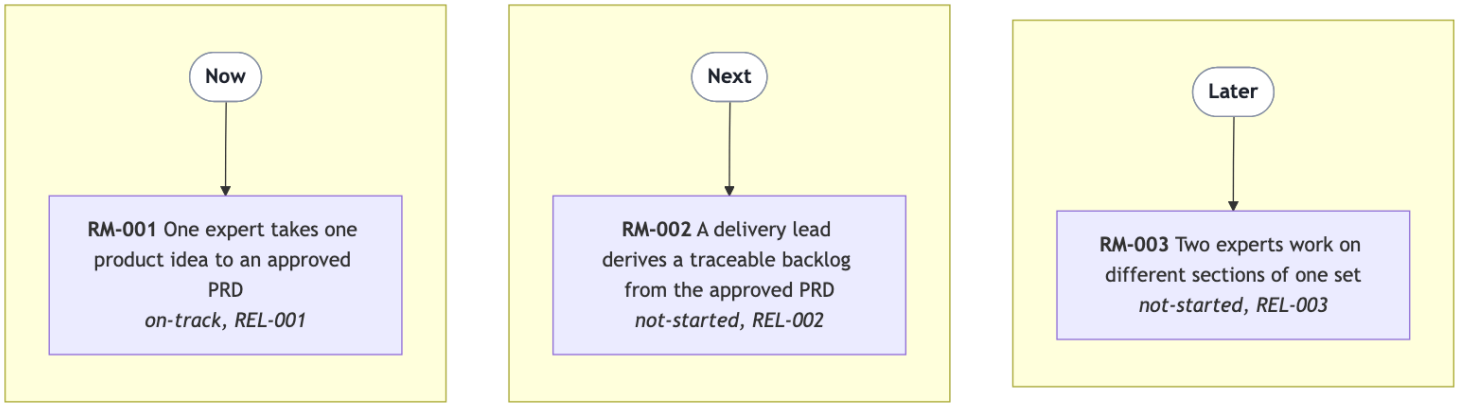
DEPENDS ON

- A-001
- A-002

SECTION 13

Roadmap

DRAFT 92/100 NO BLOCKING DEFECTS



The roadmap by horizon

The roadmap follows the release plan in section 13, one item for each release. RM-001 is now and on track: the parties behind D-002 and D-003 committed in writing to report by 2027-01-31, which answers Q-019. Q-022 is the only open decision left on REL-001's date. D-002 and D-003 are still to arrive by 2027-01-31, and the mitigations for R-001 and R-002 must be in place by the same date. RM-002 is next and waits on Q-023, which tests the delivery lead it serves. RM-003 is later, so it has no date.

The order follows "Why this order" in the release plan. RM-002 comes before RM-003 because REL-001's PRD already carries what a backlog needs, while two authors would build on A-001 and A-002 before either is tested. RM-002 starts when the first pilot PRD written with REL-001 is approved. REL-002 has no requirements yet. RM-003 starts only when OUT-002's condition holds.

RM-002 links to O-001 and RM-003 to O-002, but neither can move those numbers as section 3 measures them. The exit criteria of REL-002 and REL-003 hold their targets until each becomes an objective, as the release plan says.

Each item uses MoSCoW. No item has a RICE or ICE score, because the set holds no reach or effort figures yet.

Items

RM-001 · One expert takes one product idea to an approved PRD

HORIZON	now
STATUS	on-track
RELEASES	• REL-001
OBJECTIVES	• O-001 • O-002 • O-003
THEME	Single-expert authoring
OWNER	Richard Hightower
DEPENDS ON	• D-002 • D-003 • Q-022 • R-001 • R-002
TARGET DATE	2027-03-31
DATE STATUS	indicative
PRIORITY	

METHOD	moscow
MOSCOW	must

RM-002 · A delivery lead derives a traceable backlog from the approved PRD

HORIZON	next
STATUS	not-started
RELEASES	• REL-002
OBJECTIVES	• O-001
THEME	Backlog derivation
OWNER	Richard Hightower
DEPENDS ON	• Q-023
TARGET DATE	2027-06-30
DATE STATUS	indicative
PRIORITY	

METHOD	moscow
MOSCOW	should

RM-003 · Two experts work on different sections of one set

HORIZON	later
STATUS	not-started
RELEASES	• REL-003
OBJECTIVES	• O-002
THEME	More than one author
OWNER	Richard Hightower
DEPENDS ON	• A-001 • A-002
PRIORITY	

METHOD	moscow
MOSCOW	could

Changes

Roadmap created from the release plan at 0.13.0.

DATE	2026-11-01
REASON	Issue 24 adds a roadmap to the set. Each item follows one release, and no release changed.

RM-001 moved from at-risk to on-track.

DATE	2026-11-02
REASON	The parties behind D-002 and D-003 committed in writing to report by 2027-01-31, which answers Q-019 yes. Q-022 is now the only open decision on REL-001's date. D-002 and D-003 are still to arrive by 2027-01-31.

RM-001 also waits on R-001 and R-002, and RM-002 waits on Q-023.

DATE	2026-11-02
REASON	The mitigations for R-001 and R-002 must be in place by 2027-01-31, before REL-001. Q-023 tests PER-004, the delivery lead RM-002 serves. Richard Hightower decided both on 2026-11-02.

SECTION 14

Glossary

DRAFT 85/100 NO BLOCKING DEFECTS

TERM	MOST OFTEN CONFUSED WITH
PRD	A design document
Golden thread	The most important journey
Release line	A date
Stage	A screen
Rubric	A checklist
Blocking defect	A low score
Score	A judgement on the product idea
Spike	A piece of the product
Open question	A risk
Non-goal	Out of scope
Grade	Score
State	Threshold
Assumption	An open question
Candidate architecture decision	A requirement
Buildability review	A grading pass
Finding	A blocking defect
Root document	Section 1
Approver	The expert
Condition	A constraint
Constraint	A requirement
Objective	A requirement
Baseline	A target
Version	A release of the product
Dependency	A requirement's wait on an open question
Backlog	The set
Feature	A requirement

TERM	MOST OFTEN CONFUSED WITH
Story	A requirement
Expert	The approver
Sponsor	The expert
Engineering lead	The expert
Delivery lead	The engineering lead
Item	A line or a field of an entry
Pass	Meeting the gate
Scenario	A journey
Objection	A refuted finding the document does not support
Refute	Showing a finding wrong
Gate	The threshold
Bet	An objective
Buildable	A pass of the buildability review
Product manager	The sponsor
Out of scope	A non-goal
Outcome	A feeling
Ready	Offered for approval
Offered for approval	Ready
Sitting	A session
Experiment	A spike
Agent	A person
Exported copy	The set
Escalation	An open question
Satisfaction	A measure of the product
Pain point	The problem
Review comment	A finding
Capability	A capability the system needs
Delegated	A person deciding on someone's behalf
Verdict	A grade
Reviewer	A grading pass

TERM	MOST OFTEN CONFUSED WITH
Approval	An override
Engagement	A unit of work
Handoff	The handover
Journey	A use case
Step	A stage
Override	An approval
Release	The release line
Non-functional requirement	A requirement in section 7
Acceptance criteria	Exit criteria
Epic	A feature
Vocabulary	The glossary
Persona	A user
Buyer	The user
Risk	An assumption
Decider	The approver
Review	The buildability review
Exit criteria	Acceptance criteria
Unit of work	An engagement
Gap	A blocking defect
Wireframe	A design
Session	A sitting

The "Most often confused with" column above is the one that earns its keep. Three of them, golden thread, blocking defect and score, carry a cost when confused: a team picks the wrong journey for release 1, ships a section with a known fault because the number looked good, or believes a well-written PRD means a sound product.

References

Eight terms come from established practice, and each entry's `source` names where. Where this set departs from the source, the entry says so.

- MoSCoW: "What is MoSCoW Prioritization?", Agile Business Consortium, accessed 2026-10-27, https://www.agilebusiness.org/?p=1907
- Given/When/Then: "Introducing BDD", Dan North, 20 Sep 2006, https://dannorth.net/blog/introducing-bdd/

- Story, epic, and feature: "Epics, Features and User Stories", Mountain Goat Software, last updated 28 Apr 2025, https://www.mountaingoatsoftware.com/blog/stories-epics-and-themes
- Story map and backbone: "The New User Story Backlog is a Map", Jeff Patton & Associates, 8 Oct 2008, updated 12 Apr 2023, https://www.jpattonassociates.com/the-new-backlog/
- Spike: "Spike Solutions", in The Art of Agile Development, 1st ed., James Shore and Shane Warden, 3 Jun 2010, https://www.jamesshore.com/v2/books/aoad1/spike_solutions

PRD

MEANING A written statement of what to build, for whom, why, and how you know it works, plus what not to build. It does not say how to build it.

SEE ALSO
- section
- document set

NOT A design document. A PRD that names components or technologies has stopped being a PRD and has started being an architecture sketch.

Document set

MEANING The documents that together are one PRD: a root document and one document per section, each linking to the others.

ABBREVIATION set

SEE ALSO
- PRD

Golden thread

MEANING The one user journey that release 1 must deliver from end to end.

SEE ALSO
- journey
- release line

NOT The most important journey in general. It is the journey chosen for release 1, which is often the simplest rather than the most valuable.

Release line

MEANING The smallest set of journey stages that completes the golden thread, and so the boundary of what release 1 contains.

SEE ALSO
- golden thread
- scope

NOT A date or a deadline. It is a line across the journey, not across a calendar.

Stage

MEANING One part of a journey, naming a person, an input, an output and a decision, and holding its steps.

SEE ALSO
- journey
- step

NOT A screen or a system step. A stage survives a redesign.

The five tests

MEANING

The five properties a statement must have to be a requirement, numbered 1 to 5 in this order: a named user, an outcome a person can see, a pass-or-fail acceptance test, a priority, and no named component or technology.

SEE ALSO

- blocking defect

Rubric

MEANING

The stated standard a section is scored against: a set of dimensions, each with written anchors for what a score of 0, 2 and 4 look like.

SEE ALSO

- score
- blocking defect

NOT

A checklist. A checklist asks whether something is present; a rubric asks how good it is, against written anchors.

Blocking defect

MEANING

A named fault that stops a section advancing however high its score is, such as a requirement that names a technology or an id that refers to nothing.

SEE ALSO

- rubric
- score

NOT

A low score. Score and blocking defects are counted separately and a high score never clears a blocking defect.

Score

MEANING

A section's total against its rubric, 0 to 100, written by a grading pass and never by hand.

SEE ALSO

- rubric
- blocking defect

NOT

A measure of whether the product is a good idea. It measures how well the section is written, nothing else.

Spike

MEANING

A piece of work, usually time-boxed, whose only output is an answer to an open question, or a recorded decision. This set widens the source because some of its open questions ask what people will accept, as Q-005 does, and only work with those people answers them.

SEE ALSO

- open question

NOT

A piece of the product. A spike produces knowledge, not a feature. This set departs from the source, where a spike is a technical investigation. Here a spike may be research with people, as in Q-005, and some run inside a dependency's observed sittings, as Q-006 and Q-007 run in D-002's.

SOURCE

Spike Solutions, The Art of Agile Development, 1st ed., James Shore and Shane Warden, 3 Jun 2010, https://www.jamesshore.com/v2/books/aoad1/spike_solutions

Open question

MEANING Something not yet decided, recorded with the one person who decides it and the date by which they must.

SEE ALSO
- spike
- assumption

NOT A risk. A risk is something that might happen and nobody can simply decide; an open question has a decider.

Non-goal

MEANING Something a reasonable reader would expect the product to do, that it will not do, with no plan to.

SEE ALSO
- scope

NOT Out of scope. Out of scope means not yet, with a point at which it is revisited. A non-goal means not ever, with a reason.

Journey

MEANING The stages one user goes through to reach one goal with the product, from start to end, in order.

SEE ALSO
- stage
- golden thread

NOT A use case. A journey covers the whole path to a goal, including steps outside the product.

Step

MEANING What a user does, sees or decides at one point inside a stage.

SEE ALSO
- stage
- journey

NOT A stage. A stage holds many steps.

Section

MEANING One of the fifteen numbered parts of a PRD, each held in its own document. Section 5 holds two documents, journeys and scenarios, and section 7 holds two, functional requirements and wireframes, so a set has seventeen section documents across fifteen sections. Cross-references in a set are by section number.

SEE ALSO
- document set

Scope

MEANING What a release contains and what it does not, set against the journey stages it completes.

SEE ALSO
- release line
- non-goal

NOT The release line. The release line is the path of stages a release completes, and scope is everything the release includes and leaves out along it.

Assumption

MEANING Something taken on faith that the product depends on, recorded with what would break if it turned out false.

SEE ALSO
- open question
- risk

NOT An open question. Nobody can simply decide an assumption true; it is validated or falsified.

Risk

MEANING Something that may happen and would damage the product, recorded with a likelihood, an impact, a mitigation and one named owner.

SEE ALSO
- assumption
- spike

NOT An assumption. An assumption is believed true and must be tested. A risk is an event that may happen, and it is mitigated.

Readable form

MEANING The part of a section a person reads: its prose, its headings and its tables, written to be judged by the expert who gave the answers.

SEE ALSO
- section

Threshold

MEANING The score at or above which a section may advance: 70. A score at the threshold is not enough on its own, because the gate also needs no blocking defect.

SEE ALSO
- score
- blocking defect

Grade

MEANING A section's score together with its blocking-defect count and its weakest part, the rubric dimension that scored lowest. A grade is what an expert is shown, and a score is one number inside it.

SEE ALSO
- score
- grading pass

NOT A score. The score is part of the grade, not the whole of it.

Grading pass

MEANING One run of a section or a set against its rubric, producing a grade. It is the only thing that writes a score, and it never changes the document.

SEE ALSO
- grade
- rubric

State

MEANING Where a section stands: the state pending, shown to the expert as not started; `draft`, written; `review`, ready for someone else to read, reached through the gate; or `final`, set by the expert under TR-002. LC-001 in section 5 shows each move. The set as a whole reads as not yet approved, approved, approved with conditions, or rejected, by the approval recorded in section 15.

NOT A threshold. The threshold is the score a gate needs. The state is where the section stands.

SEE ALSO

- threshold
- blocking defect

Override

MEANING An expert's decision to advance a section past a gate it has not met, recorded in the change log with the score at the time.

SEE ALSO

- state
- threshold

NOT An approval. An override moves one section past its gate, and an approval covers the whole set.

Verdict

MEANING What a buildability review concludes, one of three: `buildable`, `buildable with gaps`, `not buildable`. In ST-009 the engineering lead may record their own verdict, and where it differs, theirs is the verdict and the buildability review's is kept beside it.

SEE ALSO

- derivability checks

NOT A grade. A grade scores one section against its rubric, and a verdict is the buildability review's conclusion on the whole set.

Derivability checks

MEANING The eleven checks a buildability review runs to decide whether a backlog can be derived from a set. Q-009 asks which edition of them release 1 grades against.

SEE ALSO

- verdict
- rubric

Candidate architecture decision

MEANING A technology or design preference the expert stated during the interview, kept out of the requirements and handed to the engineering lead as a topic for an architecture decision record.

SEE ALSO

- the five tests

NOT A requirement. It is deliberately not one, because it would fail test 5.

Unit of work

MEANING What the expert's domain counts one piece of work as: a load, a claim, a case, a patient. Every question the expert is asked uses their word for it.

NOT An engagement. An engagement is the writing of one set with this product, and the unit of work is what the expert's own domain counts.

Unverified

MEANING

A citation that was checked and could not be confirmed against its source. It is marked as such and never stated as fact.

SEE ALSO

- expert

MoSCoW

MEANING

The priority scale this set uses: `must`, `should`, `could`, `wont`. This set uses MoSCoW and not RFC 2119's MUST and SHOULD. `wont` keeps the source's meaning, won't have this time: a `wont` item is out of scope and revisited later, not a non-goal.

SEE ALSO

- the five tests
- out of scope

SOURCE

What is MoSCoW Prioritization?, Agile Business Consortium, accessed 2026-10-27, https://www.agilebusiness.org/?p=1907

Wireframe

MEANING

A sketch of one screen a person works in or one report a person reads. It shows what the person sees and does there, and names the requirements it makes concrete. Every requirement shown in a screen or a report has one, in section 7.

SEE ALSO

- the five tests

NOT

A design. A wireframe is labeled an illustration, and the development team decides how the screen looks.

Buildability review

MEANING

A reading of the whole set as the engineering lead who has to build it. It runs the eleven derivability checks and reports a verdict and, for each release-line stage, whether a requirement covers it or it is a gap.

SEE ALSO

- verdict
- derivability checks
- finding

NOT

A grading pass. A grading pass scores one section against its rubric; a buildability review reads the whole set and asks whether a backlog can be derived from it.

Finding

MEANING

One fault that a grading pass or a buildability review reports, with the document it is in quoted. A finding is checked against the document before the expert sees it, and is dropped when the document does not support it.

SEE ALSO

- blocking defect
- grading pass
- buildability review

NOT

A blocking defect. A blocking defect is one kind of finding, the kind that stops a section advancing. Every blocking defect is reported as a finding, but most findings only lower a score or name a weak dimension.

Root document

MEANING

The one document of a set that every other document names as its parent. It carries the summary, the scope, every section's state and grade, and every open question with its decider and date, and it is what an approver reads.

SEE ALSO

- document set
- approver

NOT

Section 1. The root document has no section number; section 1 is the problem and opportunity.

Approver

MEANING

A person the set names, by name and role under FR-021, who approves, approves with conditions, or rejects the whole set for a release under TR-001. A set names one or more approvers: the sponsor, or the expert when no sponsor exists. Only a named approver's decision counts. A set approved with conditions reads that way until the approver who stated each condition clears it.

SEE ALSO

- condition
- root document
- trust rule

NOT

The expert. The expert decides the state of each section; the approver decides on the whole set. They are one person only when no sponsor exists.

Condition

MEANING

What an approver who approves with conditions requires before the approval is unconditional. Each condition names the open question, requirement, or review comment it depends on and a date, and only the approver who stated it clears it. When the sponsor approves while a review comment has no response, the sponsor states a condition on that comment, and the expert records it (FR-014).

SEE ALSO

- approver
- open question

NOT

A constraint. A condition belongs to one approval and is cleared; a constraint holds for the product and is never cleared.

Constraint

MEANING

A rule the product must obey that the team cannot change by agreeing among themselves, recorded in section 9 with its category, the authority that imposes it by name, and how compliance is shown. A regulation or a standard appears in a constraint only when the expert named it.

SEE ALSO

- condition

NOT

A requirement. A requirement states what a user can do and has a priority; a constraint limits how every requirement it applies to may be met, and is not prioritized or deferred.

Objective

MEANING

A business result the product is expected to produce, recorded in section 3 with the problems it addresses, one metric, a baseline, a target with a date, and one owner. If the product ships and the metric does not move, the product did not work.

SEE ALSO

- baseline

NOT

A requirement. An objective says what the business gains; a requirement says what a user can do, and names the objective it serves.

Baseline

MEANING

The value an objective's metric has today, before the product. When it is not known, the baseline says unknown and names the open question that will establish it, and the target is an intention rather than a commitment until the question is answered.

SEE ALSO

- objective
- open question

NOT

A target. The baseline is where the metric stands now; the target is where the objective says it will be by a date.

Requirement

MEANING

A statement of what one named user can do with the product and the outcome they can see, recorded in section 7 with its acceptance criteria, a priority, the journey stage it serves, and the objective it serves. A statement is a requirement only when it meets the five tests.

SEE ALSO

- the five tests
- acceptance criteria
- objective

NOT

A story or a feature. A requirement is stated in the PRD, and stories and features are derived from requirements in the backlog.

Acceptance criteria

MEANING

The pass-or-fail statements on a requirement or a journey stage. Each is written as Given a situation, When the person acts, Then what they see, and each is one an engineer who was not in the interview could write a test from. A story derived from a requirement takes the requirement's criteria as its acceptance criteria.

SEE ALSO

- requirement
- stage
- the five tests

NOT

Exit criteria. Acceptance criteria decide whether one requirement or stage passes, and exit criteria decide whether a whole release is done. This set also departs from the source: North writes acceptance criteria as scenarios, and in this set a scenario is a different thing, a path through journey stages in section 5. The set departs because it writes Given/When/Then on one stage or one requirement, and it needs a separate word for a walk across several stages with real values, which section 5 calls a scenario.

SOURCE

Introducing BDD, Dan North, 20 Sep 2006, https://dannorth.net/blog/introducing-bdd/

Gap

MEANING

Something the set does not say that a team needs in order to build: most often a release-line stage that no requirement covers, as a buildability review reports it. A gap is closed with a requirement, or owned as an open question with a decider and a date.

SEE ALSO

- buildability review
- release line
- open question

NOT

A blocking defect. A gap is a thing the set leaves unsaid, and the buildability review reports each gap as a finding. A blocking defect is a finding that stops a section moving on.

Version

MEANING

A numbered state of the set, recorded in section 15 with its date, its author, and the sections it changed. An approval names the one version it covers, and a change written after an approval records a new version that the approval does not cover.

SEE ALSO

- approver
- condition

NOT

A release of the product. The product's releases are numbered too, and a set written by one release is read by later ones; a version of the set counts changes to the set.

Domain

MEANING

The field of work the expert specifies software for, such as freight. The expert names it before the first section. It sets the unit of work, the vocabulary every question uses, and the examples each section is taught with. The expert supplies knowledge of the domain; the product does not.

SEE ALSO

- unit of work

Reading order

MEANING

The order in which the documents of a set are read: the root document, then the section documents by section number, from section 1 to section 15. Where a section holds two documents, journeys come before scenarios and functional requirements come before wireframes. The next and the previous section of any document follow this order.

SEE ALSO

- section
- root document

Interview

MEANING

The questions the expert answers about one section in stage ST-002, after reading what the section is for. It ends as the question list says. Only the expert answers in it, so an engineer who reads the set later was not in the interview.

SEE ALSO

- section
- candidate architecture decision

Handover

MEANING

The point at which the approved set reaches the engineering lead who has to build from it. Journey J-002 starts there: its trigger is that the PRD has been approved for a release, and its first stage, ST-009, takes the approved set and the report of its buildability review. A set is handed over only once it is approved.

SEE ALSO

- approver
- candidate architecture decision

Dependency

MEANING

Something outside the product that the product relies on, recorded in section 11 with what happens if it is not there, such as D-001, an assistant session the expert is already working in.

SEE ALSO

- assumption
- risk

NOT

A requirement's wait on an open question. A requirement that waits on a question names the question, and the question lists the requirement in what it blocks.

Backlog

MEANING

The features, epics and stories a delivery lead derives from an approved set, in stages ST-007 and ST-008 of journey J-002. Every level of it names the id it came from.

SEE ALSO

- feature
- epic
- story

NOT

The set. The set says what the product must do; the backlog is derived from it once it is approved.

Feature

MEANING

The backlog item for one journey stage, carrying the ids of the requirements on that stage. The set departs from the source here: the source makes a feature any story big enough to be released, and this set makes it one journey stage, so every story traces to the backbone, which in this set is the journey stages.

SEE ALSO

- backlog
- stage

NOT

A requirement. A feature holds every requirement of one stage. Also not the everyday sense in section 1, where a "feature list" is a list of features with no named user.

SOURCE

Epics, Features and User Stories, Mountain Goat Software, updated 28 Apr 2025, https://www.mountaingoatsoftware.com/blog/stories-epics-and-themes

Epic

MEANING

A group of backlog items that one release delivers. The source defines an epic as a large user story. This set keeps that it is large, and fixes its size by the release rather than by judgment, so that each release's backlog is one group a delivery lead can plan, as ST-008's epics scoped by the release line need.

SEE ALSO

- backlog
- release line

NOT

A feature. A feature covers one journey stage, and an epic covers the stages one release delivers.

SOURCE

Epics, Features and User Stories, Mountain Goat Software, updated 28 Apr 2025, https://www.mountaingoatsoftware.com/blog/stories-epics-and-themes

Story

MEANING

The backlog item for one requirement, in the persona's own words, with the requirement's acceptance criteria as its own.

SEE ALSO

- backlog
- requirement
- acceptance criteria

NOT

A requirement. The requirement is in the set; the story is derived from it. This set departs from the source, where a story is anything a user wants. Here each story is bound to exactly one requirement, so every story carries the requirement id it came from, which is J-002's need.

SOURCE

Epics, Features and User Stories, Mountain Goat Software, updated 28 Apr 2025, https://www.mountaingoatsoftware.com/blog/stories-epics-and-themes

Expert

MEANING

The person who holds the domain knowledge in the interview, persona PER-001: the domain expert, or a product manager who speaks for one. The expert answers the interviews, owns the set, decides what the product must do, the domain's own terms, and the state of each section, and is the only one who may state a domain fact as true.

SEE ALSO

- domain
- interview
- sponsor

NOT

The approver. The expert approves their own set only when no sponsor exists, under TR-001.

Sponsor

MEANING

The person who commits budget and a date, persona PER-003, and who approves the set for a release under TR-001. The sponsor may approve, approve with conditions, or reject.

SEE ALSO

- approver
- expert

NOT

The expert. Where no sponsor exists, the expert approves their own set under TR-001, and the change log records that no separate approver existed.

Engineering lead

MEANING

The person who has to build from the approved set, persona PER-002. The engineering lead receives the set at handover, records their own verdict on its buildability review in ST-009, and receives each candidate architecture decision.

SEE ALSO

- handover
- buildability review
- candidate architecture decision

NOT

The expert. The engineering lead was not in the interview and builds from what the set says, not from what the expert meant.

Delivery lead

MEANING

The person who turns an approved set into a backlog and runs the story-mapping workshop, persona PER-004, in stages ST-007 and ST-008.

SEE ALSO

- backlog
- engineering lead

NOT

The engineering lead. The delivery lead plans the work from the set, and the engineering lead builds it.

Persona

MEANING

A named role from section 4, such as PER-001, the expert. Every requirement, journey, and trust rule names the persona it serves or who acts.

SEE ALSO

- expert
- journey

NOT

A user. Anyone who uses the product is a user, and a persona is one named role with its own goals.

Item

MEANING

One numbered entry in a section with its own id, such as a problem, a persona, a stage, a requirement, or an open question. A section is made of its items and the prose around them.

SEE ALSO

- id
- section

NOT

A line or a field of an entry. A requirement and all its acceptance criteria are one item.

Id

MEANING

The label that names one item, such as FR-008 or Q-001. Each id is unique in the set, and ids are never renumbered or reused.

SEE ALSO

- item
- link

Link

MEANING

A reference from one item to another by its id, such as the journey stage a requirement serves or the requirements an open question blocks.

SEE ALSO

- id

Priority

MEANING

How much a requirement or a scenario matters to a release, stated in MoSCoW terms. Only the expert gives a priority.

SEE ALSO

- MoSCoW

Teaching

MEANING

What the expert is told about a section before its interview: what it is, why it matters, and what belongs in it, in no more than ten lines and with an example from the expert's domain.

SEE ALSO

- interview
- domain

Decider

MEANING

The one named person who makes one decision: who answers an open question by its date, or who decides under one trust rule in section 10, whose table names that person in its Decider column.

SEE ALSO

- open question
- trust rule

NOT

The approver. The approver decides only whether the whole set is approved for a release, under TR-001, and is the decider of that one rule.

Release

MEANING

A numbered release, such as REL-001, with its outcomes, its exit criteria, and the journey stages it completes when it has any. The first release holds exactly the release line.

SEE ALSO

- release line
- scope
- exit criteria

NOT

The release line. The release line is the smallest set of stages that completes the golden thread, and a release is what ships. Release 1 ships exactly the release line.

Vocabulary

MEANING

The unit of work and the words the expert says an outsider would misread, each with the expert's own meaning. The expert states it once, before the first section. Every later question uses it, and it does not change when the expert resumes.

SEE ALSO

- unit of work
- domain

NOT

The glossary. The vocabulary is the few words the expert states at the start, and the glossary holds every domain term of the set.

Engineer

MEANING

Anyone on the team who builds from the set and was not in the interview. The engineering lead is one. An acceptance criterion is good only when an engineer could write a test from it.

SEE ALSO

- engineering lead
- acceptance criteria

AVOID

- builder

Scenario

MEANING

One concrete path a persona takes through one or more journey stages, recorded in section 5 with its steps and the outcome they see. A scenario shows a stage working, or failing, with real content.

SEE ALSO

- journey
- stage

NOT

A journey. The journey names the stages, and a scenario walks one path through them.

Trust rule

MEANING

A rule in section 10 that names who decides one decision, at which stages, on what evidence, and what happens when that person is away. A person decides every rule. A rule marked never delegated also has no agent proposal behind it.

SEE ALSO

- approver
- override

Story-mapping workshop

MEANING

The meeting in which the delivery lead and the team map a backlog against the journeys of an approved set. It needs a backbone agreed before it starts.

SEE ALSO

- backbone
- backlog
- delivery lead

SOURCE

The New User Story Backlog is a Map, Jeff Patton & Associates, 8 Oct 2008, updated 12 Apr 2023, https://www.jpattonassociates.com/the-new-backlog/

Backbone

MEANING

The ordered journey stages that a story-mapping workshop maps features against, because a story map needs an order of work to hang its stories on. A set with no journey gives the workshop no backbone, so the workshop invents one in the room. This set departs from the source, where the backbone is the capabilities the system needs. In this set the backbone is the journey stages, including stages that happen outside the product.

SEE ALSO

- journey
- story-mapping workshop

SOURCE

The New User Story Backlog is a Map, Jeff Patton & Associates, 8 Oct 2008, updated 12 Apr 2023, https://www.jpattonassociates.com/the-new-backlog/

Question list

MEANING

The questions an interview may put for one section. The interview ends when the list is finished, when the expert says they have nothing more to add, or earlier, when the expert's answers would already meet the gate.

SEE ALSO

- interview

Gate

MEANING

What a section must meet to move to its next state. For `review`, a score at or above the threshold with no blocking defect, in a set that checks clean. The expert may override a gate, and the override is recorded.

SEE ALSO

- state
- override
- threshold

NOT

The threshold. The threshold is the score alone, and a gate also needs no blocking defect. Also not the everyday sense in sections 1 and 10, where "gates an engineer would have to invent" means decision points the set left unstated.

Out of scope

MEANING

Something the product will not do in a release, with a point at which it is revisited, recorded in section 6.

SEE ALSO

- non-goal
- scope

NOT

A non-goal. A non-goal is never done, and an out-of-scope item is not done yet.

Objection

MEANING

A finding that the document supports and that the expert refutes. Its reason is stated to the expert once, and the expert's decision is recorded. The decision is reaffirmed when the expert keeps it after hearing the reason. A finding that names the same item for the same reason is the same objection, and a later grading pass does not raise it again as new.

SEE ALSO

- finding
- grading pass

NOT

A refuted finding the document does not support. A finding of this kind is dropped, not kept as an objection.

Refute

MEANING

To challenge a finding, or a score through its findings. The finding is then checked against the document. If the document does not support it, it is dropped and the expert is told. If the document supports it, it stays as an objection, and the expert's decision is recorded.

SEE ALSO

- finding
- objection

AVOID

- dispute

NOT

Showing a finding wrong. A refuted finding can still stand, when the document supports it.

Session

MEANING

The assistant conversation the expert is already working in, which the product works inside, as dependency D-001 records. The session shows what the product shows, and it carries any number of people working at once. A session ends when that conversation ends, and the set stays where it was written.

SEE ALSO

- dependency

NOT

A sitting, which is a stretch of time spent on the work and can span more than one session.

Change log

MEANING

Section 15, which records each version of the set with its date, its author, and the sections it changed, and every approval and override.

SEE ALSO

- version
- approver
- override

Target

MEANING

The value a measure must reach: where an objective says it will be, or what a non-functional requirement must meet, such as under 20 seconds at p95.

SEE ALSO

- objective
- baseline

Exit criteria

MEANING What must be true before a release counts as done, stated with each release in section 13.

SEE ALSO
- release

NOT Acceptance criteria. Acceptance criteria decide one requirement or stage, and exit criteria decide the whole release.

Non-functional requirement

MEANING A quality the product must have, such as speed, accessibility, or retention, recorded in section 8 with a measure and a target. It applies across requirements, not to one user's capability.

SEE ALSO
- requirement
- target

NOT A requirement in section 7. A requirement names a user and what they can do.

Bet

MEANING The one outcome the product stakes itself on, stated in the root document's summary: what the expert reaches with the product. The summary's objective and measure say how anyone will know it happened.

SEE ALSO
- summary
- objective

NOT An objective. An objective is measured, and the bet is the outcome the objective measures.

Summary

MEANING The opening of the root document, which states in a few sentences the problem, the users, why now, the bet, the one objective that proves it, and how it is measured. An approver reads it first.

SEE ALSO
- root document
- objective

Glossary

MEANING Section 14, this section. It gives each domain term one meaning, and every requirement and acceptance criterion uses those words.

SEE ALSO
- vocabulary

View

MEANING What the expert has in front of them at one moment, such as one section or its grade.

SEE ALSO
- section

Topic

MEANING One subject inside a section's interview, such as one requirement or one stage. A follow-up question finishes a topic before the next one opens.

SEE ALSO
- interview
- question list

Ready

MEANING

The condition a section meets before it moves to the state `review`: it meets the gate and the expert accepts it, or the expert overrides the gate and the override is recorded.

SEE ALSO

- state
- gate

NOT

Offered for approval. Ready describes one section. The expert decides, under TR-010, whether the whole set may be offered for approval. Ready is also not a verdict of the buildability review.

Capability

MEANING

The first part of a requirement: what one named user can do with the product, in their own words.

SEE ALSO

- requirement
- outcome

NOT

A capability the system needs, as the backbone's source uses the word. In this set a capability belongs to one user, and the backbone is the journey stages.

Outcome

MEANING

The second part of a requirement: what the user sees or has once the capability exists, stated so that someone else could check it.

SEE ALSO

- requirement
- capability

NOT

A feeling. "They feel confident" is not an outcome a tester can check.

Problem

MEANING

A condition a named persona suffers today, recorded in section 1 with counted evidence, why it matters now, and what it costs to leave it.

SEE ALSO

- persona
- objective

Approval

MEANING

The approver's recorded decision on one version of the set: approve, approve with conditions, or reject, with the approver, their role, the date, and the version. A rejection is an approval record too. An approval covers only the version it names, and a later change records a new version that it does not cover.

SEE ALSO

- approver
- condition
- version

NOT

An override. An override moves one section past its gate, and an approval is a decision on the whole set.

Engagement

MEANING

One piece of work in which one expert writes one set with this product, from the first question to the approval. TR-001 covers an engagement with no sponsor.

SEE ALSO

- expert
- sponsor

NOT

A unit of work. The unit of work is what the expert's own domain counts.

Lifecycle

MEANING The states one thing passes through and the moves between them, with the person who makes each move, recorded in section 5. LC-001 is the lifecycle of a section of the set.

SEE ALSO
- state
- scenario

Handoff

MEANING The point at which one stage's work passes to the next person or the next stage, named on each journey stage.

SEE ALSO
- stage
- handover

NOT The handover. The handover is the one handoff at which the approved set reaches the engineering lead.

Segment

MEANING A group of organizations that would buy the product for the same reason, recorded in section 2 with its buyer and its user, such as S-001.

SEE ALSO
- buyer

Buyer

MEANING The person in a segment who decides to pay for the product, such as a principal or a practice lead.

SEE ALSO
- segment

NOT The user. The buyer pays, and the user is the person who works with the product, such as the consultant with the domain knowledge.

Status quo

MEANING What a buyer does today with no product: a smart practitioner, a document, and a willing engineer. Section 2 records it as a competitor, because every buyer already has it.

SEE ALSO
- segment

Delegated

MEANING Given to an agent to propose. Section 10 marks for each trust rule whether an agent may propose its outcome, and the person the rule names still decides. A rule that is never delegated has no agent proposal at all.

SEE ALSO
- trust rule
- agent

NOT Decided by a person on someone else's behalf. A person who decides while the decider is away does so under the rule's escalation, and that is not delegation.

Definition of done

MEANING The conditions every backlog item must meet before the team counts it as done. The delivery lead needs the set's constraints to arrive as items of it, not as surprises, in J-002.

SEE ALSO
- backlog
- constraint

Checks clean

MEANING

Every id and every link in the set resolves to an item or a document that exists, and every document has the parts its section needs.

SEE ALSO

- gate
- link

Pilot PRD

MEANING

One of the first 5 PRDs written with release 1 by people other than the product's author. O-001 and REL-001's exit criteria count them.

SEE ALSO

- first submission
- release

First submission

MEANING

The first time a PRD goes through the buildability review. Only that first verdict counts, so no gap fixed after it changes the result.

SEE ALSO

- buildability review
- pilot PRD

Reviewer

MEANING

A person who comments on a section inside the set without being able to change it, as REL-003 plans. No persona in release 1 is a reviewer. In release 1 a sponsor comments on an exported copy instead, as a review comment.

SEE ALSO

- review

NOT

A grading pass. The grading pass that scores a section is not a person.

Iteration

MEANING

One fixed period of the development team's work after the PRD is approved. "Iteration three" means late in the build, after the plan is made.

SEE ALSO

- release

Review

MEANING

The state `review`, which a section enters once it is ready. In this state, someone other than the expert may read it.

SEE ALSO

- state
- buildability review
- reviewer

NOT

The buildability review, which is its own term, or the expert reading a draft, which the set calls reading.

Pass

MEANING

A PRD passes the buildability review when its verdict is `buildable`, or `buildable with gaps` with every gap closed or owned, as O-001 defines.

SEE ALSO

- buildability review
- verdict
- first submission

NOT

Meeting the gate, or a grading pass. A section meets the gate, a grading pass is one run of the rubric, and a PRD passes the buildability review.

Indicative

MEANING

Said of a release date the plan aims at and has not committed to, marked indicative in section 13.

SEE ALSO

- release

Follow-up

MEANING

A further question about the answer just given, asked before the next topic opens, when the answer leaves something the section needs unanswered or could be written into the section in two ways.

SEE ALSO

- interview
- topic

Specification

MEANING

In the general sense only: a requirements document written without this method, as section 1 and Q-001 use the word. This set calls itself a PRD, and calls the documents that make it up the set.

SEE ALSO

- prd
- document set

Development team

MEANING

The people who build from an approved set: the engineering lead, the delivery lead, and the engineers, as section 4 names them.

SEE ALSO

- engineering lead
- delivery lead
- engineer

Offered for approval

MEANING

Said of a set the expert has decided, under TR-010 in stage ST-005, may go to the approver, as its root document or an exported copy. A set offered for approval reads as not approved until the approver decides.

SEE ALSO

- ready
- approver
- exported copy

NOT

Ready. Ready describes one section, and a set is offered for approval as a whole.

Sitting

MEANING A stretch of time an expert spends on the work, such as one of the 6 to 9 working sittings per PRD in section 4, or one of D-002's ten observed sittings. One sitting can span more than one session.

SEE ALSO
- session

NOT A session. A session is the assistant conversation the product works inside.

Loop-back

MEANING A return to an earlier section because a later one showed it was incomplete, such as writing the requirements sending the expert back to the journeys. ST-004 allows it, and Q-007 counts loop-backs in D-002's observed sittings.

SEE ALSO
- section
- sitting

Experiment

MEANING A release shipped while assumptions it depends on are still untested, and said to be so. Section 13 names this as REL-001's course if D-002 or D-003 reports late, and Q-022 decides it.

SEE ALSO
- release
- assumption

NOT A spike. A spike produces an answer, and an experiment ships a release.

Agent

MEANING Software that proposes an outcome under a trust rule section 10 marks as delegable to an agent, such as TR-003's score and findings. The person the rule names accepts or refutes what it proposes.

SEE ALSO
- delegated
- trust rule

NOT A person. Every decider in section 10 is a person.

Exported copy

MEANING The set in a form the sponsor reads outside the product, in place of the root document, as in ST-006.

SEE ALSO
- root document
- offered for approval

NOT The set. The set is its documents, and an exported copy is one form of them for reading.

Review comment

MEANING

A comment a person reading an exported copy leaves in that copy. In release 1 that person is the sponsor (FR-022). The expert brings it back into the set and records a response with one of three statuses: accepted, when the set changed because of it; declined, when the set did not change, with the reason; or answered, when it asked a question and the set did not need to change (FR-023). Until it has a response, it is open, which is not an open question in section 12.

SEE ALSO

- exported copy
- sponsor
- finding
- reviewer

NOT

A finding. A finding comes from a grading pass or a buildability review, and a review comment comes from a person reading a copy.

Buildable

MEANING

Said of a set a development team can estimate and build from without re-deriving it, which a buildability review reports as the verdict `buildable`.

SEE ALSO

- verdict
- buildability review

NOT

A pass. A set with gaps can pass the buildability review when every gap is closed or owned, and it is still not buildable while any gap is left.

Product manager

MEANING

A person who writes product requirements for a living. PER-001 is a product manager who is, or speaks for, the domain expert, and the set calls that person the expert.

SEE ALSO

- expert

NOT

The sponsor. A product manager writes the set, and the sponsor approves it.

Re-grade

MEANING

A later grading pass on the same section, against the same rubric, after the section changed.

SEE ALSO

- grading pass
- grade

Escalation

MEANING

What happens to a decision when its decider is away or disagrees, as each trust rule in section 10 states.

SEE ALSO

- trust rule
- decider

NOT

An open question. An open question waits for an answer. An escalation names who acts when the decider cannot.

Satisfaction

MEANING How one journey stage feels to the person today, from 1 (worst) to 5 (best), as the expert scores it in section 5. Only the expert gives it.

SEE ALSO

- stage
- journey
- expert

NOT A measure of the product. The product never estimates it.

Pain point

MEANING What is hard for a person at one journey stage today, recorded in section 5 with the need that answers it.

SEE ALSO

- stage
- problem

NOT The problem in section 1. The problem is the whole cost. A pain point is one part of it at one stage.

SECTION 15

Change log and approvals

DRAFT 75/100 NO BLOCKING DEFECTS

VERSION	DATE	AUTHOR	WHAT CHANGED
0.13.0	2026-11-01	Richard Hightower	The set gains a roadmap beside the release plan, one item for each release; the release plan's roadmap_link points to it; no release, requirement, or objective changed. On 2026-11-02, Q-019 is answered yes, D-002 and D-003 are committed, and RM-001 moves to on track; REL-001 and RM-001 also wait on R-001 and R-002, RM-002 waits on Q-023, REL-002 and REL-003 gain exit criteria, and five open questions gain counted rules
0.12.0	2026-10-30	Richard Hightower	The root names its approver by name and role (FR-021) and shows the approver the scope, each section's state and grade, the open questions, and any blocking defect (FR-013); TR-015 keeps the domain check's research on by default and lets the expert turn it off for a set, PER-001's authority includes that right, and an organization's own policy can forbid it; NFR-008 and A-005 state that default
0.11.0	2026-10-29	Richard Hightower	FR-023 split out of FR-022 for the expert's import and response; the open-comment approval rule moved into FR-014, and the sponsor sets each condition's date; OUT-005 defers comments inside the set to REL-003; WF-011 shows an open comment; NFR-009 counts product releases; ST-009 hands off to the development team; section 10 says "delegable to an agent"; the glossary adds escalation, satisfaction, and pain point, and separates ready from review
0.10.0	2026-10-28	Richard Hightower	FR-022 added as a should: a sponsor's comments on an exported copy come back, each with a response, and an open one at approval becomes a condition; WF-011 shows the answered copy; IN-007, REL-001, TR-001, and the glossary include it; REL-003 keeps comments inside the set
0.9.0	2026-10-27	Richard Hightower	Counted evidence and sourced why-now claims; persona roles and frequencies; satisfaction scores; ST-006 outside the product; FR-019 to FR-021 added; FR-014 rejection; TR-014 for accepting a score; LC-001 back to draft; C-004 out of section 9; single questions due before REL-001; sourced glossary terms with one meaning each; the versioning policy
0.8.1	2026-10-27	Richard Hightower	Glossary terms, observable scenario outcomes, open questions named in section 7, competitor edges; C-001 to C-003 and S-003 moved; ST-009 for the engineering lead
0.8.0	2026-10-27	Richard Hightower	Requirement and scenario priorities, SC-007, release exit criteria, single deciders
0.7.0	2026-10-27	Richard Hightower	Root objective, primary persona and approver; release id in scope; LC-001 rule
0.6.0	2026-10-27	Richard Hightower	Links between problems, personas, objectives, competitors, trust-rule stages, and scenario views
0.5.0	2026-10-26	Richard Hightower	Given/When/Then acceptance on every stage; LC-001; scenario wireframe links; persona short names
0.4.0	2026-10-26	Richard Hightower	Navigation and refresh requirements; wireframes rewritten from their grading
0.3.0	2026-10-26	Richard Hightower	Wireframes for the requirements shown on a screen or in a report, and the term wireframe

VERSION	DATE	AUTHOR	WHAT CHANGED
0.2.0	2026-10-26	Richard Hightower	J-001's goal without the engineering lead's confirmation; FR-003 and FR-004 outcomes; FR-015 and FR-016 split out; the threshold stated as 70; NFR-010 to NFR-012; TR-007 to TR-012
0.1.0	2026-10-23	Richard Hightower	First draft of sixteen section documents across fifteen sections

0.9.0, what changed

- **Section 1:** each problem has one counted case with its source and date, and each why-now claim cites the General Assembly survey of product managers. Desired outcomes name no mechanism, and Q-018 asks for one counted cost per problem.
- **Section 2:** ChatPRD is the direct competitor. Why-we-win is one sentence we expect the sponsor to say, each buyer is PER-003, and "client contract" names the contract where "engagement" did.
- **Section 3:** O-001 counts the 5 pilot PRDs and defines a pass, O-002 counts the pilot PRDs with one journey, and O-003 counts a requirement testable when two of three engineering leads agree.
- **Section 4:** on the expert's decision, PER-001 is a senior product manager in a consultancy or an in-house product team who is, or speaks for, the domain expert. Each persona has a population, a frequency per organization, a basis, a context, and its problems, and no persona is external.
- **Section 5:** each stage carries the expert's satisfaction score, and each pain point names its stage. ST-006 happens outside the product. LC-001 makes final the expert's decision under TR-002, and lets a section at review or final go back to draft. SC-004 splits into SC-004 and SC-008.
- **Section 6:** ST-006 is outside the product, and teams new to a domain are covered by the first non-goal.
- **Section 7:** FR-019 and FR-020, both `should`, split out of FR-004 and FR-008. FR-014 records a rejection. FR-015 moves to ST-004. WF-001 draws FR-020, and the wireframes name no tool.
- **Section 8:** the NFRs name what they time, and NFR-010 and NFR-012 are measured in a test organization.
- **Section 9:** on the expert's decision, C-004 leaves, and the section may be empty. Its rule now lives in TR-005 in section 10 and FR-006 in section 7.
- **Section 10:** TR-003 only produces a score, and TR-014 has the expert accept or refute it. TR-001 covers a rejection. The three delegable rules propose and never decide.
- **Section 11:** R-001 and R-002 are mitigated by 2027-01-31, and D-002 and D-003 report by 2027-01-31.
- **Section 12:** each question is single, and those that block REL-001, including the baselines Q-001 to Q-003, are due by 2027-02-28. Q-011 is answered. Q-019 splits into Q-019 and Q-022, and Q-023 tests PER-004.
- **Section 13:** REL-001 depends on D-002, D-003, Q-019, and Q-022.
- **Section 14:** eight terms cite a verified source, and review, pass, refute, approval, and engagement each have one meaning. Session, buildable, development team, and the general sense of specification are defined. Ready describes one section, and offered for approval describes the set. Sitting, loop-back, experiment, agent, and exported copy are defined.
- **Section 15:** the versioning policy, a 0.5.0 entry, and the date the set is offered for approval, 2027-03-15.
- **Section 7 and section 10:** on the expert's decision, FR-021, a `must`, lets the expert name each approver by name and role before offering the set. TR-001's decider stays the sponsor's role.

None of this is structural under the policy. Richard Hightower decided this on 2026-10-27. ST-006 stays on the release line with the same goal and end state, and no stage is added, removed, or reordered. PER-001's new wording changes no

stage. The release line's text now quotes the existing end state, and ST-006 is done outside the product.

Versioning policy before 1.0

The set and the product share one version number. The number changes by these rules:

- A patch, such as 0.9.0 to 0.9.1, is for a correction that changes what a sentence means. A typo fix gets no version.
- A minor version, such as 0.8 to 0.9, is for new content in the set or in the product.
- A change to the release line or to the golden thread takes the set to the next major version. From 0.x, the next major version is 1.0.

A change to the golden thread or to the release line is structural: adding, removing, or reordering a stage, or changing the journey's goal or end state. Rewording a stage's text or its acceptance is not such a change, and it takes a minor version at most.

The expert set this policy at 0.9.0, and it applies from 0.9.0 onward. From 0.9.0, a reader can rely on the version number to follow it. Earlier version numbers stay as recorded. Two earlier versions do not follow it, and neither is reissued:

- 0.8.1 added stage ST-009 and moved items between sections. Under this policy it would have been the next minor version.
- 0.2.0 changed the golden thread. Under this policy it takes the set to 1.0.

Product 0.5.1 changed only the product, so the log has no entry for it.

0.2.0, what the grading changed

The reviewers that graded 0.1.0 were the product's own, run against its own PRD. They found twelve blocking defects and four failed derivability checks. The largest single finding was that FR-009, FR-010 and FR-011 were written against "its rubric" and "its threshold" with neither defined anywhere in the set: the threshold is now stated as 70 in the requirements themselves, and the rubric, the section states and the derivability checks are defined in section 14 as the product content they are.

Two findings were checked and **rejected**: the claim that "see section 6" pointed at the wrong document, and the claim that the section count was contradicted. Both came from reading file numbers as section numbers. The references were right; the numbering was unexplained, which is now fixed in the root document.

0.1.0, what to know

Written as prd-forge's own first exercise. Three things a reader should carry forward:

1. **No baselines.** All three objectives carry an explicit unknown with a matching open question (Q-001 to Q-003). The targets are intentions, not commitments.
2. **C-001 is the tension.** The product is delivered inside an assistant session. Naming that in a requirement would fail test 5, so it is recorded as a constraint. A reviewer may fairly call it architecture. At 0.8.1 it moved to section 6 as a non-goal, and the session it needs is dependency D-001 in section 11.
3. **R-003 is not closed.** Nothing in this product checks whether the product being specified should exist at all.

Approvals

None sought. This set is at `draft` and has not been offered for approval. It cannot be approved for a release without an approval recorded here, under TR-001. A section reaches `final` when the expert sets it final, under TR-002.

The set is offered for approval by 2027-03-15. Under FR-021 its one named approver is Richard Hightower, the owner of the place where this PRD is kept, in the sponsor's role, PER-003. Richard Hightower is also its author and owner, so this is a self-approval, and the entry will say so. When an outside reader reviews the set, their review is recorded beside the approval.

Each approval is recorded with the version, the approver's name and role, the date, and the decision. An approval with conditions records each condition as "conditional on <id> by <date>", where the id is the open question, requirement, or review comment the condition depends on, as FR-014 and FR-022 require. Only the approver who stated a condition clears it. A rejection records its reason.

At 0.9.0 the expert holds every other section to a score over 80, except section 2 until Q-004 and Q-005 are answered, and accepts this section below 80 until the set is offered for approval on 2027-03-15, because its rubric cannot score approvals and conditions higher before an approval exists.

> **The set gains a roadmap, 18-roadmap.md, beside the release plan in section 13. It holds on**
>
> VERSION 0.13.0
>
> DATE 2026-11-01
>
> AUTHOR Richard Hightower
>
> SECTIONS
> - 11
> - 12
> - 13
>
> RESOLVES
> - Q-019

> **The root names its approver by name and role under FR-021, and its details block now carri**
>
> VERSION 0.12.0
>
> DATE 2026-10-30
>
> AUTHOR Richard Hightower
>
> SECTIONS
> - 4
> - 8
> - 10
> - 11

> **FR-023 split out of FR-022: FR-022 is what the sponsor sees in the next copy, and FR-023 i**
>
> VERSION 0.11.0
>
> DATE 2026-10-29
>
> AUTHOR Richard Hightower
>
> SECTIONS
> - 5
> - 6
> - 7
> - 8
> - 10
> - 13
> - 14
> - 15

FR-022 added as a should in REL-001: a sponsor comments on an exported copy without using

VERSION 0.10.0

DATE 2026-10-28

AUTHOR Richard Hightower

SECTIONS
- 6
- 7
- 10
- 13
- 14
- 15

Problems carry counted, sourced evidence. Personas have named roles and frequencies, and e

VERSION 0.9.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 1
- 2
- 3
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11
- 12
- 13
- 14
- 15

RESOLVES
- Q-011

The glossary defines buildability review, finding, root document, approver, condition, con

VERSION 0.8.1

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 2
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11
- 12
- 13
- 14

Sets the expert's priorities. Ten requirements stay must and eight become should; SC-002 a

VERSION 0.8.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 5
- 7
- 12
- 13

Names the root's objective, primary persona, and approver, from section 3, the personas do

VERSION 0.7.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 4
- 5
- 6

Each problem names its persona, and each objective names the problems it addresses and the

VERSION 0.6.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 1
- 2
- 3
- 5
- 10
- 11

Every journey stage gains Given/When/Then acceptance. The scenarios document gains LC-001,

VERSION	0.5.0
DATE	2026-10-26
AUTHOR	Richard Hightower
SECTIONS	• 4 • 5

Adds FR-017, moving between sections without losing one's place, and FR-018, seeing the se

VERSION	0.4.0
DATE	2026-10-26
AUTHOR	Richard Hightower
SECTIONS	• 6 • 7 • 8 • 13 • 14

Adds wireframes to section 7. Each functional requirement now says where its user meets th

VERSION	0.3.0
DATE	2026-10-26
AUTHOR	Richard Hightower
SECTIONS	• 7 • 14

Fixes every defect the adversarial grading and the buildability review found in 0.1.0. J-0

VERSION	0.2.0
DATE	2026-10-26
AUTHOR	Richard Hightower
SECTIONS	• 4 • 5 • 6 • 7 • 8 • 10 • 13 • 14

First draft of the whole set, written with prd-forge's own conventions as the product's fi

VERSION	0.1.0
DATE	2026-10-23
AUTHOR	Richard Hightower
SECTIONS	1 2 3 4 5 6 7 8 9 10 11 12 13 14 15