PRODUCT REQUIREMENTS DOCUMENT

# prd-forge

A domain expert, with no product management training, can reach a PRD that an engineering lead confirms is buildable, and can see for themselves where it is still weak before anyone starts work.

| | |
|---|---|
| VERSION | 0.9.0 |
| STATUS | draft |
| OWNER | Richard Hightower |
| DOMAIN | product management tooling |
| CREATED | 2026-10-23 |
| UPDATED | 2026-10-27 |
| GOLDEN THREAD | J-001 |
| STAKEHOLDERS | Domain experts who must specify software for their own field; Engineering leads who build from the resulting PRD; Delivery leads who run story-mapping workshops |

## Executive summary

| | |
|---|---|
| **PROBLEM** | A domain expert knows exactly what their product must do, but what they can write down is a feature list with no named users, no pass-or-fail criteria and the technology already chosen, so the team that has to build it cannot. |
| **USERS** | Domain experts specifying software for their own field, the engineering leads who build from what they write, and the sponsors who approve it. |
| **WHY NOW** | Teams now expect to work from a written specification an assistant helped produce, and a specification that reads well but cannot be built from costs more than no specification at all, because it is believed. |
| **BET** | A domain expert, with no product management training, can reach a PRD that an engineering lead confirms is buildable, and can see for themselves where it is still weak before anyone starts work. |
| **OBJECTIVE** | O-001 |
| **MEASURE** | Share of PRDs that pass an engineering buildability review on the first submission: 80% by 2027-09-30, from a baseline Q-001 will establish. |

## Contents


| § | SECTION | STATUS | GRADE | BLOCKING |
|---|---|---|---|---|
| 1 | Problem and opportunity | draft | 71/100 | — |
| 2 | Target customer and market | draft | 55/100 | — |
| 3 | Business objectives and success metrics | draft | 75/100 | — |
| 4 | Users and personas | draft | 68/100 | — |
| 5 | User journeys | draft | 79/100 | — |
| 5 | Scenarios | draft | 78/100 | — |
| 6 | Scope | draft | 79/100 | — |
| 7 | Functional requirements | draft | 81/100 | — |
| 7 | Wireframes | draft | 78/100 | 2 |
| 8 | Non-functional requirements | draft | 82/100 | — |
| 9 | Constraints | draft | 95/100 | — |
| 10 | Decision rights and trust rules | draft | 88/100 | — |
| 11 | Assumptions, risks and dependencies | draft | 75/100 | — |

| § | SECTION | STATUS | GRADE | BLOCKING |
|---|---|---|---|---|
| 12 | Open questions | draft | 71/100 | — |
| 13 | Release plan | draft | 83/100 | — |
| 14 | Glossary | draft | 75/100 | — |
| 15 | Change log and approvals | draft | 79/100 | — |

## About this document

A domain expert knows exactly what their product must do, but what they can write down is a feature list with no named users, no pass-or-fail criteria and the technology already chosen, so the team that has to build it cannot. The users are domain experts specifying software for their own field, the engineering leads who build from what they write, and the sponsors who approve it. Teams now expect to work from a written specification an assistant helped produce, and one that reads well but cannot be built from costs more than no specification at all, because it is believed. The bet is that a domain expert with no product management training can reach a PRD an engineering lead confirms is buildable, and can see for themselves where it is still weak before anyone starts work. We will know it worked from the share of PRDs that pass an engineering buildability review on the first submission.

### Contents


| SECTION | DOCUMENT |
|---|---|
| 1 | Problem and opportunity |
| 2 | Target customer and market |
| 3 | Business objectives and success metrics |
| 4 | Users and personas |
| 5 | User journeys |
| 5 | Scenarios |
| 6 | Scope |
| 7 | Functional requirements |
| 7 | Wireframes |
| 8 | Non-functional requirements |
| 9 | Constraints |
| 10 | Decision rights and trust rules |
| 11 | Assumptions, risks and dependencies |
| 12 | Open questions |
| 13 | Release plan |
| 14 | Glossary |

| SECTION | DOCUMENT |
|---|---|
| 15 | Change log and approvals |

### How this document set is numbered

A PRD has **fifteen numbered sections** and this set has **seventeen section documents**, because section 5 holds two of them, journeys and scenarios, and section 7 holds two, functional requirements and wireframes. The two section 7 files share the prefix `08`. From section 5 onward, the file prefix and the section number differ: `07-scope.md` is section 6, and `16-change-log.md` is section 15. **Every cross-reference in this set is by section number, never by file number.** Three independent reviewers misread this, so it is stated here rather than left to be inferred.

### A note on this PRD

This is prd-forge's PRD, written with prd-forge's own conventions. Writing it surfaced one real tension: the product is delivered inside an assistant the expert already uses, and naming that delivery mechanism in a requirement would fail test 5. It was recorded as constraint C-001 until 0.8.1. It is now the separate-application non-goal in section 6 and dependency D-001 in section 11, and the requirements in section 7 describe only what a person does and sees.

Grading it with its own adversarial reviewers found twelve blocking defects across four sections, and a buildability review found four failed derivability checks. Version 0.2.0 fixes them: every graded section is now at zero blocking defects, and nine of the eleven derivability checks pass. What that grading found, and what it changed, is in section 15.

**On the scores in this set.** Each graded section carries the score from its last grading pass, with the date. Sections 5, 7 and 8 were graded before the second round of corrections, and section 10 before its last two, so each score is the last one measured and not a reading of the words as they now stand. Treat them as a floor. A score never rises without a change to the document, and never rises without a grading pass; the two gaps here are in the other direction.

SECTION 1

# Problem and opportunity

DRAFT 71/100 NO BLOCKING DEFECTS

| ID | PROBLEM | WHO |
|---|---|---|
| P-001 | What the expert writes is a feature list with no user and no testable outcome | Domain expert |
| P-002 | The steps between requirements are never written, so paths have gaps | Engineering lead |
| P-003 | Buildability is invisible until the team starts building | Sponsor or approver |
| P-004 | No source of PRD craft at the moment the expert needs it | Domain expert |

The four compound. P-004 means the expert does not know what good looks like; P-001 is what they produce as a result; P-002 is the specific gap that costs the most; and P-003 is why none of it is caught before it is expensive.

**Evidence note.** Each problem now has one counted case, with its source and date. The counts for P-001 to P-003 come from one engagement, this product's own PRD, written and graded with prd-forge. The count for P-004 comes from two public templates. One engagement shows that each problem occurs. It does not show how often. The other evidence items describe a pattern and are not measured. Three of the objectives in section 3 still have no baseline. See Q-001, Q-002, and Q-003.

### P-001

**STATEMENT** A domain expert can describe what their product must do out loud, but what they write down is a feature list: no named user, no outcome anyone could check, and a technology already chosen for reasons they cannot defend.

**WHO** Domain expert specifying software for their own field

**PERSONA** PER-001

**EVIDENCE**

- {'claim': 'In this PRD at version 0.1.0, the adversarial grade found 2 of 14 functional requirements, FR-003 and FR-004, with no outcome a person could see. Both were fixed at version 0.2.0.', 'source': 'prd-forge repository, commits 5ba59cb (version 0.1.0) and 5d4d936 (version 0.2.0)', 'as_of': datetime.date(2026, 10, 26)}
- Specifications arriving at delivery teams are commonly feature lists; teams report re-deriving the user and the acceptance criteria themselves before they can estimate.
- Experts asked to write requirements reach for the words of the software they have used, because that is the only requirements language they have been exposed to.

**WHY NOW** Specifications are now commonly drafted with an assistant, which makes them fluent. Fluent and unbuildable is a worse failure than rough and unbuildable, because nobody queries it.

**COST OF INACTION** The gap is closed by an engineer guessing, during build, at what the expert meant, and the guess is discovered at review or in production.

### P-002

**STATEMENT** Requirements are written one at a time, so the steps between them are never written at all, and nobody notices until the team tries to assemble a working path through the product.

**WHO** Engineering lead building from the specification

**PERSONA** PER-002

EVIDENCE

- {'claim': 'At version 0.1.0, the buildability review of this PRD failed 4 of 11 derivability checks and named 4 gates an engineer would have to invent.', 'source': 'prd-forge repository, the commit messages of 5ba59cb (version 0.1.0) and 5d4d936 (version 0.2.0)', 'as_of': datetime.date(2026, 10, 23)}
- A specification with no user journey gives a story-mapping workshop no backbone, so features and epics cannot be derived from it at all.
- Gaps surface as mid-iteration questions about what happens between two requirements that both looked complete.

WHY NOW

Teams are committing to outcomes per release, which needs a complete path through the product, not a complete list of features.

COST OF INACTION

The release line is drawn across a feature list rather than a journey, so a release ships features and no user can complete anything with them.

## P-003

STATEMENT

Nobody can tell a buildable specification from a readable one until the team has already started building from it.

WHO

Sponsor or approver signing the specification off

PERSONA

PER-003

EVIDENCE

- {'claim': 'At version 0.8.0, this PRD passed every schema and cross-reference check with 0 errors and 0 warnings. The adversarial grade found 23 blocking defects in 8 of 18 documents.', 'source': 'prd-forge repository, commit 6f77bc0 (version 0.8.0)', 'as_of': datetime.date(2026, 10, 26)}
- Approval is given on readability and completeness of headings, because those are the only properties an approver can assess unaided.
- The first real assessment of buildability happens when an engineer estimates, which is after approval and after the plan was made.

WHY NOW

Where the document was drafted with an assistant, readability is no longer any evidence at all of buildability.

COST OF INACTION

Plans and commitments are made on an approved document whose defects are found by the people asked to deliver against them.

## P-004

STATEMENT

An expert has nowhere to learn what a section of a PRD is for at the moment they are being asked to fill it in.

WHO

Domain expert specifying software for their own field

PERSONA

PER-001

EVIDENCE

- {'claim': 'Of 20 sections in two public PRD templates, all 20 say what to write, and 0 say what makes an entry good or how a reader checks it. The Aha! Agile PRD template has 6 sections, and the Product School PRD template has 14.', 'source': 'Aha! Agile PRD template, https://www.aha.io/roadmapping/guide/requirements-management/what-is-a-good-product-requirements-document-template, and Product School PRD template, https://productschool.com/blog/product-strategy/product-template-requirements-document-prd', 'as_of': datetime.date(2026, 10, 26)}
- The craft exists in books and in experienced product managers, neither of which is available at the moment the expert is typing.

WHY NOW

Organisations are asking domain experts to specify directly, without routing the work through a product manager first.

COST OF INACTION

Each expert reinvents the form badly, and the organisation's specifications cannot be compared to one another or improved as a body of work.

SECTION 2

# Target customer and market

DRAFT 55/100 NO BLOCKING DEFECTS

| ID | SEGMENT | BUYER | USER |
|---|---|---|---|
| S-001 | Consultancies specifying from client domain knowledge | Principal or practice lead | The consultant |
| S-002 | Regulated-industry organisations where practitioners specify | Head of product or delivery | The practitioner |

Product teams entering a domain they do not know were segment S-003 until 0.8.1. They already employ the product skill and want the domain, which this product does not supply, so section 6 now records them as a non-goal for release 1. The id S-003 is not reused.

### Why we win

The expert can see, before anyone starts building, which parts of their own specification a development team will not be able to build from, and what to answer to fix it. No competitor above tells them that.

**Honest weakness.** Against a contracted product manager we have no quality edge, only availability and cost. A sceptical buyer is right to test that.

### Segments

> **S-001 · Consultancies and advisory firms turning client domain knowledge into specifications**
>
> BUYER — Principal or practice lead
>
> USER — The consultant with the domain knowledge
>
> SIZE — `<count>` firms, per `<a named register, with its year>`, not yet measured; see Q-004.
>
> WILLINGNESS TO PAY — Already pays for the consultant's time; the comparison is billable hours saved, not a software budget.

> **S-002 · Regulated-industry organisations where practitioners specify their own internal software**
>
> BUYER — Head of product, delivery lead, or head of transformation
>
> USER — The practitioner, plus the delivery lead who reviews what they write
>
> SIZE — Not yet measured; see Q-004.
>
> WILLINGNESS TO PAY — Compared against contracting a product manager for the engagement.

### Competitors

> **A blank PRD template in the team's document tool**
>
> KIND — indirect
>
> OFFER — The section headings, immediately, free, in a tool everyone already has.
>
> OUR EDGE — After one engagement, each fact the expert could not supply is an open question with a named decider and a date, and the set holds no figure or citation that nobody can stand behind.

### A general-purpose assistant conversation with no structure

| | |
|---|---|
| KIND | indirect |
| OFFER | Fluent prose on any section, instantly, with no setup at all. |
| OUR EDGE | Fluency is the failure mode in P-003. After one engagement, the buyer holds a report, made before anyone has estimated, that names each release-line stage no requirement covers and quotes each requirement an engineer could not write a test from. |

### Contracting a product manager for the engagement

| | |
|---|---|
| KIND | indirect |
| OFFER | Genuine product judgement, and someone accountable for the outcome. |
| OUR EDGE | None on quality. The edge is availability and cost at the moment the expert is typing, and that the expert keeps the understanding afterwards rather than renting it. |

### The status quo - a smart practitioner, a document, and a willing engineer

| | |
|---|---|
| KIND | status-quo |
| OFFER | No tool, no cost, no adoption, and it does sometimes work. |
| OUR EDGE | At handover, an engineer who was not in the room receives a set that names each release-line stage no requirement covers and holds each fact the expert could not supply as an open question with a named decider and a date, so they can estimate without first interviewing the author. |

SECTION 3

# Business objectives and success metrics

DRAFT 75/100 NO BLOCKING DEFECTS

| ID | OBJECTIVE | METRIC | BASELINE | TARGET |
|---|---|---|---|---|
| O-001 | Buildable without re-derivation | Pass rate on first buildability review | unknown (Q-001) | 80% |
| O-002 | Days, not weeks, to an approved spec | Median working days, first question to approval | unknown (Q-002) | under 5 |
| O-003 | Requirements testable as written | Share of `must` requirements an outside engineer can test | unknown (Q-003) | 95% |

### Why there are no baselines

All three baselines are unknown, and each has a matching open question with a decider and a date. This is deliberate and it is the honest state: this product has not run against a real engagement yet. The targets above are therefore intentions, not commitments, and section 15 should record the moment each baseline arrives and each target is either confirmed or changed.

A reviewer should push on this. An objective with an invented baseline is worse than one with an admitted gap, because the invented one gets planned against.

### What would move these without this product

Worth asking, and recorded here because the answer is not nothing: a single experienced product manager embedded with the expert would move all three. The bet is on availability and cost, not on beating that person. See section 2.

### Objectives

#### O-001

OBJECTIVE: A specification handed to a development team can be built from without the team re-deriving it first.

PROBLEMS:
- P-001
- P-002
- P-003

METRIC: Share of PRDs that pass an engineering buildability review on first submission, where the review is the eleven derivability checks.

BASELINE: unknown - not measured; see Q-001. Do not set a target against this until Q-001 is answered.

BASELINE QUESTION: Q-001

TARGET: 80% of PRDs pass on first submission.

BY: 2027-09-30

OWNER: Richard Hightower

### O-002

| | |
|---|---|
| OBJECTIVE | A domain expert reaches an approved specification in days rather than the weeks a product manager's discovery cycle takes. |
| PROBLEMS | • P-004 |
| METRIC | Median elapsed working days from the first interview question to an approval recorded in section 15. |
| BASELINE | unknown - not measured; see Q-002. |
| BASELINE QUESTION | Q-002 |
| TARGET | Under 5 working days for a single-journey product. |
| BY | 2027-09-30 |
| OWNER | Richard Hightower |

### O-003

| | |
|---|---|
| OBJECTIVE | Requirements are testable as written, so a team can close them without asking the author what they meant. |
| PROBLEMS | • P-001 |
| METRIC | Share of functional requirements whose acceptance criteria an engineer who was not in the interview judges sufficient to write a test. |
| BASELINE | unknown - not measured; see Q-003. |
| BASELINE QUESTION | Q-003 |
| TARGET | 95% of `must` requirements. |
| BY | 2027-09-30 |
| OWNER | Richard Hightower |

SECTION 4

# Users and personas

DRAFT 68/100 NO BLOCKING DEFECTS

| ID | ROLE | FREQUENCY | MAY APPROVE |
|---|---|---|---|
| PER-001 | Domain expert specifying their own field | Weeks per product, then rarely | Scope, terms, drafts, section states, facts, and offering the set; not the PRD itself |
| PER-002 | Engineering lead building from it | A handful of intensive reads | May reject as not buildable |
| PER-003 | Sponsor or approver | Once per version | The PRD for a release |
| PER-004 | Delivery lead running story mapping | Once per release | Backlog shape and order |

### PER-001 is the primary persona

Everything is designed around someone with deep domain knowledge, no product training, and a burst of intense use followed by months away. Two consequences run through the rest of this PRD: nothing may assume prior knowledge of product vocabulary, and the work must survive being put down for a week and picked back up. See FR-015 and NFR-004.

### Who is deliberately not a persona

The assistant that delivers the teaching, the interview and the grading is not a user of this product. It is how the product reaches PER-001. Listing it here would put the delivery mechanism into the PRD, which the set keeps out of the requirements: section 6 makes a separate application a non-goal, and section 11 records the session as dependency D-001. The same goes for the grading reviewers: a reviewer that grades a section is a mechanism, and what matters to this PRD is that PER-001 sees a grade they did not produce themselves.

> **PER-001 · Domain expert specifying software for their own field**
>
> SHORT NAME: DomainExpert
>
> PRIMARY: yes
>
> GOALS:
> - Get what is in my head into a form the development team can build from
> - Stop being asked the same question by four different engineers
> - Not be the reason the project is late
>
> PAINS:
> - I know what it must do, and I cannot write it the way they need it
> - I am told my requirements are not testable and I do not know what that means
> - I reach for the names of software I have used, because that is the only requirements language I have
>
> FREQUENCY: Intensively for one to three weeks per product, then almost never until the next one.

| | |
|---|---|
| AUTHORITY | Decides what the product must do and what is out of scope; decides the domain's own terms; decides when a section's interview is finished and when a section may be left at draft; accepts or corrects a drafted section; sets a section to final; accepts or refutes a grade; decides the set is ready to be offered for approval; and is the only one who may state a domain fact as true. May not approve the PRD for a release on their own where a sponsor exists. |
| PROFICIENCY | Deep in their domain, often decades. No product management training. Assume no familiarity with MoSCoW, Given/When/Then, or the word "persona". |

### PER-002 · Engineering lead who must build from the PRD

| | |
|---|---|
| SHORT NAME | EngineerLead |
| GOALS | • Estimate without first interviewing the author<br>• Know which decisions are mine and which have been made for me<br>• Find out now, not in iteration three, what the document does not say |
| PAINS | • Specifications that name a technology and so pre-empt my decision badly<br>• Requirements I cannot write a test for<br>• Gaps between two requirements that both looked complete |
| FREQUENCY | Reads a given PRD a handful of times, intensively, at estimation and at each release planning. |
| AUTHORITY | Decides how it is built; may reject a PRD as not buildable. |
| PROFICIENCY | Expert in delivery, usually knows nothing about the expert's domain. |

### PER-003 · Sponsor or approver

| | |
|---|---|
| SHORT NAME | Sponsor |
| GOALS | • Be confident enough to commit budget and a date<br>• Know what is still undecided, and who is deciding it |
| PAINS | • I can only judge whether it reads well<br>• I find out it was not buildable after I have committed to the plan |
| FREQUENCY | Once per version, plus conditions. |
| AUTHORITY | Approves the PRD for a release; may approve with conditions; may reject. |
| PROFICIENCY | Senior in the business, not in product or engineering. |

### PER-004 · Delivery lead running the story-mapping workshop

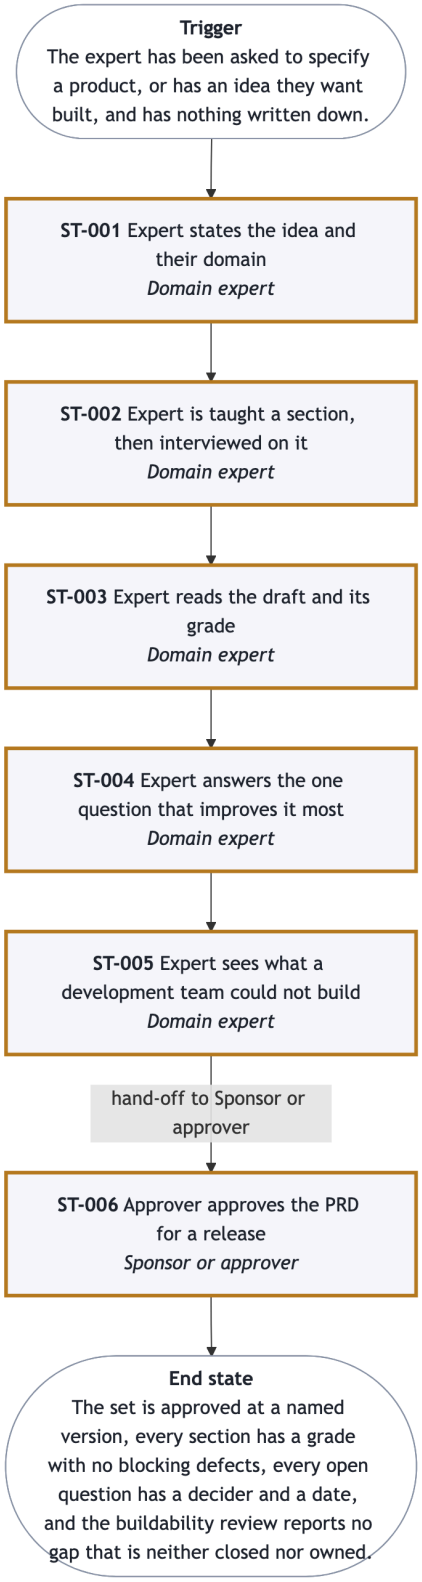
| | |
|---|---|
| SHORT NAME | DeliveryLead |
| GOALS | • Walk into the workshop with a backbone already agreed<br>• Trace every story back to something the expert actually said |
| PAINS | • No journey means no backbone, so the workshop invents one in the room<br>• Stories that trace to nothing, so nobody can say why they exist |
| FREQUENCY | Once per release, plus refinement. |
| AUTHORITY | Decides the backlog's shape and order, and decides whether a backlog can be mapped from a set at all; does not decide scope. |
| PROFICIENCY | Expert in agile delivery; depends entirely on the PRD for the domain. |

SECTION 5

# User journeys

DRAFT 79/100 NO BLOCKING DEFECTS

## J-001 - A domain expert reaches an approved PRD whose gaps are known (golden thread)




J-001 journey diagram

Persona PER-001. Trigger: they have been asked to specify a product and have nothing written. Stages ST-001 to ST-006.

| STAGE | ACTOR | DECISION |
|---|---|---|
| ST-001 States the idea and the domain | Domain expert | What this is, and the domain's own words |

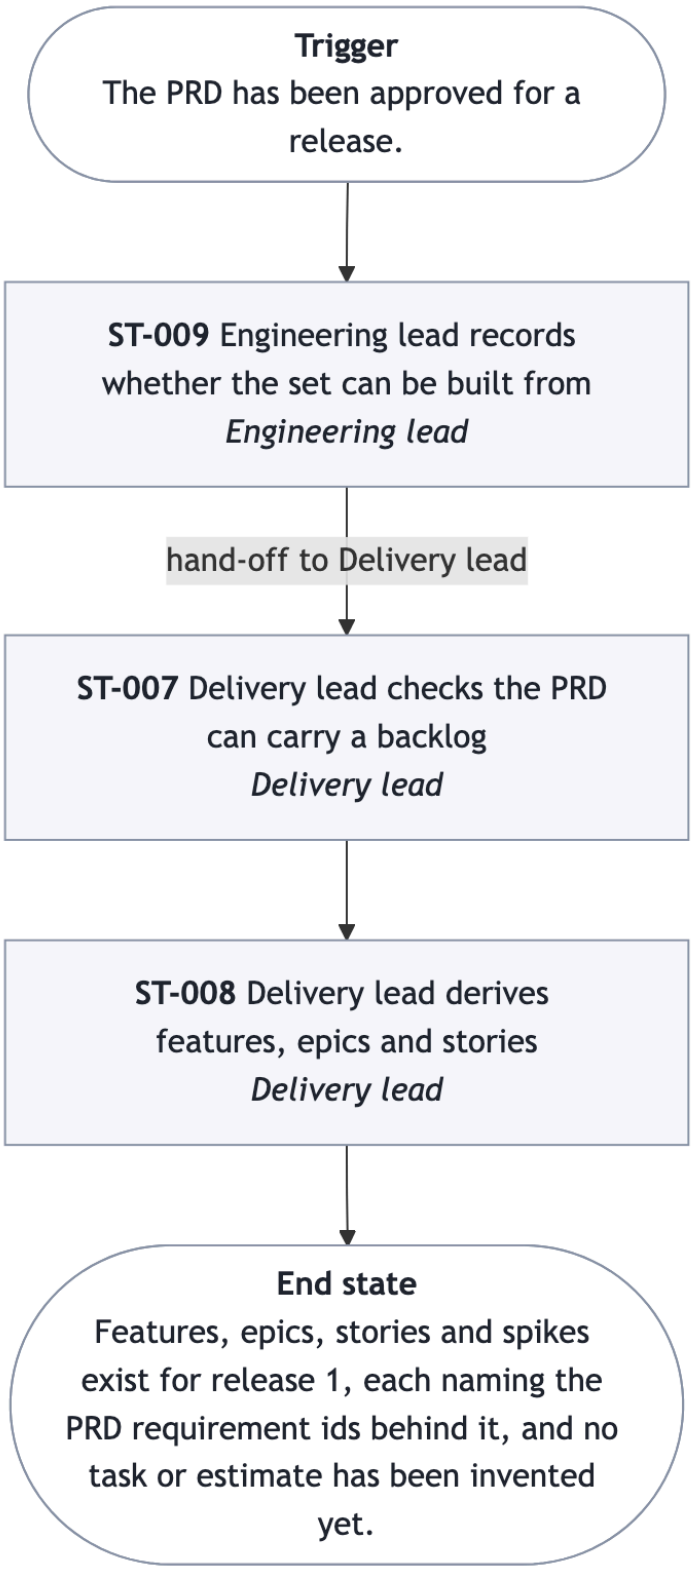
| STAGE | ACTOR | DECISION |
|---|---|---|
| ST-002 Taught a section, then interviewed | Domain expert | Have the questions reached the end of what they know |
| ST-003 Reads the draft and its grade | Domain expert | Accept, correct, or dispute the grade |
| ST-004 Answers the one question that matters most | Domain expert | Is this section done enough to leave |
| ST-005 Sees what a team could not build | Domain expert | Is it ready for someone who will commit |
| ST-006 Approver approves for a release | Sponsor or approver | Approve, approve with conditions, or reject |

ST-002 to ST-004 are a loop, run once per section. The loop is the whole product: the expert is taught, asked, shown a draft, shown a grade, and asked one more thing.

End state: approved at a named version, every section graded with nothing blocking, every open question owned and dated, and no buildability gap left that is neither closed nor owned.

The goal deliberately does not claim an engineering lead's confirmation. No stage in ST-001 to ST-006 puts one in the actor seat, so claiming it would promise something the release line cannot deliver. A lead's own recorded confirmation is OUT-004, deferred to release 2 with J-002, where it is stage ST-009.

**J-002 - A delivery lead derives a backlog (not the golden thread)**




J-002 journey diagram

Persona PER-004, stages ST-009, ST-007, and ST-008. In ST-009 the engineering lead, PER-002, reads the report of the buildability review and records their own verdict under TR-006, before the delivery lead takes over. ST-009 was added at 0.8.1, after ST-008, and ids are never renumbered. Release 1 does not have to complete this journey (see section 6), but the PRD must carry what it needs, which is what ST-007 checks.

### What these journeys deliberately do not say

No stage says where the teaching text comes from, how a grade is computed, or what reviews the set in ST-005. Those are mechanisms. Every stage above names a person, what they do and see, and what they decide.

#### J-001

| | |
|---|---|
| GOAL | An approved PRD whose buildability gaps have been found and either closed or given an owner and a date, before anyone estimates it. |
| PERSONA | PER-001 |
| TRIGGER | The expert has been asked to specify a product, or has an idea they want built, and has nothing written down. |
| GOLDEN THREAD | yes |

STAGES

### Expert states the idea and their domain

| | |
|---|---|
| ID | ST-001 |
| ACTOR | Domain expert |
| ACCEPTANCE | • Given an expert with an idea and nothing written down, when they state the idea, the domain, and three words an outsider would misread, then they see a document set with every section listed and marked as not yet started. |
| STEPS | • Says what they want built, in their own words, however roughly<br>• Is asked what the business domain is and what the unit of work is called - an engagement, a load, a claim, a case<br>• Names three words their team uses that an outsider would misread<br>• Sees that a set of documents now exists, every section listed and marked as not yet started |
| INPUTS | • The idea, the domain, and the vocabulary |
| OUTPUTS | • A document set with every section visible and honestly marked empty |
| DECISION | What this product is, and what the domain's own words are |

### Expert is taught a section, then interviewed on it

| | |
|---|---|
| ID | ST-002 |
| ACTOR | Domain expert |
| ACCEPTANCE | • Given a section not yet interviewed, when the expert reads what it is for and answers its questions one or two at a time, then the section has enough material to be drafted, or the expert has said they have nothing more to add. |
| STEPS | • Reads what the section is, why it matters, and what belongs in it, in under ten lines, with an example from their own domain<br>• Is asked one or two questions at a time, in their own vocabulary<br>• Is asked follow-ups where an answer was thin<br>• Says when they have nothing more to add |
| INPUTS | • Their knowledge of the domain, in answer to each question |
| OUTPUTS | • Enough material for the section to be drafted |
| DECISION | Whether the questions have reached the end of what they know |

### Expert reads the draft and its grade

| | |
|---|---|
| ID | ST-003 |
| ACTOR | Domain expert |
| ACCEPTANCE | • Given a drafted and graded section, when the expert opens it, then they read it as prose and see its score, each blocking defect, and its weakest part. |
| STEPS | • Reads the drafted section as prose, not as a form<br>• Sees a score, how many blocking defects there are, and what each one is<br>• Sees which part of the section is weakest, and why<br>• Decides whether the grade is fair |

| | |
|---|---|
| INPUTS | • Nothing; this is where they read |
| OUTPUTS | • A judgement on whether the draft says what they meant |
| DECISION | Accept the draft, correct it, or dispute the grade |

**Expert answers the one question that improves it most**

| | |
|---|---|
| ID | ST-004 |
| ACTOR | Domain expert |
| ACCEPTANCE | • Given one question about the weakest part of a section, when the expert answers it, then they see the section re-graded against the same rubric. |
| STEPS | • Reads a single question, not a list of twelve<br>• Answers it, or says why the premise is wrong<br>• Sees the section re-graded against the same rubric<br>• Moves to the next section, or goes back to an earlier one when this section showed it was incomplete |
| INPUTS | • One answer |
| OUTPUTS | • A section at a grade the expert accepts, and a set that is one section further on |
| DECISION | Whether this section is done enough to leave |
| HANDOFF | Back into ST-002 for the next section, until every section has been through this loop. |

**Expert sees what a development team could not build**

| | |
|---|---|
| ID | ST-005 |
| ACTOR | Domain expert |
| ACCEPTANCE | • Given a set with every section written at least once, when the expert asks for a buildability review, then they see what an engineer could not build and which release-line stages have no requirement behind them. |
| STEPS | • Asks for the whole set to be reviewed as an engineer would read it<br>• Sees what an engineer could not build, and what they would have to guess<br>• Sees which release-line stages have no requirement behind them<br>• Fixes the holes, or records them as open questions with a decider |
| INPUTS | • Every section written at least once, so there is something to review |
| OUTPUTS | • A list of what is not buildable, and a set with those holes closed |
| DECISION | Is this ready to be read by someone who will commit to it |
| HANDOFF | Sponsor or approver |

**Approver approves the PRD for a release**

| | |
|---|---|
| ID | ST-006 |
| ACTOR | Sponsor or approver |
| ACCEPTANCE | • Given a set offered for approval, when the approver approves, approves with conditions, or rejects it, then they see the decision recorded against the set's version. |

| | |
|---|---|
| STEPS | • Reads the summary, the scope, and what is still open<br>• Sees each section's grade and whether anything is still blocking<br>• Approves, approves with conditions naming what the condition is, or rejects<br>• Sees the approval recorded against this version |
| INPUTS | • Their decision and any conditions |
| OUTPUTS | • An approved PRD at a named version, with the approval recorded |
| DECISION | Approve, approve with conditions, or reject |
| HANDOFF | Delivery lead, for the story-mapping workshop |

**PAIN POINTS**

- I do not know what the section is for, so I fill the heading with adjectives
- Nobody tells me my requirements are untestable until a team tries to build them
- I reach for the names of software I have used
- I put this down for a week and cannot remember where I was
- I cannot tell whether what I have written is good

**NEEDS**

- Teach me the section before you ask me about it, in my own domain
- Tell me where it is weak, against a stated standard, and ask me one thing
- Convert my technology talk into the need underneath it, without arguing
- Let me stop and resume without starting over
- Show me, before anyone builds, what an engineer cannot build

**END STATE**

The set is approved at a named version, every section has a grade with no blocking defects, every open question has a decider and a date, and the buildability review reports no gap that is neither closed nor owned.

### J-002 · A backlog whose every story traces to a requirement the expert actually stated.

| | |
|---|---|
| PERSONA | PER-004 |
| TRIGGER | The PRD has been approved for a release. |

STAGES

### Engineering lead records whether the set can be built from

| | |
|---|---|
| ID | ST-009 |
| ACTOR | Engineering lead |
| ACCEPTANCE | • Given an approved set and the report of its buildability review, when the engineering lead records a verdict that differs from the report's, then their judgement is recorded as the verdict and the report is kept alongside it.<br>• Given a gap the engineering lead records, when it is stated, then the document is quoted for it. |
| STEPS | • Reads the report of the buildability review, its verdict, and the gaps it quotes<br>• Checks each reported gap against the document it quotes<br>• Records their own verdict on whether the set can be built from |
| INPUTS | • An approved set and the report of its buildability review |
| OUTPUTS | • The engineering lead's own recorded verdict, with the report kept alongside it |
| DECISION | Can a team build from this set |
| HANDOFF | The delivery lead, to check the PRD can carry a backlog |

### Delivery lead checks the PRD can carry a backlog

| | |
|---|---|
| ID | ST-007 |
| ACTOR | Delivery lead |
| ACCEPTANCE | • Given an approved set and the release it is mapped for, when the delivery lead checks it, then they see whether a golden-thread journey exists and what would block deriving a backlog. |
| STEPS | • Sees whether a journey exists and which one is the golden thread<br>• Sees which release-line stages have requirements behind them<br>• Sees what would block deriving a backlog, before the workshop |
| INPUTS | • An approved set, and the release it is being mapped for |
| OUTPUTS | • A go or no-go on holding the story-mapping workshop |
| DECISION | Is there a backbone to map against |

### Delivery lead derives features, epics and stories

| | |
|---|---|
| ID | ST-008 |
| ACTOR | Delivery lead |
| ACCEPTANCE | • Given an approved set, when the delivery lead derives a backlog from it, then they see one feature per journey stage, epics scoped by the release line, and a story per requirement, each naming the id it came from. |
| STEPS | • Sees one feature per journey stage, with the requirement ids it carries<br>• Sees epics scoped by the release line<br>• Sees a story per requirement, in the persona's own words, with the requirement's criteria as its acceptance criteria<br>• Sees spikes for the open questions that block the release line |

| | |
|---|---|
| INPUTS | • The approved set, and the team's sizing of the stories |
| OUTPUTS | • A backlog whose every level names what it came from |
| DECISION | Which epics come first, and why |
| HANDOFF | The engineering team, for iteration planning |

**PAIN POINTS**

- The workshop invents a backbone in the room because the PRD has none
- Stories trace to nothing, so nobody can say why they exist
- Non-functional requirements and constraints are forgotten until late

**NEEDS**

- A journey to map against, agreed before the workshop
- Every story carrying the requirement id it came from
- Constraints arriving as definition-of-done items, not as surprises

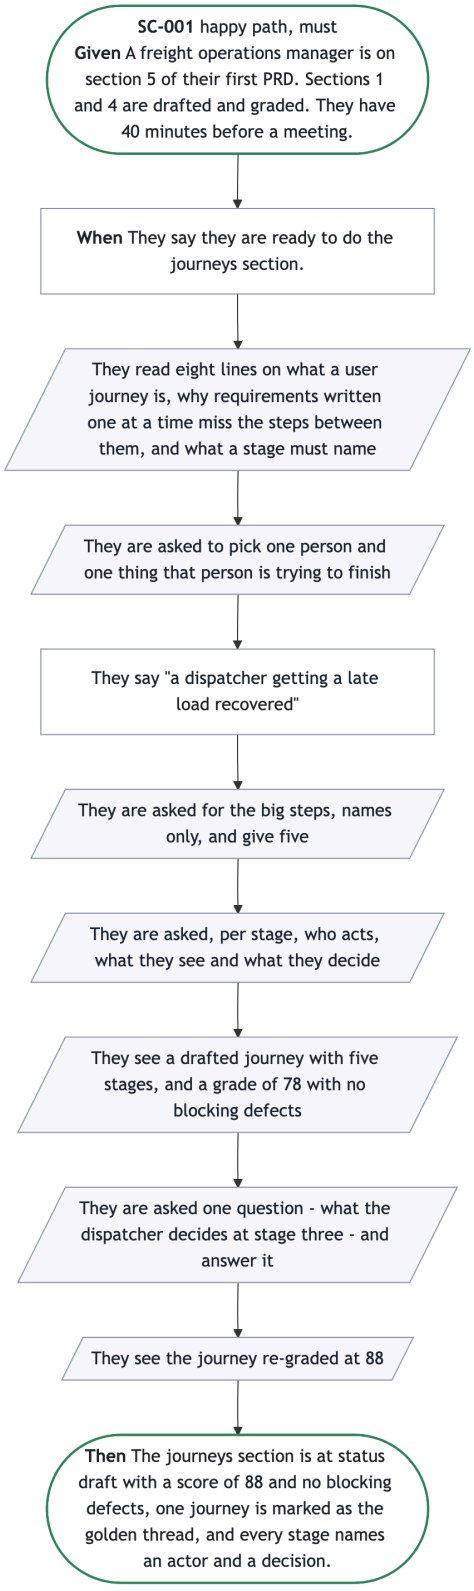
**END STATE**

Features, epics, stories and spikes exist for release 1, each naming the PRD requirement ids behind it, and no task or estimate has been invented yet.

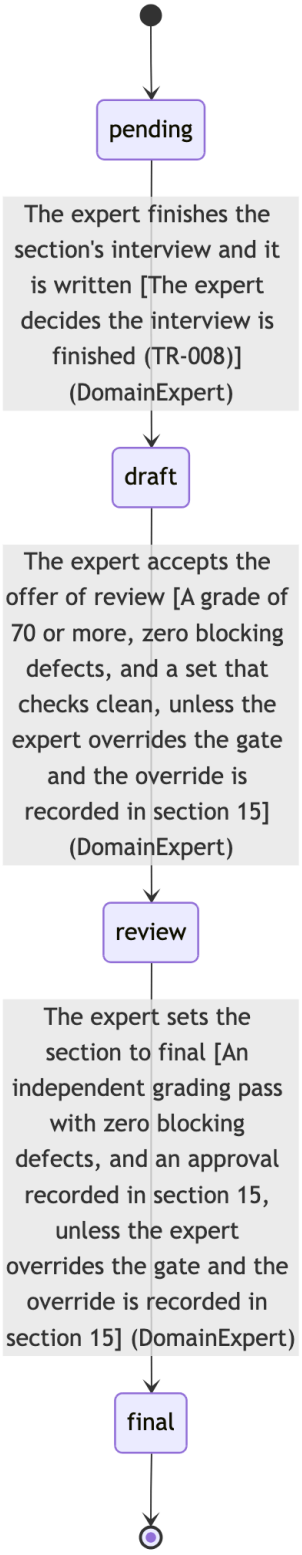
SECTION 5

# Scenarios

DRAFT 78/100 NO BLOCKING DEFECTS

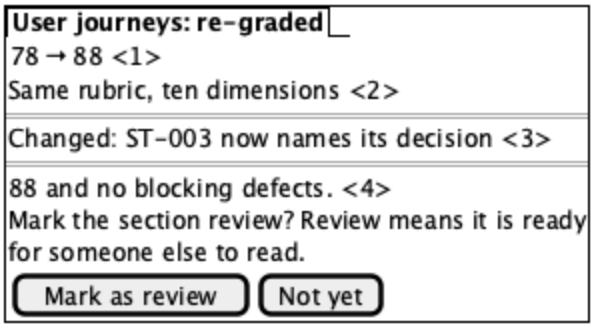



SC-001 scenario

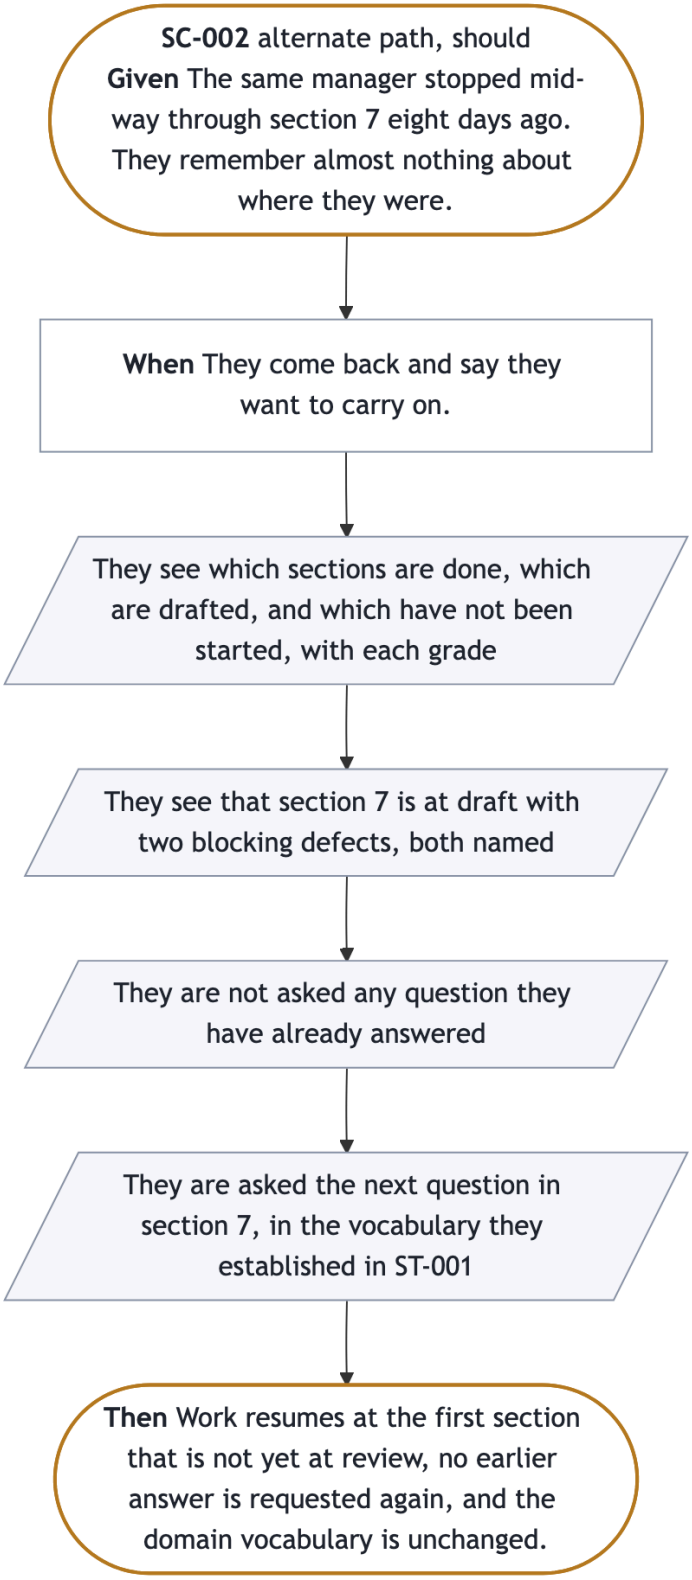



LC-001 lifecycle of a section of the PRD

## WF-008 A section re-graded after one answer

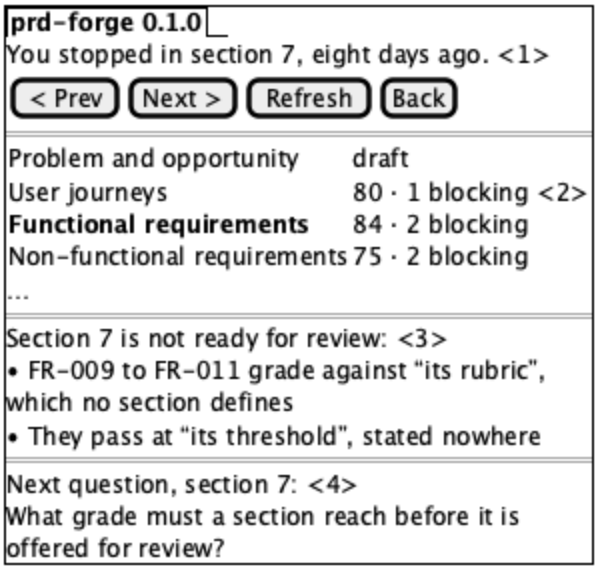



WF-008 drawing

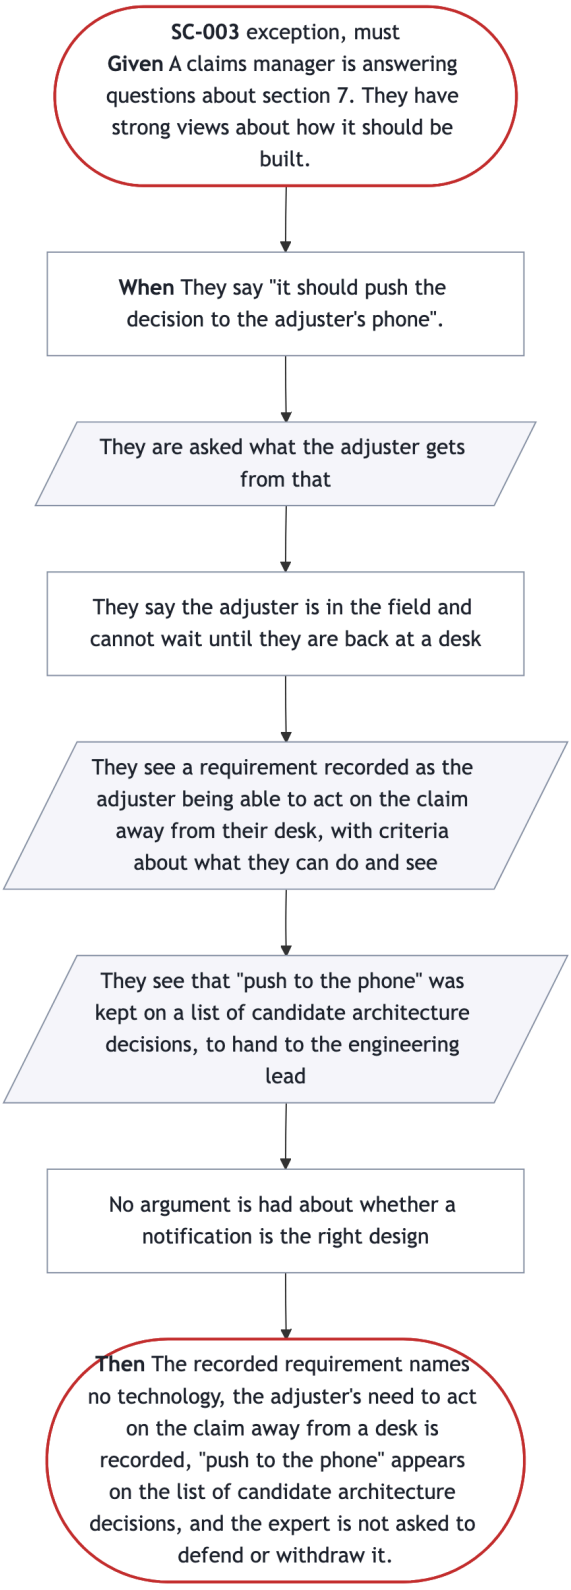



SC-002 scenario

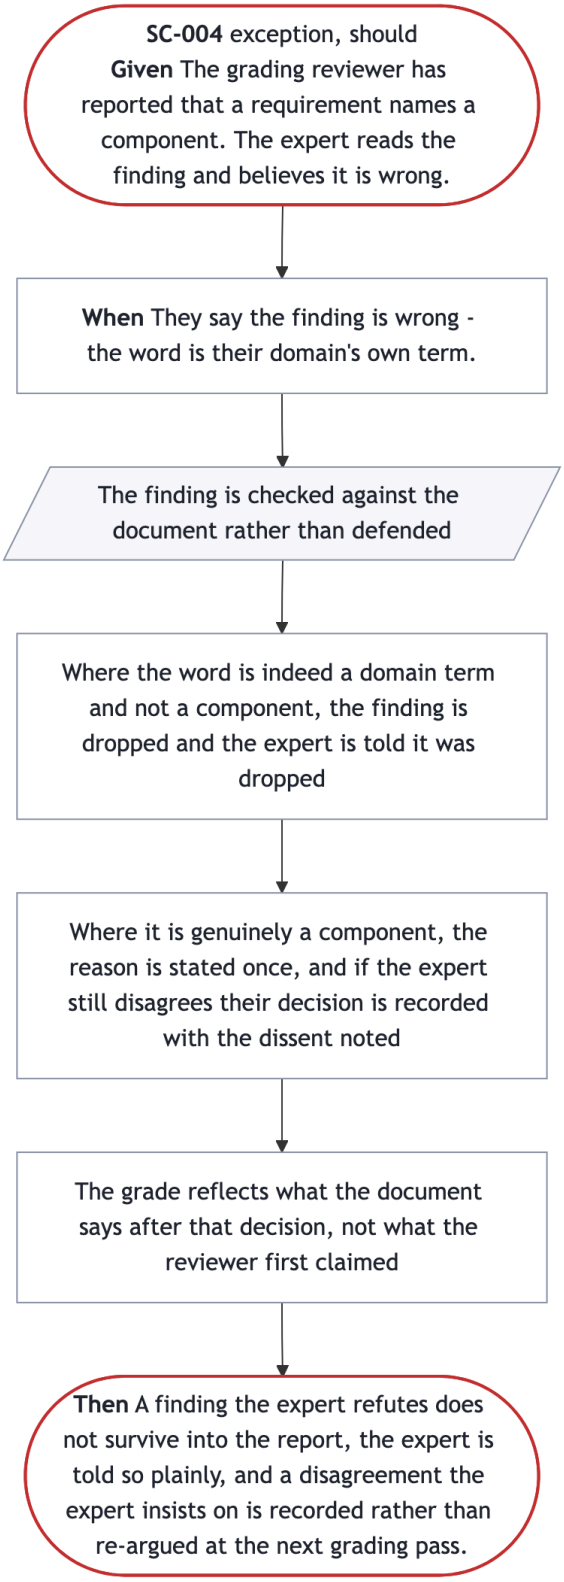
## WF-010 Resuming a set after a week away



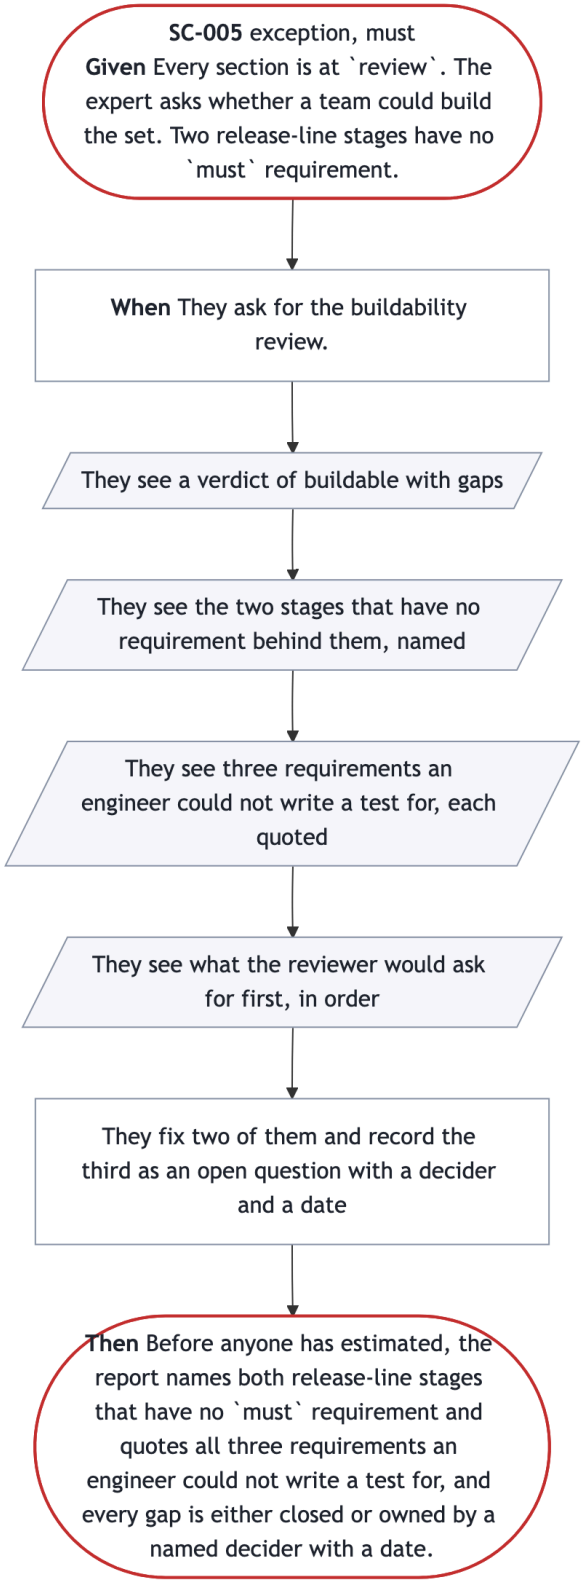

WF-010 drawing




SC-003 scenario




SC-004 scenario

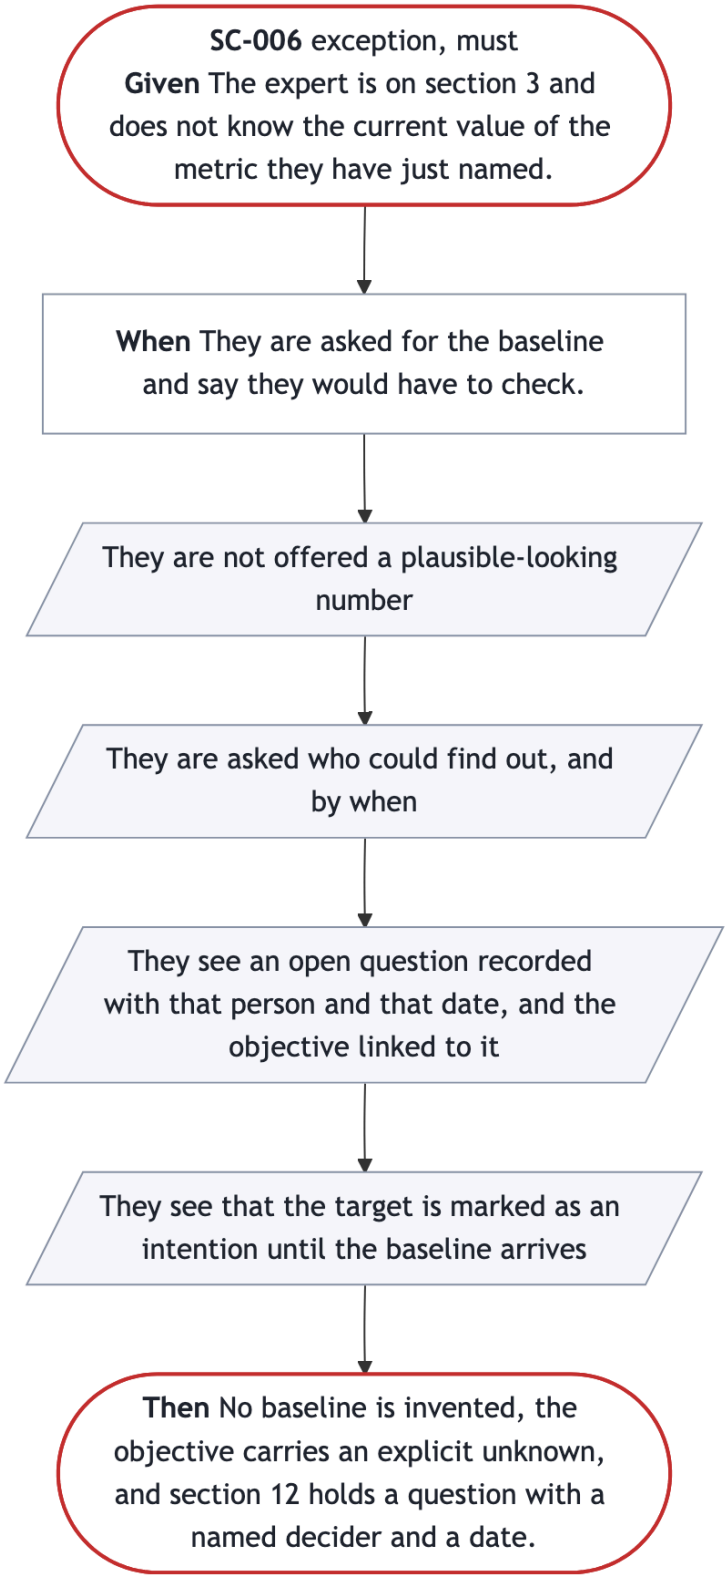



SC-005 scenario

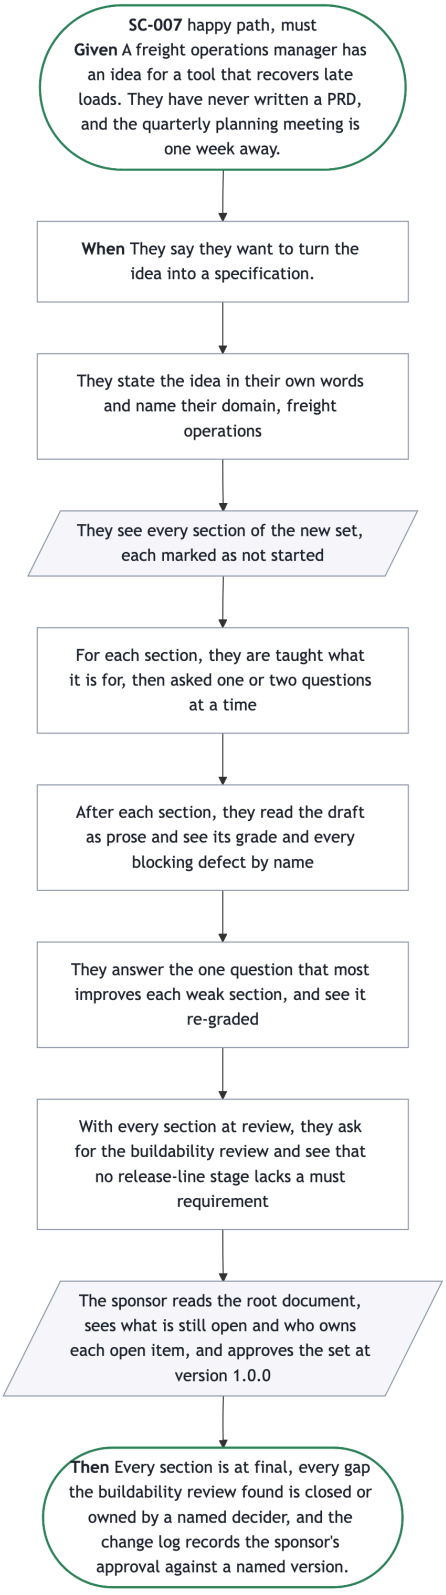
## WF-009 A buildability review that finds gaps

| What a team could not build from this set |
| --- |
| Buildable with gaps. <1> |
| ✗ ST-004 no must requirement <2><br>✗ ST-005 no must requirement<br>✓ ST-001, ST-002, ST-003, ST-006 |
| Each gap becomes a requirement, or an open question with a decider and a date. <3> |

WF-009 drawing

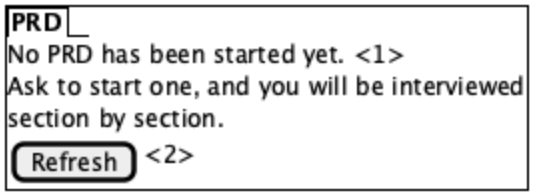



SC-006 scenario



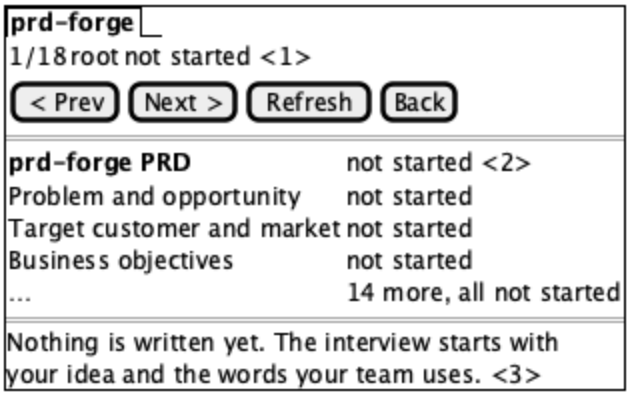

SC-007 scenario

### WF-007 Before any set exists



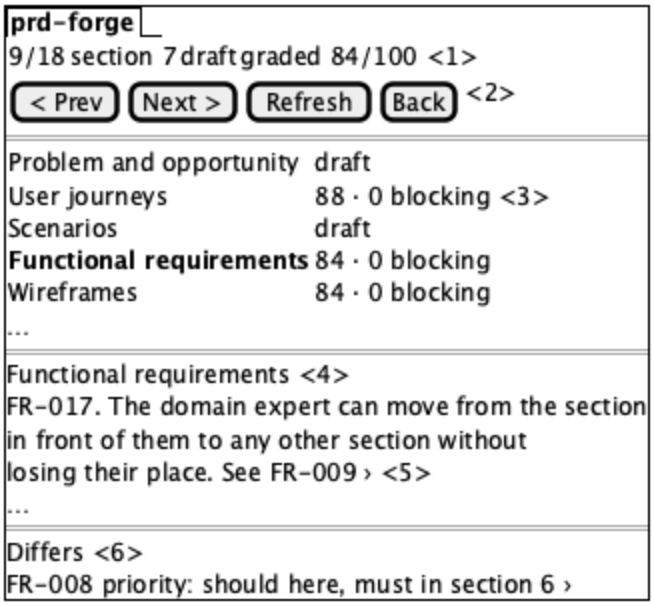

WF-007 drawing

### WF-002 A new set, before the first interview



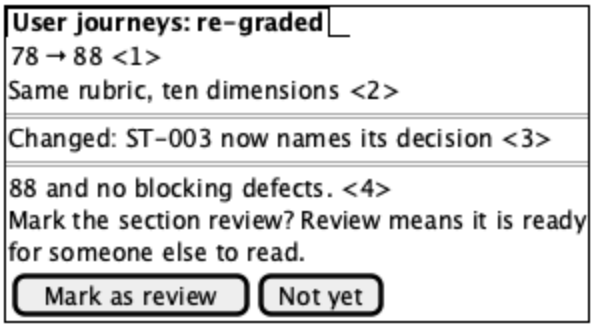

WF-002 drawing

### WF-001 The set with one section open



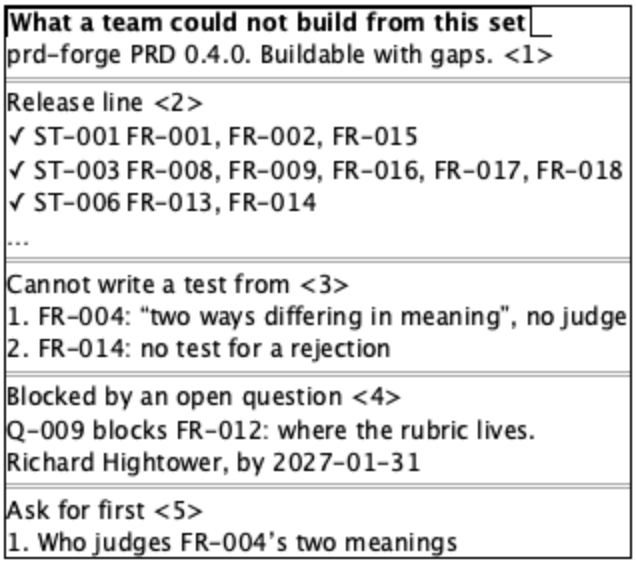

WF-001 drawing

## WF-008 A section re-graded after one answer

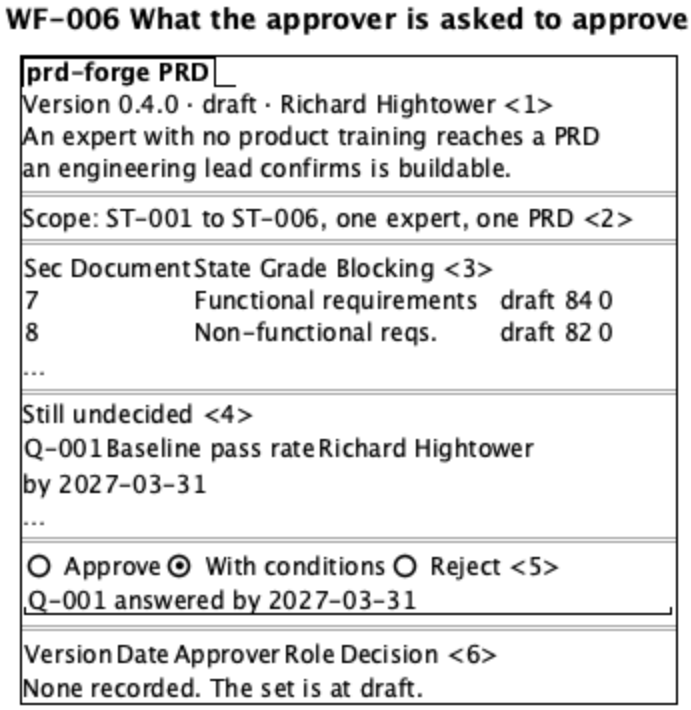



WF-008 drawing

## WF-005 What a team could not build from this set




WF-005 drawing




WF-006 drawing

| ID | STAGE | VARIANT | PRIORITY | WHAT IT COVERS |
|---|---|---|---|---|
| SC-001 | ST-002 | happy | must | A section taught, interviewed, drafted, graded, improved |
| SC-002 | ST-004 | alternate | should | Resuming after eight days away |
| SC-003 | ST-002 | exception | must | The expert states a technology |
| SC-004 | ST-003 | exception | should | The expert refutes a grading finding |
| SC-005 | ST-005 | exception | must | Buildability review finds holes |
| SC-006 | ST-002 | exception | must | The expert has no baseline for a metric |
| SC-007 | ST-001 to ST-006 | happy | must | One idea, from the first sentence to an approved set |

SC-007 is the end-to-end acceptance test for release 1: it follows the golden thread through every release-line stage. SC-001 is the same path inside one section. SC-005 is the expert's own run of the buildability review, which is why its persona is PER-001 and not the engineering lead: no engineering lead is in release 1 at all, per OUT-004.

The priorities follow the requirements. SC-002 and SC-004 are `should`, because FR-015 and FR-016, which they exercise, are `should`.

The four exceptions are the ones that decide whether this product is trustworthy. SC-003 and SC-006 are where a specification quietly acquires a technology choice or an invented fact. SC-004 is where the grading either earns the expert's trust or loses it: a reviewer that cannot be refuted is a reviewer nobody believes. SC-005 is the moment the product is supposed to pay for itself.

## Scenarios

### SC-001

| | |
|---|---|
| CONTEXT | A freight operations manager is on section 5 of their first PRD. Sections 1 and 4 are drafted and graded. They have 40 minutes before a meeting. |
| JOURNEY | J-001 |
| STAGE | ST-002 |
| STAGES | • ST-002<br>• ST-003<br>• ST-004 |
| PERSONA | PER-001 |
| VARIANT | happy |
| PRIORITY | must |
| SHOWN IN | screen |
| LIFECYCLES | • LC-001 |
| WIREFRAMES | • WF-008 |
| TRIGGER | They say they are ready to do the journeys section. |
| COURSE | • They read eight lines on what a user journey is, why requirements written one at a time miss the steps between them, and what a stage must name<br>• They are asked to pick one person and one thing that person is trying to finish<br>• They say "a dispatcher getting a late load recovered"<br>• They are asked for the big steps, names only, and give five<br>• They are asked, per stage, who acts, what they see and what they decide<br>• They see a drafted journey with five stages, and a grade of 78 with no blocking defects<br>• They are asked one question - what the dispatcher decides at stage three - and answer it<br>• They see the journey re-graded at 88 |
| EXPECTED OUTCOME | The journeys section is at status draft with a score of 88 and no blocking defects, one journey is marked as the golden thread, and every stage names an actor and a decision. |

### SC-002

**CONTEXT**
The same manager stopped mid-way through section 7 eight days ago. They remember almost nothing about where they were.

**JOURNEY**
J-001

**STAGE**
ST-004

**PERSONA**
PER-001

**VARIANT**
alternate

**PRIORITY**
should

**SHOWN IN**
screen

**LIFECYCLES**
- LC-001

**WIREFRAMES**
- WF-010

**TRIGGER**
They come back and say they want to carry on.

**COURSE**
- They see which sections are done, which are drafted, and which have not been started, with each grade
- They see that section 7 is at draft with two blocking defects, both named
- They are not asked any question they have already answered
- They are asked the next question in section 7, in the vocabulary they established in ST-001

**EXPECTED OUTCOME**
Work resumes at the first section that is not yet at review, no earlier answer is requested again, and the domain vocabulary is unchanged.

### SC-003

**CONTEXT**
A claims manager is answering questions about section 7. They have strong views about how it should be built.

**JOURNEY**
J-001

**STAGE**
ST-002

**PERSONA**
PER-001

**VARIANT**
exception

**PRIORITY**
must

**SHOWN IN**
conversation

**TRIGGER**
They say "it should push the decision to the adjuster's phone".

**COURSE**
- They are asked what the adjuster gets from that
- They say the adjuster is in the field and cannot wait until they are back at a desk
- They see a requirement recorded as the adjuster being able to act on the claim away from their desk, with criteria about what they can do and see
- They see that "push to the phone" was kept on a list of candidate architecture decisions, to hand to the engineering lead
- No argument is had about whether a notification is the right design

**EXPECTED OUTCOME**
The recorded requirement names no technology, the adjuster's need to act on the claim away from a desk is recorded, "push to the phone" appears on the list of candidate architecture decisions, and the expert is not asked to defend or withdraw it.

### SC-004

**CONTEXT**
The grading reviewer has reported that a requirement names a component. The expert reads the finding and believes it is wrong.

JOURNEY J-001

STAGE ST-003

PERSONA PER-001

VARIANT exception

PRIORITY should

SHOWN IN conversation

TRIGGER They say the finding is wrong - the word is their domain's own term.

COURSE
- The finding is checked against the document rather than defended
- Where the word is indeed a domain term and not a component, the finding is dropped and the expert is told it was dropped
- Where it is genuinely a component, the reason is stated once, and if the expert still disagrees their decision is recorded with the dissent noted
- The grade reflects what the document says after that decision, not what the reviewer first claimed

EXPECTED OUTCOME A finding the expert refutes does not survive into the report, the expert is told so plainly, and a disagreement the expert insists on is recorded rather than re-argued at the next grading pass.

### SC-005

CONTEXT Every section is at `review`. The expert asks whether a team could build the set. Two release-line stages have no `must` requirement.

JOURNEY J-001

STAGE ST-005

PERSONA PER-001

VARIANT exception

PRIORITY must

SHOWN IN report

WIREFRAMES
- WF-009

TRIGGER They ask for the buildability review.

COURSE
- They see a verdict of buildable with gaps
- They see the two stages that have no requirement behind them, named
- They see three requirements an engineer could not write a test for, each quoted
- They see what the reviewer would ask for first, in order
- They fix two of them and record the third as an open question with a decider and a date

EXPECTED OUTCOME Before anyone has estimated, the report names both release-line stages that have no `must` requirement and quotes all three requirements an engineer could not write a test for, and every gap is either closed or owned by a named decider with a date.

## SC-006

**CONTEXT**
The expert is on section 3 and does not know the current value of the metric they have just named.

**JOURNEY**
J-001

**STAGE**
ST-002

**PERSONA**
PER-001

**VARIANT**
exception

**PRIORITY**
must

**SHOWN IN**
conversation

**TRIGGER**
They are asked for the baseline and say they would have to check.

**COURSE**
- They are not offered a plausible-looking number
- They are asked who could find out, and by when
- They see an open question recorded with that person and that date, and the objective linked to it
- They see that the target is marked as an intention until the baseline arrives

**EXPECTED OUTCOME**
No baseline is invented, the objective carries an explicit unknown, and section 12 holds a question with a named decider and a date.

## SC-007

**CONTEXT**
A freight operations manager has an idea for a tool that recovers late loads. They have never written a PRD, and the quarterly planning meeting is one week away.

**JOURNEY**
J-001

**STAGE**
ST-001

**STAGES**
- ST-001
- ST-002
- ST-003
- ST-004
- ST-005
- ST-006

**PERSONA**
PER-001

**VARIANT**
happy

**PRIORITY**
must

**SHOWN IN**
screen

**LIFECYCLES**
- LC-001

**WIREFRAMES**
- WF-007
- WF-002
- WF-001
- WF-008
- WF-005
- WF-006

**TRIGGER**
They say they want to turn the idea into a specification.

COURSE

- They state the idea in their own words and name their domain, freight operations
- They see every section of the new set, each marked as not started
- For each section, they are taught what it is for, then asked one or two questions at a time
- After each section, they read the draft as prose and see its grade and every blocking defect by name
- They answer the one question that most improves each weak section, and see it re-graded
- With every section at review, they ask for the buildability review and see that no release-line stage lacks a must requirement
- The sponsor reads the root document, sees what is still open and who owns each open item, and approves the set at version 1.0.0

EXPECTED OUTCOME

Every section is at final, every gap the buildability review found is closed or owned by a named decider, and the change log records the sponsor's approval against a named version.

## Lifecycles

### LC-001 · a section of the PRD

STATES

**pending**

MEANING — Listed in the set and not yet interviewed.

**draft**

MEANING — Interviewed and written, and graded or waiting for a grade.

**review**

MEANING — Ready for someone other than the expert to read.

**final**

MEANING — Approved, with the approval recorded in section 15.

INITIAL

pending

FINAL

- final

TRANSITIONS

| | |
|---|---|
| FROM | pending |
| TO | draft |
| WHEN | The expert finishes the section's interview and it is written |
| RULE | The expert decides the interview is finished (TR-008) |
| BY | PER-001 |

| | |
|---|---|
| FROM | draft |
| TO | review |
| WHEN | The expert accepts the offer of review |
| RULE | A grade of 70 or more, zero blocking defects, and a set that checks clean, unless the expert overrides the gate and the override is recorded in section 15 |
| BY | PER-001 |

| | |
|---|---|
| FROM | review |
| TO | final |
| WHEN | The expert sets the section to final |
| RULE | An independent grading pass with zero blocking defects, and an approval recorded in section 15, unless the expert overrides the gate and the override is recorded in section 15 |
| BY | PER-001 |

SECTION 6

# Scope

DRAFT 79/100 NO BLOCKING DEFECTS

## Release line

ST-001 through ST-006 for one expert, one PRD, one golden-thread journey.

| ID | IN SCOPE | STAGE | REQUIREMENTS |
|---|---|---|---|
| IN-001 | Idea, domain, vocabulary, a visible empty set, and resumption | ST-001 | FR-001, FR-002, FR-015 |
| IN-002 | Teach then interview, in their vocabulary | ST-002 | FR-003, FR-004, FR-005, FR-019 |
| IN-003 | No invented facts; convert technology talk | ST-002 | FR-006, FR-007 |
| IN-004 | Draft as prose, with grade, blocking defects, and refutable findings, reached from any section | ST-003 | FR-008, FR-009, FR-016, FR-017, FR-018, FR-020 |
| IN-005 | One question, then re-grade | ST-004 | FR-010, FR-011 |
| IN-006 | Buildability review of the whole set | ST-005 | FR-012 |
| IN-007 | Approval recorded against a version | ST-006 | FR-013, FR-014 |

## Out of scope for R1

| ID | OUT OF SCOPE | REVISIT |
|---|---|---|
| OUT-001 | Deriving the backlog (J-002) | R2; the PRD already carries what it needs |
| OUT-002 | Two experts on one set at once | After three real engagements |
| OUT-003 | Any view beyond reading the documents and stepping through them beside the conversation | When an approver asks |
| OUT-004 | An engineering lead's own recorded confirmation | R2, with J-002 |

## Non-goals

| NON-GOAL | WHY NOT |
|---|---|
| Supplying domain knowledge | The expert has it; the product supplies product management |

| NON-GOAL | WHY NOT |
|---|---|
| Deciding scope, priority, or whether to build | The expert decides; the product shows the consequences |
| A backlog's tasks, subtasks, estimates, or sizing | They come from engineering design, not from a PRD |
| Replacing the document store, issue tracker, or approval system, or being the system of record | Product decision, 0.1.0 (was C-002) |
| A separate application to open, install, or log into | Product decision, 0.1.0 (was C-001); the session is D-001 |
| Product teams entering a domain they do not know, in release 1 | They want the domain, which prd-forge does not supply (was S-003) |
| Certifying a PRD | A grade is evidence for the expert's judgement, never a warrant |

The first is the one that defines the product: **we do not supply domain knowledge.** Someone without the domain is not served by this and should not buy it. The last is the one that protects its users: **a grade is evidence, not a warrant.** A high score on a PRD for the wrong product is still the wrong product.

Three non-goals moved here at 0.8.1. The system-of-record and separate-application non-goals were constraints C-002 and C-001 in section 9, and the non-goal on product teams new to a domain was segment S-003 in section 2.

### In scope

#### IN-001

**SUMMARY** State the idea, establish the domain and vocabulary, and create a set with every section visible and honestly marked as not started

**JOURNEY STAGE** ST-001

**REQUIREMENTS**

- FR-001
- FR-002
- FR-015

#### IN-002

**SUMMARY** Teach a section, then interview on it one or two questions at a time, in the expert's own vocabulary

**JOURNEY STAGE** ST-002

**REQUIREMENTS**

- FR-003
- FR-004
- FR-005
- FR-019

#### IN-003

**SUMMARY** Never record an invented domain fact, and convert a stated technology into the need beneath it

**JOURNEY STAGE** ST-002

**REQUIREMENTS**

- FR-006
- FR-007

### IN-004 · Show the draft as prose, with its grade and every blocking defect named

JOURNEY STAGE: ST-003

REQUIREMENTS:
- FR-008
- FR-009
- FR-016
- FR-017
- FR-018
- FR-020

### IN-005 · Ask one question that most improves the section, and re-grade after it

JOURNEY STAGE: ST-004

REQUIREMENTS:
- FR-010
- FR-011

### IN-006 · Review the whole set as an engineer would, and report what cannot be built

JOURNEY STAGE: ST-005

REQUIREMENTS:
- FR-012

### IN-007 · Record an approval, with conditions, against a named version

JOURNEY STAGE: ST-006

REQUIREMENTS:
- FR-013
- FR-014

## Out of scope

### OUT-001 · Deriving features, epics and stories from the approved PRD - journey J-002

REVISIT: Release 2. The PRD must already carry what J-002 needs, which FR-012 checks, so nothing is lost by deferring the derivation itself.

### OUT-002 · More than one expert working on the same set at the same time

REVISIT: Once a single expert's path is proven on three real engagements.

### OUT-004

SUMMARY: An engineering lead reading the set and recording their own confirmation that it can be built from

REVISIT: Release 2, with J-002. Release 1 gives the expert the buildability review (FR-012) so the gaps are found before anyone estimates; a named lead's own confirmation is a second person's time, which R1 does not ask for.

### OUT-003

SUMMARY: Any view of the document set other than reading the documents themselves and stepping through them beside the conversation

REVISIT: When an approver asks for one; PER-003 reads a handful of pages, which the documents already are.

SECTION 7

# Functional requirements

DRAFT 81/100 NO BLOCKING DEFECTS

| ID | USER | CAPABILITY | PRIORITY | STAGE |
|---|---|---|---|---|
| FR-001 | PER-001 | State the idea; establish the vocabulary once | must | ST-001 |
| FR-002 | PER-001 | See the whole set's state at a glance | should | ST-001 |
| FR-003 | PER-001 | Be taught the section before being asked | must | ST-002 |
| FR-004 | PER-001 | Be asked one or two questions at a time | must | ST-002 |
| FR-005 | PER-001 | End a section and see what was left behind | should | ST-002 |
| FR-006 | PER-001 | Unknown facts become open questions, never inventions | must | ST-002 |
| FR-007 | PER-001 | Technology talk converted; preference preserved | must | ST-002 |
| FR-008 | PER-001 | Read the draft as prose | should | ST-003 |
| FR-009 | PER-001 | See score, blocking defects, weakest part | must | ST-003 |
| FR-010 | PER-001 | Be asked one question, not twelve | must | ST-004 |
| FR-011 | PER-001 | See the re-grade against the same standard | should | ST-004 |
| FR-012 | PER-001 | See what a team could not build from it | must | ST-005 |
| FR-013 | PER-003 | See what is approved and what is open, on one page | must | ST-006 |
| FR-014 | PER-003 | Approve, with conditions, against a version | must | ST-006 |
| FR-015 | PER-001 | Stop and resume without being asked anything twice | should | ST-001 |
| FR-016 | PER-001 | Refute a finding and have it withdrawn | should | ST-003 |
| FR-017 | PER-001 | Move to any section without losing their place | should | ST-003 |
| FR-018 | PER-001 | See the set as it now stands | should | ST-003 |
| FR-019 | PER-001 | Stop being asked once the section would pass | should | ST-002 |
| FR-020 | PER-001 | Be told when the set states one item two ways | should | ST-003 |

Every release-line stage has at least one `must` requirement. Ten of the 20 are `must`: without any one of them, REL-001 does not do what it promises. The ten `should` requirements make the work faster or easier to follow, and the expert can still reach an approved set without them. The expert set this split at version 0.8.0.

Six of the `must` requirements wait on an open question in section 12. FR-004, FR-009, FR-010, and FR-012 wait on Q-009, where the rubric and the derivability checks live and who owns them. FR-010 also waits on Q-007, whether the working order is right and how looping back to the journeys feels, because Q-007 blocks ST-004. FR-013 waits on Q-005,

whether a sponsor accepts a specification that carries a grade. FR-014 waits on Q-011, how an approver records a decision.

FR-015 and FR-016 were split out of FR-002 and FR-009 at version 0.2.0, because each has a failure mode of its own: a build could deliver perfect set-state visibility and still re-ask three answered questions, and could report a grade perfectly and still defend a finding the document does not support. Ids are never renumbered, so they sit at the end rather than in stage order.

FR-019 and FR-020 were split out of FR-004 and FR-008 at version 0.9.0, for the same reason. The early stop and the check across the set each have a failure mode of their own, and each needs its own priority. The expert set both to `should`. FR-019 also waits on Q-009, because the early stop applies the rubric.

### A judgement call worth challenging

FR-002, FR-008 and FR-013 refer to "the set", "the section", "a table row" and "the root document". A reviewer may read those as naming an output format and therefore failing test 5. The judgement taken here is that a written specification is the product's subject, not its delivery technology, so these words name the artifact rather than the mechanism. What is deliberately absent everywhere: the assistant the product runs inside, how grading is performed, who or what reviews the set, and what the documents are stored as. Those are in section 6 as non-goals, in section 11 as dependencies, or nowhere.

FR-001, FR-009, FR-013 and FR-014 each read as more than one capability, and a reviewer has called them out for it. They are left whole deliberately: each is one act with several parts, and a build delivering half of one would be visibly unfinished rather than arguably done. FR-015 and FR-016 were split precisely because they were not like that, each had a failure mode the other half could hide.

> **FR-001 · State the idea in their own words and establish the domain's vocabulary once**
>
> | | |
> |---|---|
> | USER | PER-001 |
> | OUTCOME | Every question they are asked afterwards uses their unit of work and their terms, so they are never translating. |
> | ACCEPTANCE | • Given an expert who says the unit of work is a "load", when any later question refers to that unit, then it says "load" and not "item", "record" or "entity".<br>• Given three words the expert says an outsider would misread, when the first section is drafted, then each of those words appears in the glossary with the expert's own meaning.<br>• Given an expert who has not stated a domain, when they ask to begin a section, then they are asked for the domain and the unit of work first. |
> | PRIORITY | must |
> | SHOWN IN | conversation |
> | JOURNEY STAGE | ST-001 |
> | OBJECTIVE | O-003 |

### FR-002 · See every section of the set and what state each one is in, at any point

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | They can see at a glance what is done, what is weak and what has not been started, without opening any section. |
| ACCEPTANCE | • Given a set where four section documents are drafted and thirteen are pending, when the expert asks where things stand, then they see all seventeen section documents with each one's state, and the pending ones are shown as not started rather than omitted.<br>• Given a section that has been graded, when the expert sees the set's state, then that section shows its score and how many blocking defects it has.<br>• Given a set where no section has been started, when the expert asks where things stand, then every section is listed and shown as not started. |
| PRIORITY | should |
| SHOWN IN | screen |
| JOURNEY STAGE | ST-001 |
| OBJECTIVE | O-002 |

### FR-003

| | |
|---|---|
| CAPABILITY | Be taught what a section is, why it matters and what belongs in it, before being asked anything about it |
| USER | PER-001 |
| OUTCOME | Before the first question of a section reaches them, what the section is for and what belongs in it is in front of them, in their own domain's terms. |
| ACCEPTANCE | • Given a section the expert has not done before, when they begin it, then they are first told what it is, why it matters and what belongs in it, in no more than ten lines.<br>• Given a stated domain, when a section is taught, then the example used is in that domain and not in another.<br>• Given an expert who says they already know the section, when they ask to skip the teaching, then the questions begin and the teaching is not forced. |
| PRIORITY | must |
| SHOWN IN | conversation |
| JOURNEY STAGE | ST-002 |
| OBJECTIVE | O-003 |

### FR-004 · Be asked one or two questions at a time, never a questionnaire

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | No message puts more than two questions to them, so each answer they give addresses one thing. |
| ACCEPTANCE | • Given any section's interview, when the expert is asked for input, then at most two questions are put to them at once.<br>• Given an answer that leaves something the section needs still unanswered, or that could be written into the section in two ways differing in meaning, when it is received, then a follow-up is asked about that answer before the next topic is opened. |
| PRIORITY | must |
| SHOWN IN | conversation |
| JOURNEY STAGE | ST-002 |
| OBJECTIVE | O-003 |

### FR-005 · Say that they have nothing more to add, and have that end the section's interview

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | They are never trapped in an interview, and what they chose to leave incomplete is recorded rather than hidden. |
| ACCEPTANCE | • Given an expert who says they are done with a section, when the section does not yet pass its rubric, then it is written anyway at draft, and what is missing is recorded as an open question or named plainly to them.<br>• Given an expert who ends a section early, when they are told what was left behind, then the statement names the specific gap, not just a score. |
| PRIORITY | should |
| SHOWN IN | conversation |
| JOURNEY STAGE | ST-002 |
| OBJECTIVE | O-002 |

### FR-006

| | |
|---|---|
| CAPABILITY | Have a fact they do not know recorded as an open question with a decider and a date, rather than filled in for them |
| USER | PER-001 |
| OUTCOME | Nothing in their PRD is a number or a citation that nobody can stand behind. |
| ACCEPTANCE | • Given an expert who says they would have to check a figure, when the section is drafted, then no figure is written in its place and an open question records who could find out and by when.<br>• Given an objective whose baseline is unknown, when the section is drafted, then the baseline says so explicitly and an open question exists for it.<br>• Given a regulation or standard the expert has not named, when a constraint is drafted, then no regulation number appears that the expert did not supply. |
| PRIORITY | must |
| SHOWN IN | conversation |
| JOURNEY STAGE | ST-002 |
| OBJECTIVE | O-001 |

### FR-007

| | |
|---|---|
| CAPABILITY | State a technology they want and have the need beneath it recorded instead, with their preference preserved for the engineering lead |
| USER | PER-001 |
| OUTCOME | Their requirement names no technology, their reasoning is not lost, and they do not feel overruled. |
| ACCEPTANCE | • Given an expert who names a technology, when they are asked what the user gets from it and answer, then the recorded requirement states that outcome and names no technology, component, vendor or screen.<br>• Given a technology the expert stated, when the set is handed over, then that preference appears on a list of candidate architecture decisions for the engineering lead.<br>• Given an expert who insists the technology belongs in the requirement, when they have been told once why it does not, then their decision is recorded and is not argued again at the next grading pass. |
| PRIORITY | must |
| SHOWN IN | conversation |

| | |
|---|---|
| JOURNEY STAGE | ST-002 |
| OBJECTIVE | O-001 |

### FR-008 · Read the drafted section as prose rather than as a filled-in form

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | They can judge from the prose alone whether it says what they meant. |
| ACCEPTANCE | • Given a drafted section, when the expert reviews it, then every item is shown in a sentence or in a table row under named column headings, and no item is shown only as a bare list of names and values. |
| PRIORITY | should |
| SHOWN IN | screen |
| JOURNEY STAGE | ST-003 |
| OBJECTIVE | O-003 |

### FR-009 · See a score, every blocking defect by name, and which part of the section is weakest

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | They know how good the section is against a stated standard, and what specifically is wrong, rather than being told it looks good. |
| ACCEPTANCE | • Given a graded section, when the expert reads the grade, then they see a score, a count of blocking defects, and each blocking defect stated with where it is.<br>• Given a section scoring 70 or more with one blocking defect, when the grade is reported, then the section is not offered as ready, and the defect is named.<br>• Given a graded section, when the expert reads the grade, then the weakest part is named and why it is weak is stated in no more than three lines. |
| PRIORITY | must |
| SHOWN IN | screen |
| JOURNEY STAGE | ST-003 |
| OBJECTIVE | O-001 |

### FR-010

| | |
|---|---|
| CAPABILITY | Be asked the single question that would most improve the section, not a list of everything wrong with it |
| USER | PER-001 |
| OUTCOME | They can act immediately, and the section improves one answer at a time instead of stalling under a list. |
| ACCEPTANCE | • Given a graded section scoring under 70, or carrying any blocking defect, when the expert is asked to improve it, then exactly one question is put to them.<br>• Given that one question, when it is asked, then it addresses the weakest part that was just named, and not a different part. |
| PRIORITY | must |
| SHOWN IN | conversation |
| JOURNEY STAGE | ST-004 |
| OBJECTIVE | O-003 |

### FR-011 · See the section re-graded against the same standard after answering

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | They can see whether their answer actually improved the section, and by how much. |
| ACCEPTANCE | • Given an answered question, when the section is re-graded, then it is scored against the same rubric as before and both scores are visible.<br>• Given a re-grade, when the score has risen, then the document has changed; a score never rises without a change to the document.<br>• Given a section scoring 70 or more with no blocking defects, when it is re-graded, then the expert is offered the state `review`, and is told that `review` means the section is ready for someone else to read. |
| PRIORITY | should |
| SHOWN IN | screen |
| JOURNEY STAGE | ST-004 |
| OBJECTIVE | O-003 |

### FR-012

| | |
|---|---|
| CAPABILITY | See what a development team would not be able to build from the set, before anyone has estimated it |
| USER | PER-001 |
| OUTCOME | They learn where their own specification fails while it is still cheap to fix, rather than at iteration three. |
| ACCEPTANCE | • Given a set in which every section has been written at least once, when the expert asks for a buildability review, then they see a verdict and, for each release-line stage, whether a requirement covers it or it is a gap.<br>• Given a requirement whose acceptance criteria an engineer who was not in the interview judges insufficient to write a test from, when the review reports it, then the requirement is named and the document is quoted.<br>• Given an open question that blocks a `must` requirement, when the review runs, then that requirement's wait on the question is reported.<br>• Given a review finding, when it is reported to the expert, then it has been checked against the document and anything not supported by the document has been dropped. |
| PRIORITY | must |
| SHOWN IN | report |
| JOURNEY STAGE | ST-005 |
| OBJECTIVE | O-001 |

### FR-013

| | |
|---|---|
| CAPABILITY | See, in the set's own root document, what is being approved, what is still undecided and who owns each undecided thing |
| USER | PER-003 |
| OUTCOME | They can commit budget and a date knowing what they are committing to and what is still open. |
| ACCEPTANCE | • Given a set offered for approval, when the approver opens its root document, then that one document carries the summary, the scope, every section's state and grade, and every open question with its decider and date, without the approver opening any other document.<br>• Given a set with any blocking defect anywhere, when the approver opens it, then they see that, with the sections it is in. |
| PRIORITY | must |
| SHOWN IN | report |
| JOURNEY STAGE | ST-006 |
| OBJECTIVE | O-001 |

### FR-014

| | |
|---|---|
| CAPABILITY | Approve, approve with conditions, or reject, and have that recorded against a named version |
| USER | PER-003 |
| OUTCOME | What was approved, when, by whom and on what conditions is answerable later without asking anyone. |
| ACCEPTANCE | • Given an approver who approves with conditions, when they state the conditions, then each condition names the open question or requirement it depends on and a date.<br>• Given a recorded approval, when the set is read later, then the version approved, the approver, their role, the date and the decision are all present.<br>• Given an approved set that is then changed, when the change is written, then a new version is recorded and the earlier approval is not represented as covering it. |
| PRIORITY | must |
| SHOWN IN | report |
| JOURNEY STAGE | ST-006 |
| OBJECTIVE | O-001 |

### FR-015 · Stop in the middle of a set and resume days later without being asked anything twice

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | They pick up at the first unfinished section, and nothing they have already told the product is asked for again. |
| ACCEPTANCE | • Given an expert returning after eight days, when they say they want to carry on, then work resumes at the first section not yet at `review`.<br>• Given an expert who has already stated their domain and vocabulary, when they resume, then they are not asked for either again and the vocabulary in use is unchanged.<br>• Given an expert who answered a question in a section before stopping, when they resume that section, then that question is not put to them again. |
| PRIORITY | should |
| SHOWN IN | conversation |
| JOURNEY STAGE | ST-001 |
| OBJECTIVE | O-002 |

### FR-016 · Refute a grading finding by pointing at the document, and have it withdrawn

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | They are told the finding was dropped, and it does not reappear at the next grading pass. |
| ACCEPTANCE | • Given a finding the expert refutes, when it is checked against the document and the document does not support it, then it is dropped and the expert is told it was dropped.<br>• Given a finding the document does support, when the expert disputes it, then the reason is stated once and their decision is recorded.<br>• Given a decision the expert reaffirmed after an objection, when the section is graded again, then the objection is not raised as a new finding. |
| PRIORITY | should |
| SHOWN IN | conversation |
| JOURNEY STAGE | ST-003 |
| OBJECTIVE | O-001 |

### FR-017

| | |
|---|---|
| CAPABILITY | Move from the section in front of them to any other section of the set without losing their place |
| USER | PER-001 |
| OUTCOME | They reach the section they want in one step, and can always return to the view they were on before. |
| ACCEPTANCE | • Given a section in front of them, when they ask for the next or the previous section, then the section after or before it in reading order is shown, and at the first or the last section the view does not change.<br>• Given a section that cites an id another section defines, such as Q-001, when they follow that id, then the section that defines it is shown and the id is named above it.<br>• Given they have moved from a section to three other things in the set, one after another, when they ask to go back three times, then they see each earlier thing again, the most recent first, and end at the section they started from. |
| PRIORITY | should |
| SHOWN IN | screen |
| JOURNEY STAGE | ST-003 |
| OBJECTIVE | O-002 |

### FR-018 · See the set as it now stands after it was changed outside the view in front of them

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | After they ask to refresh, what they read is the set as it now stands. |
| ACCEPTANCE | • Given a section changed after they opened the set, when they ask to refresh, then within 10 seconds that section shows its new text and its new state.<br>• Given a section that is graded again while the set is open, when they ask to refresh, then within 10 seconds the new grade is shown.<br>• Given no set existed when they looked, when a set is written and they ask to refresh, then every section of it is listed. |
| PRIORITY | should |
| SHOWN IN | screen |
| JOURNEY STAGE | ST-003 |
| OBJECTIVE | O-002 |

### FR-019 · Stop being asked once the section would already pass

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | The interview ends when the section has what it needs, so they are not asked questions it no longer needs. |
| ACCEPTANCE | • Given an expert whose answers would already score the section 70 or more with no blocking defects, when further questions remain unasked on the question list, then the interview stops rather than continuing to the end of it.<br>• Given an expert whose answers would score the section under 70, or with a blocking defect, when questions remain on the question list, then the next question is put to them. |
| PRIORITY | should |
| SHOWN IN | conversation |
| JOURNEY STAGE | ST-002 |
| OBJECTIVE | O-003 |

### FR-020 · Be told when the set states one item two different ways

| | |
|---|---|
| USER | PER-001 |
| OUTCOME | They see every id, priority or link that two places in the set state differently, before anyone builds from either. |
| ACCEPTANCE | • Given an id, priority or link that two places in the set state differently, when the section is read, then the expert is told of the difference rather than it being silently resolved.<br>• Given a set in which every id, priority and link is stated the same way wherever it appears, when the section is read, then no difference is reported. |
| PRIORITY | should |
| SHOWN IN | screen |
| JOURNEY STAGE | ST-003 |
| OBJECTIVE | O-003 |

SECTION 7

# Wireframes

DRAFT 78/100 2 BLOCKING

Every requirement in section 7 whose outcome a person meets on a screen or in a report is drawn here as a text sketch. Each sketch shows what the person sees and what they can do there, with the set's own content. None of them decides how anything looks or how it is built.

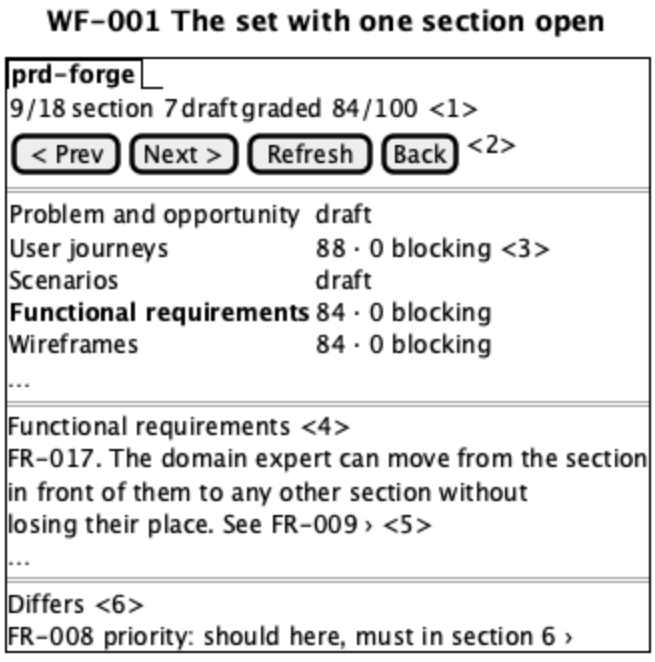
Five sketches follow the domain expert through the golden thread: a new set with nothing started, the set with a section open, why a section got its grade, the rubric behind it, and a re-grade. Two reports decide whether the set is buildable: the buildability review and the approval pack. Three smaller sketches draw the other states: no set at all, which FR-018's third criterion names; the review that finds gaps, which SC-005 names; and the return after a week away, which SC-002 names.

The numbered marks in each sketch are explained, in order, by its `shows` list. Each sketch is also drawn for the printed PRD.

### WF-001 The set with one section open

Illustration, not a design. Illustrates FR-002, FR-008, FR-017, FR-018, and FR-020. PER-001 sees it at ST-003.

The expert reads a drafted section with the whole set beside it. Every graded section shows its score and its blocking count, so what is weak is visible without opening anything. When two places in the set state one id, priority, or link differently, mark 6 lists each difference with both values. The difference drawn is an example, not a fact of this set.



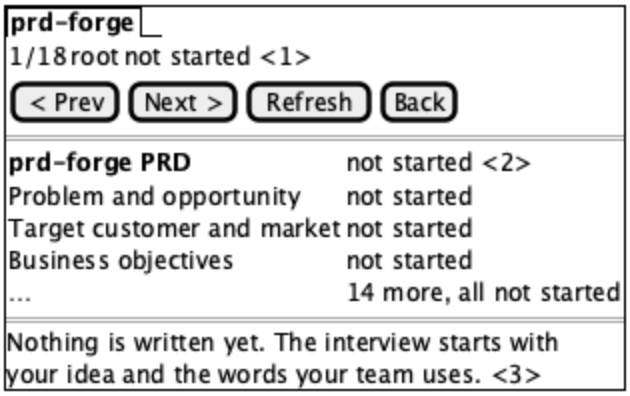

WF-001 drawing

### WF-002 A new set, before the first interview

Illustration, not a design. Illustrates FR-002, FR-017, and FR-018. PER-001 sees it at ST-001.

At ST-001 the set exists and nothing in it is written. Every section is listed and marked not started, so the expert sees the whole shape of the work before the first question.

**WF-002 A new set, before the first interview**



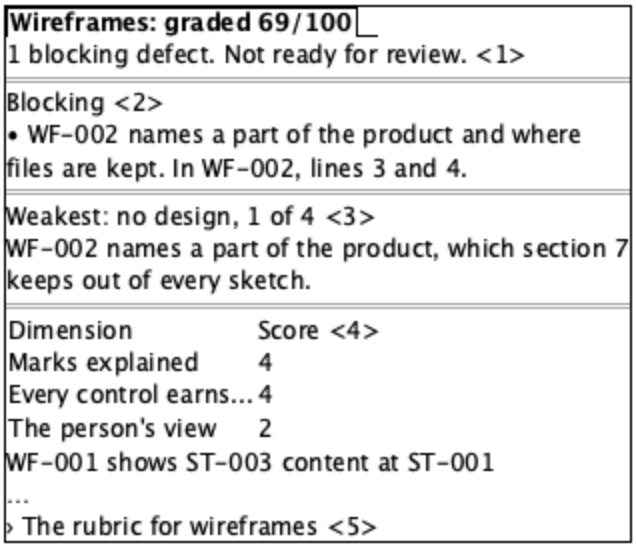

WF-002 drawing

### WF-003 Why a section got its grade

Illustration, not a design. Illustrates FR-009 and FR-017. PER-001 sees it at ST-003.

A score alone does not tell the expert what to fix. This view names each blocking defect and where it is, and names the weakest part, quoting the section.

**WF-003 Why a section got its grade**




WF-003 drawing

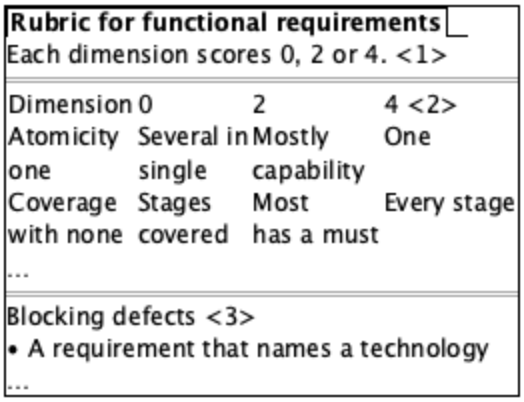
### WF-004 The standard a section is graded against

Illustration, not a design. Illustrates FR-009 and FR-017. PER-001 sees it at ST-003.

A grade against a stated standard is only fair if the expert can read the standard. This is the same rubric the graders used.

**WF-004 The standard a section is graded against**

| Rubric for functional requirements<br>Each dimension scores 0, 2 or 4. <1> | | | |
|---|---|---|---|
| Dimension | 0 | 2 | 4 <2> |
| Atomicity | Several in one | Mostly single capability | One |
| Coverage | Stages with none | Most covered | Every stage has a must |
| ... | | | |
| Blocking defects <3><br>• A requirement that names a technology<br>... | | | |



WF-004 drawing

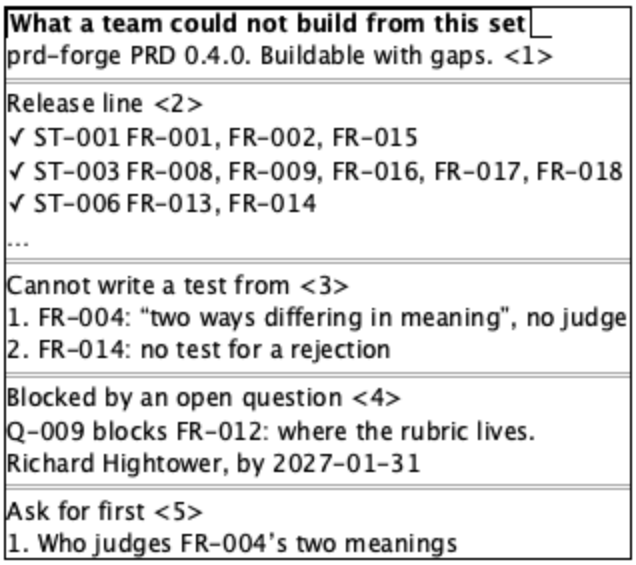
### WF-005 What a team could not build from this set

Illustration, not a design. Illustrates FR-012. PER-001 sees it at ST-005.

The expert reads this before anyone has estimated anything, while a fix is still a sentence. Every finding quotes the set, so the expert can check it rather than trust it.

**WF-005 What a team could not build from this set**

| What a team could not build from this set<br>prd-forge PRD 0.4.0. Buildable with gaps. <1> |
|---|
| Release line <2><br>✓ ST-001 FR-001, FR-002, FR-015<br>✓ ST-003 FR-008, FR-009, FR-016, FR-017, FR-018<br>✓ ST-006 FR-013, FR-014<br>... |
| Cannot write a test from <3><br>1. FR-004: "two ways differing in meaning", no judge<br>2. FR-014: no test for a rejection |
| Blocked by an open question <4><br>Q-009 blocks FR-012: where the rubric lives.<br>Richard Hightower, by 2027-01-31 |
| Ask for first <5><br>1. Who judges FR-004's two meanings |

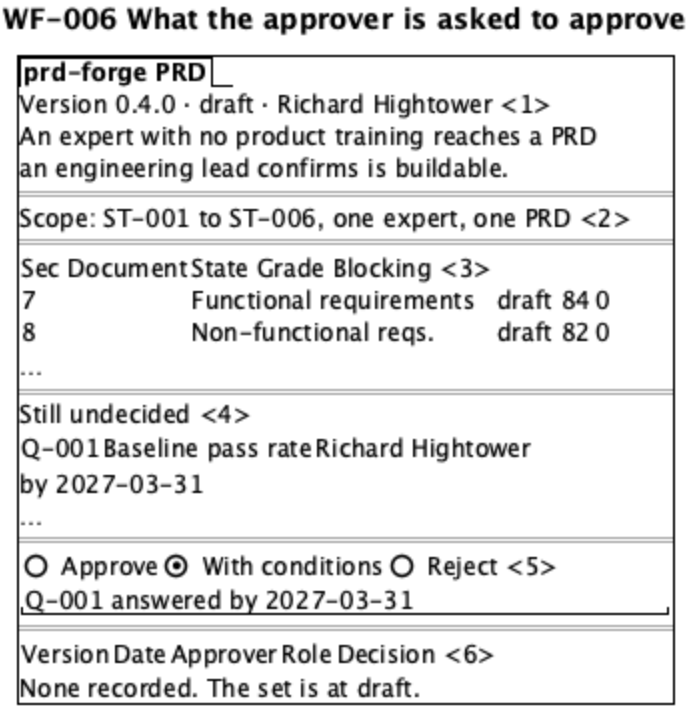


WF-005 drawing

### WF-006 What the approver is asked to approve

Illustration, not a design. Illustrates FR-013 and FR-014. PER-003 sees it at ST-006.

The approver reads one document, not eighteen. It says what they are committing to, what is still open and who owns it, and where their decision is recorded.

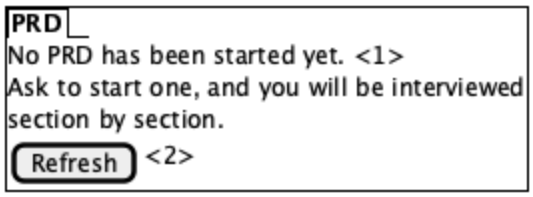



WF-006 drawing

### WF-007 Before any set exists

Illustration, not a design. Illustrates FR-018. PER-001 sees it at ST-001.

Before the first interview there is no set to list. The view says so plainly, and a refresh lists every section once the set is written.

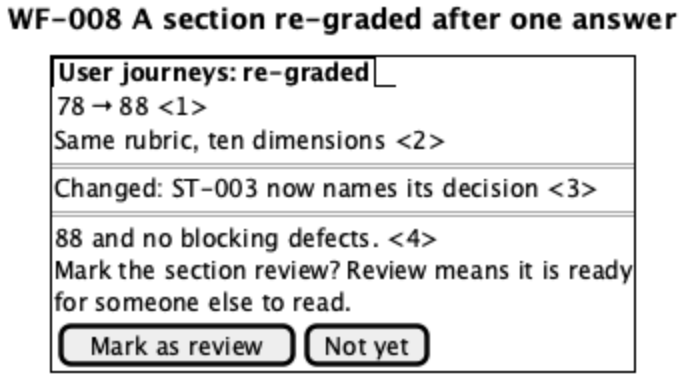



WF-007 drawing

### WF-008 A section re-graded after one answer

Illustration, not a design. Illustrates FR-011. PER-001 sees it at ST-004.

After one answer the section is graded again against the same rubric. Both scores are visible, and at 70 or more with no blocking defects the expert is offered review. The values are SC-001's.

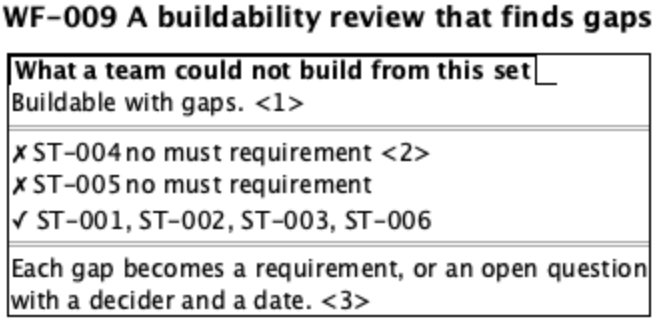



WF-008 drawing

### WF-009 A buildability review that finds gaps

Illustration, not a design. Illustrates FR-012. PER-001 sees it at ST-005.

SC-005 is the review that finds work missing: two release-line stages have no must requirement behind them. The gaps are named first, before anything else in the review.

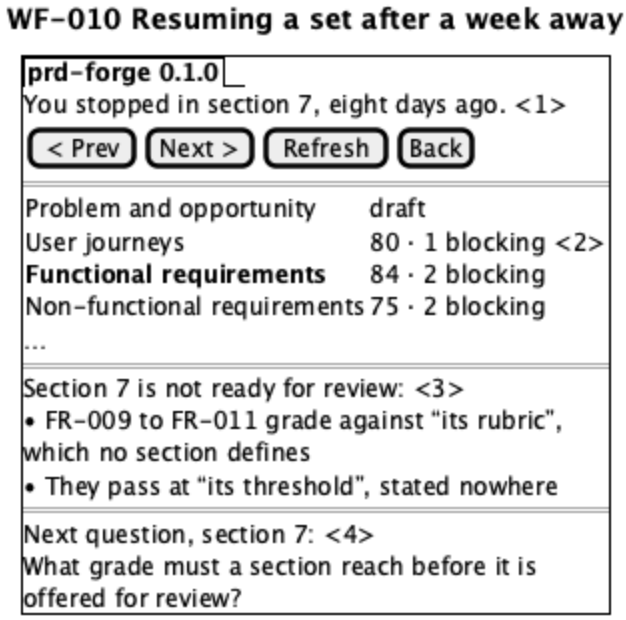



WF-009 drawing

### WF-010 Resuming a set after a week away

Illustration, not a design. Illustrates FR-002 and FR-009. PER-001 sees it at ST-004.

SC-002 is the expert who returns after eight days and remembers nothing. The values are prd-forge's own at version 0.1.0, when section 7 scored 84 with two blocking defects: a grade of 70 or more that is still not offered as ready.




WF-010 drawing

### WF-001 · The set with one section open

**TYPE** screen

**PERSONA** PER-001

**JOURNEY STAGE** ST-003

**ILLUSTRATES**

- FR-002
- FR-008
- FR-017
- FR-018
- FR-020

**SHOWS**

- Where they are in the set, the section's state, and its grade
- Next, previous, refresh, and back to the view before
- Each section's grade and blocking count; opening it shows why
- The section itself, as prose
- An ID that another section defines opens that section
- Each place where the set states one id, priority, or link two ways, with both values, and each place can be followed

### WF-002 · A new set, before the first interview

**TYPE** screen

**PERSONA** PER-001

**JOURNEY STAGE** ST-001

**ILLUSTRATES**
- FR-002
- FR-017
- FR-018

**SHOWS**
- Where they are, and that nothing is started yet
- Every section listed, each marked not started
- What happens first

### WF-003 · Why a section got its grade

**TYPE** screen

**PERSONA** PER-001

**JOURNEY STAGE** ST-003

**ILLUSTRATES**
- FR-009
- FR-017

**SHOWS**
- The score, and that the section is not ready
- Each blocking defect, and where it is
- The weakest part, and why, in three lines or fewer
- Each dimension's score, with the evidence for it
- The standard the grade was given against

### WF-004 · The standard a section is graded against

**TYPE** screen

**PERSONA** PER-001

**JOURNEY STAGE** ST-003

**ILLUSTRATES**
- FR-009
- FR-017

**SHOWS**
- How a dimension becomes a grade
- What 0, 2 and 4 look like for each dimension
- What stops a section at any grade

### WF-005 · What a team could not build from this set

| | |
|---|---|
| TYPE | report |
| PERSONA | PER-001 |
| JOURNEY STAGE | ST-005 |
| ILLUSTRATES | • FR-012 |
| SHOWS | • Which version of the set was reviewed, and the verdict<br>• Each release-line stage, covered or a gap<br>• Each requirement an engineer could not test, quoted<br>• An open question that blocks a must, its decider and date<br>• What the reviewer would ask for first |

### WF-006 · What the approver is asked to approve

| | |
|---|---|
| TYPE | report |
| PERSONA | PER-003 |
| JOURNEY STAGE | ST-006 |
| ILLUSTRATES | • FR-013<br>• FR-014 |
| SHOWS | • The version, its state, its owner, and what it is for<br>• What release 1 covers<br>• Every section with its state and grade, and what blocks<br>• What is still undecided, who decides it, and by when<br>• The decision, with any condition tied to a question<br>• Who decided what, in which role, on which version |

### WF-007 · Before any set exists

| | |
|---|---|
| TYPE | screen |
| PERSONA | PER-001 |
| JOURNEY STAGE | ST-001 |
| ILLUSTRATES | • FR-018 |
| SHOWS | • That nothing has been started<br>• Look again once a set exists |

### WF-008 · A section re-graded after one answer

| | |
|---|---|
| TYPE | screen |
| PERSONA | PER-001 |
| JOURNEY STAGE | ST-004 |
| ILLUSTRATES | • FR-011 |
| SHOWS | • The new score beside the old one<br>• The same standard both times<br>• What changed in the section<br>• The offer of review, and what review means |

### WF-009 · A buildability review that finds gaps

| | |
|---|---|
| TYPE | report |
| PERSONA | PER-001 |
| JOURNEY STAGE | ST-005 |
| ILLUSTRATES | • FR-012 |
| SHOWS | • The verdict<br>• The release-line stages with no requirement, named<br>• That each gap must become a requirement or a question |

### WF-010 · Resuming a set after a week away

| | |
|---|---|
| TYPE | screen |
| PERSONA | PER-001 |
| JOURNEY STAGE | ST-004 |
| ILLUSTRATES | • FR-002<br>• FR-009 |
| SHOWS | • Where they stopped<br>• Every section's grade and blocking count<br>• Each blocking defect in the section, named<br>• The next question, not one already answered |

SECTION 8

# Non-functional requirements

DRAFT 82/100 NO BLOCKING DEFECTS

| ID | CATEGORY | TARGET |
|---|---|---|
| NFR-001 | performance | Next question under 20s at p95 |
| NFR-002 | performance | Section grade under 90s at p95 |
| NFR-003 | scale | 120 requirements, 10 journeys x 10 stages, 200 terms |
| NFR-004 | durability | 100% of a set readable with no part of the product running |
| NFR-005 | usability | 90% of sections completed without asking what a term means |
| NFR-006 | accessibility | Zero WCAG 2.2 AA violations on the golden-thread path |
| NFR-007 | observability | 100% of grades reconstructable after the session |
| NFR-008 | privacy | Zero destinations beyond the session the expert already uses |
| NFR-009 | supportability | 100% of older sets readable within a major version |
| NFR-010 | retention | 100% retrievable from the organisation's own store, no expiry imposed by us |
| NFR-011 | availability | Zero answers lost per interruption; resume within 1 minute |
| NFR-012 | security | Zero read paths the organisation's own access rules do not govern |

### Where these came from

NFR-001 and NFR-002 come from PER-001's frequency: a burst of intense use. A 20-second wait is tolerable once; across 60 questions it is 20 minutes of dead time and the expert stops.

NFR-004 is the one that matters most and is cheapest to get wrong. PER-001 works in bursts weeks apart. If the set is only readable through this product, then a version change, a tool change or an expired session destroys weeks of an expert's thinking. The target is deliberately absolute.

NFR-005 is the only requirement here that needs people rather than instruments to measure, and it is the one most likely to be skipped. It is the direct test of whether P-004 was solved.

NFR-006 covers the approver too, not only the expert: PER-003 reads this once and must not need a mouse to do it.

There is no concurrency target, and that is a decision rather than an omission. Under D-001 the product reaches one expert inside a session they are already in, so simultaneous users are the host's load and not ours; NFR-003 therefore reads scale as the size of one set, which is the dimension we do control. If D-001 or the separate-application non-goal in section 6 ever changes, concurrency comes back and this paragraph goes.

NFR-006 and NFR-012 set targets on things the non-goals in section 6 say we do not own: what the host renders, and where the customer keeps its documents. They are kept as `must` because a target nobody owns is how accessibility and access control get lost; who owns them is Q-011.

NFR-010, NFR-011 and NFR-012 were added at version 0.2.0. Retention, availability and security had simply not been reached, and C-003, now assumption A-005, stated a retention rule that no measurable requirement stood behind. NFR-010 is the one that answers A-005's second half: the product retains nothing of its own, so the only retention question is how long the organisation keeps the documents, and the answer must be that we impose no expiry on them.

## Requirements

### NFR-001

| | |
|---|---|
| STATEMENT | An expert answering a question does not lose their train of thought waiting for the next one. |
| CATEGORY | performance |
| METRIC | 95th percentile elapsed time from the expert submitting an answer to the next question or the drafted section appearing. |
| TARGET | under 20 seconds at p95, and under 60 seconds at p99 |
| MEASUREMENT | At least 300 timed observations per release, taken across five scripted runs of all seventeen section documents on a set already holding 60 requirements, from the expert's own side of the conversation rather than from inside the product. |
| PRIORITY | must |
| APPLIES TO | • FR-004<br>• FR-010 |

### NFR-002 · A grading pass does not make the expert wait so long that they stop asking for one.

| | |
|---|---|
| CATEGORY | performance |
| METRIC | 95th percentile elapsed time for one section's grade to be reported, measured per grading pass on the section just drafted. |
| TARGET | under 90 seconds at p95 for a single section |
| MEASUREMENT | At least 200 grading passes per release, spread across all seventeen section documents and across sets of 20, 60 and 120 requirements. |
| PRIORITY | must |
| APPLIES TO | • FR-009<br>• FR-011 |

### NFR-003 · A set stays workable at the size real products reach.

| | |
|---|---|
| CATEGORY | scale |
| METRIC | The size of one set at which every requirement in section 7 still passes its acceptance criteria and every target in this section still holds. |
| TARGET | 120 functional requirements, 10 journeys of 10 stages, and 200 glossary terms |
| MEASUREMENT | A generated set at those sizes, run through every function in section 7, per release. |
| PRIORITY | must |

### NFR-004 · Work survives the expert stopping, and survives the tool they were using going away.

| | |
|---|---|
| CATEGORY | durability |
| METRIC | Share of a set's content recoverable after the authoring session ends, with no part of the product running. |
| TARGET | 100%, readable without this product |
| MEASUREMENT | End a session mid-section, then read the whole set with nothing but a text editor, once per release. |
| PRIORITY | must |
| APPLIES TO | • FR-015 |

### NFR-005

| | |
|---|---|
| STATEMENT | An expert with no product management training can complete a section without asking anyone what a word means. |
| CATEGORY | usability |
| METRIC | Share of sections a first-time expert completes without asking for the meaning of a product management term. |
| TARGET | 90% of sections, measured over 10 first-time experts |
| MEASUREMENT | Observed sessions with 10 first-time experts before release, questions to the observer counted and categorised. |
| PRIORITY | must |

### NFR-006

| | |
|---|---|
| STATEMENT | Everything the expert and the approver read is usable with a screen reader and without a mouse. |
| CATEGORY | accessibility |
| METRIC | Conformance to WCAG 2.2 AA of everything the product puts in front of the expert or the approver along the golden-thread journey. |
| TARGET | zero level-A or level-AA violations on the golden-thread path |
| MEASUREMENT | One manual screen-reader pass of the golden-thread journey per release, plus an automated check in the release pipeline of everything that journey shows. |
| PRIORITY | must |

### NFR-007 · An expert can tell why a section received the score it did, after the fact.

| | |
|---|---|
| CATEGORY | observability |
| METRIC | Share of grades for which the score per rubric dimension and the quotation behind each can be recovered after the session. |
| TARGET | 100% for any section at review or better |
| MEASUREMENT | Re-read a graded set a week later and reconstruct every score, once per release. |
| PRIORITY | should |
| APPLIES TO | • FR-009 |

### NFR-008

| | |
|---|---|
| STATEMENT | What the expert says about their domain stays within whatever boundary their organisation already works in; this product adds no new place for it to go. |
| CATEGORY | privacy |
| METRIC | Number of destinations set content reaches beyond the boundary A-005 names. |
| TARGET | zero |
| MEASUREMENT | Review of every outbound destination per release, recorded against A-005. |
| PRIORITY | must |

### NFR-009 · A set written by an earlier version is still readable and still gradable by a later one.

| | |
|---|---|
| CATEGORY | supportability |
| METRIC | Share of sets written by version N that version N+1 reads and grades without the expert rewriting anything. |
| TARGET | 100% within a major version |
| MEASUREMENT | Keep one set per released version and run the current version against all of them, per release. |
| PRIORITY | should |

### NFR-010

| | |
|---|---|
| STATEMENT | A set remains readable and retrievable for as long as the organisation keeps the specifications it builds from, without this product being involved. |
| CATEGORY | retention |
| METRIC | Share of a set retrievable and readable from the organisation's own document store, with no part of this product present. |
| TARGET | 100%, for as long as the organisation retains the set; no expiry imposed by us |
| MEASUREMENT | Retrieve a set written one release ago from the organisation's own store and read it with a text editor, once per release. |
| PRIORITY | must |
| APPLIES TO | • FR-014 |

### NFR-011

| | |
|---|---|
| STATEMENT | An interview in progress survives the product being unavailable; no session loses an answer the expert has already given. |
| CATEGORY | availability |
| METRIC | Answers lost per interruption, and time to resume after one. |
| TARGET | Zero answers lost, on every interruption; resume within 1 minute at p95 and within 3 minutes at p99 |
| MEASUREMENT | Interrupt a session mid-section 100 times per release, spread across all seventeen section documents, and count answers lost and seconds to resume. |
| PRIORITY | must |
| APPLIES TO | • FR-015 |

### NFR-012

| | |
|---|---|
| STATEMENT | Nobody but the people the organisation has already given access to its documents can read a set. |
| CATEGORY | security |
| METRIC | Number of ways to read a set that the organisation's own document access rules do not govern. |
| TARGET | zero |
| MEASUREMENT | Enumerate every read path to a set per release and check each against the organisation's own access rules; recorded against the system-of-record non-goal in section 6 and A-005. |
| PRIORITY | must |

SECTION 9

# Constraints

DRAFT 95/100 NO BLOCKING DEFECTS

This section holds no constraints. No rule imposed from outside the team binds prd-forge itself. Each rule that this section once held was a decision of our own, a belief about customers, or the product's chosen response to a risk. The team could change each one by agreeing among themselves, so each one now lives where that kind of item belongs. The ids C-001 to C-004 are not reused.

An expert who uses prd-forge can still be bound by a professional code, a contract, or a firm policy. Those constraints belong in that expert's own PRD, not in this one.

### Where the rules live now

At version 0.8.1, C-001 to C-003 left this section.

- C-001, no separate application for the expert to open, install, or log into, is a non-goal in section 6. The session it needs is dependency D-001 in section 11.
- C-002, the output is plain documents and the product is not the system of record, is a non-goal in section 6.
- C-003, domain content stays in the expert's session and in the documents, is assumption A-005 in section 11.

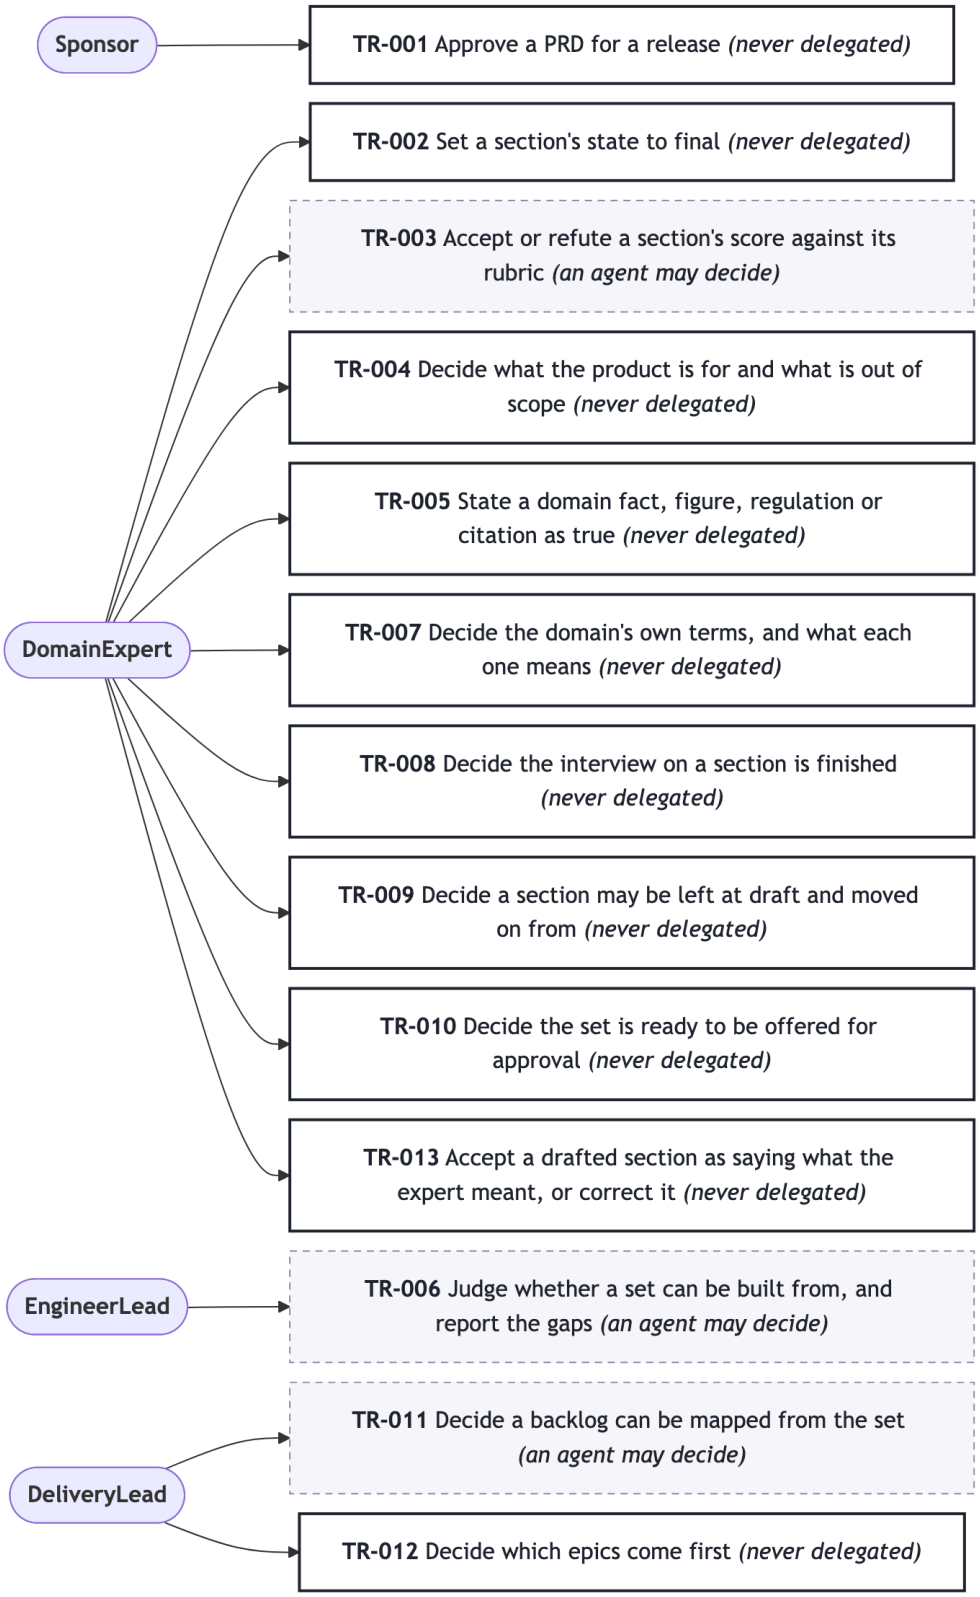
At version 0.9.0, C-004 left this section.

- C-004, no regulation, standard, clause, or statistic is stated as fact unless the expert supplied it or it was checked against its source, is trust rule TR-005 in section 10 and requirement FR-006 in section 7. TR-005 makes the expert the only one who can state such a fact as true, and it is never delegated. FR-006 records a fact the expert does not know as an open question with a decider and a date.

SECTION 10

# Decision rights and trust rules

DRAFT 88/100 NO BLOCKING DEFECTS




Who decides what

| ID | DECISION | DECIDER | AGENT MAY DECIDE | NEVER DELEGATED |
|---|---|---|---|---|
| TR-001 | Approve a PRD for a release | PER-003 | no | yes |
| TR-002 | Set a section to final | PER-001 | no | yes |
| TR-003 | Accept or refute a section's score | PER-001 | **yes** | no |
| TR-004 | Decide scope | PER-001 | no | yes |
| TR-005 | State a domain fact as true | PER-001 | no | yes |
| TR-006 | Judge buildability, and report the gaps | PER-002 | **yes** | no |
| TR-007 | Decide the domain's own terms | PER-001 | no | yes |
| TR-008 | Decide a section's interview is finished | PER-001 | no | yes |
| TR-009 | Decide a section may be left at draft | PER-001 | no | yes |
| TR-010 | Decide the set is ready to offer for approval | PER-001 | no | yes |
| TR-011 | Decide a backlog can be mapped | PER-004 | **yes** | no |
| TR-012 | Decide which epics come first | PER-004 | no | yes |
| TR-013 | Accept or correct a drafted section | PER-001 | no | yes |

Every decision named in a stage of section 5 appears here. TR-007 to TR-013 were added at version 0.2.0 to close that: five release-line stage decisions and both of J-002's had no rule, which meant gates an engineering team would have had to invent. TR-013 covers the one in the middle of the loop, the expert accepting that a draft says what they meant, which is distinct from scoring it (TR-003), from leaving it at draft (TR-009) and from setting it final (TR-002).

### The shape of the trust boundary

Only three of thirteen are delegable, and all three are **assessments**, not decisions: TR-003 produces a document's score, TR-006 judges whether it can be built from, and TR-011 judges whether a backlog can be mapped from it. All three produce evidence for a person. None changes a document, and none can set a state.

Everything that commits anyone to anything, such as approving, finalising, choosing scope, or asserting a fact, is never delegated. TR-005 is the strictest: a fact nobody can stand behind is not written at all, by anyone.

The delegable assessments all carry the same two obligations: **quote the document for every claim**, and **drop any finding the document does not support**. An assessment that cannot be refuted by pointing at the text is not evidence, it is an assertion, and the expert will correctly stop believing it. See R-002 in section 11.

> **TR-001 · Approve a PRD for a release**
>
> | | |
> |---|---|
> | STAGES | • ST-006 |
> | APPROVER | PER-003 |
> | AN AGENT MAY DECIDE | no |
> | NEVER DELEGATED | yes |
> | EVIDENCE REQUIRED | The summary, the scope, every section's state and grade, every blocking defect, and every open question with its decider and date. |

| | |
|---|---|
| ESCALATION | No substitute approver, and approval waits. Where no sponsor exists at all, such as a solo expert or an internal tool with no PER-003, the expert approves their own set under this rule, shown the same evidence, and the change log records that no separate approver existed. A set is never treated as approved because its score is high. |
| ACCEPTANCE | • Given a set with any blocking defect, when approval is sought, then the approver is shown the defects and the set is not presented as ready.<br>• Given a set with no recorded approval, when anyone reads its state, then it reads as not approved, whatever its scores are.<br>• Given an engagement with no sponsor, when the expert approves their own set, then the approval records that no separate approver existed.<br>• Given an approver who approves with conditions, when the approval is recorded, then each condition names the open question or requirement it depends on and a date, and the set reads as approved-with-conditions until every condition is cleared by that same approver. |

### TR-002 · Set a section's state to final

| | |
|---|---|
| STAGES | • ST-004 |
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | The section's grade, its blocking defect count, and the result of validating it. |
| ESCALATION | The expert may override the gate; the override and the score at the time are recorded in the change log. Nobody may set a section final on the expert's behalf while they are away: it waits. |
| ACCEPTANCE | • Given a section with one or more blocking defects, when final is requested, then it is not set without an explicit override from the expert.<br>• Given an override, when it is applied, then the change log records that it was an override and the score at the time. |

### TR-003 · Accept or refute a section's score against its rubric

| | |
|---|---|
| STAGES | • ST-003 |
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | yes |
| NEVER DELEGATED | no |
| EVIDENCE REQUIRED | The rubric's dimensions and a quotation from the document for every score below full marks. |
| ESCALATION | The expert may refute any finding by pointing at the document; a finding the document does not support is dropped (FR-016). While the expert is away a grade stands unaccepted, and the section does not advance. |
| ACCEPTANCE | • Given a score below full marks on a dimension, when it is reported, then a quotation from the document is given for it.<br>• Given a finding the expert refutes and the document does not support, when it is checked, then it is dropped and the expert is told it was dropped. |

### TR-004 · Decide what the product is for and what is out of scope

| | |
|---|---|
| STAGES | • ST-001 |
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | What the consequence of the choice is, including any release-line stage it would leave uncovered. |
| ESCALATION | None, and no substitute. The expert decides; while they are away, scope does not change. Where the choice weakens the PRD the objection is stated once and the decision is recorded with the dissent noted. |
| ACCEPTANCE | • Given an expert decision that leaves a release-line stage uncovered, when it is recorded, then the consequence is stated to the expert once and the decision stands.<br>• Given a decision the expert reaffirmed after an objection, when later grading passes run, then the objection is not raised again as a new finding. |

### TR-005 · State a domain fact, figure, regulation or citation as true

| | |
|---|---|
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | The expert's own statement of it, or a source that has been checked. |
| ESCALATION | Where neither exists the fact is not written at all; it becomes an open question with a decider and a date. Nobody may supply the fact on the expert's behalf while they are away. |
| ACCEPTANCE | • Given a figure or citation the expert did not supply and that has not been verified, when a section is drafted, then it does not appear in the document.<br>• Given a citation that was checked and could not be confirmed, when it is reported, then it is marked unverified rather than stated as fact. |

### TR-006 · Judge whether a set can be built from, and report the gaps

| | |
|---|---|
| STAGES | • ST-009 |
| APPROVER | PER-002 |
| AN AGENT MAY DECIDE | yes |
| NEVER DELEGATED | no |
| EVIDENCE REQUIRED | The release line, the requirements covering each stage, and a quotation for every gap claimed. |
| ESCALATION | The engineering lead's own judgement overrides any reported verdict; the report is evidence for them, not a substitute for them. In release 1 no engineering lead is in the loop, so the report goes to the expert, who decides what to do with it under TR-010. |
| ACCEPTANCE | • Given a reported verdict, when the engineering lead disagrees, then their judgement is recorded as the verdict and the report is kept alongside it.<br>• Given a reported gap, when it is stated, then the document is quoted for it. |

### TR-007 · Decide the domain's own terms, and what each one means

| | |
|---|---|
| STAGES | • ST-001 |
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | Where the term is already used in the set, and any standard-setter's definition that differs from the expert's. |
| ESCALATION | None, and no substitute. Terms are the expert's; while they are away the vocabulary does not change. |
| ACCEPTANCE | • Given a term the expert has defined, when any later question or document uses that idea, then it uses the expert's word and not a synonym.<br>• Given a standard-setter's definition that differs, when the term is recorded, then both are stated and the expert says which the set uses. |

### TR-008 · Decide the interview on a section is finished

| | |
|---|---|
| STAGES | • ST-002 |
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | What the section would score as it stands, and what is still missing from it. |
| ESCALATION | None. The expert ends the interview; nobody ends it for them, and no question list runs past the point they call it done. |
| ACCEPTANCE | • Given an expert who says they have nothing more to add, when they say it, then no further question is put to them in that section.<br>• Given an interview ended early, when the section is written, then what is missing is named to the expert, not hidden behind a score. |

### TR-009 · Decide a section may be left at draft and moved on from

| | |
|---|---|
| STAGES | • ST-004 |
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | The section's score, its blocking defects, and which later sections depend on what is missing. |
| ESCALATION | None, and no substitute. While the expert is away the section stays where it is. |
| ACCEPTANCE | • Given a section below 70 or carrying a blocking defect, when the expert chooses to move on, then it is written at `draft` and what was left behind is recorded as an open question or named plainly to them. |

### TR-010 · Decide the set is ready to be offered for approval

| | |
|---|---|
| STAGES | • ST-005 |
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | Every section's state and grade, every blocking defect, and the buildability report's gaps with, for each, whether it is closed or owned. |
| ESCALATION | None, and no substitute. The expert offers the set; nobody offers it for them. |
| ACCEPTANCE | • Given a buildability gap that is neither closed nor given an owner and a date, when the expert offers the set for approval, then that gap is stated to them first.<br>• Given a set offered for approval, when the approver opens it, then what the expert was shown under this rule is what the approver sees. |

### TR-011 · Decide a backlog can be mapped from the set

| | |
|---|---|
| STAGES | • ST-007 |
| APPROVER | PER-004 |
| AN AGENT MAY DECIDE | yes |
| NEVER DELEGATED | no |
| EVIDENCE REQUIRED | Whether a golden thread exists, and which release-line stages have requirements behind them. |
| ESCALATION | The delivery lead's own judgement stands over any reported verdict; where they are unavailable the workshop waits rather than proceeding without a backbone. |
| ACCEPTANCE | • Given a set with no journey, when the check runs, then it reports that a backlog cannot be derived, and names that as the reason. |

### TR-012 · Decide which epics come first

| | |
|---|---|
| STAGES | • ST-008 |
| APPROVER | PER-004 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | The release line, the dependencies in section 11, and the objective each epic moves. |
| ESCALATION | The delivery lead may re-order within the release line at any time; order never changes scope, which stays with TR-004. While they are away the order stands as last recorded, and nobody re-orders on their behalf. |
| ACCEPTANCE | • Given an ordering, when it is recorded, then each decision names whether value, dependency or the release line drove it. |

### TR-013 · Accept a drafted section as saying what the expert meant, or correct it

| | |
|---|---|
| STAGES | • ST-003 |
| APPROVER | PER-001 |
| AN AGENT MAY DECIDE | no |
| NEVER DELEGATED | yes |
| EVIDENCE REQUIRED | The drafted section in readable form, and its grade, so the expert judges the words and not the score. |
| ESCALATION | None, and no substitute. An unaccepted draft stays at `draft`; nobody accepts a section's wording on the expert's behalf while they are away. |
| ACCEPTANCE | • Given a drafted section the expert has not accepted, when the set's state is read, then that section is not shown as ready for anyone else to read.<br>• Given an expert who corrects a drafted section, when the correction is applied, then it is their words that are kept, and the section is graded again before it advances. |

SECTION 11

# Assumptions, risks and dependencies

DRAFT 75/100 NO BLOCKING DEFECTS

## Assumptions

| ID | ASSUMPTION | IF FALSE |
|---|---|---|
| A-001 | An expert will answer 50-100 questions across a few sittings | The method fails; sections 5-7 change completely |
| A-002 | An expert will act on a grade they did not produce | The grading is decoration |
| A-003 | An engineering lead accepts a PRD naming no technology | Test 5 becomes a liability |
| A-004 | Organisations permit drafting inside an assistant session | Section 6's separate-application non-goal and A-005 have to change |
| A-005 | Organizations accept that domain content stays in the session and the documents | NFR-008, NFR-010, and NFR-012 have to change |

A-001 and A-002 are the two the product stands on, and neither has been tested. Both are validated by the same ten observed sessions (D-002), which makes D-002 the single most load-bearing dependency in this PRD.

A-005 was constraint C-003 in section 9 until 0.8.1. It moved here because it is a belief about what customers accept, which a customer can falsify, and not a rule imposed from outside.

## Risks

| ID | RISK | L | I | OWNER |
|---|---|---|---|---|
| R-001 | High score becomes false assurance | high | high | Richard Hightower |
| R-002 | Unsupported findings destroy trust in all findings | high | high | Richard Hightower |
| R-003 | A beautifully graded PRD for the wrong product | medium | high | Richard Hightower |
| R-004 | Interview too long; sets abandoned | medium | medium | Richard Hightower |
| R-005 | A version change strands an expert's set | low | high | Richard Hightower |

**R-003 is not closed.** The mitigations listed reduce it and do not solve it. Nothing in this product checks whether the product being specified should exist. This section states it rather than hide it behind a mitigation that sounds adequate.

## Dependencies

| ID | NEED | FROM | BY | STATUS |
|---|---|---|---|---|
| D-001 | An assistant session able to carry the interview | The customer's platform | 2027-01-31 | committed |
| D-002 | 10 observed first-time experts | Early customers | 2027-03-31 | requested |

| ID | NEED | FROM | BY | STATUS |
|---|---|---|---|---|
| D-003 | 3 engineering leads to read a set cold | Early customers | 2027-03-31 | requested |

D-002 and D-003 are both merely requested, and both are dated on R1's target day in section 13. This is a real schedule problem: R1 ships on the day the assumptions it rests on are due to be tested, so it cannot use the results. Recorded, not resolved.

## Assumptions

### A-001

| | |
|---|---|
| ASSUMPTION | A domain expert will answer fifty to a hundred questions about their own field across a few sittings, if each one is clearly worth answering. |
| IF FALSE | The whole method fails. The product would have to produce a usable partial PRD from perhaps fifteen answers, which changes sections 5, 6 and 7 completely and makes the grading thresholds meaningless. |
| VALIDATED BY | Observed sessions with 10 first-time experts, counting where they stop - the same sessions NFR-005 uses. |
| OWNER | Richard Hightower |

### A-002 · An expert will act on a grade they did not produce, rather than dismissing it.

| | |
|---|---|
| IF FALSE | The grading is decoration and sections 3 and 7 are measuring something nobody uses. The product reduces to a structured interview. |
| VALIDATED BY | In the same observed sessions, whether the expert answers the one question or moves on regardless. |
| OWNER | Richard Hightower |

### A-003

| | |
|---|---|
| ASSUMPTION | An engineering lead will accept a PRD that names no technology, rather than reading the absence as incompleteness. |
| IF FALSE | Test 5 becomes a liability rather than a discipline, and the candidate architecture decisions list in FR-007 has to become a first-class output rather than a hand-over note. |
| VALIDATED BY | Three engineering leads reading a completed set cold and saying what they would ask for first. |
| OWNER | Richard Hightower |

### A-004

| | |
|---|---|
| ASSUMPTION | Organisations will let a specification for their product be drafted inside an assistant session. |
| IF FALSE | The separate-application non-goal in section 6 and A-005 are unacceptable as written and the delivery model has to change, which invalidates most of section 7's assumptions about immediacy. |
| VALIDATED BY | The information security review of the first three customer engagements. |
| OWNER | Richard Hightower |

### A-005

| | |
|---|---|
| ASSUMPTION | Organizations accept that the domain content an expert enters stays in the expert's own session and in the documents. This product sends it nowhere else and keeps it nowhere else. |
| IF FALSE | An organization whose information policy needs more than this cannot adopt the product as written, and NFR-008, NFR-010, and NFR-012, which are measured against this boundary, have to change. |
| OWNER | Richard Hightower |
| CATEGORY | viability |
| EVIDENCE | none |
| VALIDATE BY | 2027-03-15 |

## Risks

### R-001

| | |
|---|---|
| RISK | The expert accepts a weak PRD because the score was high. The grade becomes false assurance, which is worse than no grade. |
| LIKELIHOOD | high |
| IMPACT | high |
| MITIGATION | Blocking defects are counted and reported separately from the score and no score clears them (FR-009); the non-goal in section 6 says explicitly that a grade is evidence, not a warrant; FR-012 puts a buildability verdict in front of the expert before approval. |
| OWNER | Richard Hightower |

### R-002

| | |
|---|---|
| RISK | The grading reports findings that are not in the document. The expert catches one, stops trusting all of them, and the grading becomes noise they skip. |
| LIKELIHOOD | high |
| IMPACT | high |
| MITIGATION | Every finding must quote the document, and anything unsupported is dropped before the expert sees it (FR-016, TR-003, TR-006); a refuted finding is acknowledged as dropped rather than defended (SC-004). |
| OWNER | Richard Hightower |

### R-003

| | |
|---|---|
| RISK | A fluent, well-structured, well-graded PRD is produced for the wrong product, and the structure makes the error harder to see rather than easier. |
| LIKELIHOOD | medium |
| IMPACT | high |
| MITIGATION | Section 3 asks what would move the objectives without this product at all; section 1 requires evidence per problem and names what is unmeasured. Neither is sufficient - this risk is not closed, and a reviewer should say so. |
| OWNER | Richard Hightower |

### R-004

RISK
The interview is too long to finish, so sets are abandoned half-written and the product is remembered as the thing that asked a hundred questions.

LIKELIHOOD
medium

IMPACT
medium

MITIGATION
FR-005 lets the expert end any section and records what was left; FR-015 makes resumption cheap; the working order in the method puts the sections that unlock the others first, so an abandoned set is still useful.

OWNER
Richard Hightower

### R-005

RISK
A set written by one version cannot be read or graded by a later one, and an expert loses weeks of thinking to an upgrade.

LIKELIHOOD
low

IMPACT
high

MITIGATION
NFR-004 makes every set readable with no part of the product running; NFR-009 requires older sets to stay gradable within a major version.

OWNER
Richard Hightower

## Dependencies

### D-001

NEED
An assistant session the expert is already working in, able to carry a long interview and read and write the documents.

FROM
The assistant platform the customer already uses

BY
2027-01-31

STATUS
committed

BLOCKS
- ST-001
- ST-002

### D-002 · Ten first-time domain experts willing to be observed completing a PRD.

FROM
Early customers and the professional network

BY
2027-03-31

STATUS
requested

BLOCKS
- A-001
- A-002
- NFR-005

### D-003

| | |
|---|---|
| NEED | Three engineering leads willing to read a completed set cold and report what they would ask for first. |
| FROM | Early customers |
| BY | 2027-03-31 |
| STATUS | requested |
| BLOCKS | • A-003<br>• O-001 |

SECTION 12

# Open questions

DRAFT 71/100 NO BLOCKING DEFECTS

| ID | QUESTION | DECIDER | BY | BLOCKS |
|---|---|---|---|---|
| Q-001 | Today's first-pass buildability rate | R. Hightower | 2027-03-31 | O-001 |
| Q-002 | Today's idea-to-approval elapsed time | R. Hightower | 2027-03-31 | O-002 |
| Q-003 | Today's share of testable requirements | R. Hightower | 2027-03-31 | O-003 |
| Q-004 | Segment sizes, and which converts | R. Hightower | 2027-06-30 | S-001, S-002 |
| Q-005 | Will a sponsor accept a graded specification | R. Hightower | 2027-01-31 | FR-013, TR-001 |
| Q-006 | Sections completed per sitting, and where they stop | R. Hightower | 2027-03-31 | R-004, NFR-005 |
| Q-007 | Is the working order right, and how does looping back feel | R. Hightower | 2027-03-31 | ST-004 |
| Q-008 | Disciplined or incomplete, to an engineering lead | R. Hightower | 2027-03-31 | A-003 |
| Q-009 | Where the rubric and the checks live, and who owns them | R. Hightower | 2027-01-31 | FR-004, FR-009..FR-012, FR-019, NFR-007 |
| Q-010 | Where an answer lives before its section is drafted | R. Hightower | 2027-01-31 | FR-015, NFR-011 |
| Q-011 | How an approver records a decision, and who owns NFR-006/012 | R. Hightower | 2027-02-28 | FR-014, NFR-006, NFR-012 |

### The three that matter most

**Q-001 to Q-003 are the baselines for all three objectives.** Until they are answered, section 3's targets are intentions. Nothing should be committed to a plan on the strength of them.

**Q-009, Q-010 and Q-011 came from the buildability review at 0.2.0**, and all three block `must` requirements on the release line. Q-009 is the largest: six requirements are written against a rubric whose shape the set defines but whose content it does not contain, so two of the six release-line stages can be scaffolded and not finished until it is answered. Q-010 is the one where guessing breaks assumption A-005 rather than costing rework.

**Q-005 is the one that could invalidate the product's central idea.** It was dated after the release it blocks until 0.2.0 moved it to 2027-01-31; it is a two-day question and there was no defence for asking it late. If showing a sponsor a

grade reduces their confidence rather than raising it, then FR-013 is wrong and the grading becomes an author's private tool rather than part of the approval. Q-005 is a cheap two-day question, and it should be asked early.

**Q-008 tests test 5.** If engineering leads read the absence of technology as incompleteness, the discipline this whole product enforces is a liability, and A-003 fails.

### Honest note

Every decider here is the same person. For a one-person product that is accurate; for anything larger it is a defect, and the first thing to fix when there is a second person.

#### Q-001

| | |
|---|---|
| QUESTION | What share of specifications reaching a development team today pass a buildability review on first submission? |
| DECIDER | Richard Hightower |
| BY | 2027-03-31 |
| BLOCKS | • O-001 |
| SPIKE | Run the eleven derivability checks against ten specifications collected from real engagements, and count. Time-boxed to 3 days. |

#### Q-002

| | |
|---|---|
| QUESTION | How long does it take today, in working days, from a product idea to an approved specification? |
| DECIDER | Richard Hightower |
| BY | 2027-03-31 |
| BLOCKS | • O-002 |
| SPIKE | Reconstruct the timeline from first discussion to approval for five recent engagements. Time-boxed to 2 days. |

#### Q-003

| | |
|---|---|
| QUESTION | What share of functional requirements in today's specifications can an engineer who was not in the interview write a test from? |
| DECIDER | Richard Hightower |
| BY | 2027-03-31 |
| BLOCKS | • O-003 |
| SPIKE | Three engineering leads score the requirements of the same ten specifications from Q-001, independently. Time-boxed to 3 days. |

#### Q-004

| | |
|---|---|
| QUESTION | How many organisations are in each segment in section 2, and which segment actually converts? |
| DECIDER | Richard Hightower |
| BY | 2027-06-30 |
| BLOCKS | • S-001<br>• S-002 |
| SPIKE | Name the register or source per segment and count; then compare against the first ten conversations. Time-boxed to 5 days. |

### Q-005

**QUESTION** Will a sponsor accept a specification that carries a machine-produced grade, or does the grade reduce their confidence rather than raise it?

**DECIDER** Richard Hightower

**CONSULTED** the first three sponsors

**BY** 2027-01-31

**BLOCKS**

- FR-013
- TR-001

**SPIKE** Show three sponsors the same set with and without grades and ask which they would approve. Time-boxed to 2 days.

### Q-009

**QUESTION** Where do the rubric and the eleven derivability checks live, who owns them, and how is a set told which version it was graded against?

**DECIDER** Richard Hightower

**BY** 2027-01-31

**BLOCKS**

- FR-004
- FR-009
- FR-010
- FR-011
- FR-012
- FR-019
- NFR-007

**SPIKE** Decide whether they ship as documents of the product, name their owner, and decide whether a grade records the version it was made against. Time-boxed to 2 days. Until it is answered, ST-003 and ST-004 can be scaffolded and not finished.

### Q-010

**QUESTION** Where does an answer live between the moment the expert gives it and the moment the section is drafted, given A-005, that content is retained nowhere outside the documents, and given that NFR-004 requires the set readable with no part of the product running?

**DECIDER** Richard Hightower

**BY** 2027-01-31

**BLOCKS**

- FR-015
- NFR-011

**SPIKE** Decide whether an in-progress interview is itself a document of the set, or whether answers before a draft are not retained at all and FR-015 holds only within a session. Time-boxed to 2 days. Guessing this one breaks assumption A-005, which customers rely on, rather than merely costing rework.

### Q-011

**QUESTION** How does an approver record a decision, given that D-001 puts the product inside the expert's session and OUT-003 excludes any view of our own, and who owns NFR-006's accessibility and NFR-012's access control, which describe things the non-goals in section 6 say we do not own?

**DECIDER** Richard Hightower

**BY** 2027-02-28

**BLOCKS**

- FR-014
- NFR-006
- NFR-012

**SPIKE** Walk one approval end to end with a real sponsor and record every step they take that the product does not provide. Time-boxed to 3 days.

### Q-006

**QUESTION** How many sections will an expert complete in one sitting before stopping, and which section do they stop at?

**DECIDER** Richard Hightower

**BY** 2027-03-31

**BLOCKS**

- R-004
- NFR-005

**SPIKE** Count in the ten observed sessions from D-002. No extra work.

### Q-007

**QUESTION** Is the working order in the method right - does writing the requirements really send an expert back to the journeys, and does that feel like progress or like failure to them?

**DECIDER** Richard Hightower

**BY** 2027-03-31

**BLOCKS**

- ST-004

**SPIKE** Count loop-backs and ask the expert how it felt, in the ten observed sessions. No extra work.

### Q-008

**QUESTION** Does an engineering lead read a PRD that names no technology as disciplined or as incomplete?

**DECIDER** Richard Hightower

**CONSULTED** the three engineering leads from D-003

**BY** 2027-03-31

**BLOCKS**

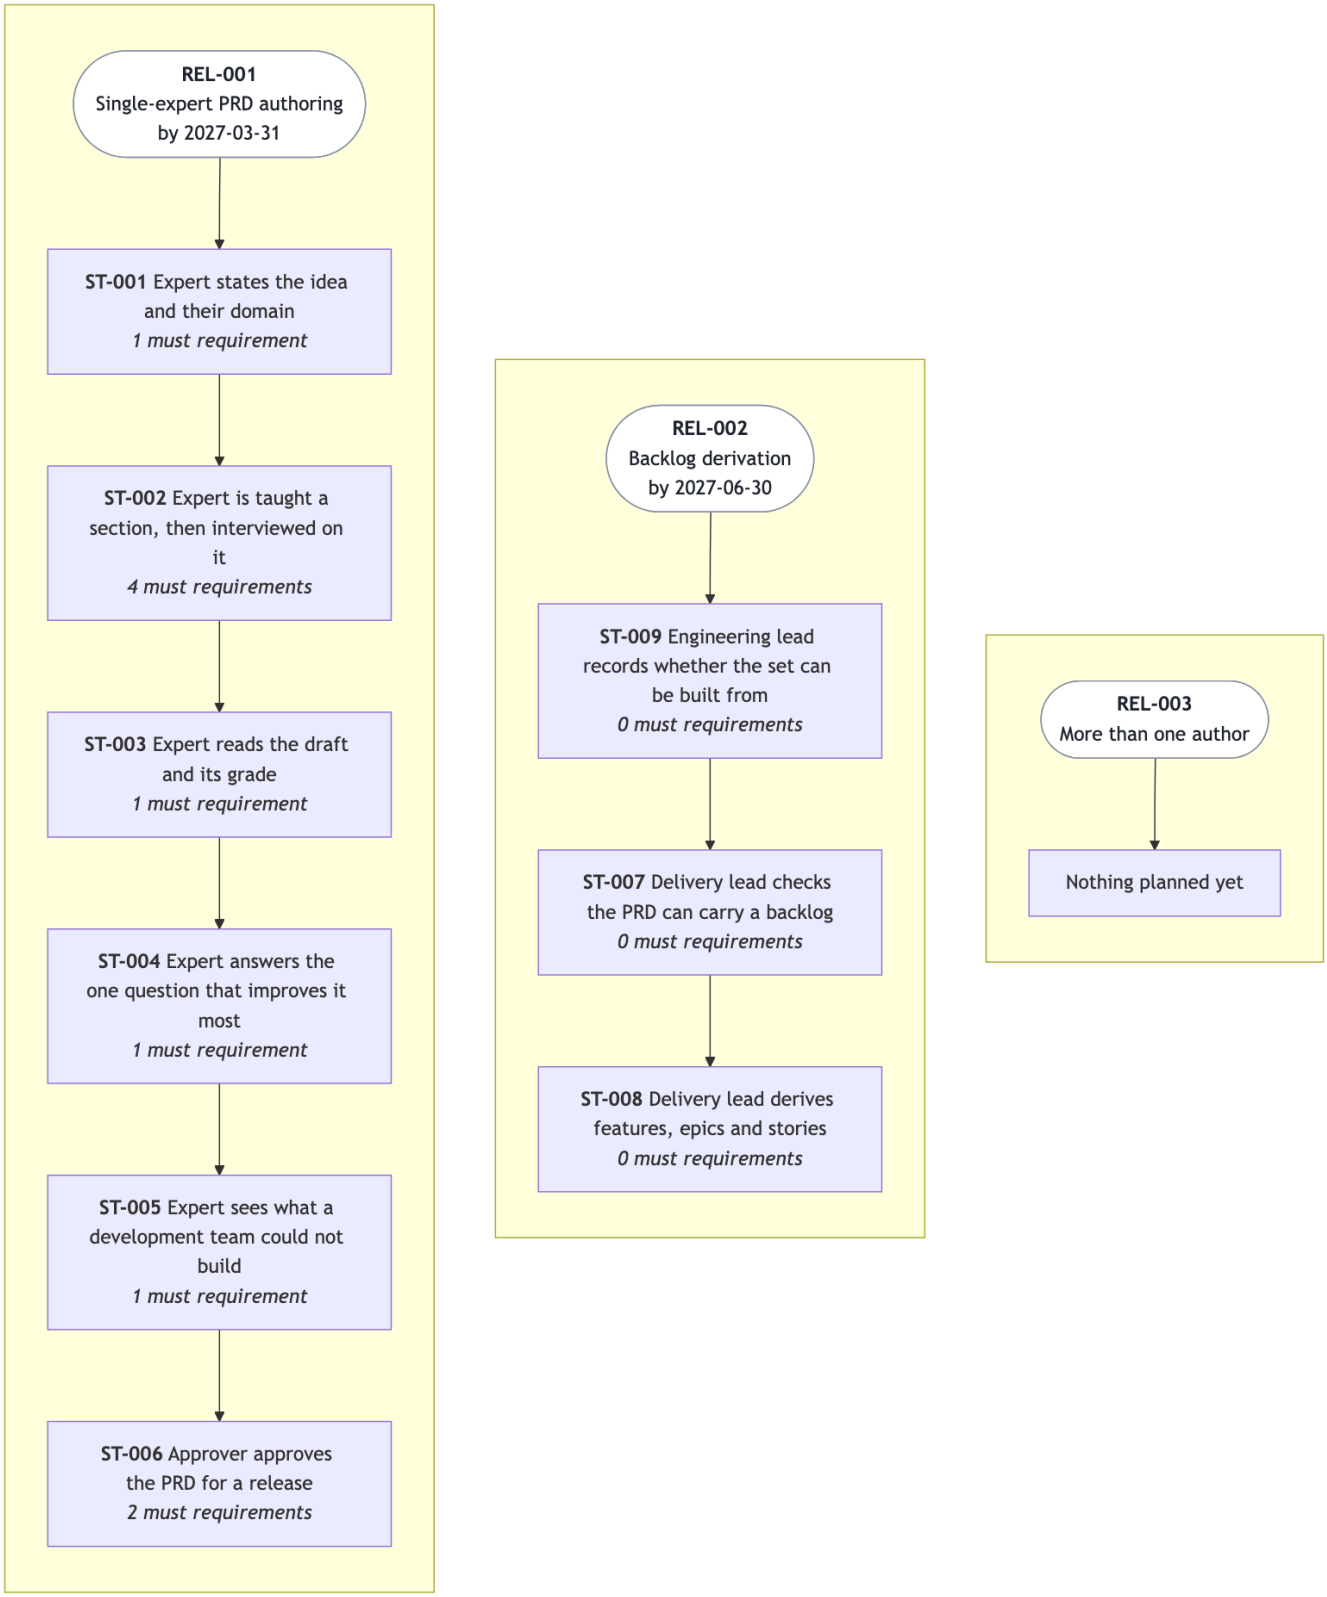
- A-003

**SPIKE** Ask each lead, after reading a set cold, whether anything was missing and whether they wanted a technology named. Time-boxed to 1 day.

SECTION 13

# Release plan

DRAFT 83/100 NO BLOCKING DEFECTS




The release line

| ID | RELEASE | STAGES | TARGET |
|---|---|---|---|
| REL-001 | Single-expert PRD authoring | ST-001..ST-006 | 2027-03-31 |
| REL-002 | Backlog derivation | ST-009, ST-007, ST-008 | 2027-06-30 |
| REL-003 | More than one author | none | none |

### Why this order

REL-001 is exactly the release line in section 6. It is the smallest release in which anyone gets anything: an expert reaching an approved, reviewed PRD.

REL-002 is deferred rather than cut because the PRD REL-001 produces already has to carry everything a backlog needs, and FR-012 checks precisely that. Deferring the derivation costs nothing; deferring what it needs would have cost a rewrite.

REL-003 waits for evidence that one author's path works at all. Building for two authors before one is proven would be building on A-001 and A-002 before either is tested.

### How you know each release worked

REL-001 exits when the pilot PRDs meet the three objective targets in section 3: 80% pass the buildability review on first submission, the median reaches approval in under 5 working days, and 95% of `must` requirements are testable as written. REL-002 exits when every derived story names its requirement and a delivery lead runs a workshop from the backlog without re-deriving it.

Both target dates are indicative, not committed. The schedule problem below is the reason.

### The schedule problem, stated plainly

REL-001 targets 2027-03-31. D-002 and D-003 in section 11, the ten observed experts and the three engineering leads, are also dated 2027-03-31 and are only *requested*. As a result, REL-001 ships at the same moment the assumptions beneath it are first tested, not after. Either REL-001 moves, or it ships knowingly as an experiment. The sponsor makes this decision under TR-001, and it is not made here.

### Not in this section, deliberately

No tasks, no estimates, no sizing, no team assignments, and no technology milestone. Those come from engineering design and iteration planning.

> **REL-001 · Single-expert PRD authoring**
>
> OUTCOMES
>
> - A domain expert with no product training can take one product idea to an approved PRD without leaving the session they started in.
> - The expert can see, before anyone estimates, which parts of their own specification a development team could not build from.
> - A sponsor can see what they are approving, what is still open and who owns each open thing, on one page.

**JOURNEY STAGES**

- ST-001
- ST-002
- ST-003
- ST-004
- ST-005
- ST-006

**REQUIREMENTS**

- FR-001
- FR-002
- FR-003
- FR-004
- FR-005
- FR-006
- FR-007
- FR-008
- FR-009
- FR-010
- FR-011
- FR-012
- FR-013
- FR-014
- FR-015
- FR-016
- FR-017
- FR-018
- FR-019
- FR-020

**OBJECTIVES**

- O-001
- O-002
- O-003

**EXIT CRITERIA**

- At least 80% of the pilot PRDs pass the buildability review on first submission (O-001).
- The median pilot PRD reaches a recorded approval in under 5 working days from the first interview question (O-002).
- An engineer who was not in the interview judges 95% of the pilot PRDs' `must` requirements testable as written (O-003).

**TARGET DATE** 2027-03-31

**DATE STATUS** indicative

### REL-002 · Backlog derivation

**OUTCOMES**

- A delivery lead can walk into a story-mapping workshop with features and a release line already derived from the approved PRD.
- Every story names the requirement id it came from, so nobody has to ask why it exists.

**JOURNEY STAGES**

- ST-009
- ST-007
- ST-008

**OBJECTIVES**

- O-001

**EXIT CRITERIA**

- Every story derived from a pilot PRD names the requirement id it came from.
- A delivery lead runs a story-mapping workshop from a derived backlog without re-deriving any feature from the PRD.

**TARGET DATE** 2027-06-30

**DATE STATUS** indicative

### REL-003 · More than one author

**OUTCOMES**

- Two experts can work on different sections of the same specification without overwriting each other.
- A reviewer can comment on a section without being able to change it.

**OBJECTIVES**

- O-002

SECTION 14

# Glossary

DRAFT 75/100 NO BLOCKING DEFECTS

| TERM | MOST OFTEN CONFUSED WITH |
|---|---|
| PRD | A design document |
| Golden thread | The most important journey |
| Release line | A date |
| Stage | A screen |
| Rubric | A checklist |
| Blocking defect | A low score |
| Score | A judgement on the product idea |
| Spike | A piece of the product |
| Open question | A risk |
| Non-goal | Out of scope |
| Grade | Score |
| State | Threshold |
| Assumption | An open question |
| Candidate architecture decision | A requirement |
| Buildability review | A grading pass |
| Finding | A blocking defect |
| Root document | Section 1 |
| Approver | The expert |
| Condition | A constraint |
| Constraint | A requirement |
| Objective | A requirement |
| Baseline | A target |
| Version | A release of the product |
| Dependency | A requirement's wait on an open question |
| Backlog | The set |
| Feature | A requirement |

| TERM | MOST OFTEN CONFUSED WITH |
|---|---|
| Story | A requirement |
| Expert | The approver |
| Sponsor | The expert |
| Engineering lead | The expert |
| Delivery lead | The engineering lead |
| Item | A line or a field of an entry |
| Scenario | A journey |
| Objection | A refuted finding |
| Gate | The threshold |
| Out of scope | A non-goal |
| Outcome | A feeling |
| Approval | An override |
| Engagement | A unit of work |
| Handoff | The handover |
| Journey | A use case |
| Step | A stage |
| Override | An approval |
| Release | The release line |
| Non-functional requirement | A requirement in section 7 |
| Acceptance criteria | Exit criteria |
| Epic | A feature |
| Vocabulary | The glossary |
| Persona | A user |
| Buyer | The user |

The "Most often confused with" column above is the one that earns its keep. Every one of these has been confused in practice, and three of them, golden thread, blocking defect and score, carry a cost when confused: a team picks the wrong journey for release 1, ships a section with a known fault because the number looked good, or believes a well-written specification means a sound product.

## References

Eight terms come from established practice, and each entry's `source` names where. Where this set departs from the source, the entry says so.

- MoSCoW: "What is MoSCoW Prioritization?", Agile Business Consortium, accessed 2026-10-27, https://www.agilebusiness.org/?p=1907
- Given/When/Then: "Introducing BDD", Dan North, 20 Sep 2006, https://dannorth.net/blog/introducing-bdd/
- Story, epic, and feature: "Epics, Features and User Stories", Mountain Goat Software, last updated 28 Apr 2025, https://www.mountaingoatsoftware.com/blog/stories-epics-and-themes
- Story map and backbone: "The New User Story Backlog is a Map", Jeff Patton & Associates, 8 Oct 2008, updated 12 Apr 2023, https://www.jpattonassociates.com/the-new-backlog/
- Spike: "Spike Solutions", in The Art of Agile Development, 1st ed., James Shore and Shane Warden, 3 Jun 2010, https://www.jamesshore.com/v2/books/aoad1/spike_solutions


### PRD

**MEANING** A written statement of what to build, for whom, why, and how you know it works, plus what not to build. It does not say how to build it.

**SEE ALSO**
- section
- document set

**NOT** A design document. A PRD that names components or technologies has stopped being a PRD and has started being an architecture sketch.

### Document set

**MEANING** The documents that together are one PRD: a root document and one document per section, each linking to the others.

**SEE ALSO**
- PRD

### Golden thread

**MEANING** The one user journey that release 1 must deliver from end to end.

**SEE ALSO**
- journey
- release line

**NOT** The most important journey in general. It is the journey chosen for release 1, which is often the simplest rather than the most valuable.

### Release line

**MEANING** The smallest set of journey stages that completes the golden thread, and so the boundary of what release 1 contains.

**SEE ALSO**
- golden thread
- scope

**NOT** A date or a deadline. It is a line across the journey, not across a calendar.

### Stage

**MEANING** One large step of a journey, naming a person, an input, an output and a decision.

**SEE ALSO**
- journey
- step

**NOT** A screen or a system step. A stage survives a redesign.

### The five tests

**MEANING** The five numbered properties a statement must have to be a product requirement. Test 1, it names a user or a stakeholder. Test 2, it states an outcome a person can see. Test 3, it has a pass-or-fail acceptance test. Test 4, it has a priority. Test 5, it does not name the component or the technology that delivers it. The numbers are part of the term: the set refers to "test 5" by number.

**SEE ALSO**

- blocking defect

### Rubric

**MEANING** The stated standard a section is scored against: a set of dimensions, each with written anchors for what a score of 0, 2 and 4 look like.

**SEE ALSO**

- score
- blocking defect

**NOT** A checklist. A checklist asks whether something is present; a rubric asks how good it is, against written anchors.

### Blocking defect

**MEANING** A named fault that stops a section advancing however high its score is, such as a requirement that names a technology or an id that refers to nothing.

**SEE ALSO**

- rubric
- score

**NOT** A low score. Score and blocking defects are counted separately and a high score never clears a blocking defect.

### Score

**MEANING** A section's total against its rubric, 0 to 100, written by a grading pass and never by hand.

**SEE ALSO**

- rubric
- blocking defect

**NOT** A measure of whether the product is a good idea. It measures how well the section is written, nothing else.

### Spike

**MEANING** A time-boxed piece of work whose only output is an answer to an open question, or a recorded decision.

**SEE ALSO**

- open question

**NOT** A piece of the product. A spike produces knowledge, not a feature.

**SOURCE** Spike Solutions, The Art of Agile Development, 1st ed., James Shore and Shane Warden, 3 Jun 2010, https://www.jamesshore.com/v2/books/aoad1/spike_solutions

### Open question

**MEANING** Something not yet decided, recorded with the one person who decides it and the date by which they must.

**SEE ALSO**

- spike
- assumption

**NOT** A risk. A risk is something that might happen and nobody can simply decide; an open question has a decider.

### Non-goal

**MEANING** Something a reasonable reader would expect the product to do, that it will not do, with no plan to.

**SEE ALSO**
- scope

**NOT** Out of scope. Out of scope means not yet, with a point at which it is revisited. A non-goal means not ever, with a reason.

### Journey

**MEANING** The steps one user takes to reach one goal with the product, from start to end, in order.

**SEE ALSO**
- stage
- golden thread

**NOT** A use case. A journey covers the whole path to a goal, including steps outside the product.

### Step

**MEANING** What a user does, sees or decides at one point inside a stage.

**SEE ALSO**
- stage
- journey

**NOT** A stage. A stage holds many steps.

### Section

**MEANING** One of the fifteen numbered parts of a PRD, each held in its own document. Section 5 holds two documents, journeys and scenarios, and section 7 holds two, functional requirements and wireframes, so a set has seventeen section documents across fifteen sections. Cross-references in a set are by section number.

**SEE ALSO**
- document set

### Scope

**MEANING** What a release contains and what it does not, set against the journey stages it completes.

**SEE ALSO**
- release line
- non-goal

### Assumption

**MEANING** Something taken on faith that the product depends on, recorded with what would break if it turned out false.

**SEE ALSO**
- open question
- risk

**NOT** An open question. Nobody can simply decide an assumption true; it is validated or falsified.

### Risk

**MEANING** Something that may happen and would damage the product, recorded with a likelihood, an impact, a mitigation and one named owner.

**SEE ALSO**
- assumption
- spike

### Readable form

MEANING

The part of a section a person reads: its prose, its headings and its tables, written to be judged by the expert who gave the answers.

SEE ALSO

- section

### Threshold

MEANING

The score at or above which a section may advance: 70. A score at the threshold is not enough on its own, because the gate also needs no blocking defect.

SEE ALSO

- score
- blocking defect

### Grade

MEANING

A section's score together with its blocking-defect count and its weakest dimension, the rubric dimension that scored lowest, which the set also calls the section's weakest part. A grade is what an expert is shown; a score is one number inside it.

SEE ALSO

- score
- grading pass

NOT

A score. The score is part of the grade, not the whole of it.

### Grading pass

MEANING

One run of a section or a set against its rubric, producing a grade. It is the only thing that writes a score, and it never changes the document.

SEE ALSO

- grade
- rubric

### State

MEANING

Where a section stands, one of four in this order: `pending` (not yet interviewed, and shown to the expert as not started), `draft` (written), `review` (ready for someone else to read), `final` (decided by the expert). A section is offered `review` at a score of 70 or more with no blocking defects, and moves there when the expert accepts. It reaches `final` when the expert sets it final under TR-002. Approval under TR-001 covers the whole set, not one section.

NOT

A threshold. The threshold is the score a gate needs. The state is where the section stands.

SEE ALSO

- threshold
- blocking defect

### Override

MEANING

An expert's decision to advance a section past a gate it has not met, recorded in the change log with the score at the time.

SEE ALSO

- state
- threshold

NOT

An approval. An override moves one section past its gate, and an approval covers the whole set.

### Verdict

**MEANING**

What a buildability review concludes, one of three: `buildable`, `buildable with gaps`, `not buildable`. In ST-009 the engineering lead may record their own verdict, and where it differs, theirs is the verdict and the review's is kept beside it.

**SEE ALSO**

- derivability checks

### Derivability checks

**MEANING**

The eleven checks a buildability review runs to decide whether a backlog can be derived from a set. They are product content, shipped and versioned with the product, and the set states their count and their verdict vocabulary rather than restating the list.

**SEE ALSO**

- verdict
- rubric

### Candidate architecture decision

**MEANING**

A technology or design preference the expert stated during the interview, kept out of the requirements and handed to the engineering lead as a topic for an architecture decision record.

**SEE ALSO**

- the five tests

**NOT**

A requirement. It is deliberately not one, because it would fail test 5.

### Unit of work

**MEANING**

What the expert's domain counts one piece of work as: a load, a claim, a case, a patient. Every question the expert is asked uses their word for it.

### Unverified

**MEANING**

A citation that was checked and could not be confirmed against its source. It is marked as such and never stated as fact.

**SEE ALSO**

- the five tests

### MoSCoW

**MEANING**

The priority scale this set uses: `must`, `should`, `could`, `wont`. This set uses MoSCoW and not RFC 2119's MUST and SHOULD. `wont` keeps the source's meaning, won't have this time: a `wont` item is out of scope and revisited later, not a non-goal.

**SEE ALSO**

- the five tests
- out of scope

**SOURCE**

What is MoSCoW Prioritization?, Agile Business Consortium, accessed 2026-10-27, https://www.agilebusiness.org/?p=1907

### Wireframe

**MEANING**

A sketch of one screen a person works in or one report a person reads, labeled an illustration, not a design. It shows what the person sees and does there, and names the requirements it makes concrete. Every requirement shown in a screen or a report has one, in section 7.

**SEE ALSO**

- the five tests

### Buildability review

**MEANING**

A reading of the whole set as the engineering lead who has to build it. It runs the eleven derivability checks and reports a verdict and, for each release-line stage, whether a requirement covers it or it is a gap.

**SEE ALSO**

- verdict
- derivability checks
- finding

**NOT**

A grading pass. A grading pass scores one section against its rubric; a buildability review reads the whole set and asks whether a backlog can be derived from it.

### Finding

**MEANING**

One fault that a grading pass or a buildability review reports, with the document it is in quoted. A finding is checked against the document before the expert sees it, and is dropped when the document does not support it.

**SEE ALSO**

- blocking defect
- grading pass
- buildability review

**NOT**

A blocking defect. A blocking defect is one kind of finding, the kind that stops a section advancing. Every blocking defect is reported as a finding, but most findings only lower a score or name a weak dimension.

### Root document

**MEANING**

The one document of a set that every other document names as its parent. It carries the summary, the scope, every section's state and grade, and every open question with its decider and date, and it is what an approver reads.

**SEE ALSO**

- document set
- approver

**NOT**

Section 1. The root document has no section number; section 1 is the problem and opportunity.

### Approver

**MEANING**

The one person who decides whether a set is approved for a release: the sponsor, or the expert when no sponsor exists. They approve, approve with conditions, or reject, and the decision is recorded in section 15 against a named version. An approver who approves with conditions approves the set on the conditions they state, and the set reads as approved with conditions until that same approver clears every one.

**SEE ALSO**

- condition
- root document
- trust rule

**NOT**

The expert. The expert decides the state of each section; the approver decides on the whole set. They are one person only when no sponsor exists.

### Condition

**MEANING**

What an approver who approves with conditions requires before the approval is unconditional. Each condition names the open question or requirement it depends on and a date, and only the approver who stated it clears it.

**SEE ALSO**

- approver
- open question

**NOT**

A constraint. A condition belongs to one approval and is cleared; a constraint holds for the product and is never cleared.

### Constraint

**MEANING**

A rule the product must obey that the team cannot change by agreeing among themselves, recorded in section 9 with its category, the authority that imposes it by name, and how compliance is shown. A regulation or a standard appears in a constraint only when the expert named it.

**SEE ALSO**

- condition

**NOT**

A requirement. A requirement states what a user can do and has a priority; a constraint limits how every requirement it applies to may be met, and is not prioritized or deferred.

### Objective

**MEANING**

A business result the product is expected to produce, recorded in section 3 with the problems it addresses, one metric, a baseline, a target with a date, and one owner. If the product ships and the metric does not move, the product did not work.

**SEE ALSO**

- baseline

**NOT**

A requirement. An objective says what the business gains; a requirement says what a user can do, and names the objective it serves.

### Baseline

**MEANING**

The value an objective's metric has today, before the product. When it is not known, the baseline says unknown and names the open question that will establish it, and the target is an intention rather than a commitment until the question is answered.

**SEE ALSO**

- objective
- open question

**NOT**

A target. The baseline is where the metric stands now; the target is where the objective says it will be by a date.

### Requirement

**MEANING**

A statement of what one named user can do with the product and the outcome they can see, recorded in section 7 with its acceptance criteria, a priority, the journey stage it serves, and the objective it serves. A statement is a requirement only when it passes the five tests.

**SEE ALSO**

- the five tests
- acceptance criteria
- objective

### Acceptance criteria

**MEANING**

The pass-or-fail statements on a requirement or a journey stage. Each is written as Given a situation, When the person acts, Then what they see, and each is one an engineer who was not in the interview could write a test from. A story derived from a requirement takes the requirement's criteria as its acceptance criteria.

**SEE ALSO**

- requirement
- stage
- the five tests

**NOT**

Exit criteria. Acceptance criteria decide whether one requirement or stage passes, and exit criteria decide whether a whole release is done.

**SOURCE**

Introducing BDD, Dan North, 20 Sep 2006, https://dannorth.net/blog/introducing-bdd/

### Gap

**MEANING**

Something the set does not say that a team needs in order to build: most often a release-line stage that no requirement covers, as a buildability review reports it. A gap is closed with a requirement, or owned as an open question with a decider and a date.

**SEE ALSO**

- buildability review
- release line
- open question

### Version

**MEANING**

A numbered state of the set, recorded in section 15 with its date, its author, and the sections it changed. An approval names the one version it covers, and a change written after an approval records a new version that the approval does not cover.

**SEE ALSO**

- approver
- condition

**NOT**

A release of the product. The product's releases are numbered too, and a set written by one release is read by later ones; a version of the set counts changes to the set.

### Domain

**MEANING**

The field of work the expert specifies software for, such as freight. The expert names it before the first section. It sets the unit of work, the vocabulary every question uses, and the examples each section is taught with. The expert supplies knowledge of the domain; the product does not.

**SEE ALSO**

- unit of work

### Reading order

**MEANING**

The order in which the documents of a set are read: the root document, then the section documents by section number, from section 1 to section 15. Where a section holds two documents, journeys come before scenarios and functional requirements come before wireframes. The next and the previous section of any document follow this order.

**SEE ALSO**

- section
- root document

### Interview

**MEANING**

The questions the expert answers about one section, in stage ST-002. It starts after the expert reads what the section is for, puts one or two questions at a time, and ends when the section has enough material to be drafted, when the expert says they have nothing more to add, or when their answers would already score the section 70 or more with no blocking defect. Only the expert answers in it. An engineer who reads the set later was not in the interview.

**SEE ALSO**

- section
- candidate architecture decision

### Handover

**MEANING**

The point at which the approved set reaches the engineering lead who has to build from it. Journey J-002 starts there: its trigger is that the PRD has been approved for a release, and its first stage, ST-009, takes the approved set and the report of its buildability review. A set is handed over only once it is approved.

**SEE ALSO**

- approver
- candidate architecture decision

### Dependency

**MEANING**

Something outside the product that the product relies on, recorded in section 11 with what happens if it is not there, such as D-001, an assistant session the expert is already working in.

**SEE ALSO**

- assumption
- risk

**NOT**

A requirement's wait on an open question. A requirement that waits on a question names the question, and the question lists the requirement in what it blocks.

### Backlog

**MEANING**

The features, epics and stories a delivery lead derives from an approved set, in stages ST-007 and ST-008 of journey J-002. Every level of it names the id it came from.

**SEE ALSO**

- feature
- epic
- story

**NOT**

The set. The set says what the product must do; the backlog is derived from it once it is approved.

### Feature

**MEANING**

The backlog item for one journey stage, carrying the ids of the requirements on that stage. The set departs from the source here: the source makes a feature any story big enough to be released, and this set makes it one journey stage.

**SEE ALSO**

- backlog
- stage

**NOT**

A requirement. A feature holds every requirement of one stage.

**SOURCE**

Epics, Features and User Stories, Mountain Goat Software, updated 28 Apr 2025, https://www.mountaingoatsoftware.com/blog/stories-epics-and-themes

### Epic

**MEANING**

A group of backlog items scoped by the release line, so one epic holds the items that one release delivers. The source defines an epic as a large user story. This set keeps that it is large, and fixes its size by the release line rather than by judgment.

**SEE ALSO**

- backlog
- release line

**NOT**

A feature. A feature covers one journey stage, and an epic covers the stages one release delivers.

**SOURCE**

Epics, Features and User Stories, Mountain Goat Software, updated 28 Apr 2025, https://www.mountaingoatsoftware.com/blog/stories-epics-and-themes

### Story

**MEANING**

The backlog item for one requirement, in the persona's own words, with the requirement's acceptance criteria as its own.

**SEE ALSO**

- backlog
- requirement
- acceptance criteria

**NOT**

A requirement. The requirement is in the set; the story is derived from it.

**SOURCE**

Epics, Features and User Stories, Mountain Goat Software, updated 28 Apr 2025, https://www.mountaingoatsoftware.com/blog/stories-epics-and-themes

### Expert

**MEANING**

The person with the domain knowledge, persona PER-001, who answers the interviews and owns the set. The expert decides what the product must do, the domain's own terms and the state of each section, and is the only one who may state a domain fact as true.

**SEE ALSO**

- domain
- interview
- sponsor

**NOT**

The approver. The expert approves their own set only when no sponsor exists, under TR-001.

### Sponsor

**MEANING**

The person who commits budget and a date, persona PER-003, and who approves the set for a release under TR-001. The sponsor may approve, approve with conditions, or reject.

**SEE ALSO**

- approver
- expert

**NOT**

The expert. Where no sponsor exists, the expert approves their own set under TR-001, and the change log records that no separate approver existed.

### Engineering lead

MEANING

The person who has to build from the approved set, persona PER-002. The engineering lead receives the set at handover, records their own verdict on its buildability review in ST-009, and receives each candidate architecture decision.

SEE ALSO

- handover
- buildability review
- candidate architecture decision

NOT

The expert. The engineering lead was not in the interview and builds from what the set says, not from what the expert meant.

### Delivery lead

MEANING

The person who turns an approved set into a backlog and runs the story-mapping workshop, persona PER-004, in stages ST-007 and ST-008.

SEE ALSO

- backlog
- engineering lead

NOT

The engineering lead. The delivery lead plans the work from the set, and the engineering lead builds it.

### Persona

MEANING

A named role from section 4, such as PER-001, the expert. Every requirement, journey, and trust rule names the persona it serves or who acts.

SEE ALSO

- expert
- journey

NOT

A user. Anyone who uses the product is a user, and a persona is one named role with its own goals.

### Item

MEANING

One numbered entry in a section with its own id, such as a problem, a persona, a stage, a requirement, or an open question. A section is made of its items and the prose around them.

SEE ALSO

- id
- section

NOT

A line or a field of an entry. A requirement and all its acceptance criteria are one item.

### Id

MEANING

The label that names one item, such as FR-008 or Q-001. Each id is unique in the set, and ids are never renumbered or reused.

SEE ALSO

- item
- link

### Link

MEANING

A reference from one item to another by its id, such as the journey stage a requirement serves or the requirements an open question blocks.

SEE ALSO

- id

### Priority

**MEANING**

How much a requirement or a scenario matters to a release, stated in MoSCoW terms. Only the expert gives a priority.

**SEE ALSO**

- moscow

### Teaching

**MEANING**

What the expert is told about a section before its interview: what it is, why it matters, and what belongs in it, in no more than ten lines and with an example from the expert's domain.

**SEE ALSO**

- interview
- domain

### Decider

**MEANING**

The one named person who makes one decision: who answers an open question by its date, or who decides under one trust rule in section 10, whose table names that person in its Decider column.

**SEE ALSO**

- open question
- trust rule

**NOT**

The approver. The approver decides only whether the whole set is approved for a release, under TR-001, and is the decider of that one rule.

### Release

**MEANING**

A numbered group of journey stages delivered together, such as REL-001, with its exit criteria. The first release holds exactly the release line.

**SEE ALSO**

- release line
- scope
- exit criteria

**NOT**

The release line. The release line is the smallest set of stages that completes the golden thread, and a release is what ships. Release 1 ships exactly the release line.

### Vocabulary

**MEANING**

The unit of work and the words the expert says an outsider would misread, each with the expert's own meaning. The expert states it once, before the first section. Every later question uses it, and it does not change when the expert resumes.

**SEE ALSO**

- unit of work
- domain

**NOT**

The glossary. The vocabulary is the few words the expert states at the start, and the glossary holds every domain term of the set.

### Engineer

**MEANING**

Anyone on the team who builds from the set and was not in the interview. The engineering lead is one. An acceptance criterion passes the five tests only when an engineer could write a test from it.

**SEE ALSO**

- engineering lead
- acceptance criteria

**AVOID**

- builder

### Scenario

**MEANING** One concrete path a persona takes through one or more journey stages, recorded in section 5 with its steps and the outcome they see. A scenario shows a stage working, or failing, with real content.

**SEE ALSO**
- journey
- stage

**NOT** A journey. The journey names the stages, and a scenario walks one path through them.

### Trust rule

**MEANING** A rule in section 10 that names who decides one decision, at which stages, on what evidence, and what happens when that person is away. A rule marked never delegated is always decided by a person.

**SEE ALSO**
- approver
- override

### Story-mapping workshop

**MEANING** The meeting in which the delivery lead and the team map a backlog against the journeys of an approved set. It needs a backbone agreed before it starts.

**SEE ALSO**
- backbone
- backlog
- delivery lead

**SOURCE** The New User Story Backlog is a Map, Jeff Patton & Associates, 8 Oct 2008, updated 12 Apr 2023, https://www.jpattonassociates.com/the-new-backlog/

### Backbone

**MEANING** The ordered journey stages that a story-mapping workshop maps features against. A set with no journey gives the workshop no backbone, so the workshop invents one in the room.

**SEE ALSO**
- journey
- story-mapping workshop

**SOURCE** The New User Story Backlog is a Map, Jeff Patton & Associates, 8 Oct 2008, updated 12 Apr 2023, https://www.jpattonassociates.com/the-new-backlog/

### Question list

**MEANING** The questions an interview may put for one section. The interview ends before the list does when the expert's answers would already pass the section.

**SEE ALSO**
- interview

### Gate

**MEANING** What a section must pass to move to its next state, such as a score of 70 or more with no blocking defect for `review`. The expert may override a gate, and the override is recorded.

**SEE ALSO**
- state
- override
- threshold

**NOT** The threshold. The threshold is the score alone, and a gate also needs no blocking defect.

### Out of scope

**MEANING** Something the product will not do in a release, with a point at which it is revisited, recorded in section 6.

**SEE ALSO**
- non-goal
- scope

**NOT** A non-goal. A non-goal is never done, and an out-of-scope item is not done yet.

### Objection

**MEANING** A finding that the document supports and that the expert disputes. Its reason is stated to the expert once, and the expert's decision is recorded. The decision is reaffirmed when the expert keeps it after hearing the reason. A finding that names the same item for the same reason is the same objection, and a later grading pass does not raise it again as new.

**SEE ALSO**
- finding
- grading pass

**NOT** A refuted finding. A finding the document does not support is dropped, not kept as an objection.

### Session

**MEANING** The assistant conversation the expert is already working in, which the product works inside, as dependency D-001 records. A session ends when that conversation ends, and the set stays where it was written.

**SEE ALSO**
- dependency

### Change log

**MEANING** Section 15, which records each version of the set with its date, its author, and the sections it changed, and every approval and override.

**SEE ALSO**
- version
- approver
- override

### Target

**MEANING** The value a measure must reach: where an objective says it will be, or what a non-functional requirement must meet, such as under 20 seconds at p95.

**SEE ALSO**
- objective
- baseline

### Exit criteria

**MEANING** What must be true before a release counts as done, stated with each release in section 13.

**SEE ALSO**
- release

**NOT** Acceptance criteria. Acceptance criteria decide one requirement or stage, and exit criteria decide the whole release.

### Non-functional requirement

MEANING

A quality the product must have, such as speed, accessibility, or retention, recorded in section 8 with a measure and a target. It applies across requirements, not to one user's capability.

SEE ALSO

- requirement
- target

NOT

A requirement in section 7. A requirement names a user and what they can do.

### Summary

MEANING

The opening of the root document, which states in a few sentences the problem, the users, why now, the bet, the one objective that proves it, and how it is measured. An approver reads it first.

SEE ALSO

- root document
- objective

### Glossary

MEANING

Section 14, this section. It gives each domain term one meaning, and every requirement and acceptance criterion uses those words.

SEE ALSO

- vocabulary

### View

MEANING

What the expert has in front of them at one moment, such as one section or its grade.

SEE ALSO

- section

### Topic

MEANING

One subject inside a section's interview, such as one requirement or one stage. A follow-up question finishes a topic before the next one opens.

SEE ALSO

- interview
- question list

### Ready

MEANING

At `review`: fit for someone else to read. A section is ready when it passes the gate and the expert accepts, or when the expert overrides the gate and the override is recorded.

SEE ALSO

- state
- gate

### Capability

MEANING

The first part of a requirement: what one named user can do with the product, in their own words.

SEE ALSO

- requirement
- outcome

### Outcome

**MEANING** The second part of a requirement: what the user sees or has once the capability exists, stated so that someone else could check it.

**SEE ALSO**
- requirement
- capability

**NOT** A feeling. "They feel confident" is not an outcome a tester can check.

### Problem

**MEANING** A condition a named persona suffers today, recorded in section 1 with counted evidence, why it matters now, and what it costs to leave it.

**SEE ALSO**
- persona
- objective

### Approval

**MEANING** The approver's recorded decision on one version of the set: approve, approve with conditions, or reject, with the approver, their role, the date, and the version. A rejection is an approval record too. An approval covers only the version it names, and a later change records a new version that it does not cover.

**SEE ALSO**
- approver
- condition
- version

**NOT** An override. An override moves one section past its gate, and an approval is a decision on the whole set.

### Engagement

**MEANING** One piece of work in which one expert writes one set, from the first question to the approval. TR-001 covers an engagement with no sponsor.

**SEE ALSO**
- expert
- sponsor

**NOT** A unit of work. The unit of work is what the expert's own domain counts.

### Lifecycle

**MEANING** The states one thing passes through and the moves between them, with the person who makes each move, recorded in section 5. LC-001 is the lifecycle of a section of the set.

**SEE ALSO**
- state
- scenario

### Handoff

**MEANING** The point at which one stage's work passes to the next person or the next stage, named on each journey stage.

**SEE ALSO**
- stage
- handover

**NOT** The handover. The handover is the one handoff at which the approved set reaches the engineering lead.

### Segment

**MEANING** A group of organizations that would buy the product for the same reason, recorded in section 2 with its buyer and its user, such as S-001.

**SEE ALSO**
- buyer

### Buyer

**MEANING** The person in a segment who decides to pay for the product, such as a principal or a practice lead.

**SEE ALSO**
- segment

**NOT** The user. The buyer pays, and the user is the person who works with the product, such as the consultant with the domain knowledge.

### Status quo

**MEANING** What a buyer does today with no product: a smart practitioner, a document, and a willing engineer. Section 2 records it as a competitor, because every buyer already has it.

**SEE ALSO**
- segment

### Delegated

**MEANING** Given to an agent to decide. Section 10 marks for each trust rule whether an agent may decide it, and a rule that is never delegated is always decided by the person it names.

**SEE ALSO**
- trust rule

### Definition of done

**MEANING** The conditions every backlog item must meet before the team counts it as done. The delivery lead needs the set's constraints to arrive as items of it, not as surprises, in J-002.

**SEE ALSO**
- backlog
- constraint

SECTION 15

# Change log and approvals

DRAFT 79/100 NO BLOCKING DEFECTS

| VERSION | DATE | AUTHOR | WHAT CHANGED |
|---|---|---|---|
| 0.9.0 | 2026-10-27 | Richard Hightower | Six glossary terms, FR-017 without view names, FR-010's wait on Q-007, the status-quo edge; counted evidence for P-001 to P-004; C-004 moved out of section 9; FR-008 and FR-012 restated; interview, handover, and dependency defined; FR-018 on request within 10 seconds; backlog terms, expert, and sponsor defined; FR-008 checks the set for one item stated two ways; final set by the expert; ten more glossary terms; FR-019 and FR-020 split out; sources cited for standard terms; the versioning policy; FR-020 drawn in WF-001; Decider for trust rules |
| 0.8.1 | 2026-10-27 | Richard Hightower | Glossary terms, observable scenario outcomes, open questions named in section 7, competitor edges; C-001 to C-003 and S-003 moved; ST-009 for the engineering lead |
| 0.8.0 | 2026-10-27 | Richard Hightower | Requirement and scenario priorities, SC-007, release exit criteria, single deciders |
| 0.7.0 | 2026-10-27 | Richard Hightower | Root objective, primary persona and approver; release id in scope; LC-001 rule |
| 0.6.0 | 2026-10-27 | Richard Hightower | Links from the specification review: problems, objectives, competitors, decision stages, and scenario views |
| 0.4.0 | 2026-10-26 | Richard Hightower | Navigation and refresh requirements; wireframes rewritten from their grading |
| 0.3.0 | 2026-10-26 | Richard Hightower | Wireframes for every screen and report, and recorded reasons for the grades |
| 0.2.0 | 2026-10-26 | Richard Hightower | Every defect the grading and the buildability review found, over two rounds |
| 0.1.0 | 2026-10-23 | Richard Hightower | First draft of all sixteen sections |

### Versioning policy before 1.0

The set and the product share one version number. The number changes by these rules:

- A patch, such as 0.9.0 to 0.9.1, is for fixes that add no content.
- A minor version, such as 0.8 to 0.9, is for new content in the set or in the product.
- A change to the release line or to the golden thread takes the set to the next major version. From 0.x, the next major version is 1.0.

The expert set this policy at 0.9.0, and it applies from 0.9.0 onward. Earlier version numbers stay as recorded. Two earlier versions do not follow it, and neither is reissued:

- 0.8.1 added stage ST-009 and moved items between sections. Under this policy it would have been the next minor version.
- 0.2.0 changed the golden thread. Under this policy it takes the set to 1.0.

Version 0.5.0 changed the product and not this PRD, so the log has no entry for it.

### 0.2.0, what the grading changed

The reviewers that graded 0.1.0 were the product's own, run against its own PRD. They found twelve blocking defects and four failed derivability checks. The largest single finding was that FR-009, FR-010 and FR-011 were written against "its rubric" and "its threshold" with neither defined anywhere in the set: the threshold is now stated as 70 in the requirements themselves, and the rubric, the section states and the derivability checks are defined in section 14 as the product content they are.

Two findings were checked and **rejected**: the claim that "see section 6" pointed at the wrong document, and the claim that the section count was contradicted. Both came from reading file numbers as section numbers. The references were right; the numbering was unexplained, which is now fixed in the root document.

### 0.1.0, what to know

Written as prd-forge's own first exercise. Three things a reader should carry forward:

1. **No baselines.** All three objectives carry an explicit unknown with a matching open question (Q-001 to Q-003). The targets are intentions, not commitments.
2. **C-001 is the tension.** The product is delivered inside an assistant session. Naming that in a requirement would fail test 5, so it is recorded as a constraint. A reviewer may fairly call it architecture. At 0.8.1 it moved to section 6 as a non-goal, and the session it needs is dependency D-001 in section 11.
3. **R-003 is not closed.** Nothing in this product checks whether the product being specified should exist at all.

### Approvals

None sought. This set is at `draft` and has not been offered for approval. It cannot be approved for a release without an approval recorded here, under TR-001. A section reaches `final` when the expert sets it final, under TR-002.

> **Fixes the blocking defects the 0.8.1 grading found that do not wait on the expert. The glo**
>
> VERSION 0.9.0
>
> DATE 2026-10-27
>
> AUTHOR Richard Hightower
>
> SECTIONS
> - 1
> - 2
> - 5
> - 6
> - 7
> - 9
> - 12
> - 13
> - 14
> - 15

> **Fixes the blocking defects that do not wait on evidence. The glossary defines buildability**
>
> VERSION 0.8.1
>
> DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 2
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11
- 12
- 13
- 14

**Sets the expert's priorities. Ten requirements stay must and eight become should; SC-002 a**

VERSION 0.8.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 5
- 7
- 12
- 13

**Names the root's objective, primary persona, and approver, from section 3, the personas do**

VERSION 0.7.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 4
- 5
- 6

**Adds the links the specification review asked for. Each problem names its persona, and eac**

VERSION 0.6.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS
- 1
- 2
- 3
- 5
- 10
- 11

**Adds FR-017, moving between sections without losing one's place, and FR-018, seeing the se**

VERSION 0.4.0

DATE 2026-10-26

AUTHOR Richard Hightower

SECTIONS
- 6
- 7
- 8
- 13
- 14

**Adds wireframes to section 7. Each functional requirement now says where its user meets th**

VERSION 0.3.0

DATE 2026-10-26

AUTHOR Richard Hightower

SECTIONS
- 5
- 7
- 8
- 14

**Fixes every defect the adversarial grading and the buildability review found in 0.1.0. J-0**

VERSION 0.2.0

DATE 2026-10-26

AUTHOR Richard Hightower

SECTIONS
- 4
- 5
- 6
- 7
- 8
- 10
- 13
- 14

**First draft of the whole set, written with prd-forge's own conventions as the product's fi**

| | |
|---|---|
| VERSION | 0.1.0 |
| DATE | 2026-10-23 |
| AUTHOR | Richard Hightower |
| SECTIONS | 1<br>2<br>3<br>4<br>5<br>6<br>7<br>8<br>9<br>10<br>11<br>12<br>13<br>14<br>15 |