PRODUCT REQUIREMENTS DOCUMENT

prd-forge

A domain expert, with no product management training, can reach a PRD that an engineering lead confirms is buildable, and can see for themselves where it is still weak before anyone starts work.

VERSION	0.9.0
STATUS	draft
OWNER	Richard Hightower
DOMAIN	product management tooling
CREATED	2026-10-23
UPDATED	2026-10-27
GOLDEN THREAD	J-001
STAKEHOLDERS	Domain experts who must specify software for their own field; Engineering leads who build from the resulting PRD; Delivery leads who run story-mapping workshops

Executive summary

PROBLEM	A domain expert knows exactly what their product must do, but what they can write down is a feature list with no named users, no pass-or-fail criteria and the technology already chosen, so the team that has to build it cannot.
USERS	Domain experts specifying software for their own field, the engineering leads who build from what they write, and the sponsors who approve it.
WHY NOW	Teams now expect to work from a written specification an assistant helped produce, and a specification that reads well but cannot be built from costs more than no specification at all, because it is believed.
BET	A domain expert, with no product management training, can reach a PRD that an engineering lead confirms is buildable, and can see for themselves where it is still weak before anyone starts work.
OBJECTIVE	O-001
MEASURE	Share of PRDs that pass an engineering buildability review on the first submission: 80% by 2027-09-30, from a baseline Q-001 will establish.

Contents

§	SECTION	STATUS	GRADE	BLOCKING
1	Problem and opportunity	draft	71/100	—
2	Target customer and market	draft	55/100	—
3	Business objectives and success metrics	draft	75/100	—
4	Users and personas	draft	68/100	—
5	User journeys	draft	77/100	—
5	Scenarios	draft	78/100	—
6	Scope	draft	79/100	—
7	Functional requirements	draft	78/100	1
7	Wireframes	draft	88/100	—
8	Non-functional requirements	draft	82/100	—
9	Constraints	draft	100/100	—
10	Decision rights and trust rules	draft	88/100	—
11	Assumptions, risks and dependencies	draft	75/100	—

§	SECTION	STATUS	GRADE	BLOCKING
12	Open questions	draft	79/100	—
13	Release plan	draft	88/100	—
14	Glossary	draft	65/100	4
15	Change log and approvals	draft	63/100	1

About this document

A domain expert knows exactly what their product must do, but what they can write down is a feature list with no named users, no pass-or-fail criteria and the technology already chosen, so the team that has to build it cannot. The users are domain experts specifying software for their own field, the engineering leads who build from what they write, and the sponsors who approve it. Teams now expect to work from a written specification an assistant helped produce, and one that reads well but cannot be built from costs more than no specification at all, because it is believed. The bet is that a domain expert with no product management training can reach a PRD an engineering lead confirms is buildable, and can see for themselves where it is still weak before anyone starts work. We will know it worked from the share of PRDs that pass an engineering buildability review on the first submission.

Contents

SECTION	DOCUMENT
1	Problem and opportunity
2	Target customer and market
3	Business objectives and success metrics
4	Users and personas
5	User journeys
5	Scenarios
6	Scope
7	Functional requirements
7	Wireframes
8	Non-functional requirements
9	Constraints
10	Decision rights and trust rules
11	Assumptions, risks and dependencies
12	Open questions
13	Release plan
14	Glossary

SECTION	DOCUMENT
15	Change log and approvals

How this document set is numbered

A PRD has **fifteen numbered sections** and this set has **seventeen section documents**, because section 5 holds two of them, journeys and scenarios, and section 7 holds two, functional requirements and wireframes. The two section 7 files share the prefix `08`. So from section 5 onward the file prefix and the section number differ — `07-scope.md` is section 6, `16-change-log.md` is section 15. **Every cross-reference in this set is by section number, never by file number.** Three independent reviewers misread this, so it is stated here rather than left to be inferred.

A note on this PRD

This is prd-forge's PRD, written with prd-forge's own conventions. Writing it surfaced one real tension: the product is delivered inside an assistant the expert already uses, and naming that delivery mechanism in a requirement would fail test 5. It was recorded as constraint C-001 until 0.8.1. It is now the separate-application non-goal in section 6 and dependency D-001 in section 11, and the requirements in section 7 describe only what a person does and sees.

Grading it with its own adversarial reviewers found twelve blocking defects across four sections, and a buildability review found four failed derivability checks. Version 0.2.0 fixes them: every graded section is now at zero blocking defects, and nine of the eleven derivability checks pass. What that grading found, and what it changed, is in section 15.

On the scores in this set. Each graded section carries the score from its last grading pass, with the date. Sections 5, 7 and 8 were graded before the second round of corrections, and section 10 before its last two, so each score is the last one measured and not a reading of the words as they now stand. Treat them as a floor. A score never rises without a change to the document, and never rises without a grading pass; the two gaps here are in the other direction.

SECTION 1

Problem and opportunity

DRAFT 71/100 NO BLOCKING DEFECTS

ID	PROBLEM	WHO
P-001	What the expert writes is a feature list with no user and no testable outcome	Domain expert
P-002	The steps between requirements are never written, so paths have gaps	Engineering lead
P-003	Buildability is invisible until the team starts building	Sponsor or approver
P-004	No source of PRD craft at the moment the expert needs it	Domain expert

The four compound. P-004 means the expert does not know what good looks like; P-001 is what they produce as a result; P-002 is the specific gap that costs the most; and P-003 is why none of it is caught before it is expensive.

Evidence note. Each problem now has one counted case, with its source and date. The counts for P-001 to P-003 come from one engagement, this product's own PRD, written and graded with prd-forge. The count for P-004 comes from two public templates. One engagement shows that each problem occurs. It does not show how often. The other evidence items describe a pattern and are not measured. Three of the objectives in section 3 still have no baseline. See Q-001, Q-002, and Q-003.

P-001

STATEMENT A domain expert can describe what their product must do out loud, but what they write down is a feature list: no named user, no outcome anyone could check, and a technology already chosen for reasons they cannot defend.

WHO Domain expert specifying software for their own field

PERSONA PER-001

EVIDENCE

- {'claim': 'In this PRD at version 0.1.0, the adversarial grade found 2 of 14 functional requirements, FR-003 and FR-004, with no outcome a person could see. Both were fixed at version 0.2.0.', 'source': 'prd-forge repository, commits 5ba59cb (version 0.1.0) and 5d4d936 (version 0.2.0)', 'as_of': datetime.date(2026, 10, 26)}
- Specifications arriving at delivery teams are commonly feature lists; teams report re-deriving the user and the acceptance criteria themselves before they can estimate.
- Experts asked to write requirements reach for the words of the software they have used, because that is the only requirements language they have been exposed to.

WHY NOW Specifications are now commonly drafted with an assistant, which makes them fluent. Fluent and unbuildable is a worse failure than rough and unbuildable, because nobody queries it.

COST OF INACTION The gap is closed by an engineer guessing, during build, at what the expert meant, and the guess is discovered at review or in production.

P-002

STATEMENT Requirements are written one at a time, so the steps between them are never written at all, and nobody notices until the team tries to assemble a working path through the product.

WHO Engineering lead building from the specification

PERSONA PER-002

EVIDENCE

- {'claim': 'At version 0.1.0, the buildability review of this PRD failed 4 of 11 derivability checks and named 4 gates an engineer would have to invent.', 'source': 'prd-forge repository, the commit messages of 5ba59cb (version 0.1.0) and 5d4d936 (version 0.2.0)', 'as_of': datetime.date(2026, 10, 23)}
- A specification with no user journey gives a story-mapping workshop no backbone, so features and epics cannot be derived from it at all.
- Gaps surface as mid-iteration questions about what happens between two requirements that both looked complete.

WHY NOW

Teams are committing to outcomes per release, which needs a complete path through the product, not a complete list of features.

COST OF INACTION

The release line is drawn across a feature list rather than a journey, so a release ships features and no user can complete anything with them.

P-003

STATEMENT

Nobody can tell a buildable specification from a readable one until the team has already started building from it.

WHO

Sponsor or approver signing the specification off

PERSONA

PER-003

EVIDENCE

- {'claim': 'At version 0.8.0, this PRD passed every schema and cross-reference check with 0 errors and 0 warnings. The adversarial grade found 23 blocking defects in 8 of 18 documents.', 'source': 'prd-forge repository, commit 6f77bc0 (version 0.8.0)', 'as_of': datetime.date(2026, 10, 26)}
- Approval is given on readability and completeness of headings, because those are the only properties an approver can assess unaided.
- The first real assessment of buildability happens when an engineer estimates, which is after approval and after the plan was made.

WHY NOW

Where the document was drafted with an assistant, readability is no longer any evidence at all of buildability.

COST OF INACTION

Plans and commitments are made on an approved document whose defects are found by the people asked to deliver against them.

P-004

STATEMENT

An expert has nowhere to learn what a section of a PRD is for at the moment they are being asked to fill it in.

WHO

Domain expert specifying software for their own field

PERSONA

PER-001

EVIDENCE

- {'claim': 'Of 20 sections in two public PRD templates, all 20 say what to write, and 0 say what makes an entry good or how a reader checks it. The Aha! Agile PRD template has 6 sections, and the Product School PRD template has 14.', 'source': 'Aha! Agile PRD template, https://www.aha.io/roadmapping/guide/requirements-management/what-is-a-good-product-requirements-document-template, and Product School PRD template, https://productschool.com/blog/product-strategy/product-template-requirements-document-prd', 'as_of': datetime.date(2026, 10, 26)}
- The craft exists in books and in experienced product managers, neither of which is available at the moment the expert is typing.

WHY NOW

Organisations are asking domain experts to specify directly, without routing the work through a product manager first.

COST OF INACTION

Each expert reinvents the form badly, and the organisation's specifications cannot be compared to one another or improved as a body of work.

SECTION 2

Target customer and market

DRAFT 55/100 NO BLOCKING DEFECTS

ID	SEGMENT	BUYER	USER
S-001	Consultancies specifying from client domain knowledge	Principal or practice lead	The consultant
S-002	Regulated-industry organisations where practitioners specify	Head of product or delivery	The practitioner

Product teams entering a domain they do not know were segment S-003 until 0.8.1. They already employ the product skill and want the domain, which this product does not supply, so section 6 now records them as a non-goal for release 1. The id S-003 is not reused.

Why we win

The expert can see, before anyone starts building, which parts of their own specification a development team will not be able to build from, and what to answer to fix it. No competitor above tells them that.

Honest weakness. Against a contracted product manager we have no quality edge, only availability and cost. A sceptical buyer is right to test that.

Segments

> **S-001 · Consultancies and advisory firms turning client domain knowledge into specifications**
>
> BUYER — Principal or practice lead
>
> USER — The consultant with the domain knowledge
>
> SIZE — `<count>` firms, per `<a named register, with its year>`, not yet measured; see Q-004.
>
> WILLINGNESS TO PAY — Already pays for the consultant's time; the comparison is billable hours saved, not a software budget.

> **S-002 · Regulated-industry organisations where practitioners specify their own internal software**
>
> BUYER — Head of product, delivery lead, or head of transformation
>
> USER — The practitioner, plus the delivery lead who reviews what they write
>
> SIZE — Not yet measured; see Q-004.
>
> WILLINGNESS TO PAY — Compared against contracting a product manager for the engagement.

Competitors

> **A blank PRD template in the team's document tool**
>
> KIND — indirect
>
> OFFER — The section headings, immediately, free, in a tool everyone already has.
>
> OUR EDGE — After one engagement, each fact the expert could not supply is an open question with a named decider and a date, and the set holds no figure or citation that nobody can stand behind.

A general-purpose assistant conversation with no structure

KIND	indirect
OFFER	Fluent prose on any section, instantly, with no setup at all.
OUR EDGE	Fluency is the failure mode in P-003. After one engagement, the buyer holds a report, made before anyone has estimated, that names each release-line stage no requirement covers and quotes each requirement an engineer could not write a test from.

Contracting a product manager for the engagement

KIND	indirect
OFFER	Genuine product judgement, and someone accountable for the outcome.
OUR EDGE	None on quality. The edge is availability and cost at the moment the expert is typing, and that the expert keeps the understanding afterwards rather than renting it.

The status quo - a smart practitioner, a document, and a willing engineer

KIND	status-quo
OFFER	No tool, no cost, no adoption, and it does sometimes work.
OUR EDGE	At handover, an engineer who was not in the room receives a set that names each release-line stage no requirement covers and holds each fact the expert could not supply as an open question with a named decider and a date, so they can estimate without first interviewing the author.

SECTION 3

Business objectives and success metrics

DRAFT 75/100 NO BLOCKING DEFECTS

ID	OBJECTIVE	METRIC	BASELINE	TARGET
O-001	Buildable without re-derivation	Pass rate on first buildability review	unknown (Q-001)	80%
O-002	Days, not weeks, to an approved spec	Median working days, first question to approval	unknown (Q-002)	under 5
O-003	Requirements testable as written	Share of `must` requirements an outside engineer can test	unknown (Q-003)	95%

Why there are no baselines

All three baselines are unknown, and each has a matching open question with a decider and a date. That is deliberate and it is the honest state: this product has not run against a real engagement yet. The targets above are therefore intentions, not commitments, and section 15 should record the moment each baseline arrives and each target is either confirmed or changed.

A reviewer should push on this. An objective with an invented baseline is worse than one with an admitted gap, because the invented one gets planned against.

What would move these without this product

Worth asking, and recorded here because the answer is not nothing: a single experienced product manager embedded with the expert would move all three. The bet is on availability and cost, not on beating that person — see section 2.

Objectives

O-001

OBJECTIVE: A specification handed to a development team can be built from without the team re-deriving it first.

PROBLEMS:
- P-001
- P-002
- P-003

METRIC: Share of PRDs that pass an engineering buildability review on first submission, where the review is the eleven derivability checks.

BASELINE: unknown - not measured; see Q-001. Do not set a target against this until Q-001 is answered.

BASELINE QUESTION: Q-001

TARGET: 80% of PRDs pass on first submission.

BY: 2027-09-30

OWNER: Richard Hightower

O-002

OBJECTIVE	A domain expert reaches an approved specification in days rather than the weeks a product manager's discovery cycle takes.
PROBLEMS	• P-004
METRIC	Median elapsed working days from the first interview question to an approval recorded in section 15.
BASELINE	unknown - not measured; see Q-002.
BASELINE QUESTION	Q-002
TARGET	Under 5 working days for a single-journey product.
BY	2027-09-30
OWNER	Richard Hightower

O-003

OBJECTIVE	Requirements are testable as written, so a team can close them without asking the author what they meant.
PROBLEMS	• P-001
METRIC	Share of functional requirements whose acceptance criteria an engineer who was not in the interview judges sufficient to write a test.
BASELINE	unknown - not measured; see Q-003.
BASELINE QUESTION	Q-003
TARGET	95% of `must` requirements.
BY	2027-09-30
OWNER	Richard Hightower

SECTION 4

Users and personas

DRAFT 68/100 NO BLOCKING DEFECTS

ID	ROLE	FREQUENCY	MAY APPROVE
PER-001	Domain expert specifying their own field	Weeks per product, then rarely	Scope, terms, drafts, section states, facts, and offering the set; not the PRD itself
PER-002	Engineering lead building from it	A handful of intensive reads	May reject as not buildable
PER-003	Sponsor or approver	Once per version	The PRD for a release
PER-004	Delivery lead running story mapping	Once per release	Backlog shape and order

PER-001 is the primary persona

Everything is designed around someone with deep domain knowledge, no product training, and a burst of intense use followed by months away. Two consequences run through the rest of this PRD: nothing may assume prior knowledge of product vocabulary, and the work must survive being put down for a week and picked back up — see FR-015 and NFR-004.

Who is deliberately not a persona

The assistant that delivers the teaching, the interview and the grading is not a user of this product. It is how the product reaches PER-001. Listing it here would put the delivery mechanism into the PRD, which the set keeps out of the requirements: section 6 makes a separate application a non-goal, and section 11 records the session as dependency D-001. The same goes for the grading reviewers: a reviewer that grades a section is a mechanism, and what matters to this PRD is that PER-001 sees a grade they did not produce themselves.

> **PER-001 · Domain expert specifying software for their own field**
>
> SHORT NAME: DomainExpert
>
> PRIMARY: yes
>
> GOALS:
> - Get what is in my head into a form the development team can build from
> - Stop being asked the same question by four different engineers
> - Not be the reason the project is late
>
> PAINS:
> - I know what it must do, and I cannot write it the way they need it
> - I am told my requirements are not testable and I do not know what that means
> - I reach for the names of software I have used, because that is the only requirements language I have
>
> FREQUENCY: Intensively for one to three weeks per product, then almost never until the next one.

AUTHORITY	Decides what the product must do and what is out of scope; decides the domain's own terms; decides when a section's interview is finished and when a section may be left at draft; accepts or corrects a drafted section; sets a section to final; accepts or refutes a grade; decides the set is ready to be offered for approval; and is the only one who may state a domain fact as true. May not approve the PRD for a release on their own where a sponsor exists.
PROFICIENCY	Deep in their domain, often decades. No product management training. Assume no familiarity with MoSCoW, Given/When/Then, or the word "persona".

PER-002 · Engineering lead who must build from the PRD

SHORT NAME	EngineerLead
GOALS	• Estimate without first interviewing the author • Know which decisions are mine and which have been made for me • Find out now, not in iteration three, what the document does not say
PAINS	• Specifications that name a technology and so pre-empt my decision badly • Requirements I cannot write a test for • Gaps between two requirements that both looked complete
FREQUENCY	Reads a given PRD a handful of times, intensively, at estimation and at each release planning.
AUTHORITY	Decides how it is built; may reject a PRD as not buildable.
PROFICIENCY	Expert in delivery, usually knows nothing about the expert's domain.

PER-003 · Sponsor or approver

SHORT NAME	Sponsor
GOALS	• Be confident enough to commit budget and a date • Know what is still undecided, and who is deciding it
PAINS	• I can only judge whether it reads well • I find out it was not buildable after I have committed to the plan
FREQUENCY	Once per version, plus conditions.
AUTHORITY	Approves the PRD for a release; may approve with conditions; may reject.
PROFICIENCY	Senior in the business, not in product or engineering.

PER-004 · Delivery lead running the story-mapping workshop

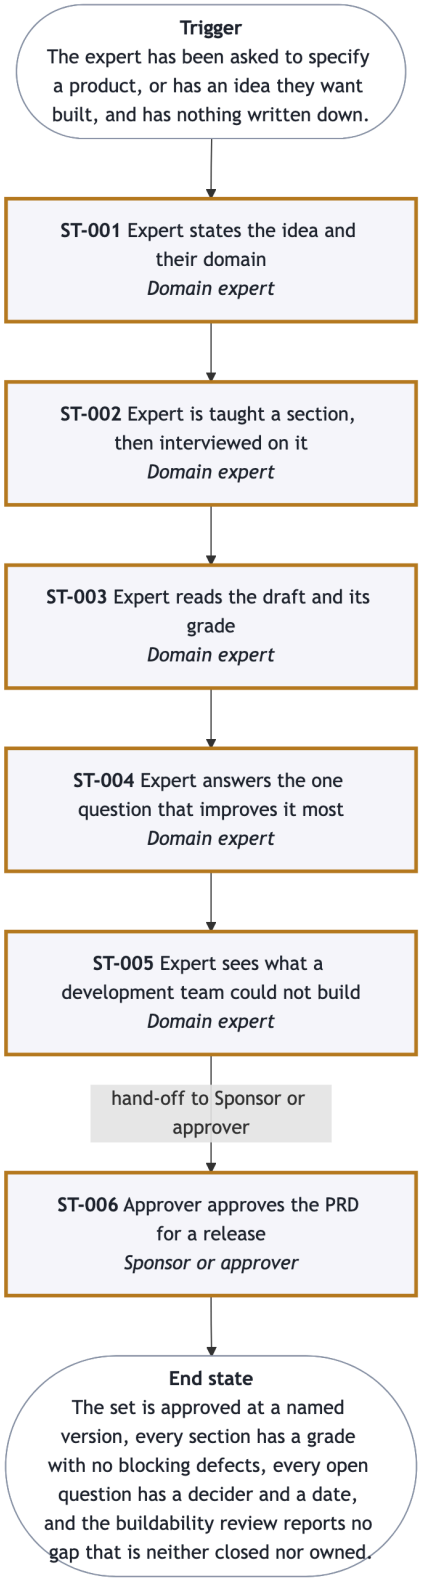
SHORT NAME	DeliveryLead
GOALS	• Walk into the workshop with a backbone already agreed • Trace every story back to something the expert actually said
PAINS	• No journey means no backbone, so the workshop invents one in the room • Stories that trace to nothing, so nobody can say why they exist
FREQUENCY	Once per release, plus refinement.
AUTHORITY	Decides the backlog's shape and order, and decides whether a backlog can be mapped from a set at all; does not decide scope.
PROFICIENCY	Expert in agile delivery; depends entirely on the PRD for the domain.

SECTION 5

User journeys

DRAFT 77/100 NO BLOCKING DEFECTS

J-001 - A domain expert reaches an approved PRD whose gaps are known (golden thread)



J-001 journey diagram

Persona PER-001. Trigger: they have been asked to specify a product and have nothing written. Stages ST-001 to ST-006.

STAGE	ACTOR	DECISION
ST-001 States the idea and the domain	Domain expert	What this is, and the domain's own words

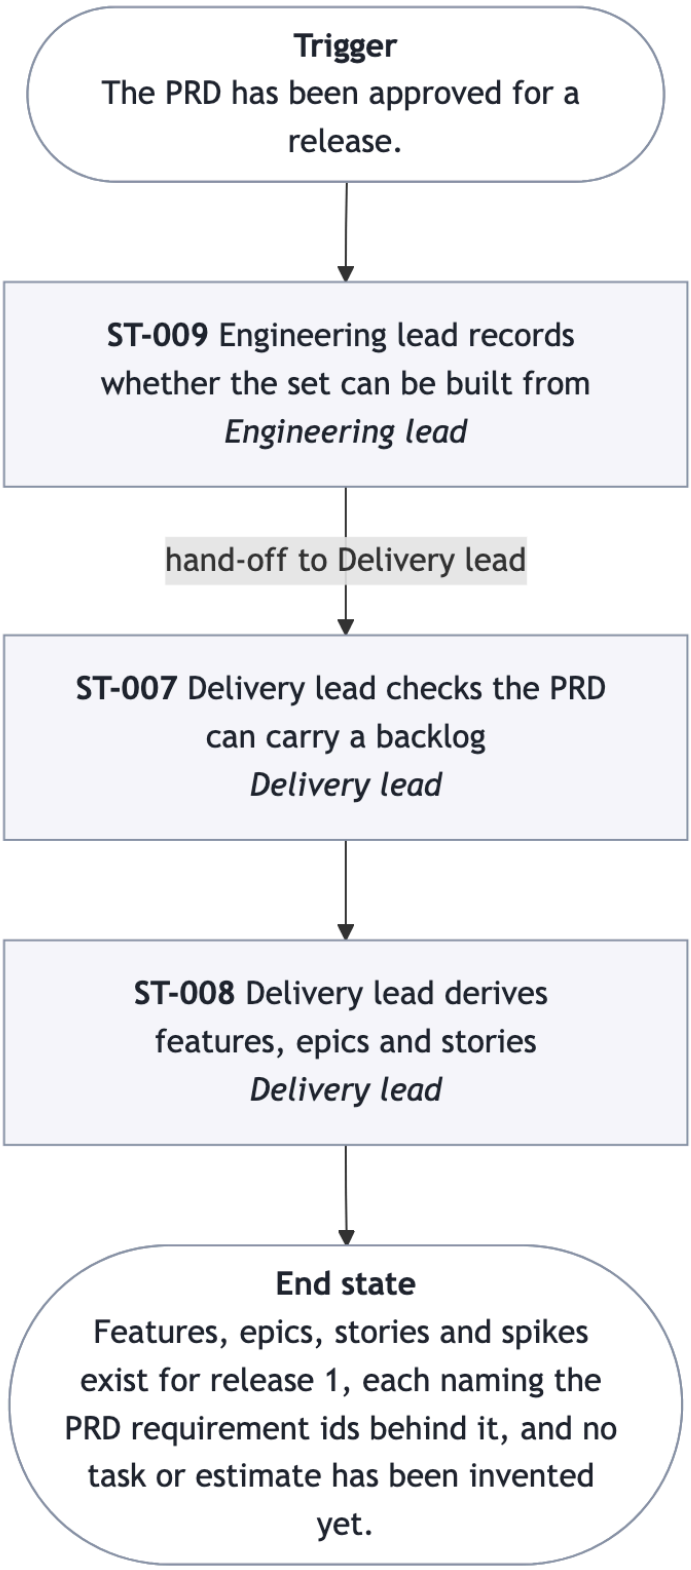
STAGE	ACTOR	DECISION
ST-002 Taught a section, then interviewed	Domain expert	Have the questions reached the end of what they know
ST-003 Reads the draft and its grade	Domain expert	Accept, correct, or dispute the grade
ST-004 Answers the one question that matters most	Domain expert	Is this section done enough to leave
ST-005 Sees what a team could not build	Domain expert	Is it ready for someone who will commit
ST-006 Approver approves for a release	Sponsor or approver	Approve, approve with conditions, or reject

ST-002 to ST-004 are a loop, run once per section. That is the whole product: the expert is taught, asked, shown a draft, shown a grade, and asked one more thing.

End state: approved at a named version, every section graded with nothing blocking, every open question owned and dated, and no buildability gap left that is neither closed nor owned.

The goal deliberately does not claim an engineering lead's confirmation. No stage in ST-001 to ST-006 puts one in the actor seat, so claiming it would promise something the release line cannot deliver. A lead's own recorded confirmation is OUT-004, deferred to release 2 with J-002, where it is stage ST-009.

J-002 - A delivery lead derives a backlog (not the golden thread)



J-002 journey diagram

Persona PER-004, stages ST-009, ST-007, and ST-008. In ST-009 the engineering lead, PER-002, reads the buildability report and records their own verdict under TR-006, before the delivery lead takes over. ST-009 was added at 0.8.1, after ST-008, and ids are never renumbered. Release 1 does not have to complete this journey (see section 6), but the PRD must carry what it needs, which is what ST-007 checks.

What these journeys deliberately do not say

No stage says where the teaching text comes from, how a grade is computed, or what reviews the set in ST-005. Those are mechanisms. Every stage above names a person, what they do and see, and what they decide.

J-001

GOAL	An approved PRD whose buildability gaps have been found and either closed or given an owner and a date, before anyone estimates it.
PERSONA	PER-001
TRIGGER	The expert has been asked to specify a product, or has an idea they want built, and has nothing written down.
GOLDEN THREAD	yes

STAGES

Expert states the idea and their domain

ID	ST-001
ACTOR	Domain expert
ACCEPTANCE	• Given an expert with an idea and nothing written down, when they state the idea, the domain, and three words an outsider would misread, then they see a document set with every section listed and marked as not yet started.
STEPS	• Says what they want built, in their own words, however roughly • Is asked what the business domain is and what the unit of work is called - an engagement, a load, a claim, a case • Names three words their team uses that an outsider would misread • Sees that a set of documents now exists, every section listed and marked as not yet started
INPUTS	• The idea, the domain, and the vocabulary
OUTPUTS	• A document set with every section visible and honestly marked empty
DECISION	What this product is, and what the domain's own words are

Expert is taught a section, then interviewed on it

ID	ST-002
ACTOR	Domain expert
ACCEPTANCE	• Given a section not yet interviewed, when the expert reads what it is for and answers its questions one or two at a time, then the section has enough material to be drafted, or the expert has said they have nothing more to add.
STEPS	• Reads what the section is, why it matters, and what belongs in it, in under ten lines, with an example from their own domain • Is asked one or two questions at a time, in their own vocabulary • Is asked follow-ups where an answer was thin • Says when they have nothing more to add
INPUTS	• Their knowledge of the domain, in answer to each question
OUTPUTS	• Enough material for the section to be drafted
DECISION	Whether the questions have reached the end of what they know

Expert reads the draft and its grade

ID	ST-003
ACTOR	Domain expert
ACCEPTANCE	• Given a drafted and graded section, when the expert opens it, then they read it as prose and see its score, each blocking defect, and its weakest part.
STEPS	• Reads the drafted section as prose, not as a form • Sees a score, how many blocking defects there are, and what each one is • Sees which part of the section is weakest, and why • Decides whether the grade is fair

INPUTS	• Nothing; this is where they read
OUTPUTS	• A judgement on whether the draft says what they meant
DECISION	Accept the draft, correct it, or dispute the grade

Expert answers the one question that improves it most

ID	ST-004
ACTOR	Domain expert
ACCEPTANCE	• Given one question about the weakest part of a section, when the expert answers it, then they see the section re-graded against the same rubric.
STEPS	• Reads a single question, not a list of twelve • Answers it, or says why the premise is wrong • Sees the section re-graded against the same rubric • Moves to the next section, or goes back to an earlier one when this section showed it was incomplete
INPUTS	• One answer
OUTPUTS	• A section at a grade the expert accepts, and a set that is one section further on
DECISION	Whether this section is done enough to leave
HANDOFF	Back into ST-002 for the next section, until every section has been through this loop.

Expert sees what a development team could not build

ID	ST-005
ACTOR	Domain expert
ACCEPTANCE	• Given a set with every section written at least once, when the expert asks for a buildability review, then they see what a builder could not build and which release-line stages have no requirement behind them.
STEPS	• Asks for the whole set to be reviewed as a builder would read it • Sees what a builder could not build, and what they would have to guess • Sees which release-line stages have no requirement behind them • Fixes the holes, or records them as open questions with a decider
INPUTS	• Every section written at least once, so there is something to review
OUTPUTS	• A list of what is not buildable, and a set with those holes closed
DECISION	Is this ready to be read by someone who will commit to it
HANDOFF	Sponsor or approver

Approver approves the PRD for a release

ID	ST-006
ACTOR	Sponsor or approver
ACCEPTANCE	• Given a set offered for approval, when the approver approves, approves with conditions, or rejects it, then they see the decision recorded against the set's version.

STEPS	• Reads the summary, the scope, and what is still open • Sees each section's grade and whether anything is still blocking • Approves, approves with conditions naming what the condition is, or rejects • Sees the approval recorded against this version
INPUTS	• Their decision and any conditions
OUTPUTS	• An approved PRD at a named version, with the approval recorded
DECISION	Approve, approve with conditions, or reject
HANDOFF	Delivery lead, for the story-mapping workshop

PAIN POINTS

- I do not know what the section is for, so I fill the heading with adjectives
- Nobody tells me my requirements are untestable until a team tries to build them
- I reach for the names of software I have used
- I put this down for a week and cannot remember where I was
- I cannot tell whether what I have written is good

NEEDS

- Teach me the section before you ask me about it, in my own domain
- Tell me where it is weak, against a stated standard, and ask me one thing
- Convert my technology talk into the need underneath it, without arguing
- Let me stop and resume without starting over
- Show me, before anyone builds, what a builder cannot build

END STATE

The set is approved at a named version, every section has a grade with no blocking defects, every open question has a decider and a date, and the buildability review reports no gap that is neither closed nor owned.

J-002 · A backlog whose every story traces to a requirement the expert actually stated.

PERSONA	PER-004
TRIGGER	The PRD has been approved for a release.

STAGES

Engineering lead records whether the set can be built from

ID	ST-009
ACTOR	Engineering lead
ACCEPTANCE	• Given an approved set and its buildability report, when the engineering lead records a verdict that differs from the report's, then their judgement is recorded as the verdict and the report is kept alongside it. • Given a gap the engineering lead records, when it is stated, then the document is quoted for it.
STEPS	• Reads the buildability report, its verdict, and the gaps it quotes • Checks each reported gap against the document it quotes • Records their own verdict on whether the set can be built from
INPUTS	• An approved set and its buildability report
OUTPUTS	• The engineering lead's own recorded verdict, with the report kept alongside it
DECISION	Can a team build from this set
HANDOFF	The delivery lead, to check the PRD can carry a backlog

Delivery lead checks the PRD can carry a backlog

ID	ST-007
ACTOR	Delivery lead
ACCEPTANCE	• Given an approved set and the release it is mapped for, when the delivery lead checks it, then they see whether a golden-thread journey exists and what would block deriving a backlog.
STEPS	• Sees whether a journey exists and which one is the golden thread • Sees which release-line stages have requirements behind them • Sees what would block deriving a backlog, before the workshop
INPUTS	• An approved set, and the release it is being mapped for
OUTPUTS	• A go or no-go on holding the story-mapping workshop
DECISION	Is there a backbone to map against

Delivery lead derives features, epics and stories

ID	ST-008
ACTOR	Delivery lead
ACCEPTANCE	• Given an approved set, when the delivery lead derives a backlog from it, then they see one feature per journey stage, epics scoped by the release line, and a story per requirement, each naming the id it came from.
STEPS	• Sees one feature per journey stage, with the requirement ids it carries • Sees epics scoped by the release line • Sees a story per requirement, in the persona's own words, with the requirement's criteria as its acceptance criteria • Sees spikes for the open questions that block the release line
INPUTS	• The approved set, and the team's sizing of the stories

OUTPUTS	• A backlog whose every level names what it came from
DECISION	Which epics come first, and why
HANDOFF	The engineering team, for iteration planning

PAIN POINTS

- The workshop invents a backbone in the room because the PRD has none
- Stories trace to nothing, so nobody can say why they exist
- Non-functional requirements and constraints are forgotten until late

NEEDS

- A journey to map against, agreed before the workshop
- Every story carrying the requirement id it came from
- Constraints arriving as definition-of-done items, not as surprises

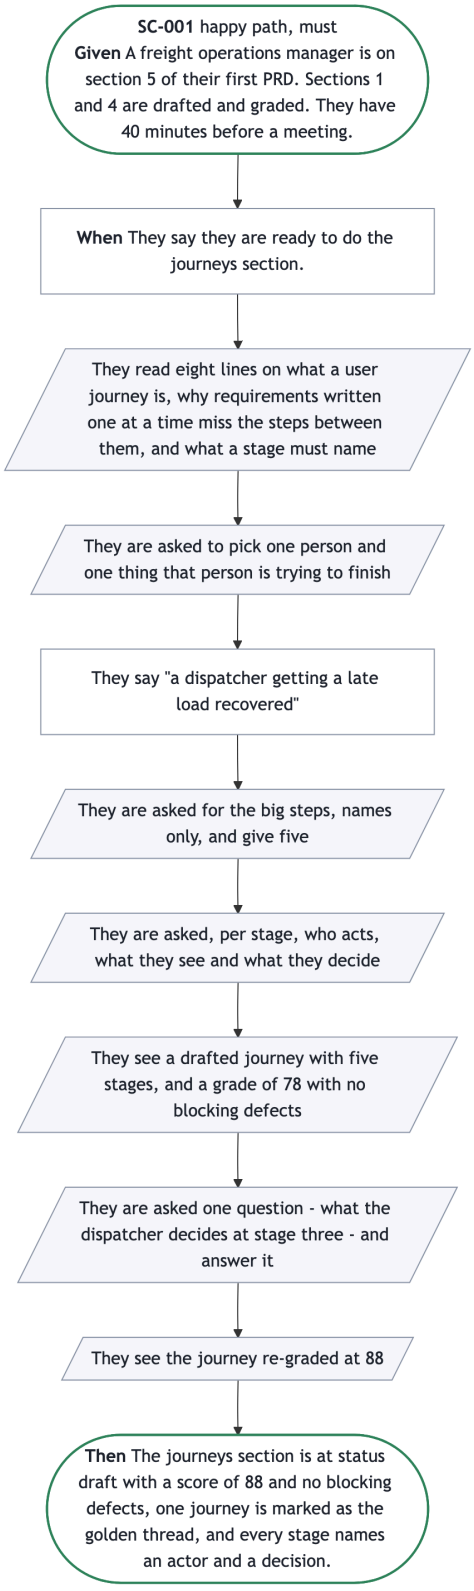
END STATE

Features, epics, stories and spikes exist for release 1, each naming the PRD requirement ids behind it, and no task or estimate has been invented yet.

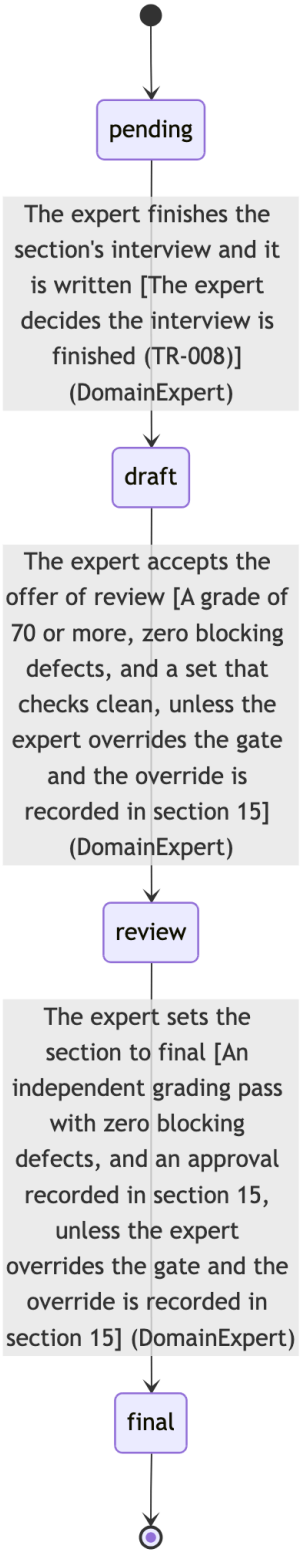
SECTION 5

Scenarios

DRAFT 78/100 NO BLOCKING DEFECTS

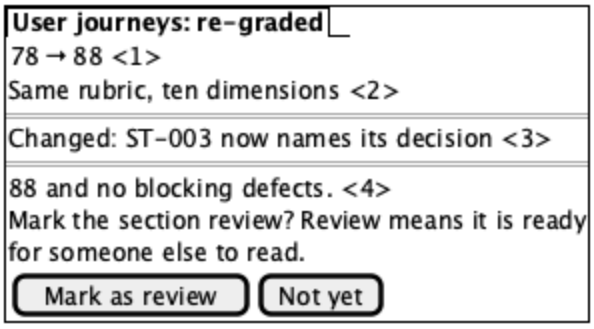


SC-001 scenario

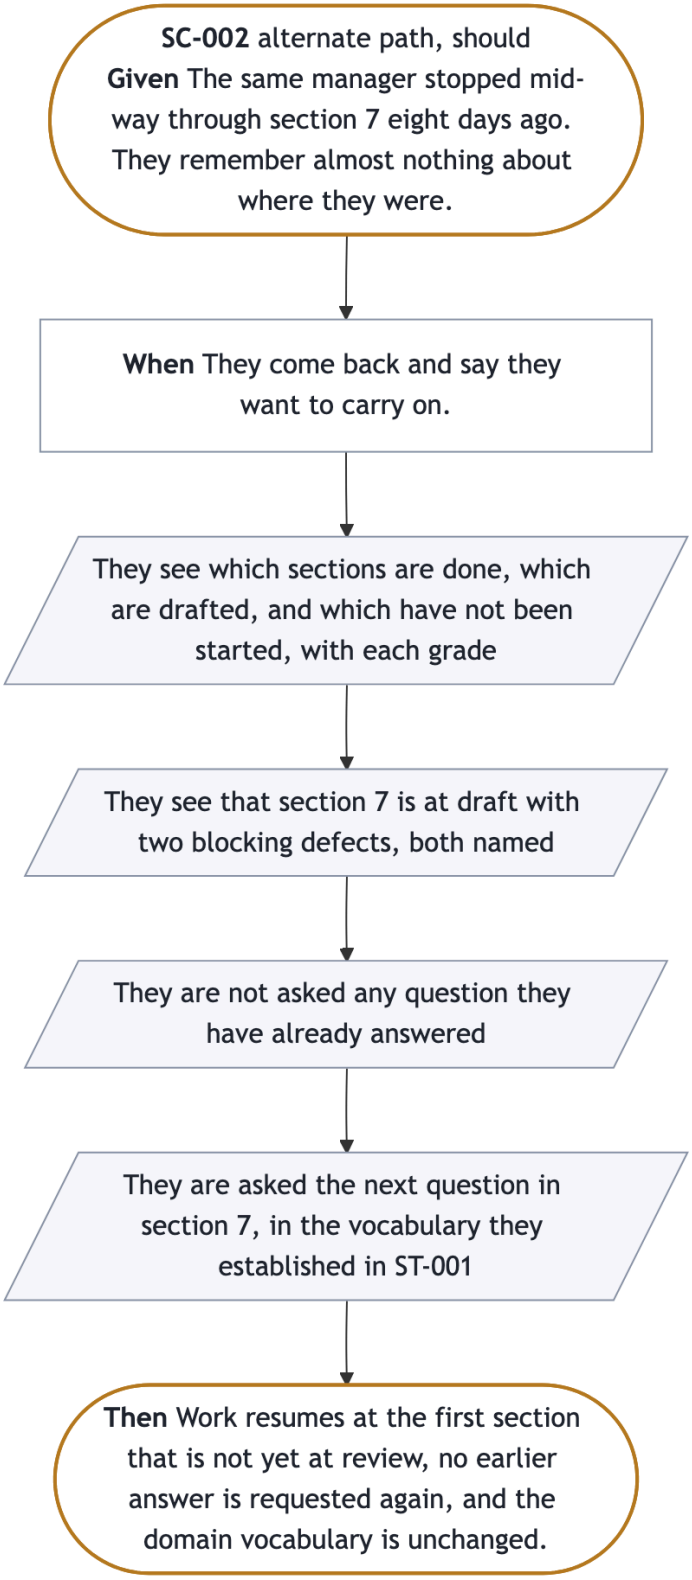


LC-001 lifecycle of a section of the PRD

WF-008 A section re-graded after one answer

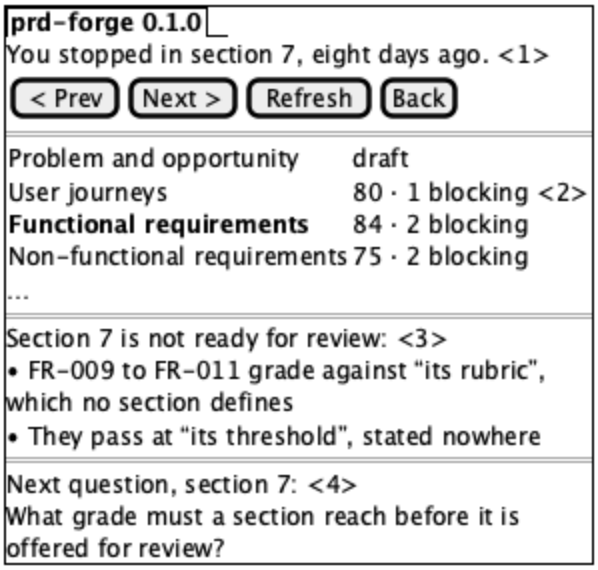


WF-008 drawing

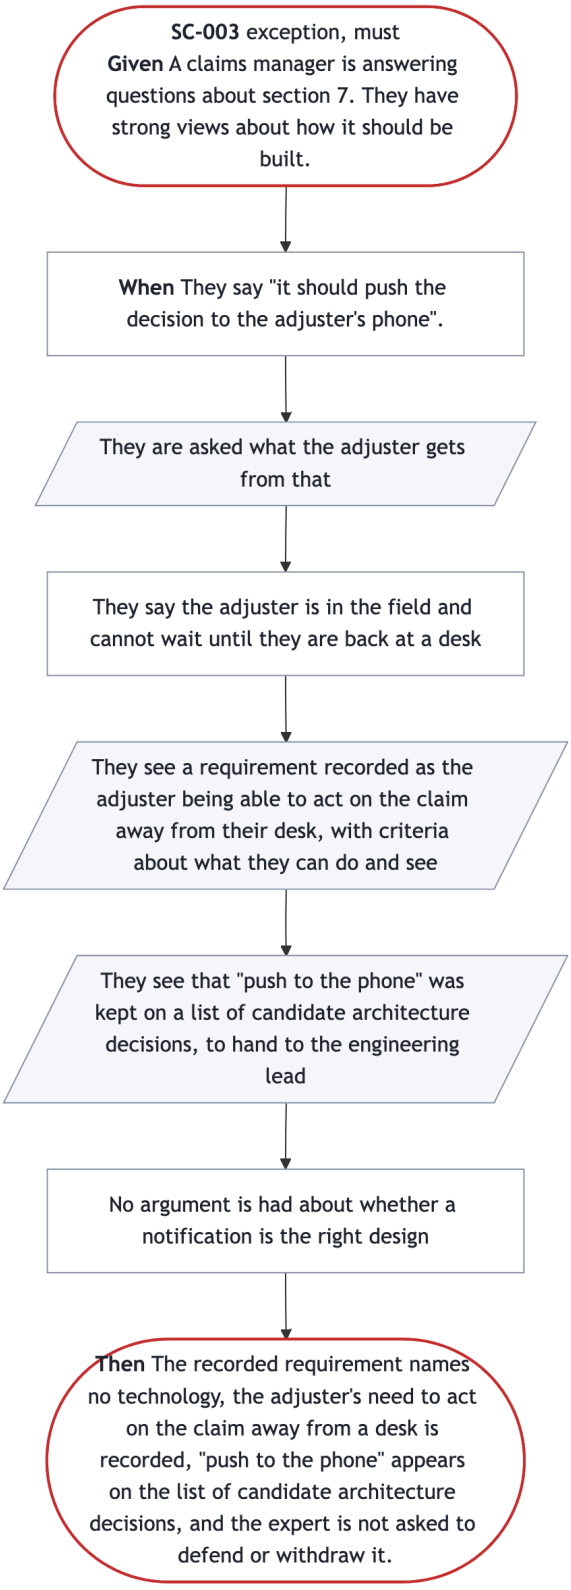


SC-002 scenario

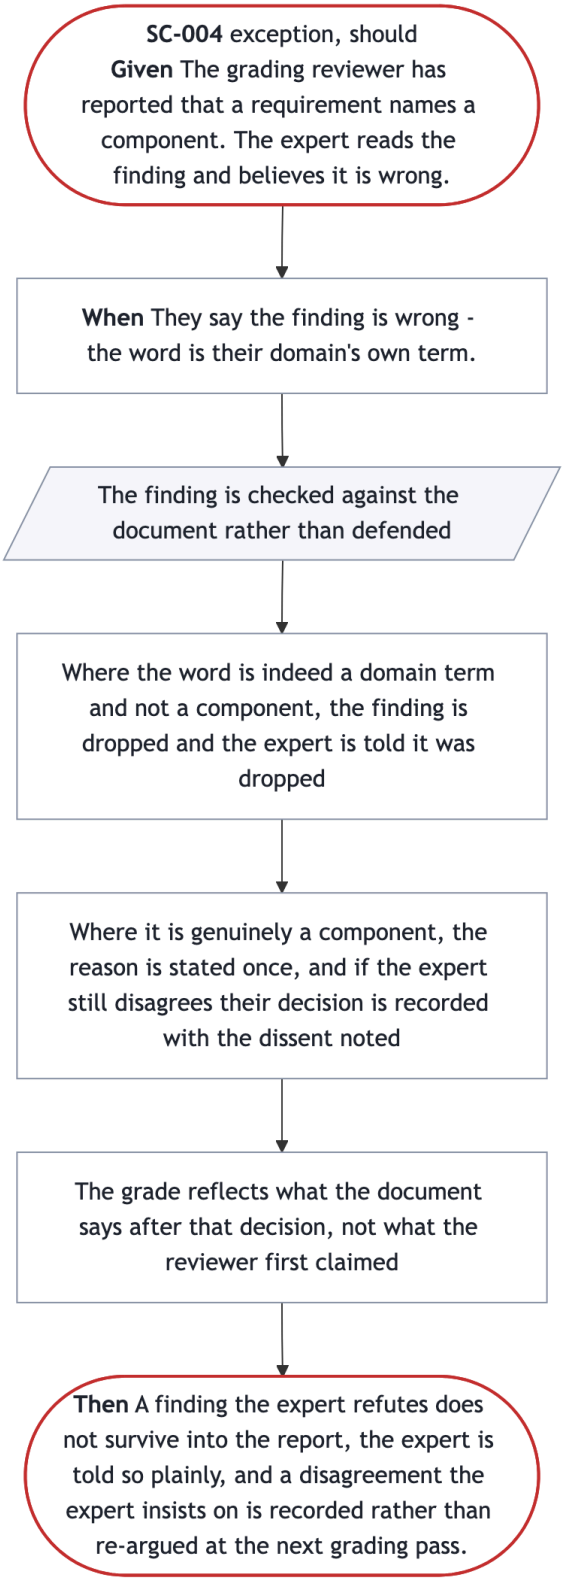
WF-010 Resuming a set after a week away



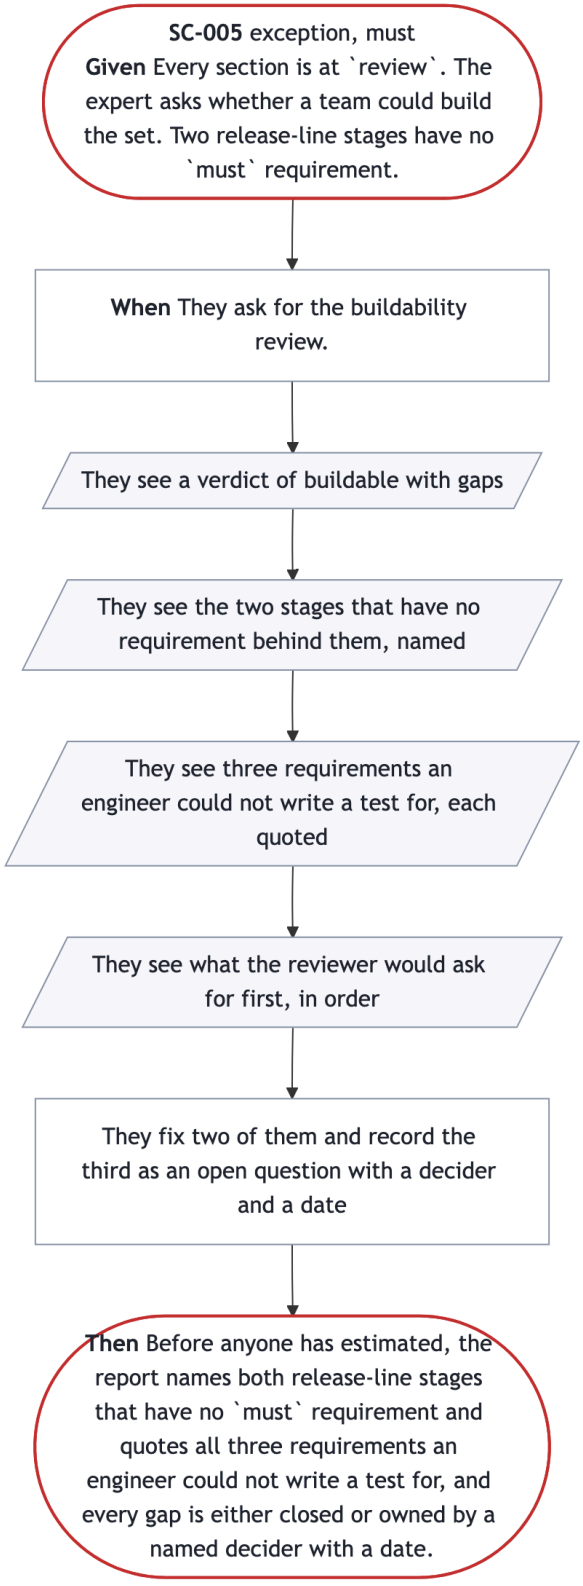
WF-010 drawing



SC-003 scenario



SC-004 scenario

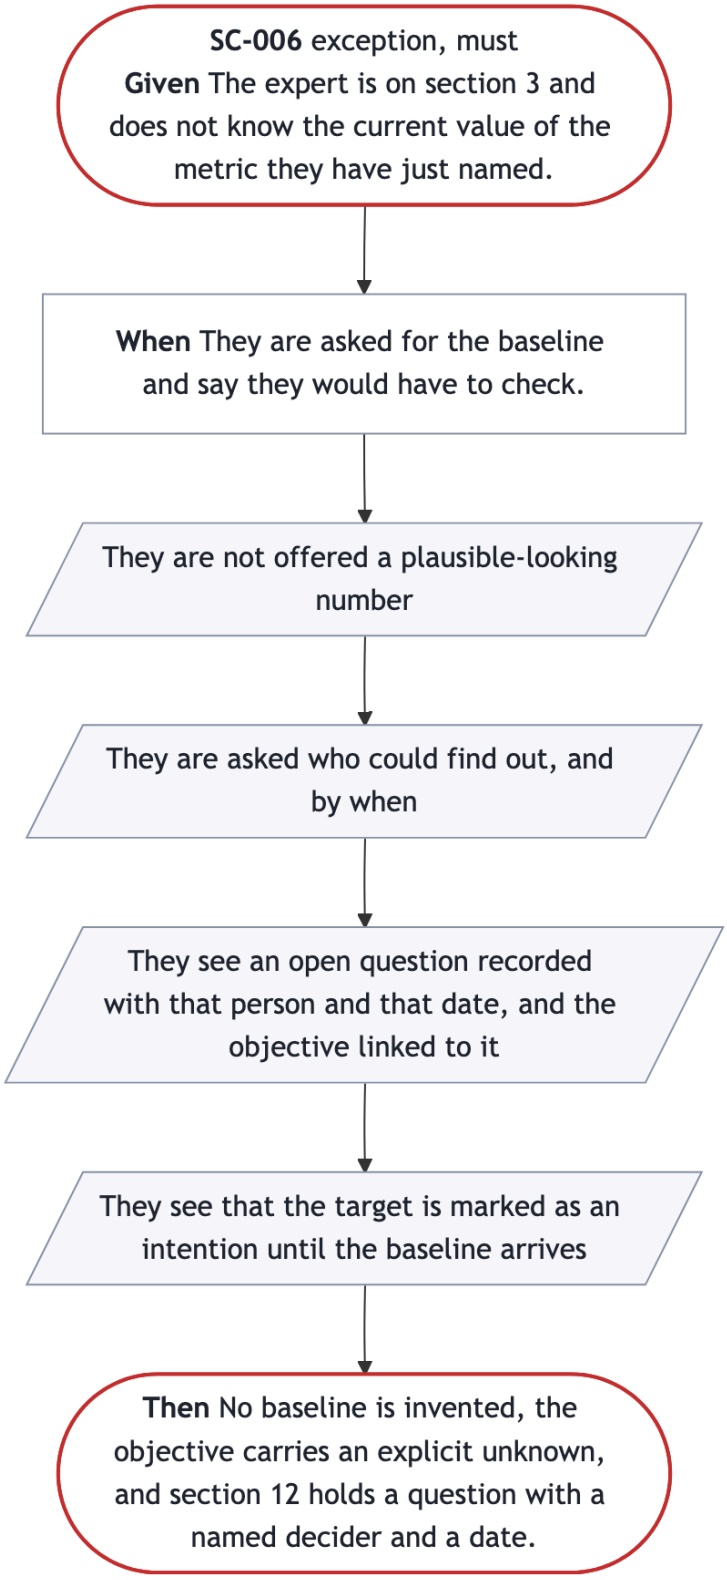


SC-005 scenario

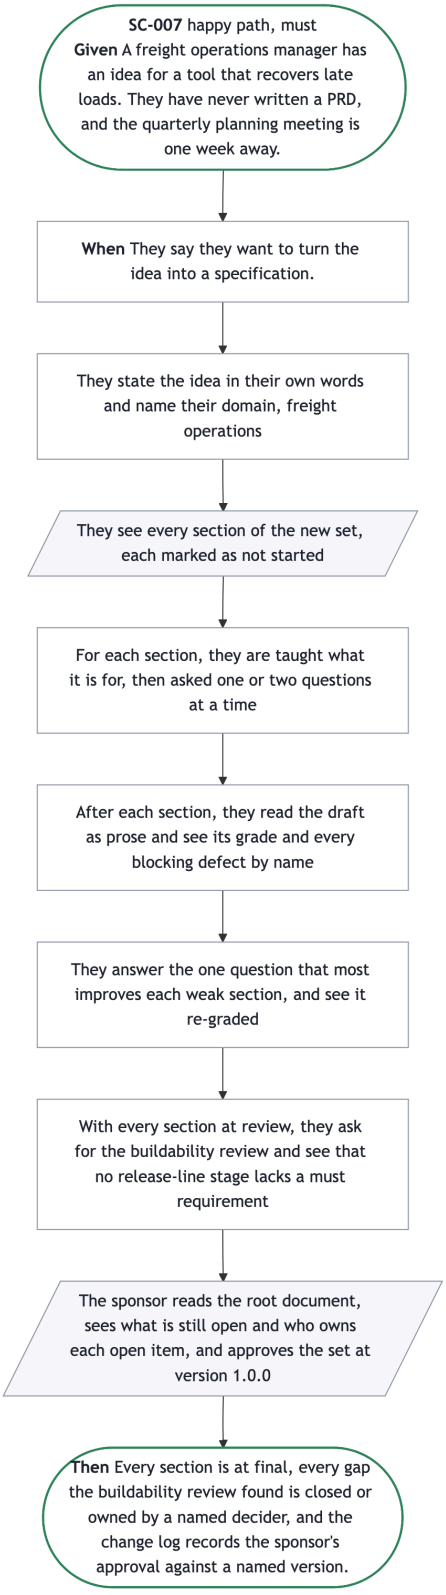
WF-009 A buildability review that finds gaps

What a team could not build from this set
Buildable with gaps. <1>
✗ ST-004 no must requirement <2> ✗ ST-005 no must requirement ✓ ST-001, ST-002, ST-003, ST-006
Each gap becomes a requirement, or an open question with a decider and a date. <3>

WF-009 drawing

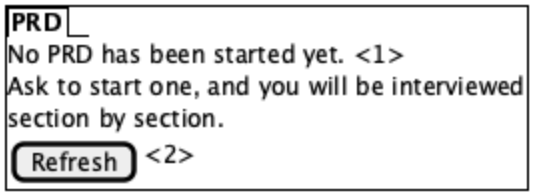


SC-006 scenario



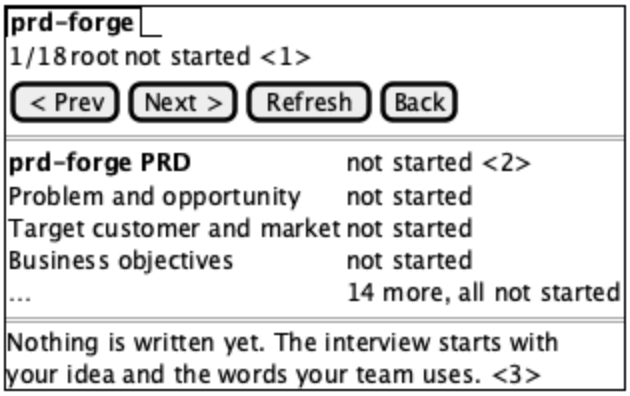
SC-007 scenario

WF-007 Before any set exists



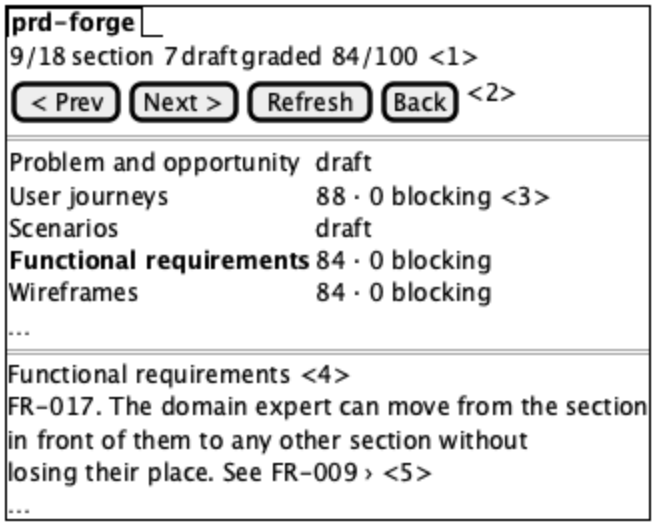
WF-007 drawing

WF-002 A new set, before the first interview



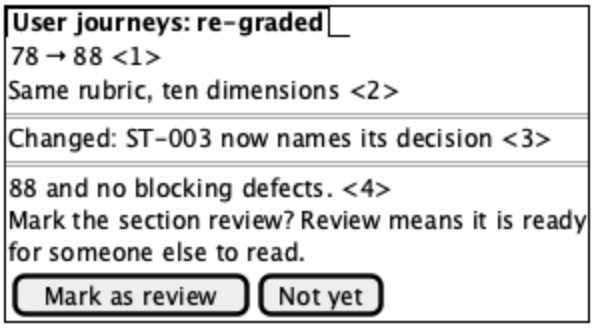
WF-002 drawing

WF-001 The set with one section open



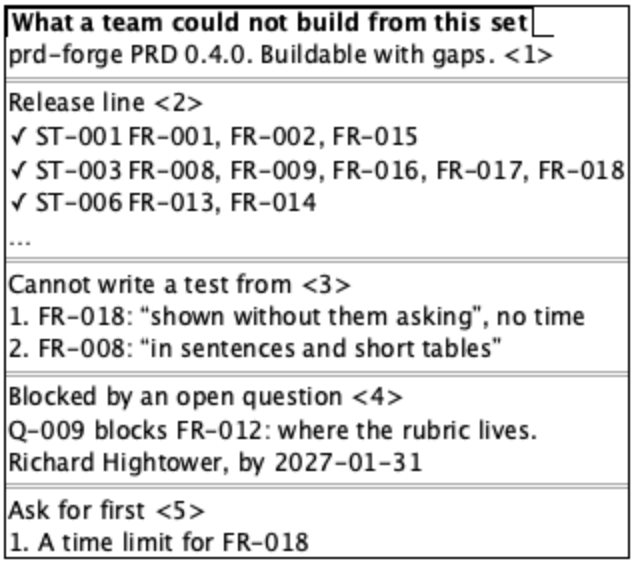
WF-001 drawing

WF-008 A section re-graded after one answer

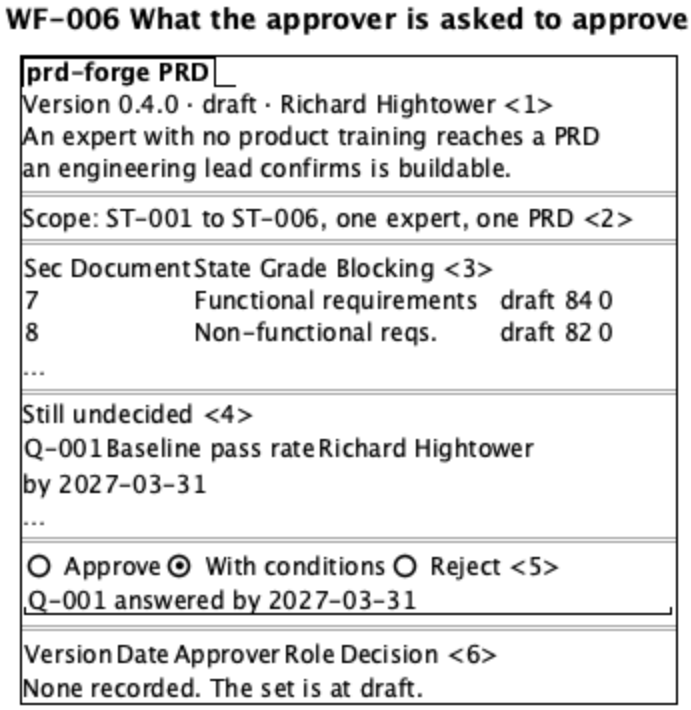


WF-008 drawing

WF-005 What a team could not build from this set



WF-005 drawing



WF-006 drawing

ID	STAGE	VARIANT	PRIORITY	WHAT IT COVERS
SC-001	ST-002	happy	must	A section taught, interviewed, drafted, graded, improved
SC-002	ST-004	alternate	should	Resuming after eight days away
SC-003	ST-002	exception	must	The expert states a technology
SC-004	ST-003	exception	should	The expert refutes a grading finding
SC-005	ST-005	exception	must	Buildability review finds holes
SC-006	ST-002	exception	must	The expert has no baseline for a metric
SC-007	ST-001 to ST-006	happy	must	One idea, from the first sentence to an approved set

SC-007 is the end-to-end acceptance test for release 1: it follows the golden thread through every release-line stage. SC-001 is the same path inside one section. SC-005 is the expert's own run of the buildability review, which is why its persona is PER-001 and not the engineering lead: no engineering lead is in release 1 at all, per OUT-004.

The priorities follow the requirements. SC-002 and SC-004 are `should`, because FR-015 and FR-016, which they exercise, are `should`.

The four exceptions are the ones that decide whether this product is trustworthy. SC-003 and SC-006 are where a specification quietly acquires a technology choice or an invented fact. SC-004 is where the grading either earns the expert's trust or loses it: a reviewer that cannot be refuted is a reviewer nobody believes. SC-005 is the moment the product is supposed to pay for itself.

Scenarios

SC-001

CONTEXT	A freight operations manager is on section 5 of their first PRD. Sections 1 and 4 are drafted and graded. They have 40 minutes before a meeting.
JOURNEY	J-001
STAGE	ST-002
STAGES	• ST-002 • ST-003 • ST-004
PERSONA	PER-001
VARIANT	happy
PRIORITY	must
SHOWN IN	screen
LIFECYCLES	• LC-001
WIREFRAMES	• WF-008
TRIGGER	They say they are ready to do the journeys section.
COURSE	• They read eight lines on what a user journey is, why requirements written one at a time miss the steps between them, and what a stage must name • They are asked to pick one person and one thing that person is trying to finish • They say "a dispatcher getting a late load recovered" • They are asked for the big steps, names only, and give five • They are asked, per stage, who acts, what they see and what they decide • They see a drafted journey with five stages, and a grade of 78 with no blocking defects • They are asked one question - what the dispatcher decides at stage three - and answer it • They see the journey re-graded at 88
EXPECTED OUTCOME	The journeys section is at status draft with a score of 88 and no blocking defects, one journey is marked as the golden thread, and every stage names an actor and a decision.

SC-002

CONTEXT
The same manager stopped mid-way through section 7 eight days ago. They remember almost nothing about where they were.

JOURNEY
J-001

STAGE
ST-004

PERSONA
PER-001

VARIANT
alternate

PRIORITY
should

SHOWN IN
screen

LIFECYCLES
- LC-001

WIREFRAMES
- WF-010

TRIGGER
They come back and say they want to carry on.

COURSE
- They see which sections are done, which are drafted, and which have not been started, with each grade
- They see that section 7 is at draft with two blocking defects, both named
- They are not asked any question they have already answered
- They are asked the next question in section 7, in the vocabulary they established in ST-001

EXPECTED OUTCOME
Work resumes at the first section that is not yet at review, no earlier answer is requested again, and the domain vocabulary is unchanged.

SC-003

CONTEXT
A claims manager is answering questions about section 7. They have strong views about how it should be built.

JOURNEY
J-001

STAGE
ST-002

PERSONA
PER-001

VARIANT
exception

PRIORITY
must

SHOWN IN
conversation

TRIGGER
They say "it should push the decision to the adjuster's phone".

COURSE
- They are asked what the adjuster gets from that
- They say the adjuster is in the field and cannot wait until they are back at a desk
- They see a requirement recorded as the adjuster being able to act on the claim away from their desk, with criteria about what they can do and see
- They see that "push to the phone" was kept on a list of candidate architecture decisions, to hand to the engineering lead
- No argument is had about whether a notification is the right design

EXPECTED OUTCOME
The recorded requirement names no technology, the adjuster's need to act on the claim away from a desk is recorded, "push to the phone" appears on the list of candidate architecture decisions, and the expert is not asked to defend or withdraw it.

SC-004

CONTEXT
The grading reviewer has reported that a requirement names a component. The expert reads the finding and believes it is wrong.

JOURNEY J-001

STAGE ST-003

PERSONA PER-001

VARIANT exception

PRIORITY should

SHOWN IN conversation

TRIGGER They say the finding is wrong - the word is their domain's own term.

COURSE
- The finding is checked against the document rather than defended
- Where the word is indeed a domain term and not a component, the finding is dropped and the expert is told it was dropped
- Where it is genuinely a component, the reason is stated once, and if the expert still disagrees their decision is recorded with the dissent noted
- The grade reflects what the document says after that decision, not what the reviewer first claimed

EXPECTED OUTCOME A finding the expert refutes does not survive into the report, the expert is told so plainly, and a disagreement the expert insists on is recorded rather than re-argued at the next grading pass.

SC-005

CONTEXT Every section is at `review`. The expert asks whether a team could build the set. Two release-line stages have no `must` requirement.

JOURNEY J-001

STAGE ST-005

PERSONA PER-001

VARIANT exception

PRIORITY must

SHOWN IN report

WIREFRAMES
- WF-009

TRIGGER They ask for the buildability review.

COURSE
- They see a verdict of buildable with gaps
- They see the two stages that have no requirement behind them, named
- They see three requirements an engineer could not write a test for, each quoted
- They see what the reviewer would ask for first, in order
- They fix two of them and record the third as an open question with a decider and a date

EXPECTED OUTCOME Before anyone has estimated, the report names both release-line stages that have no `must` requirement and quotes all three requirements an engineer could not write a test for, and every gap is either closed or owned by a named decider with a date.

SC-006

CONTEXT
The expert is on section 3 and does not know the current value of the metric they have just named.

JOURNEY
J-001

STAGE
ST-002

PERSONA
PER-001

VARIANT
exception

PRIORITY
must

SHOWN IN
conversation

TRIGGER
They are asked for the baseline and say they would have to check.

COURSE
- They are not offered a plausible-looking number
- They are asked who could find out, and by when
- They see an open question recorded with that person and that date, and the objective linked to it
- They see that the target is marked as an intention until the baseline arrives

EXPECTED OUTCOME
No baseline is invented, the objective carries an explicit unknown, and section 12 holds a question with a named decider and a date.

SC-007

CONTEXT
A freight operations manager has an idea for a tool that recovers late loads. They have never written a PRD, and the quarterly planning meeting is one week away.

JOURNEY
J-001

STAGE
ST-001

STAGES
- ST-001
- ST-002
- ST-003
- ST-004
- ST-005
- ST-006

PERSONA
PER-001

VARIANT
happy

PRIORITY
must

SHOWN IN
screen

LIFECYCLES
- LC-001

WIREFRAMES
- WF-007
- WF-002
- WF-001
- WF-008
- WF-005
- WF-006

TRIGGER
They say they want to turn the idea into a specification.

COURSE

- They state the idea in their own words and name their domain, freight operations
- They see every section of the new set, each marked as not started
- For each section, they are taught what it is for, then asked one or two questions at a time
- After each section, they read the draft as prose and see its grade and every blocking defect by name
- They answer the one question that most improves each weak section, and see it re-graded
- With every section at review, they ask for the buildability review and see that no release-line stage lacks a must requirement
- The sponsor reads the root document, sees what is still open and who owns each open item, and approves the set at version 1.0.0

EXPECTED OUTCOME

Every section is at final, every gap the buildability review found is closed or owned by a named decider, and the change log records the sponsor's approval against a named version.

Lifecycles

LC-001 · a section of the PRD

STATES

pending

MEANING — Listed in the set and not yet interviewed.

draft

MEANING — Interviewed and written, and graded or waiting for a grade.

review

MEANING — Ready for someone other than the expert to read.

final

MEANING — Approved, with the approval recorded in section 15.

INITIAL

pending

FINAL

- final

TRANSITIONS

FROM	pending
TO	draft
WHEN	The expert finishes the section's interview and it is written
RULE	The expert decides the interview is finished (TR-008)
BY	PER-001

FROM	draft
TO	review
WHEN	The expert accepts the offer of review
RULE	A grade of 70 or more, zero blocking defects, and a set that checks clean, unless the expert overrides the gate and the override is recorded in section 15
BY	PER-001

FROM	review
TO	final
WHEN	The expert sets the section to final
RULE	An independent grading pass with zero blocking defects, and an approval recorded in section 15, unless the expert overrides the gate and the override is recorded in section 15
BY	PER-001

SECTION 6

Scope

DRAFT 79/100 NO BLOCKING DEFECTS

Release line

ST-001 through ST-006 for one expert, one PRD, one golden-thread journey.

ID	IN SCOPE	STAGE	REQUIREMENTS
IN-001	Idea, domain, vocabulary, a visible empty set, and resumption	ST-001	FR-001, FR-002, FR-015
IN-002	Teach then interview, in their vocabulary	ST-002	FR-003, FR-004, FR-005
IN-003	No invented facts; convert technology talk	ST-002	FR-006, FR-007
IN-004	Draft as prose, with grade, blocking defects, and refutable findings, reached from any section	ST-003	FR-008, FR-009, FR-016, FR-017, FR-018
IN-005	One question, then re-grade	ST-004	FR-010, FR-011
IN-006	Buildability review of the whole set	ST-005	FR-012
IN-007	Approval recorded against a version	ST-006	FR-013, FR-014

Out of scope for R1

ID	OUT OF SCOPE	REVISIT
OUT-001	Deriving the backlog (J-002)	R2; the PRD already carries what it needs
OUT-002	Two experts on one set at once	After three real engagements
OUT-003	Any view beyond reading the documents and stepping through them beside the conversation	When an approver asks
OUT-004	An engineering lead's own recorded confirmation	R2, with J-002

Non-goals

NON-GOAL	WHY NOT
Supplying domain knowledge	The expert has it; the product supplies product management

NON-GOAL	WHY NOT
Deciding scope, priority, or whether to build	The expert decides; the product shows the consequences
A backlog's tasks, subtasks, estimates, or sizing	They come from engineering design, not from a PRD
Replacing the document store, issue tracker, or approval system, or being the system of record	Product decision, 0.1.0 (was C-002)
A separate application to open, install, or log into	Product decision, 0.1.0 (was C-001); the session is D-001
Product teams entering a domain they do not know, in release 1	They want the domain, which prd-forge does not supply (was S-003)
Certifying a PRD	A grade is evidence for the expert's judgement, never a warrant

The first is the one that defines the product: **we do not supply domain knowledge.** Someone without the domain is not served by this and should not buy it. The last is the one that protects its users: **a grade is evidence, not a warrant.** A high score on a PRD for the wrong product is still the wrong product.

Three non-goals moved here at 0.8.1. The system-of-record and separate-application non-goals were constraints C-002 and C-001 in section 9, and the non-goal on product teams new to a domain was segment S-003 in section 2.

In scope

IN-001

SUMMARY State the idea, establish the domain and vocabulary, and create a set with every section visible and honestly marked as not started

JOURNEY STAGE ST-001

REQUIREMENTS

- FR-001
- FR-002
- FR-015

IN-002

SUMMARY Teach a section, then interview on it one or two questions at a time, in the expert's own vocabulary

JOURNEY STAGE ST-002

REQUIREMENTS

- FR-003
- FR-004
- FR-005

IN-003

SUMMARY Never record an invented domain fact, and convert a stated technology into the need beneath it

JOURNEY STAGE ST-002

REQUIREMENTS

- FR-006
- FR-007

IN-004 · Show the draft as prose, with its grade and every blocking defect named

JOURNEY STAGE: ST-003

REQUIREMENTS:
- FR-008
- FR-009
- FR-016
- FR-017
- FR-018

IN-005 · Ask one question that most improves the section, and re-grade after it

JOURNEY STAGE: ST-004

REQUIREMENTS:
- FR-010
- FR-011

IN-006 · Review the whole set as a builder would, and report what cannot be built

JOURNEY STAGE: ST-005

REQUIREMENTS:
- FR-012

IN-007 · Record an approval, with conditions, against a named version

JOURNEY STAGE: ST-006

REQUIREMENTS:
- FR-013
- FR-014

Out of scope

OUT-001 · Deriving features, epics and stories from the approved PRD - journey J-002

REVISIT: Release 2. The PRD must already carry what J-002 needs, which FR-012 checks, so nothing is lost by deferring the derivation itself.

OUT-002 · More than one expert working on the same set at the same time

REVISIT: Once a single expert's path is proven on three real engagements.

OUT-004

SUMMARY: An engineering lead reading the set and recording their own confirmation that it can be built from

REVISIT: Release 2, with J-002. Release 1 gives the expert the buildability report (FR-012) so the gaps are found before anyone estimates; a named lead's own confirmation is a second person's time, which R1 does not ask for.

OUT-003

SUMMARY: Any view of the document set other than reading the documents themselves and stepping through them beside the conversation

REVISIT: When an approver asks for one; PER-003 reads a handful of pages, which the documents already are.

SECTION 7

Functional requirements

DRAFT 78/100 1 BLOCKING

ID	USER	CAPABILITY	PRIORITY	STAGE
FR-001	PER-001	State the idea; establish the vocabulary once	must	ST-001
FR-002	PER-001	See the whole set's state at a glance	should	ST-001
FR-003	PER-001	Be taught the section before being asked	must	ST-002
FR-004	PER-001	Be asked one or two questions at a time	must	ST-002
FR-005	PER-001	End a section and see what was left behind	should	ST-002
FR-006	PER-001	Unknown facts become open questions, never inventions	must	ST-002
FR-007	PER-001	Technology talk converted; preference preserved	must	ST-002
FR-008	PER-001	Read the draft as prose	should	ST-003
FR-009	PER-001	See score, blocking defects, weakest part	must	ST-003
FR-010	PER-001	Be asked one question, not twelve	must	ST-004
FR-011	PER-001	See the re-grade against the same standard	should	ST-004
FR-012	PER-001	See what a team could not build from it	must	ST-005
FR-013	PER-003	See what is approved and what is open, on one page	must	ST-006
FR-014	PER-003	Approve, with conditions, against a version	must	ST-006
FR-015	PER-001	Stop and resume without being asked anything twice	should	ST-001
FR-016	PER-001	Refute a finding and have it withdrawn	should	ST-003
FR-017	PER-001	Move to any section without losing their place	should	ST-003
FR-018	PER-001	See the set as it now stands	should	ST-003

Every release-line stage has at least one `must` requirement. Ten of the 18 are `must`: without any one of them, REL-001 does not do what it promises. The eight `should` requirements make the work faster or easier to follow, and the expert can still reach an approved set without them. The expert set this split at version 0.8.0.

Six of the `must` requirements wait on an open question in section 12. FR-004, FR-009, FR-010, and FR-012 wait on Q-009, where the rubric and the derivability checks live and who owns them. FR-010 also waits on Q-007, whether the working order is right and how looping back to the journeys feels, because Q-007 blocks ST-004. FR-013 waits on Q-005, whether a sponsor accepts a specification that carries a grade. FR-014 waits on Q-011, how an approver records a decision.

FR-015 and FR-016 were split out of FR-002 and FR-009 at version 0.2.0, because each has a failure mode of its own: a build could deliver perfect set-state visibility and still re-ask three answered questions, and could report a grade perfectly

and still defend a finding the document does not support. Ids are never renumbered, so they sit at the end rather than in stage order.

A judgement call worth challenging

FR-002, FR-008 and FR-013 refer to "the set", "the section", "written English" and "the root document". A reviewer may read those as naming an output format and therefore failing test 5. The judgement taken here is that a written specification is the product's subject, not its delivery technology, so these words name the artifact rather than the mechanism. What is deliberately absent everywhere: the assistant the product runs inside, how grading is performed, who or what reviews the set, and what the documents are stored as. Those are in section 6 as non-goals, in section 11 as dependencies, or nowhere.

FR-001, FR-009, FR-013 and FR-014 each read as more than one capability, and a reviewer has called them out for it. They are left whole deliberately: each is one act with several parts, and a build delivering half of one would be visibly unfinished rather than arguably done. FR-015 and FR-016 were split precisely because they were not like that, each had a failure mode the other half could hide.

> **FR-001 · State the idea in their own words and establish the domain's vocabulary once**
>
> | | |
> |---|---|
> | USER | PER-001 |
> | OUTCOME | Every question they are asked afterwards uses their unit of work and their terms, so they are never translating. |
> | ACCEPTANCE | • Given an expert who says the unit of work is a "load", when any later question refers to that unit, then it says "load" and not "item", "record" or "entity".
• Given three words the expert says an outsider would misread, when the first section is drafted, then each of those words appears in the glossary with the expert's own meaning.
• Given an expert who has not stated a domain, when they ask to begin a section, then they are asked for the domain and the unit of work first. |
> | PRIORITY | must |
> | SHOWN IN | conversation |
> | JOURNEY STAGE | ST-001 |
> | OBJECTIVE | O-003 |

FR-002 · See every section of the set and what state each one is in, at any point

USER	PER-001
OUTCOME	They can see at a glance what is done, what is weak and what has not been started, without opening any section.
ACCEPTANCE	• Given a set where four section documents are drafted and thirteen are untouched, when the expert asks where things stand, then they see all seventeen section documents with each one's state, and the untouched ones are shown as not started rather than omitted. • Given a section that has been graded, when the expert sees the set's state, then that section shows its score and how many blocking defects it has. • Given a set where no section has been started, when the expert asks where things stand, then every section is listed and shown as not started.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-001
OBJECTIVE	O-002

FR-003

CAPABILITY	Be taught what a section is, why it matters and what belongs in it, before being asked anything about it
USER	PER-001
OUTCOME	Before the first question of a section reaches them, what the section is for and what belongs in it is in front of them, in their own domain's terms.
ACCEPTANCE	• Given a section the expert has not done before, when they begin it, then they are first told what it is, why it matters and what belongs in it, in no more than ten lines. • Given a stated domain, when a section is taught, then the example used is in that domain and not in another. • Given an expert who says they already know the section, when they ask to skip the teaching, then the questions begin and the teaching is not forced.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-003

FR-004 · Be asked one or two questions at a time, never a questionnaire

USER	PER-001
OUTCOME	No message puts more than two questions to them, so each answer they give addresses one thing.
ACCEPTANCE	• Given any section's interview, when the expert is asked for input, then at most two questions are put to them at once. • Given an answer that leaves a field the section needs still unfillable, or that could be written into the section in two ways differing in meaning, when it is received, then a follow-up is asked about that answer before the next topic is opened. • Given an expert whose answers would already score the section 70 or more with no blocking defects, when further questions remain unasked on the list, then the interview stops rather than continuing to the end of it.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-003

FR-005 · Say that they have nothing more to add, and have that end the section's interview

USER	PER-001
OUTCOME	They are never trapped in an interview, and what they chose to leave incomplete is recorded rather than hidden.
ACCEPTANCE	• Given an expert who says they are done with a section, when the section does not yet pass its rubric, then it is written anyway at draft, and what is missing is recorded as an open question or named plainly to them. • Given an expert who ends a section early, when they are told what was left behind, then the statement names the specific gap, not just a score.
PRIORITY	should
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-002

FR-006

CAPABILITY	Have a fact they do not know recorded as an open question with a decider and a date, rather than filled in for them
USER	PER-001
OUTCOME	Nothing in their PRD is a number or a citation that nobody can stand behind.
ACCEPTANCE	• Given an expert who says they would have to check a figure, when the section is drafted, then no figure is written in its place and an open question records who could find out and by when. • Given an objective whose baseline is unknown, when the section is drafted, then the baseline says so explicitly and an open question exists for it. • Given a regulation or standard the expert has not named, when a constraint is drafted, then no regulation number appears that the expert did not supply.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-001

FR-007

CAPABILITY	State a technology they want and have the need beneath it recorded instead, with their preference preserved for the engineering lead
USER	PER-001
OUTCOME	Their requirement names no technology, their reasoning is not lost, and they do not feel overruled.
ACCEPTANCE	• Given an expert who names a technology, when they are asked what the user gets from it and answer, then the recorded requirement states that outcome and names no technology, component, vendor or screen. • Given a technology the expert stated, when the set is handed over, then that preference appears on a list of candidate architecture decisions for the engineering lead. • Given an expert who insists the technology belongs in the requirement, when they have been told once why it does not, then their decision is recorded and is not argued again at the next grading pass.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-002
OBJECTIVE	O-001

FR-008 · Read the drafted section as prose rather than as a filled-in form

USER	PER-001
OUTCOME	They can judge from the prose alone whether it says what they meant, and every id, priority and link the team builds from is the one they read.
ACCEPTANCE	• Given a drafted section, when the expert reviews it, then what they are shown reads as written English, in sentences and short tables, and they can judge every item in the section from it alone. • Given an item whose id, priority or link differs between what the team builds from and what the expert reads in the prose, when the section is read, then the expert is told of the difference rather than it being silently resolved.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-003

FR-009 · See a score, every blocking defect by name, and which part of the section is weakest

USER	PER-001
OUTCOME	They know how good the section is against a stated standard, and what specifically is wrong, rather than being told it looks good.
ACCEPTANCE	• Given a graded section, when the expert reads the grade, then they see a score, a count of blocking defects, and each blocking defect stated with where it is. • Given a section scoring 70 or more with one blocking defect, when the grade is reported, then the section is not offered as ready, and the defect is named. • Given a graded section, when the expert reads the grade, then the weakest part is named and why it is weak is stated in no more than three lines.
PRIORITY	must
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-001

FR-010

CAPABILITY	Be asked the single question that would most improve the section, not a list of everything wrong with it
USER	PER-001
OUTCOME	They can act immediately, and the section improves one answer at a time instead of stalling under a list.
ACCEPTANCE	• Given a graded section scoring under 70, or carrying any blocking defect, when the expert is asked to improve it, then exactly one question is put to them. • Given that one question, when it is asked, then it addresses the weakest part that was just named, and not a different part.
PRIORITY	must
SHOWN IN	conversation
JOURNEY STAGE	ST-004
OBJECTIVE	O-003

FR-011 · See the section re-graded against the same standard after answering

USER	PER-001
OUTCOME	They can see whether their answer actually improved the section, and by how much.
ACCEPTANCE	• Given an answered question, when the section is re-graded, then it is scored against the same rubric as before and both scores are visible. • Given a re-grade, when the score has risen, then the document has changed; a score never rises without a change to the document. • Given a section scoring 70 or more with no blocking defects, when it is re-graded, then the expert is offered the state `review`, and is told that `review` means the section is ready for someone else to read.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-004
OBJECTIVE	O-003

FR-012

CAPABILITY	See what a development team would not be able to build from the set, before anyone has estimated it
USER	PER-001
OUTCOME	They learn where their own specification fails while it is still cheap to fix, rather than at iteration three.
ACCEPTANCE	• Given a set in which every section has been written at least once, when the expert asks for a buildability review, then they see a verdict and, for each release-line stage, whether a requirement covers it or it is a gap. • Given a requirement whose acceptance criteria an engineer who was not in the interview judges insufficient to write a test from, when the review reports it, then the requirement is named and the document is quoted. • Given an open question that blocks a `must` requirement, when the review runs, then that requirement's wait on the question is reported. • Given a review finding, when it is reported to the expert, then it has been checked against the document and anything not supported by the document has been dropped.
PRIORITY	must
SHOWN IN	report
JOURNEY STAGE	ST-005
OBJECTIVE	O-001

FR-013

CAPABILITY	See, in the set's own root document, what is being approved, what is still undecided and who owns each undecided thing
USER	PER-003
OUTCOME	They can commit budget and a date knowing what they are committing to and what is still open.
ACCEPTANCE	• Given a set offered for approval, when the approver opens its root document, then that one document carries the summary, the scope, every section's state and grade, and every open question with its decider and date, without the approver opening any other document. • Given a set with any blocking defect anywhere, when the approver opens it, then they see that, with the sections it is in.
PRIORITY	must
SHOWN IN	report
JOURNEY STAGE	ST-006
OBJECTIVE	O-001

FR-014

CAPABILITY	Approve, approve with conditions, or reject, and have that recorded against a named version
USER	PER-003
OUTCOME	What was approved, when, by whom and on what conditions is answerable later without asking anyone.
ACCEPTANCE	• Given an approver who approves with conditions, when they state the conditions, then each condition names the open question or requirement it depends on and a date. • Given a recorded approval, when the set is read later, then the version approved, the approver, their role, the date and the decision are all present. • Given an approved set that is then changed, when the change is written, then a new version is recorded and the earlier approval is not represented as covering it.
PRIORITY	must
SHOWN IN	report
JOURNEY STAGE	ST-006
OBJECTIVE	O-001

FR-015 · Stop in the middle of a set and resume days later without being asked anything twice

USER	PER-001
OUTCOME	They pick up at the first unfinished section, and nothing they have already told the product is asked for again.
ACCEPTANCE	• Given an expert returning after eight days, when they say they want to carry on, then work resumes at the first section not yet at `review`. • Given an expert who has already stated their domain and vocabulary, when they resume, then they are not asked for either again and the vocabulary in use is unchanged. • Given an expert who answered a question in a section before stopping, when they resume that section, then that question is not put to them again.
PRIORITY	should
SHOWN IN	conversation
JOURNEY STAGE	ST-001
OBJECTIVE	O-002

FR-016 · Refute a grading finding by pointing at the document, and have it withdrawn

USER	PER-001
OUTCOME	They are told the finding was dropped, and it does not reappear at the next grading pass.
ACCEPTANCE	• Given a finding the expert refutes, when it is checked against the document and the document does not support it, then it is dropped and the expert is told it was dropped. • Given a finding the document does support, when the expert disputes it, then the reason is stated once and their decision is recorded. • Given a decision the expert reaffirmed after an objection, when the section is graded again, then the objection is not raised as a new finding.
PRIORITY	should
SHOWN IN	conversation
JOURNEY STAGE	ST-003
OBJECTIVE	O-001

FR-017

CAPABILITY	Move from the section in front of them to any other section of the set without losing their place
USER	PER-001
OUTCOME	They reach the section they want in one step, and can always return to the view they were on before.
ACCEPTANCE	• Given a section in front of them, when they ask for the next or the previous section, then the section after or before it in reading order is shown, and at the first or the last section the view does not change. • Given a section that cites an ID another section defines, such as Q-001, when they follow that ID, then the section that defines it is shown and the ID is named above it. • Given they have moved from a section to three other things in the set, one after another, when they ask to go back three times, then they see each earlier thing again, the most recent first, and end at the section they started from.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-002

FR-018 · See the set as it now stands after it was changed outside the view in front of them

USER	PER-001
OUTCOME	What they read is never older than the last change written to the set.
ACCEPTANCE	• Given a section changed after they opened the set, when they ask to refresh, then that section shows its new text and its new state. • Given a section that is graded again while the set is open, when the new grade is written, then the new grade is shown without them asking. • Given no set existed when they looked, when a set is written and they ask to refresh, then every section of it is listed.
PRIORITY	should
SHOWN IN	screen
JOURNEY STAGE	ST-003
OBJECTIVE	O-002

SECTION 7

Wireframes

DRAFT 88/100 NO BLOCKING DEFECTS

Every requirement in section 7 whose outcome a person meets on a screen or in a report is drawn here as a text sketch. Each sketch shows what the person sees and what they can do there, with the set's own content. None of them decides how anything looks or how it is built.

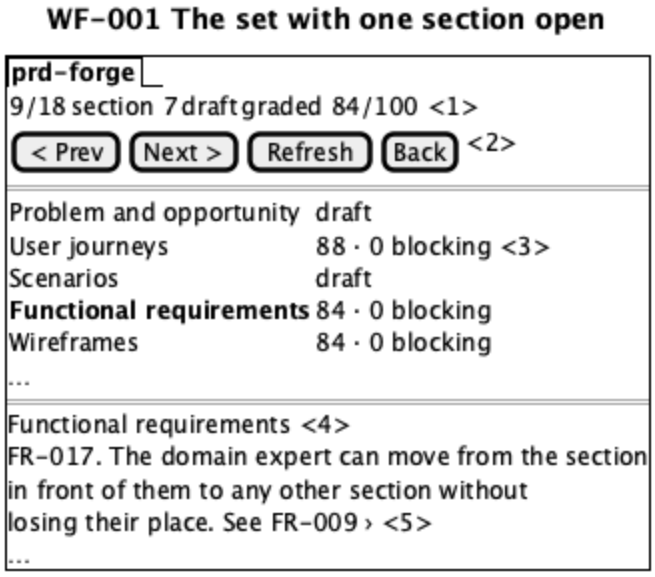
Five sketches follow the domain expert through the golden thread: a new set with nothing started, the set with a section open, why a section got its grade, the rubric behind it, and a re-grade. Two reports decide whether the set is buildable: the buildability review and the approval pack. Three smaller sketches draw the other states: no set at all, which FR-018's third criterion names; the review that finds gaps, which SC-005 names; and the return after a week away, which SC-002 names.

The numbered marks in each sketch are explained, in order, by its `shows` list. Each sketch is also drawn in PlantUML Salt for the printed PRD.

WF-001 The set with one section open

Illustration, not a design. Illustrates FR-002, FR-008, FR-017, and FR-018. PER-001 sees it at ST-003.

The expert reads a drafted section with the whole set beside it. Every graded section shows its score and its blocking count, so what is weak is visible without opening anything.



WF-001 drawing

WF-002 A new set, before the first interview

Illustration, not a design. Illustrates FR-002, FR-017, and FR-018. PER-001 sees it at ST-001.

At ST-001 the set exists and nothing in it is written. Every section is listed and marked not started, so the expert sees the whole shape of the work before the first question.

WF-002 A new set, before the first interview

prd-forge
1/18 root not started <1>
< Prev | Next > | Refresh | Back

prd-forge PRD	not started <2>
Problem and opportunity	not started
Target customer and market	not started
Business objectives	not started
...	14 more, all not started

Nothing is written yet. The interview starts with your idea and the words your team uses. <3>

WF-002 drawing

WF-003 Why a section got its grade

Illustration, not a design. Illustrates FR-009 and FR-017. PER-001 sees it at ST-003.

A score alone does not tell the expert what to fix. This view names each blocking defect and where it is, and names the weakest part, quoting the section.

WF-003 Why a section got its grade

Wireframes: graded 69/100
1 blocking defect. Not ready for review. <1>

Blocking <2>
• WF-002 names a part of the product and where files are kept. In WF-002, lines 3 and 4.

Weakest: no design, 1 of 4 <3>
WF-002 names a part of the product, which section 7 keeps out of every sketch.

Dimension	Score <4>
Marks explained	4
Every control earns...	4
The person's view	2

WF-001 shows ST-003 content at ST-001
...
› The rubric for wireframes <5>

WF-003 drawing

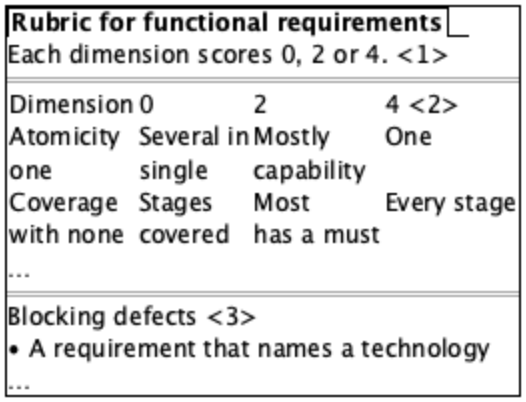
WF-004 The standard a section is graded against

Illustration, not a design. Illustrates FR-009 and FR-017. PER-001 sees it at ST-003.

A grade against a stated standard is only fair if the expert can read the standard. This is the same rubric the graders used.

WF-004 The standard a section is graded against

Rubric for functional requirements Each dimension scores 0, 2 or 4. <1>			
Dimension	0	2	4 <2>
Atomicity	Several in one	Mostly single capability	One
Coverage	Stages with none	Most covered	Every stage has a must
...			
Blocking defects <3> • A requirement that names a technology ...			



WF-004 drawing

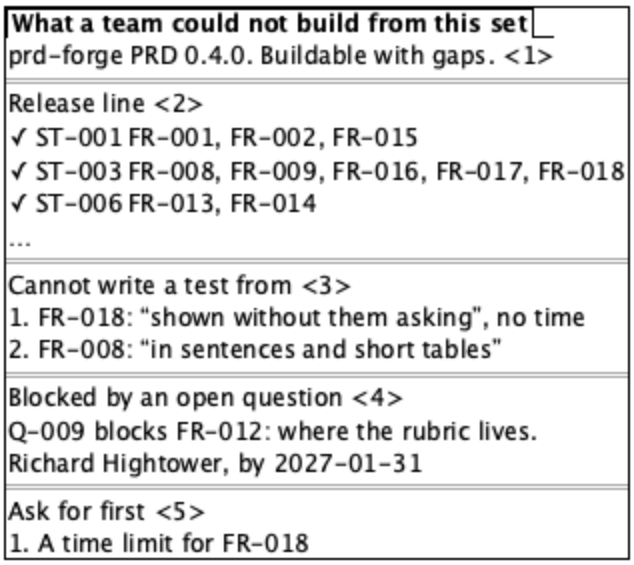
WF-005 What a team could not build from this set

Illustration, not a design. Illustrates FR-012. PER-001 sees it at ST-005.

The expert reads this before anyone has estimated anything, while a fix is still a sentence. Every finding quotes the set, so the expert can check it rather than trust it.

WF-005 What a team could not build from this set

What a team could not build from this set prd-forge PRD 0.4.0. Buildable with gaps. <1>
Release line <2> ✓ ST-001 FR-001, FR-002, FR-015 ✓ ST-003 FR-008, FR-009, FR-016, FR-017, FR-018 ✓ ST-006 FR-013, FR-014 ...
Cannot write a test from <3> 1. FR-018: "shown without them asking", no time 2. FR-008: "in sentences and short tables"
Blocked by an open question <4> Q-009 blocks FR-012: where the rubric lives. Richard Hightower, by 2027-01-31
Ask for first <5> 1. A time limit for FR-018

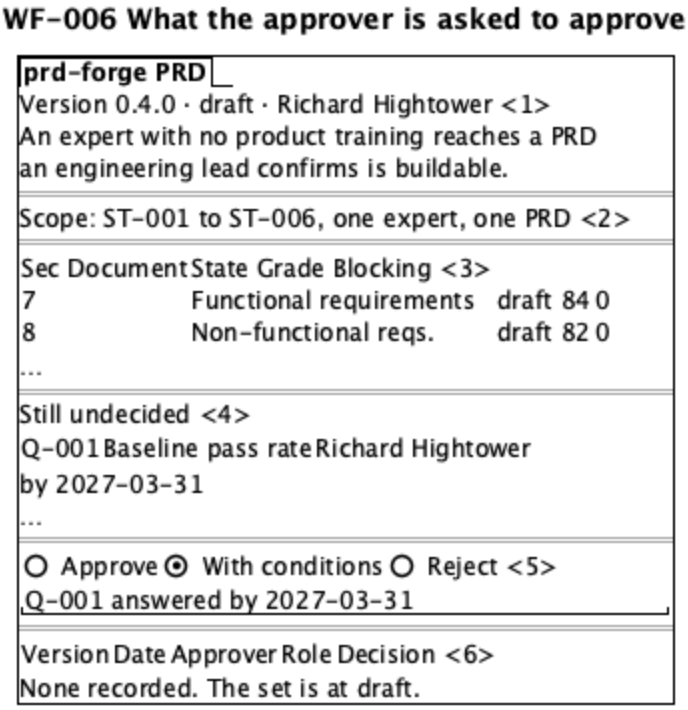


WF-005 drawing

WF-006 What the approver is asked to approve

Illustration, not a design. Illustrates FR-013 and FR-014. PER-003 sees it at ST-006.

The approver reads one document, not eighteen. It says what they are committing to, what is still open and who owns it, and where their decision is recorded.

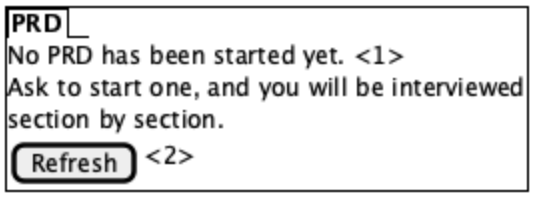


WF-006 drawing

WF-007 Before any set exists

Illustration, not a design. Illustrates FR-018. PER-001 sees it at ST-001.

Before the first interview there is no set to list. The view says so plainly, and a refresh lists every section once the set is written.

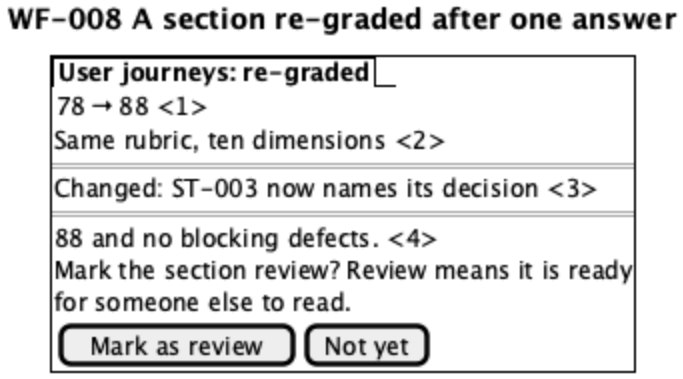


WF-007 drawing

WF-008 A section re-graded after one answer

Illustration, not a design. Illustrates FR-011. PER-001 sees it at ST-004.

After one answer the section is graded again against the same rubric. Both scores are visible, and at 70 or more with no blocking defects the expert is offered review. The values are SC-001's.

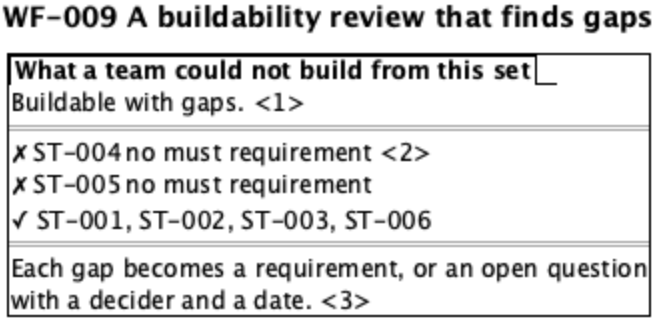


WF-008 drawing

WF-009 A buildability review that finds gaps

Illustration, not a design. Illustrates FR-012. PER-001 sees it at ST-005.

SC-005 is the review that finds work missing: two release-line stages have no must requirement behind them. The gaps are named first, before anything else in the review.

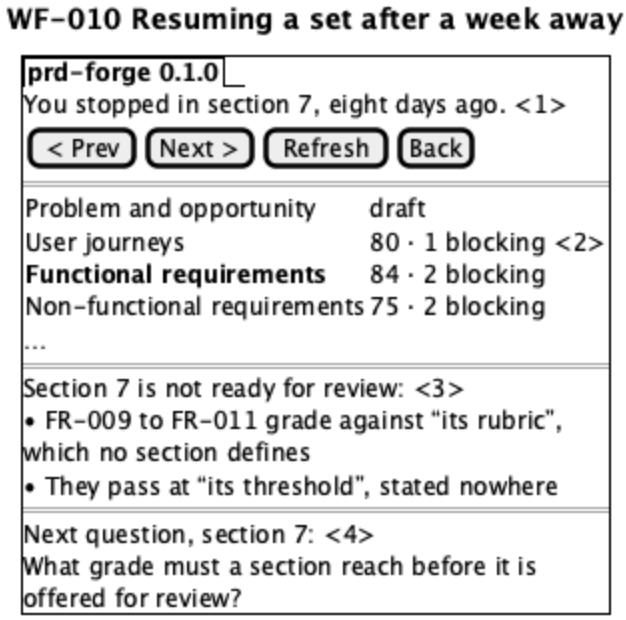


WF-009 drawing

WF-010 Resuming a set after a week away

Illustration, not a design. Illustrates FR-002 and FR-009. PER-001 sees it at ST-004.

SC-002 is the expert who returns after eight days and remembers nothing. The values are prd-forge's own at version 0.1.0, when section 7 scored 84 with two blocking defects: a grade of 70 or more that is still not offered as ready.



WF-010 drawing

WF-001 · The set with one section open

TYPE screen

PERSONA PER-001

JOURNEY STAGE ST-003

ILLUSTRATES

- FR-002
- FR-008
- FR-017
- FR-018

SHOWS

- Where they are in the set, the section's state, and its grade
- Next, previous, refresh, and back to the view before
- Each section's grade and blocking count; opening it shows why
- The section itself, as prose
- An ID that another section defines opens that section

WF-002 · A new set, before the first interview

TYPE screen

PERSONA PER-001

JOURNEY STAGE ST-001

ILLUSTRATES

- FR-002
- FR-017
- FR-018

SHOWS

- Where they are, and that nothing is started yet
- Every section listed, each marked not started
- What happens first

WF-003 · Why a section got its grade

TYPE	screen
PERSONA	PER-001
JOURNEY STAGE	ST-003
ILLUSTRATES	• FR-009 • FR-017
SHOWS	• The score, and that the section is not ready • Each blocking defect, and where it is • The weakest part, and why, in three lines or fewer • Each dimension's score, with the evidence for it • The standard the grade was given against

WF-004 · The standard a section is graded against

TYPE	screen
PERSONA	PER-001
JOURNEY STAGE	ST-003
ILLUSTRATES	• FR-009 • FR-017
SHOWS	• How a dimension becomes a grade • What 0, 2 and 4 look like for each dimension • What stops a section at any grade

WF-005 · What a team could not build from this set

TYPE	report
PERSONA	PER-001
JOURNEY STAGE	ST-005
ILLUSTRATES	• FR-012
SHOWS	• Which version of the set was reviewed, and the verdict • Each release-line stage, covered or a gap • Each requirement an engineer could not test, quoted • An open question that blocks a must, its decider and date • What the reviewer would ask for first

WF-006 · What the approver is asked to approve

TYPE report

PERSONA PER-003

JOURNEY STAGE ST-006

ILLUSTRATES
- FR-013
- FR-014

SHOWS
- The version, its state, its owner, and what it is for
- What release 1 covers
- Every section with its state and grade, and what blocks
- What is still undecided, who decides it, and by when
- The decision, with any condition tied to a question
- Who decided what, in which role, on which version

WF-007 · Before any set exists

TYPE screen

PERSONA PER-001

JOURNEY STAGE ST-001

ILLUSTRATES
- FR-018

SHOWS
- That nothing has been started
- Look again once a set exists

WF-008 · A section re-graded after one answer

TYPE screen

PERSONA PER-001

JOURNEY STAGE ST-004

ILLUSTRATES
- FR-011

SHOWS
- The new score beside the old one
- The same standard both times
- What changed in the section
- The offer of review, and what review means

WF-009 · A buildability review that finds gaps

TYPE report

PERSONA PER-001

JOURNEY STAGE ST-005

ILLUSTRATES
- FR-012

SHOWS
- The verdict
- The release-line stages with no requirement, named
- That each gap must become a requirement or a question

WF-010 · Resuming a set after a week away

TYPE screen

PERSONA PER-001

JOURNEY STAGE ST-004

ILLUSTRATES

- FR-002
- FR-009

SHOWS

- Where they stopped
- Every section's grade and blocking count
- Each blocking defect in the section, named
- The next question, not one already answered

SECTION 8

Non-functional requirements

DRAFT 82/100 NO BLOCKING DEFECTS

ID	CATEGORY	TARGET
NFR-001	performance	Next question under 20s at p95
NFR-002	performance	Section grade under 90s at p95
NFR-003	scale	120 requirements, 10 journeys x 10 stages, 200 terms
NFR-004	durability	100% of a set readable with no part of the product running
NFR-005	usability	90% of sections completed without asking what a term means
NFR-006	accessibility	Zero WCAG 2.2 AA violations on the golden-thread path
NFR-007	observability	100% of grades reconstructable after the session
NFR-008	privacy	Zero destinations beyond the session the expert already uses
NFR-009	supportability	100% of older sets readable within a major version
NFR-010	retention	100% retrievable from the organisation's own store, no expiry imposed by us
NFR-011	availability	Zero answers lost per interruption; resume within 1 minute
NFR-012	security	Zero read paths the organisation's own access rules do not govern

Where these came from

NFR-001 and NFR-002 come from PER-001's frequency: a burst of intense use. A 20-second wait is tolerable once; across 60 questions it is 20 minutes of dead time and the expert stops.

NFR-004 is the one that matters most and is cheapest to get wrong. PER-001 works in bursts weeks apart. If the set is only readable through this product, then a version change, a tool change or an expired session destroys weeks of an expert's thinking. The target is deliberately absolute.

NFR-005 is the only requirement here that needs people rather than instruments to measure, and it is the one most likely to be skipped. It is the direct test of whether P-004 was solved.

NFR-006 covers the approver too, not only the expert: PER-003 reads this once and must not need a mouse to do it.

There is no concurrency target, and that is a decision rather than an omission. Under D-001 the product reaches one expert inside a session they are already in, so simultaneous users are the host's load and not ours; NFR-003 therefore reads scale as the size of one set, which is the dimension we do control. If D-001 or the separate-application non-goal in section 6 ever changes, concurrency comes back and this paragraph goes.

NFR-006 and NFR-012 set targets on things the non-goals in section 6 say we do not own: what the host renders, and where the customer keeps its documents. They are kept as `must` because a target nobody owns is how accessibility and access control get lost; who owns them is Q-011.

NFR-010, NFR-011 and NFR-012 were added at version 0.2.0. Retention, availability and security had simply not been reached, and C-003, now assumption A-005, stated a retention rule that no measurable requirement stood behind. NFR-010 is the one that answers A-005's second half: the product retains nothing of its own, so the only retention question is how long the organisation keeps the documents, and the answer must be that we impose no expiry on them.

Requirements

NFR-001

STATEMENT	An expert answering a question does not lose their train of thought waiting for the next one.
CATEGORY	performance
METRIC	95th percentile elapsed time from the expert submitting an answer to the next question or the drafted section appearing.
TARGET	under 20 seconds at p95, and under 60 seconds at p99
MEASUREMENT	At least 300 timed observations per release, taken across five scripted runs of all seventeen section documents on a set already holding 60 requirements, from the expert's own side of the conversation rather than from inside the product.
PRIORITY	must
APPLIES TO	• FR-004 • FR-010

NFR-002 · A grading pass does not make the expert wait so long that they stop asking for one.

CATEGORY	performance
METRIC	95th percentile elapsed time for one section's grade to be reported, measured per grading pass on the section just drafted.
TARGET	under 90 seconds at p95 for a single section
MEASUREMENT	At least 200 grading passes per release, spread across all seventeen section documents and across sets of 20, 60 and 120 requirements.
PRIORITY	must
APPLIES TO	• FR-009 • FR-011

NFR-003 · A set stays workable at the size real products reach.

CATEGORY	scale
METRIC	The size of one set at which every requirement in section 7 still passes its acceptance criteria and every target in this section still holds.
TARGET	120 functional requirements, 10 journeys of 10 stages, and 200 glossary terms
MEASUREMENT	A generated set at those sizes, run through every function in section 7, per release.
PRIORITY	must

NFR-004 · Work survives the expert stopping, and survives the tool they were using going away.

CATEGORY	durability
METRIC	Share of a set's content recoverable after the authoring session ends, with no part of the product running.
TARGET	100%, readable without this product
MEASUREMENT	End a session mid-section, then read the whole set with nothing but a text editor, once per release.
PRIORITY	must
APPLIES TO	• FR-015

NFR-005

STATEMENT	An expert with no product management training can complete a section without asking anyone what a word means.
CATEGORY	usability
METRIC	Share of sections a first-time expert completes without asking for the meaning of a product management term.
TARGET	90% of sections, measured over 10 first-time experts
MEASUREMENT	Observed sessions with 10 first-time experts before release, questions to the observer counted and categorised.
PRIORITY	must

NFR-006

STATEMENT	Everything the expert and the approver read is usable with a screen reader and without a mouse.
CATEGORY	accessibility
METRIC	Conformance to WCAG 2.2 AA of everything the product puts in front of the expert or the approver along the golden-thread journey.
TARGET	zero level-A or level-AA violations on the golden-thread path
MEASUREMENT	One manual screen-reader pass of the golden-thread journey per release, plus an automated check in the release pipeline of everything that journey shows.
PRIORITY	must

NFR-007 · An expert can tell why a section received the score it did, after the fact.

CATEGORY	observability
METRIC	Share of grades for which the score per rubric dimension and the quotation behind each can be recovered after the session.
TARGET	100% for any section at review or better
MEASUREMENT	Re-read a graded set a week later and reconstruct every score, once per release.
PRIORITY	should
APPLIES TO	• FR-009

NFR-008

STATEMENT	What the expert says about their domain stays within whatever boundary their organisation already works in; this product adds no new place for it to go.
CATEGORY	privacy
METRIC	Number of destinations set content reaches beyond the boundary A-005 names.
TARGET	zero
MEASUREMENT	Review of every outbound destination per release, recorded against A-005.
PRIORITY	must

NFR-009 · A set written by an earlier version is still readable and still gradable by a later one.

CATEGORY	supportability
METRIC	Share of sets written by version N that version N+1 reads and grades without the expert rewriting anything.
TARGET	100% within a major version
MEASUREMENT	Keep one set per released version and run the current version against all of them, per release.
PRIORITY	should

NFR-010

STATEMENT	A set remains readable and retrievable for as long as the organisation keeps the specifications it builds from, without this product being involved.
CATEGORY	retention
METRIC	Share of a set retrievable and readable from the organisation's own document store, with no part of this product present.
TARGET	100%, for as long as the organisation retains the set; no expiry imposed by us
MEASUREMENT	Retrieve a set written one release ago from the organisation's own store and read it with a text editor, once per release.
PRIORITY	must
APPLIES TO	• FR-014

NFR-011

STATEMENT	An interview in progress survives the product being unavailable; no session loses an answer the expert has already given.
CATEGORY	availability
METRIC	Answers lost per interruption, and time to resume after one.
TARGET	Zero answers lost, on every interruption; resume within 1 minute at p95 and within 3 minutes at p99
MEASUREMENT	Interrupt a session mid-section 100 times per release, spread across all seventeen section documents, and count answers lost and seconds to resume.
PRIORITY	must
APPLIES TO	• FR-015

NFR-012

STATEMENT	Nobody but the people the organisation has already given access to its documents can read a set.
CATEGORY	security
METRIC	Number of ways to read a set that the organisation's own document access rules do not govern.
TARGET	zero
MEASUREMENT	Enumerate every read path to a set per release and check each against the organisation's own access rules; recorded against the system-of-record non-goal in section 6 and A-005.
PRIORITY	must

SECTION 9

Constraints

DRAFT 100/100 NO BLOCKING DEFECTS

This section holds no constraints. No rule imposed from outside the team binds prd-forge itself. Each rule that this section once held was a decision of our own, a belief about customers, or the product's chosen response to a risk. The team could change each one by agreeing among themselves, so each one now lives where that kind of item belongs. The ids C-001 to C-004 are not reused.

An expert who uses prd-forge can still be bound by a professional code, a contract, or a firm policy. Those constraints belong in that expert's own PRD, not in this one.

Where the rules live now

At version 0.8.1, C-001 to C-003 left this section.

- C-001, no separate application for the expert to open, install, or log into, is a non-goal in section 6. The session it needs is dependency D-001 in section 11.
- C-002, the output is plain documents and the product is not the system of record, is a non-goal in section 6.
- C-003, domain content stays in the expert's session and in the documents, is assumption A-005 in section 11.

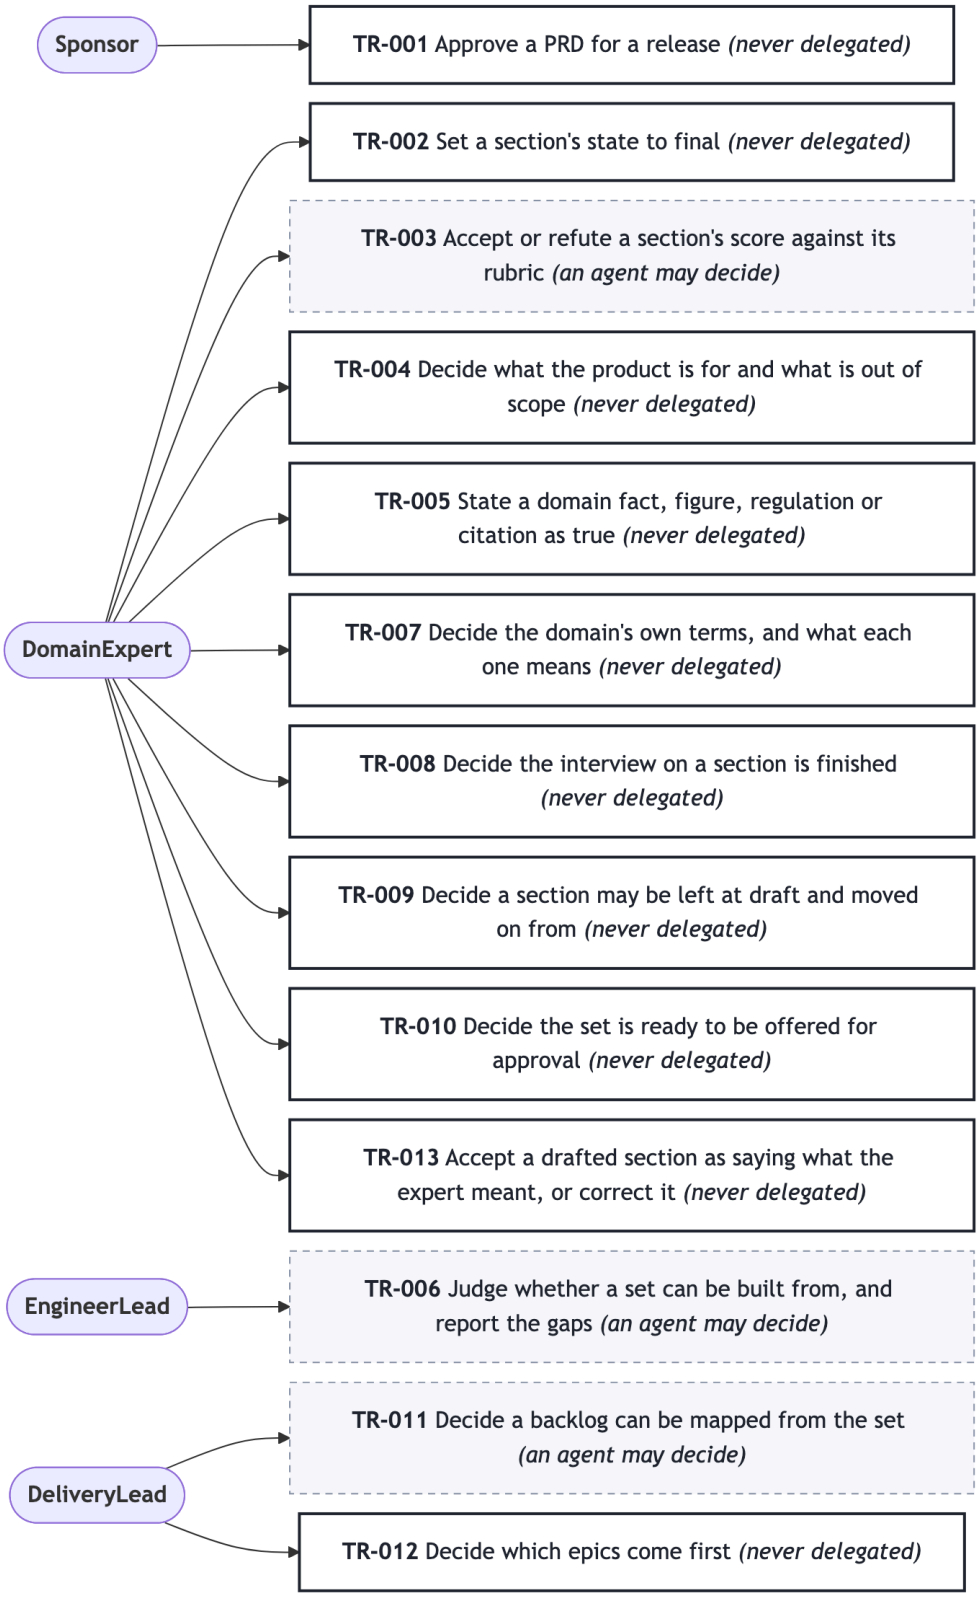
At version 0.9.0, C-004 left this section.

- C-004, no regulation, standard, clause, or statistic is stated as fact unless the expert supplied it or it was checked against its source, is decision rule TR-005 in section 10 and requirement FR-006 in section 7. TR-005 makes the expert the only one who can state such a fact as true, and it is never delegated. FR-006 records a fact the expert does not know as an open question with a decider and a date.

SECTION 10

Decision rights and trust rules

DRAFT 88/100 NO BLOCKING DEFECTS



Who decides what

ID	DECISION	APPROVER	AGENT MAY DECIDE	NEVER DELEGATED
TR-001	Approve a PRD for a release	PER-003	no	yes
TR-002	Set a section to final	PER-001	no	yes
TR-003	Accept or refute a section's score	PER-001	**yes**	no
TR-004	Decide scope	PER-001	no	yes
TR-005	State a domain fact as true	PER-001	no	yes
TR-006	Judge buildability, and report the gaps	PER-002	**yes**	no
TR-007	Decide the domain's own terms	PER-001	no	yes
TR-008	Decide a section's interview is finished	PER-001	no	yes
TR-009	Decide a section may be left at draft	PER-001	no	yes
TR-010	Decide the set is ready to offer for approval	PER-001	no	yes
TR-011	Decide a backlog can be mapped	PER-004	**yes**	no
TR-012	Decide which epics come first	PER-004	no	yes
TR-013	Accept or correct a drafted section	PER-001	no	yes

Every decision named in a stage of section 5 appears here. TR-007 to TR-013 were added at version 0.2.0 to close that: five release-line stage decisions and both of J-002's had no rule, which meant gates an engineering team would have had to invent. TR-013 covers the one in the middle of the loop — the expert accepting that a draft says what they meant — which is distinct from scoring it (TR-003), from leaving it at draft (TR-009) and from setting it final (TR-002).

The shape of the trust boundary

Only three of thirteen are delegable, and all three are **assessments**, not decisions: TR-003 produces a document's score, TR-006 judges whether it can be built from, and TR-011 judges whether a backlog can be mapped from it. All three produce evidence for a person. None changes a document, and none can set a state.

Everything that commits anyone to anything — approving, finalising, choosing scope, asserting a fact — is never delegated. TR-005 is the strictest: a fact nobody can stand behind is not written at all, by anyone.

The delegable assessments all carry the same two obligations: **quote the document for every claim**, and **drop any finding the document does not support**. An assessment that cannot be refuted by pointing at the text is not evidence, it is an assertion, and the expert will correctly stop believing it — see R-002 in section 11.

> **TR-001 · Approve a PRD for a release**
>
> **STAGES** • ST-006
>
> **APPROVER** PER-003
>
> **AN AGENT MAY DECIDE** no
>
> **NEVER DELEGATED** yes
>
> **EVIDENCE REQUIRED** The summary, the scope, every section's state and grade, every blocking defect, and every open question with its decider and date.

ESCALATION	No substitute approver, and approval waits. Where no sponsor exists at all — a solo expert, or an internal tool with no PER-003 — the expert approves their own set under this rule, shown the same evidence, and the change log records that no separate approver existed. A set is never treated as approved because its score is high.
ACCEPTANCE	• Given a set with any blocking defect, when approval is sought, then the approver is shown the defects and the set is not presented as ready. • Given a set with no recorded approval, when anyone reads its state, then it reads as not approved, whatever its scores are. • Given an engagement with no sponsor, when the expert approves their own set, then the approval records that no separate approver existed. • Given an approver who approves with conditions, when the approval is recorded, then each condition names the open question or requirement it depends on and a date, and the set reads as approved-with-conditions until every condition is cleared by that same approver.

TR-002 · Set a section's state to final

STAGES	• ST-004
APPROVER	PER-001
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The section's grade, its blocking defect count, and the result of validating it.
ESCALATION	The expert may override the gate; the override and the score at the time are recorded in the change log. Nobody may set a section final on the expert's behalf while they are away: it waits.
ACCEPTANCE	• Given a section with one or more blocking defects, when final is requested, then it is not set without an explicit override from the expert. • Given an override, when it is applied, then the change log records that it was an override and the score at the time.

TR-003 · Accept or refute a section's score against its rubric

STAGES	• ST-003
APPROVER	PER-001
AN AGENT MAY DECIDE	yes
NEVER DELEGATED	no
EVIDENCE REQUIRED	The rubric's dimensions and a quotation from the document for every score below full marks.
ESCALATION	The expert may refute any finding by pointing at the document; a finding the document does not support is dropped (FR-016). While the expert is away a grade stands unaccepted, and the section does not advance.
ACCEPTANCE	• Given a score below full marks on a dimension, when it is reported, then a quotation from the document is given for it. • Given a finding the expert refutes and the document does not support, when it is checked, then it is dropped and the expert is told it was dropped.

TR-004 · Decide what the product is for and what is out of scope

STAGES	• ST-001
APPROVER	PER-001
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	What the consequence of the choice is, including any release-line stage it would leave uncovered.
ESCALATION	None, and no substitute. The expert decides; while they are away, scope does not change. Where the choice weakens the PRD the objection is stated once and the decision is recorded with the dissent noted.
ACCEPTANCE	• Given an expert decision that leaves a release-line stage uncovered, when it is recorded, then the consequence is stated to the expert once and the decision stands. • Given a decision the expert reaffirmed after an objection, when later grading passes run, then the objection is not raised again as a new finding.

TR-005 · State a domain fact, figure, regulation or citation as true

APPROVER	PER-001
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The expert's own statement of it, or a source that has been checked.
ESCALATION	Where neither exists the fact is not written at all; it becomes an open question with a decider and a date. Nobody may supply the fact on the expert's behalf while they are away.
ACCEPTANCE	• Given a figure or citation the expert did not supply and that has not been verified, when a section is drafted, then it does not appear in the document. • Given a citation that was checked and could not be confirmed, when it is reported, then it is marked unverified rather than stated as fact.

TR-006 · Judge whether a set can be built from, and report the gaps

STAGES	• ST-009
APPROVER	PER-002
AN AGENT MAY DECIDE	yes
NEVER DELEGATED	no
EVIDENCE REQUIRED	The release line, the requirements covering each stage, and a quotation for every gap claimed.
ESCALATION	The engineering lead's own judgement overrides any reported verdict; the report is evidence for them, not a substitute for them. In release 1 no engineering lead is in the loop, so the report goes to the expert, who decides what to do with it under TR-010.
ACCEPTANCE	• Given a reported verdict, when the engineering lead disagrees, then their judgement is recorded as the verdict and the report is kept alongside it. • Given a reported gap, when it is stated, then the document is quoted for it.

TR-007 · Decide the domain's own terms, and what each one means

STAGES	• ST-001
APPROVER	PER-001
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	Where the term is already used in the set, and any standard-setter's definition that differs from the expert's.
ESCALATION	None, and no substitute. Terms are the expert's; while they are away the vocabulary does not change.
ACCEPTANCE	• Given a term the expert has defined, when any later question or document uses that idea, then it uses the expert's word and not a synonym. • Given a standard-setter's definition that differs, when the term is recorded, then both are stated and the expert says which the set uses.

TR-008 · Decide the interview on a section is finished

STAGES	• ST-002
APPROVER	PER-001
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	What the section would score as it stands, and what is still missing from it.
ESCALATION	None. The expert ends the interview; nobody ends it for them, and no question list runs past the point they call it done.
ACCEPTANCE	• Given an expert who says they have nothing more to add, when they say it, then no further question is put to them in that section. • Given an interview ended early, when the section is written, then what is missing is named to the expert, not hidden behind a score.

TR-009 · Decide a section may be left at draft and moved on from

STAGES	• ST-004
APPROVER	PER-001
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The section's score, its blocking defects, and which later sections depend on what is missing.
ESCALATION	None, and no substitute. While the expert is away the section stays where it is.
ACCEPTANCE	• Given a section below 70 or carrying a blocking defect, when the expert chooses to move on, then it is written at `draft` and what was left behind is recorded as an open question or named plainly to them.

TR-010 · Decide the set is ready to be offered for approval

STAGES	• ST-005
APPROVER	PER-001
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	Every section's state and grade, every blocking defect, and the buildability report's gaps with, for each, whether it is closed or owned.
ESCALATION	None, and no substitute. The expert offers the set; nobody offers it for them.
ACCEPTANCE	• Given a buildability gap that is neither closed nor given an owner and a date, when the expert offers the set for approval, then that gap is stated to them first. • Given a set offered for approval, when the approver opens it, then what the expert was shown under this rule is what the approver sees.

TR-011 · Decide a backlog can be mapped from the set

STAGES	• ST-007
APPROVER	PER-004
AN AGENT MAY DECIDE	yes
NEVER DELEGATED	no
EVIDENCE REQUIRED	Whether a golden thread exists, and which release-line stages have requirements behind them.
ESCALATION	The delivery lead's own judgement stands over any reported verdict; where they are unavailable the workshop waits rather than proceeding without a backbone.
ACCEPTANCE	• Given a set with no journey, when the check runs, then it reports that a backlog cannot be derived, and names that as the reason.

TR-012 · Decide which epics come first

STAGES	• ST-008
APPROVER	PER-004
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The release line, the dependencies in section 11, and the objective each epic moves.
ESCALATION	The delivery lead may re-order within the release line at any time; order never changes scope, which stays with TR-004. While they are away the order stands as last recorded, and nobody re-orders on their behalf.
ACCEPTANCE	• Given an ordering, when it is recorded, then each decision names whether value, dependency or the release line drove it.

TR-013 · Accept a drafted section as saying what the expert meant, or correct it

STAGES	• ST-003
APPROVER	PER-001
AN AGENT MAY DECIDE	no
NEVER DELEGATED	yes
EVIDENCE REQUIRED	The drafted section in readable form, and its grade, so the expert judges the words and not the score.
ESCALATION	None, and no substitute. An unaccepted draft stays at `draft`; nobody accepts a section's wording on the expert's behalf while they are away.
ACCEPTANCE	• Given a drafted section the expert has not accepted, when the set's state is read, then that section is not shown as ready for anyone else to read. • Given an expert who corrects a drafted section, when the correction is applied, then it is their words that are kept, and the section is graded again before it advances.

SECTION 11

Assumptions, risks and dependencies

DRAFT 75/100 NO BLOCKING DEFECTS

Assumptions

ID	ASSUMPTION	IF FALSE
A-001	An expert will answer 50-100 questions across a few sittings	The method fails; sections 5-7 change completely
A-002	An expert will act on a grade they did not produce	The grading is decoration
A-003	An engineering lead accepts a PRD naming no technology	Test 5 becomes a liability
A-004	Organisations permit drafting inside an assistant session	Section 6's separate-application non-goal and A-005 have to change
A-005	Organizations accept that domain content stays in the session and the documents	NFR-008, NFR-010, and NFR-012 have to change

A-001 and A-002 are the two the product stands on, and neither has been tested. Both are validated by the same ten observed sessions (D-002), which makes D-002 the single most load-bearing dependency in this PRD.

A-005 was constraint C-003 in section 9 until 0.8.1. It moved here because it is a belief about what customers accept, which a customer can falsify, and not a rule imposed from outside.

Risks

ID	RISK	L	I	OWNER
R-001	High score becomes false assurance	high	high	Richard Hightower
R-002	Unsupported findings destroy trust in all findings	high	high	Richard Hightower
R-003	A beautifully graded PRD for the wrong product	medium	high	Richard Hightower
R-004	Interview too long; sets abandoned	medium	medium	Richard Hightower
R-005	A version change strands an expert's set	low	high	Richard Hightower

R-003 is not closed. The mitigations listed reduce it and do not solve it. Nothing in this product checks whether the product being specified should exist. That is stated here rather than hidden behind a mitigation that sounds adequate.

Dependencies

ID	NEED	FROM	BY	STATUS
D-001	An assistant session able to carry the interview	The customer's platform	2027-01-31	committed
D-002	10 observed first-time experts	Early customers	2027-03-31	requested

ID	NEED	FROM	BY	STATUS
D-003	3 engineering leads to read a set cold	Early customers	2027-03-31	requested

D-002 and D-003 are both merely requested, and both are dated on R1's target day in section 13. That is a real schedule problem: R1 ships on the day the assumptions it rests on are due to be tested, so it cannot use the results. Recorded, not resolved.

Assumptions

A-001

ASSUMPTION	A domain expert will answer fifty to a hundred questions about their own field across a few sittings, if each one is clearly worth answering.
IF FALSE	The whole method fails. The product would have to produce a usable partial PRD from perhaps fifteen answers, which changes sections 5, 6 and 7 completely and makes the grading thresholds meaningless.
VALIDATED BY	Observed sessions with 10 first-time experts, counting where they stop - the same sessions NFR-005 uses.
OWNER	Richard Hightower

A-002 · An expert will act on a grade they did not produce, rather than dismissing it.

IF FALSE	The grading is decoration and sections 3 and 7 are measuring something nobody uses. The product reduces to a structured interview.
VALIDATED BY	In the same observed sessions, whether the expert answers the one question or moves on regardless.
OWNER	Richard Hightower

A-003

ASSUMPTION	An engineering lead will accept a PRD that names no technology, rather than reading the absence as incompleteness.
IF FALSE	Test 5 becomes a liability rather than a discipline, and the candidate architecture decisions list in FR-007 has to become a first-class output rather than a hand-over note.
VALIDATED BY	Three engineering leads reading a completed set cold and saying what they would ask for first.
OWNER	Richard Hightower

A-004

ASSUMPTION	Organisations will let a specification for their product be drafted inside an assistant session.
IF FALSE	The separate-application non-goal in section 6 and A-005 are unacceptable as written and the delivery model has to change, which invalidates most of section 7's assumptions about immediacy.
VALIDATED BY	The information security review of the first three customer engagements.
OWNER	Richard Hightower

A-005

ASSUMPTION	Organizations accept that the domain content an expert enters stays in the expert's own session and in the documents. This product sends it nowhere else and keeps it nowhere else.
IF FALSE	An organization whose information policy needs more than this cannot adopt the product as written, and NFR-008, NFR-010, and NFR-012, which are measured against this boundary, have to change.
OWNER	Richard Hightower
CATEGORY	viability
EVIDENCE	none
VALIDATE BY	2027-03-15

Risks

R-001

RISK	The expert accepts a weak PRD because the score was high. The grade becomes false assurance, which is worse than no grade.
LIKELIHOOD	high
IMPACT	high
MITIGATION	Blocking defects are counted and reported separately from the score and no score clears them (FR-009); the non-goal in section 6 says explicitly that a grade is evidence, not a warrant; FR-012 puts a buildability verdict in front of the expert before approval.
OWNER	Richard Hightower

R-002

RISK	The grading reports findings that are not in the document. The expert catches one, stops trusting all of them, and the grading becomes noise they skip.
LIKELIHOOD	high
IMPACT	high
MITIGATION	Every finding must quote the document, and anything unsupported is dropped before the expert sees it (FR-016, TR-003, TR-006); a refuted finding is acknowledged as dropped rather than defended (SC-004).
OWNER	Richard Hightower

R-003

RISK	A fluent, well-structured, well-graded PRD is produced for the wrong product, and the structure makes the error harder to see rather than easier.
LIKELIHOOD	medium
IMPACT	high
MITIGATION	Section 3 asks what would move the objectives without this product at all; section 1 requires evidence per problem and names what is unmeasured. Neither is sufficient - this risk is not closed, and a reviewer should say so.
OWNER	Richard Hightower

R-004

RISK
The interview is too long to finish, so sets are abandoned half-written and the product is remembered as the thing that asked a hundred questions.

LIKELIHOOD
medium

IMPACT
medium

MITIGATION
FR-005 lets the expert end any section and records what was left; FR-015 makes resumption cheap; the working order in the method puts the sections that unlock the others first, so an abandoned set is still useful.

OWNER
Richard Hightower

R-005

RISK
A set written by one version cannot be read or graded by a later one, and an expert loses weeks of thinking to an upgrade.

LIKELIHOOD
low

IMPACT
high

MITIGATION
NFR-004 makes every set readable with no part of the product running; NFR-009 requires older sets to stay gradable within a major version.

OWNER
Richard Hightower

Dependencies

D-001

NEED
An assistant session the expert is already working in, able to carry a long interview and read and write the documents.

FROM
The assistant platform the customer already uses

BY
2027-01-31

STATUS
committed

BLOCKS
- ST-001
- ST-002

D-002 · Ten first-time domain experts willing to be observed completing a PRD.

FROM
Early customers and the professional network

BY
2027-03-31

STATUS
requested

BLOCKS
- A-001
- A-002
- NFR-005

D-003

NEED	Three engineering leads willing to read a completed set cold and report what they would ask for first.
FROM	Early customers
BY	2027-03-31
STATUS	requested
BLOCKS	• A-003 • O-001

SECTION 12

Open questions

DRAFT 79/100 NO BLOCKING DEFECTS

ID	QUESTION	DECIDER	BY	BLOCKS
Q-001	Today's first-pass buildability rate	R. Hightower	2027-03-31	O-001
Q-002	Today's idea-to-approval elapsed time	R. Hightower	2027-03-31	O-002
Q-003	Today's share of testable requirements	R. Hightower	2027-03-31	O-003
Q-004	Segment sizes, and which converts	R. Hightower	2027-06-30	S-001, S-002
Q-005	Will a sponsor accept a graded specification	R. Hightower	2027-01-31	FR-013, TR-001
Q-006	Sections completed per sitting, and where they stop	R. Hightower	2027-03-31	R-004, NFR-005
Q-007	Is the working order right, and how does looping back feel	R. Hightower	2027-03-31	ST-004
Q-008	Disciplined or incomplete, to an engineering lead	R. Hightower	2027-03-31	A-003
Q-009	Where the rubric and the checks live, and who owns them	R. Hightower	2027-01-31	FR-004, FR-009..FR-012, NFR-007
Q-010	Where an answer lives before its section is drafted	R. Hightower	2027-01-31	FR-015, NFR-011
Q-011	How an approver records a decision, and who owns NFR-006/012	R. Hightower	2027-02-28	FR-014, NFR-006, NFR-012

The three that matter most

Q-001 to Q-003 are the baselines for all three objectives. Until they are answered, section 3's targets are intentions. Nothing should be committed to a plan on the strength of them.

Q-009, Q-010 and Q-011 came from the buildability review at 0.2.0, and all three block `must` requirements on the release line. Q-009 is the largest: four requirements are written against a rubric whose shape the set defines but whose content it does not contain, so two of the six release-line stages can be scaffolded and not finished until it is answered. Q-010 is the one where guessing breaks assumption A-005 rather than costing rework.

Q-005 is the one that could invalidate the product's central idea. It was dated after the release it blocks until 0.2.0 moved it to 2027-01-31; it is a two-day question and there was no defence for asking it late. If showing a sponsor a grade reduces their confidence rather than raising it, then FR-013 is wrong and the grading becomes an author's private tool rather than part of the approval. That is a cheap two-day question and it should be asked early.

Q-008 tests test 5. If engineering leads read the absence of technology as incompleteness, the discipline this whole product enforces is a liability, and A-003 fails.

Honest note

Every decider here is the same person. For a one-person product that is accurate; for anything larger it is a defect, and the first thing to fix when there is a second person.

Q-001

QUESTION	What share of specifications reaching a development team today pass a buildability review on first submission?
DECIDER	Richard Hightower
BY	2027-03-31
BLOCKS	• O-001
SPIKE	Run the eleven derivability checks against ten specifications collected from real engagements, and count. Time-boxed to 3 days.

Q-002

QUESTION	How long does it take today, in working days, from a product idea to an approved specification?
DECIDER	Richard Hightower
BY	2027-03-31
BLOCKS	• O-002
SPIKE	Reconstruct the timeline from first discussion to approval for five recent engagements. Time-boxed to 2 days.

Q-003

QUESTION	What share of functional requirements in today's specifications can an engineer who was not in the interview write a test from?
DECIDER	Richard Hightower
BY	2027-03-31
BLOCKS	• O-003
SPIKE	Three engineering leads score the requirements of the same ten specifications from Q-001, independently. Time-boxed to 3 days.

Q-004

QUESTION	How many organisations are in each segment in section 2, and which segment actually converts?
DECIDER	Richard Hightower
BY	2027-06-30
BLOCKS	• S-001 • S-002
SPIKE	Name the register or source per segment and count; then compare against the first ten conversations. Time-boxed to 5 days.

Q-005

QUESTION Will a sponsor accept a specification that carries a machine-produced grade, or does the grade reduce their confidence rather than raise it?

DECIDER Richard Hightower

CONSULTED the first three sponsors

BY 2027-01-31

BLOCKS

- FR-013
- TR-001

SPIKE Show three sponsors the same set with and without grades and ask which they would approve. Time-boxed to 2 days.

Q-009

QUESTION Where do the rubric and the eleven derivability checks live, who owns them, and how is a set told which version it was graded against?

DECIDER Richard Hightower

BY 2027-01-31

BLOCKS

- FR-004
- FR-009
- FR-010
- FR-011
- FR-012
- NFR-007

SPIKE Decide whether they ship as documents of the product, name their owner, and decide whether a grade records the version it was made against. Time-boxed to 2 days. Until it is answered, ST-003 and ST-004 can be scaffolded and not finished.

Q-010

QUESTION Where does an answer live between the moment the expert gives it and the moment the section is drafted, given A-005, that content is retained nowhere outside the documents, and given that NFR-004 requires the set readable with no part of the product running?

DECIDER Richard Hightower

BY 2027-01-31

BLOCKS

- FR-015
- NFR-011

SPIKE Decide whether an in-progress interview is itself a document of the set, or whether answers before a draft are not retained at all and FR-015 holds only within a session. Time-boxed to 2 days. Guessing this one breaks assumption A-005, which customers rely on, rather than merely costing rework.

Q-011

QUESTION How does an approver record a decision, given that D-001 puts the product inside the expert's session and OUT-003 excludes any view of our own, and who owns NFR-006's accessibility and NFR-012's access control, which describe things the non-goals in section 6 say we do not own?

DECIDER Richard Hightower

BY 2027-02-28

BLOCKS

- FR-014
- NFR-006
- NFR-012

SPIKE Walk one approval end to end with a real sponsor and record every step they take that the product does not provide. Time-boxed to 3 days.

Q-006

QUESTION How many sections will an expert complete in one sitting before stopping, and which section do they stop at?

DECIDER Richard Hightower

BY 2027-03-31

BLOCKS

- R-004
- NFR-005

SPIKE Count in the ten observed sessions from D-002. No extra work.

Q-007

QUESTION Is the working order in the method right - does writing the requirements really send an expert back to the journeys, and does that feel like progress or like failure to them?

DECIDER Richard Hightower

BY 2027-03-31

BLOCKS

- ST-004

SPIKE Count loop-backs and ask the expert how it felt, in the ten observed sessions. No extra work.

Q-008

QUESTION Does an engineering lead read a PRD that names no technology as disciplined or as incomplete?

DECIDER Richard Hightower

CONSULTED the three engineering leads from D-003

BY 2027-03-31

BLOCKS

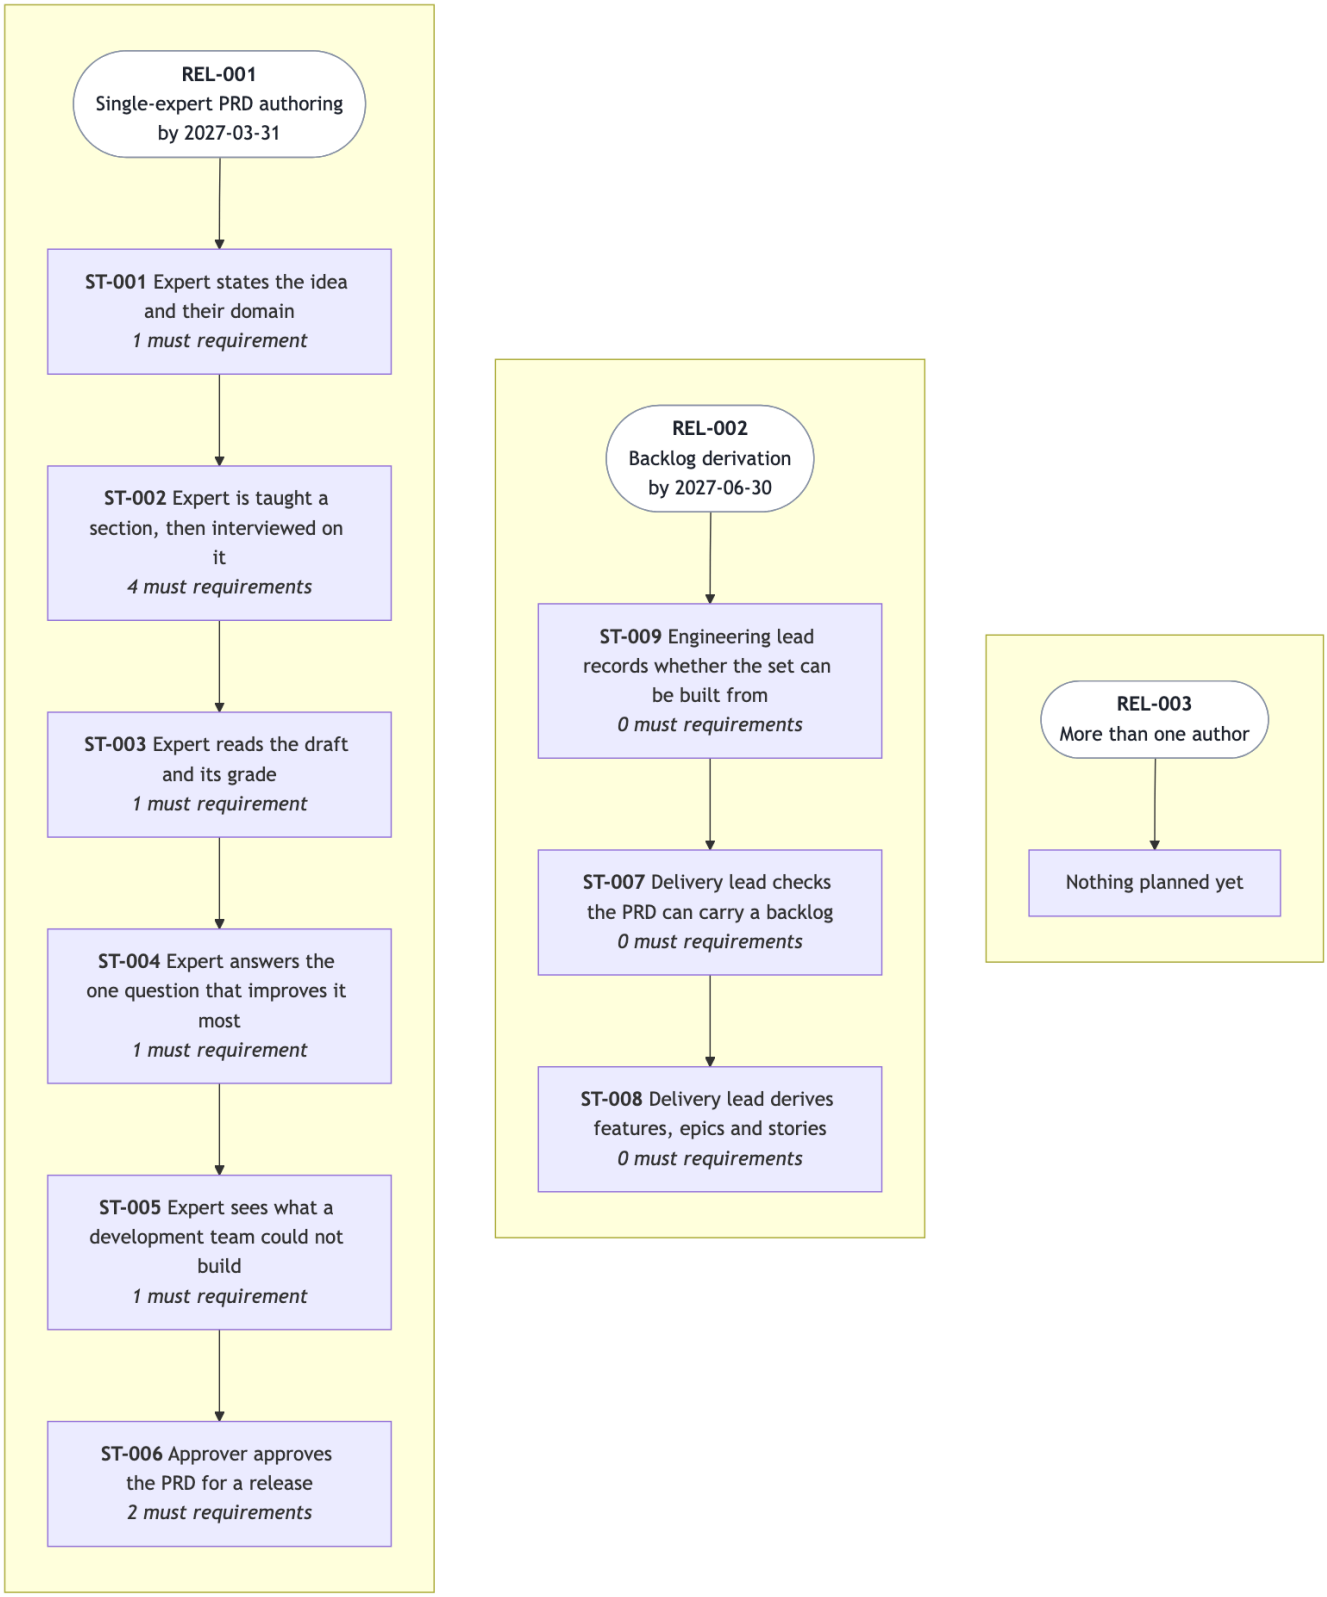
- A-003

SPIKE Ask each lead, after reading a set cold, whether anything was missing and whether they wanted a technology named. Time-boxed to 1 day.

SECTION 13

Release plan

DRAFT 88/100 NO BLOCKING DEFECTS



The release line

ID	RELEASE	STAGES	TARGET
REL-001	Single-expert PRD authoring	ST-001..ST-006	2027-03-31
REL-002	Backlog derivation	ST-009, ST-007, ST-008	2027-06-30
REL-003	More than one author	—	none

Why this order

REL-001 is exactly the release line in section 6. It is the smallest release in which anyone gets anything: an expert reaching an approved, reviewed PRD.

REL-002 is deferred rather than cut because the PRD REL-001 produces already has to carry everything a backlog needs — FR-012 checks precisely that. Deferring the derivation costs nothing; deferring what it needs would have cost a rewrite.

REL-003 waits for evidence that one author's path works at all. Building for two authors before one is proven would be building on A-001 and A-002 before either is tested.

How you know each release worked

REL-001 exits when the pilot PRDs meet the three objective targets in section 3: 80% pass the buildability review on first submission, the median reaches approval in under 5 working days, and 95% of `must` requirements are testable as written. REL-002 exits when every derived story names its requirement and a delivery lead runs a workshop from the backlog without re-deriving it.

Both target dates are indicative, not committed. The schedule problem below is the reason.

The schedule problem, stated plainly

REL-001 targets 2027-03-31. D-002 and D-003 in section 11 — the ten observed experts and the three engineering leads — are also dated 2027-03-31 and are only *requested*. So REL-001 ships at the same moment the assumptions beneath it are first tested, not after. Either REL-001 moves, or it ships knowingly as an experiment. That is a decision for the sponsor at TR-001 and it is not made here.

Not in this section, deliberately

No tasks, no estimates, no sizing, no team assignments, and no technology milestone. Those come from engineering design and iteration planning.

> **REL-001 · Single-expert PRD authoring**
>
> OUTCOMES
>
> - A domain expert with no product training can take one product idea to an approved PRD without leaving the session they started in.
> - The expert can see, before anyone estimates, which parts of their own specification a development team could not build from.
> - A sponsor can see what they are approving, what is still open and who owns each open thing, on one page.

JOURNEY STAGES

- ST-001
- ST-002
- ST-003
- ST-004
- ST-005
- ST-006

REQUIREMENTS

- FR-001
- FR-002
- FR-003
- FR-004
- FR-005
- FR-006
- FR-007
- FR-008
- FR-009
- FR-010
- FR-011
- FR-012
- FR-013
- FR-014
- FR-015
- FR-016
- FR-017
- FR-018

OBJECTIVES

- O-001
- O-002
- O-003

EXIT CRITERIA

- At least 80% of the pilot PRDs pass the buildability review on first submission (O-001).
- The median pilot PRD reaches a recorded approval in under 5 working days from the first interview question (O-002).
- An engineer who was not in the interview judges 95% of the pilot PRDs' `must` requirements testable as written (O-003).

TARGET DATE 2027-03-31

DATE STATUS indicative

REL-002 · Backlog derivation

OUTCOMES

- A delivery lead can walk into a story-mapping workshop with features and a release line already derived from the approved PRD.
- Every story names the requirement id it came from, so nobody has to ask why it exists.

JOURNEY STAGES

- ST-009
- ST-007
- ST-008

OBJECTIVES

- O-001

EXIT CRITERIA

- Every story derived from a pilot PRD names the requirement id it came from.
- A delivery lead runs a story-mapping workshop from a derived backlog without re-deriving any feature from the PRD.

TARGET DATE 2027-06-30

DATE STATUS indicative

REL-003 · More than one author

OUTCOMES

- Two experts can work on different sections of the same specification without overwriting each other.
- A reviewer can comment on a section without being able to change it.

OBJECTIVES

- O-002

SECTION 14

Glossary

DRAFT 65/100 4 BLOCKING

TERM	MOST OFTEN CONFUSED WITH
PRD	A design document
Golden thread	The most important journey
Release line	A date
Stage	A screen
Rubric	A checklist
Blocking defect	A low score
Score	A verdict on the product idea
Spike	A piece of the product
Open question	A risk
Non-goal	Out of scope
Grade	Score
State	Threshold
Assumption	Open question, risk
Candidate architecture decision	A requirement
Buildability review	A grading pass
Finding	A blocking defect
Root document	Section 1
Approver	The expert
Condition	A constraint
Constraint	A requirement
Objective	A requirement
Baseline	A target
Version	A release of the product
Dependency	A requirement's wait on an open question

Seventeen terms were added at version 0.2.0. Five of them existed only as `see_also` pointers to nothing, which is this set's own named blocking defect; the rest were load-bearing inside acceptance criteria while being defined nowhere:

threshold, grade, grading pass, state, verdict, override, unit of work and the rest. The five tests are now numbered, because the set cites "test 5" by number five times.

The `not` column above is the one that earns its keep. Every one of these has been confused in practice, and three of them, golden thread, blocking defect and score, carry a cost when confused: a team picks the wrong journey for release 1, ships a section with a known fault because the number looked good, or believes a well-written specification means a sound product.

PRD

MEANING A written statement of what to build, for whom, why, and how you know it works, plus what not to build. It does not say how to build it.

SEE ALSO
- section
- document set

NOT A design document. A PRD that names components or technologies has stopped being a PRD and has started being an architecture sketch.

Document set

MEANING The documents that together are one PRD: a root document and one document per section, each linking to the others.

SEE ALSO
- PRD

Golden thread

MEANING The one user journey that release 1 must deliver from end to end.

SEE ALSO
- journey
- release line

NOT The most important journey in general. It is the journey chosen for release 1, which is often the simplest rather than the most valuable.

Release line

MEANING The smallest set of journey stages that completes the golden thread, and so the boundary of what release 1 contains.

SEE ALSO
- golden thread
- scope

NOT A date or a deadline. It is a line across the journey, not across a calendar.

Stage

MEANING One large step of a journey, naming a person, an input, an output and a decision.

SEE ALSO
- journey
- step

NOT A screen or a system step. A stage survives a redesign.

The five tests

MEANING

The five numbered properties a statement must have to be a product requirement. Test 1, it names a user or a stakeholder. Test 2, it states an outcome a person can see. Test 3, it has a pass-or-fail acceptance test. Test 4, it has a priority. Test 5, it does not name the component or the technology that delivers it. The numbers are part of the term: the set refers to "test 5" by number.

SEE ALSO

- blocking defect

Rubric

MEANING

The stated standard a section is scored against: a set of dimensions, each with written anchors for what a score of 0, 2 and 4 look like.

SEE ALSO

- score
- blocking defect

NOT

A checklist. A checklist asks whether something is present; a rubric asks how good it is, against written anchors.

Blocking defect

MEANING

A named fault that stops a section advancing however high its score is, such as a requirement that names a technology or an id that refers to nothing.

SEE ALSO

- rubric
- score

NOT

A low score. Score and blocking defects are counted separately and a high score never clears a blocking defect.

Score

MEANING

A section's total against its rubric, 0 to 100, written by a grading pass and never by hand.

SEE ALSO

- rubric
- blocking defect

NOT

A measure of whether the product is a good idea. It measures how well the section is written, nothing else.

Spike

MEANING

A time-boxed piece of work whose only output is an answer to an open question, or a recorded decision.

SEE ALSO

- open question

NOT

A piece of the product. A spike produces knowledge, not a feature.

Open question

MEANING

Something not yet decided, recorded with the one person who decides it and the date by which they must.

SEE ALSO

- spike
- assumption

NOT

A risk. A risk is something that might happen and nobody can simply decide; an open question has a decider.

Non-goal

MEANING Something a reasonable reader would expect the product to do, that it will not do, with no plan to.

SEE ALSO
- scope

NOT Out of scope. Out of scope means not yet, with a condition for revisiting; a non-goal means not ever, with a reason.

Journey

MEANING The steps one user takes to reach one goal with the product, from start to end, in order.

SEE ALSO
- stage
- golden thread

NOT A use case. A journey covers the whole path to a goal, including steps outside the product.

Step

MEANING What a user does, sees or decides at one point inside a stage.

SEE ALSO
- stage
- journey

NOT A stage. A stage holds many steps.

Section

MEANING One of the fifteen numbered parts of a PRD, each held in its own document. Section 5 holds two documents, journeys and scenarios, and section 7 holds two, functional requirements and wireframes, so a set has seventeen section documents across fifteen sections. Cross-references in a set are by section number.

SEE ALSO
- document set

Scope

MEANING What a release contains and what it does not, set against the journey stages it completes.

SEE ALSO
- release line
- non-goal

Assumption

MEANING Something taken on faith that the product depends on, recorded with what would break if it turned out false.

SEE ALSO
- open question
- risk

NOT An open question. Nobody can simply decide an assumption true; it is validated or falsified.

Risk

MEANING Something that may happen and would damage the product, recorded with a likelihood, an impact, a mitigation and one named owner.

SEE ALSO
- assumption
- spike

Readable form

MEANING

The part of a section a person reads: its prose, its headings and its tables, written to be judged by the expert who gave the answers.

SEE ALSO

- section

Threshold

MEANING

The score at or above which a section may advance, provided it also carries no blocking defect. It is 70.

SEE ALSO

- score
- blocking defect

Grade

MEANING

A section's score together with its blocking-defect count and its weakest dimension, the rubric dimension that scored lowest, which the set also calls the section's weakest part. A grade is what an expert is shown; a score is one number inside it.

SEE ALSO

- score
- grading pass

NOT

A score. The score is part of the grade, not the whole of it.

Grading pass

MEANING

One run of a section or a set against its rubric, producing a grade. It is the only thing that writes a score, and it never changes the document.

SEE ALSO

- grade
- rubric

State

MEANING

Where a section stands, one of four in this order: `pending` (not yet interviewed), `draft` (written), `review` (ready for someone else to read), `final` (approved). A section reaches `review` at a score of 70 or more with no blocking defects, and `final` only with an approval recorded.

SEE ALSO

- threshold
- blocking defect

Override

MEANING

An expert's decision to advance a section past a gate it has not met, recorded in the change log with the score at the time.

SEE ALSO

- state
- threshold

Verdict

MEANING

What a buildability review concludes, one of three: `buildable`, `buildable with gaps`, `not buildable`.

SEE ALSO

- derivability checks

Derivability checks

MEANING

The eleven checks a buildability review runs to decide whether a backlog can be derived from a set. They are product content, shipped and versioned with the product, and the set states their count and their verdict vocabulary rather than restating the list.

SEE ALSO

- verdict
- rubric

Candidate architecture decision

MEANING

A technology or design preference the expert stated during the interview, kept out of the requirements and handed to the engineering lead as a topic for an architecture decision record.

SEE ALSO

- the five tests

NOT

A requirement. It is deliberately not one, because it would fail test 5.

Unit of work

MEANING

What the expert's domain counts one piece of work as: an engagement, a load, a claim, a case, a patient. Every question the expert is asked uses their word for it.

Unverified

MEANING

A citation that was checked and could not be confirmed against its source. It is marked as such and never stated as fact.

SEE ALSO

- the five tests

MoSCoW

MEANING

The priority scale this set uses: `must`, `should`, `could`, `wont`. This set uses MoSCoW and not RFC 2119's MUST and SHOULD.

SEE ALSO

- the five tests

Wireframe

MEANING

A sketch of one screen a person works in or one report a person reads, labeled an illustration, not a design. It shows what the person sees and does there, and names the requirements it makes concrete. Every requirement shown in a screen or a report has one, in section 7.

SEE ALSO

- the five tests

Buildability review

MEANING

A reading of the whole set as the engineering lead who has to build it. It runs the eleven derivability checks and reports a verdict and, for each release-line stage, whether a requirement covers it or it is a gap.

SEE ALSO

- verdict
- derivability checks
- finding

NOT

A grading pass. A grading pass scores one section against its rubric; a buildability review reads the whole set and asks whether a backlog can be derived from it.

Finding

MEANING

One fault that a grading pass or a buildability review reports, with the document it is in quoted. A finding is checked against the document before the expert sees it, and is dropped when the document does not support it.

SEE ALSO

- blocking defect
- grading pass
- buildability review

NOT

A blocking defect. A blocking defect is one kind of finding, the kind that stops a section advancing. Every blocking defect is reported as a finding, but most findings only lower a score or name a weak dimension.

Root document

MEANING

The one document of a set that every other document names as its parent. It carries the summary, the scope, every section's state and grade, and every open question with its decider and date, and it is what an approver reads.

SEE ALSO

- document set
- approver

NOT

Section 1. The root document has no section number; section 1 is the problem and opportunity.

Approver

MEANING

The one person who decides whether a set is approved for a release: the sponsor, or the expert when no sponsor exists. They approve, approve with conditions, or reject, and the decision is recorded in section 15 against a named version. An approver who approves with conditions approves the set on the conditions they state, and the set reads as approved with conditions until that same approver clears every one.

SEE ALSO

- condition
- root document

NOT

The expert. The expert decides the state of each section; the approver decides on the whole set. They are one person only when no sponsor exists.

Condition

MEANING

What an approver who approves with conditions requires before the approval is unconditional. Each condition names the open question or requirement it depends on and a date, and only the approver who stated it clears it.

SEE ALSO

- approver
- open question

NOT

A constraint. A condition belongs to one approval and is cleared; a constraint holds for the product and is never cleared.

Constraint

MEANING

A rule the product must obey that the team cannot change by agreeing among themselves, recorded in section 9 with its category, the authority that imposes it by name, and how compliance is shown. A regulation or a standard appears in a constraint only when the expert named it.

SEE ALSO

- condition

NOT

A requirement. A requirement states what a user can do and has a priority; a constraint limits how every requirement it applies to may be met, and is not prioritized or deferred.

Objective

MEANING

A business result the product is expected to produce, recorded in section 3 with the problems it addresses, one metric, a baseline, a target with a date, and one owner. If the product ships and the metric does not move, the product did not work.

SEE ALSO

- baseline

NOT

A requirement. An objective says what the business gains; a requirement says what a user can do, and names the objective it serves.

Baseline

MEANING

The value an objective's metric has today, before the product. When it is not known, the baseline says unknown and names the open question that will establish it, and the target is an intention rather than a commitment until the question is answered.

SEE ALSO

- objective
- open question

NOT

A target. The baseline is where the metric stands now; the target is where the objective says it will be by a date.

Requirement

MEANING

A statement of what one named user can do with the product and the outcome they can see, recorded in section 7 with its acceptance criteria, a priority, the journey stage it serves, and the objective it serves. A statement is a requirement only when it passes the five tests.

SEE ALSO

- the five tests
- acceptance criteria
- objective

Acceptance criteria

MEANING

The pass-or-fail statements on a requirement or a journey stage. Each is written as Given a situation, When the person acts, Then what they see, and each is one an engineer who was not in the interview could write a test from. A story derived from a requirement takes the requirement's criteria as its acceptance criteria.

SEE ALSO

- requirement
- stage
- the five tests

Gap

MEANING

Something the set does not say that a team needs in order to build: most often a release-line stage that no requirement covers, as a buildability review reports it. A gap is closed with a requirement, or owned as an open question with a decider and a date.

SEE ALSO

- buildability review
- release line
- open question

Version

MEANING

A numbered state of the set, recorded in section 15 with its date, its author, and the sections it changed. An approval names the one version it covers, and a change written after an approval records a new version that the approval does not cover.

SEE ALSO

- approver
- condition

NOT

A release of the product. The product's releases are numbered too, and a set written by one release is read by later ones; a version of the set counts changes to the set.

Domain

MEANING

The field of work the expert specifies software for, such as freight. The expert names it before the first section. It sets the unit of work, the vocabulary every question uses, and the examples each section is taught with. The expert supplies knowledge of the domain; the product does not.

SEE ALSO

- unit of work

Reading order

MEANING

The order in which the documents of a set are read: the root document, then the section documents by section number, from section 1 to section 15. Where a section holds two documents, journeys come before scenarios and functional requirements come before wireframes. The next and the previous section of any document follow this order.

SEE ALSO

- section
- root document

Interview

MEANING

The questions the expert answers about one section, in stage ST-002. It starts after the expert reads what the section is for, puts one or two questions at a time, and ends when the section has enough material to be drafted, when the expert says they have nothing more to add, or when their answers would already score the section 70 or more with no blocking defect. Only the expert answers in it. An engineer who reads the set later was not in the interview.

SEE ALSO

- section
- candidate architecture decision

Handover

MEANING

The point at which the approved set reaches the engineering lead who has to build from it. Journey J-002 starts there: its trigger is that the PRD has been approved for a release, and its first stage, ST-009, takes the approved set and its buildability report. A set is handed over only once it is approved.

SEE ALSO

- approver
- candidate architecture decision

Dependency

MEANING

Something outside the product that the product relies on, recorded in section 11 with what happens if it is not there, such as D-001, an assistant session the expert is already working in.

SEE ALSO

- assumption
- risk

NOT

A requirement's wait on an open question. A requirement that waits on a question names the question, and the question lists the requirement in what it blocks.

SECTION 15

Change log and approvals

DRAFT 63/100 1 BLOCKING

VERSION	DATE	AUTHOR	WHAT CHANGED
0.9.0	2026-10-27	Richard Hightower	Six glossary terms, FR-017 without view names, FR-010's wait on Q-007, the status-quo edge; counted evidence for P-001 to P-004; C-004 moved out of section 9; FR-008 and FR-012 restated; interview, handover, and dependency defined; the versioning policy
0.8.1	2026-10-27	Richard Hightower	Glossary terms, observable scenario outcomes, open questions named in section 7, competitor edges; C-001 to C-003 and S-003 moved; ST-009 for the engineering lead
0.8.0	2026-10-27	Richard Hightower	Requirement and scenario priorities, SC-007, release exit criteria, single deciders
0.7.0	2026-10-27	Richard Hightower	Root objective, primary persona and approver; release id in scope; LC-001 rule
0.6.0	2026-10-27	Richard Hightower	Links from the specification review: problems, objectives, competitors, decision stages, and scenario views
0.4.0	2026-10-26	Richard Hightower	Navigation and refresh requirements; wireframes rewritten from their grading
0.3.0	2026-10-26	Richard Hightower	Wireframes for every screen and report, and recorded reasons for the grades
0.2.0	2026-10-26	Richard Hightower	Every defect the grading and the buildability review found, over two rounds
0.1.0	2026-10-23	Richard Hightower	First draft of all sixteen sections

Versioning policy before 1.0

The set and the product share one version number. The number changes by these rules:

- A patch, such as 0.9.0 to 0.9.1, is for fixes that add no content.
- A minor version, such as 0.8 to 0.9, is for new content in the set or in the plugin.
- A change to the golden thread takes the set to the next major version. From 0.x, the next major version is 1.0.

The expert set this policy at 0.9.0. Two earlier versions do not follow it, and neither is reissued:

- 0.8.1 added stage ST-009 and moved items between sections. Under this policy it would have been the next minor version.
- 0.2.0 changed the golden thread. Under this policy it takes the set to 1.0.

0.2.0, what the grading changed

The reviewers that graded 0.1.0 were the product's own, run against its own PRD. They found twelve blocking defects and four failed derivability checks. The largest single finding was that FR-009, FR-010 and FR-011 were written against

"its rubric" and "its threshold" with neither defined anywhere in the set: the threshold is now stated as 70 in the requirements themselves, and the rubric, the section states and the derivability checks are defined in section 14 as the product content they are.

Two findings were checked and **rejected**: the claim that "see section 6" pointed at the wrong document, and the claim that the section count was contradicted. Both came from reading file numbers as section numbers. The references were right; the numbering was unexplained, which is now fixed in the root document.

0.1.0, what to know

Written as prd-forge's own first exercise. Three things a reader should carry forward:

1. **No baselines.** All three objectives carry an explicit unknown with a matching open question (Q-001 to Q-003). The targets are intentions, not commitments.
2. **C-001 is the tension.** The product is delivered inside an assistant session. Naming that in a requirement would fail test 5, so it is recorded as a constraint. A reviewer may fairly call it architecture. At 0.8.1 it moved to section 6 as a non-goal, and the session it needs is dependency D-001 in section 11.
3. **R-003 is not closed.** Nothing in this product checks whether the product being specified should exist at all.

Approvals

None sought. This set is at `draft` and has not been offered for approval. It cannot reach `final` without an approval recorded here, per TR-001.

> **Fixes the blocking defects the 0.8.1 grading found that do not wait on the expert. The glo**
>
> VERSION 0.9.0
>
> DATE 2026-10-27
>
> AUTHOR Richard Hightower
>
> SECTIONS
> - 1
> - 2
> - 7
> - 9
> - 14
> - 15

> **Fixes the blocking defects that do not wait on evidence. The glossary defines buildability**
>
> VERSION 0.8.1
>
> DATE 2026-10-27
>
> AUTHOR Richard Hightower

SECTIONS

- 2
- 4
- 5
- 6
- 7
- 8
- 9
- 10
- 11
- 12
- 13
- 14

Sets the expert's priorities. Ten requirements stay must and eight become should; SC-002 a

VERSION 0.8.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS

- 5
- 7
- 12
- 13

Names the root's objective, primary persona, and approver, from section 3, the personas do

VERSION 0.7.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS

- 4
- 5
- 6

Adds the links the specification review asked for. Each problem names its persona, and eac

VERSION 0.6.0

DATE 2026-10-27

AUTHOR Richard Hightower

SECTIONS

- 1
- 2
- 3
- 5
- 10
- 11

Adds FR-017, moving between sections without losing one's place, and FR-018, seeing the se

VERSION	0.4.0
DATE	2026-10-26
AUTHOR	Richard Hightower
SECTIONS	• 6 • 7 • 8 • 13 • 14

Adds wireframes to section 7. Each functional requirement now says where its user meets th

VERSION	0.3.0
DATE	2026-10-26
AUTHOR	Richard Hightower
SECTIONS	• 5 • 7 • 8 • 14

Fixes every defect the adversarial grading and the buildability review found in 0.1.0. J-0

VERSION	0.2.0
DATE	2026-10-26
AUTHOR	Richard Hightower
SECTIONS	• 4 • 5 • 6 • 7 • 8 • 10 • 13 • 14

First draft of the whole set, written with prd-forge's own conventions as the product's fi

VERSION	0.1.0
DATE	2026-10-23
AUTHOR	Richard Hightower
SECTIONS	1 2 3 4 5 6 7 8 9 10 11 12 13 14 15